

Combiner Function

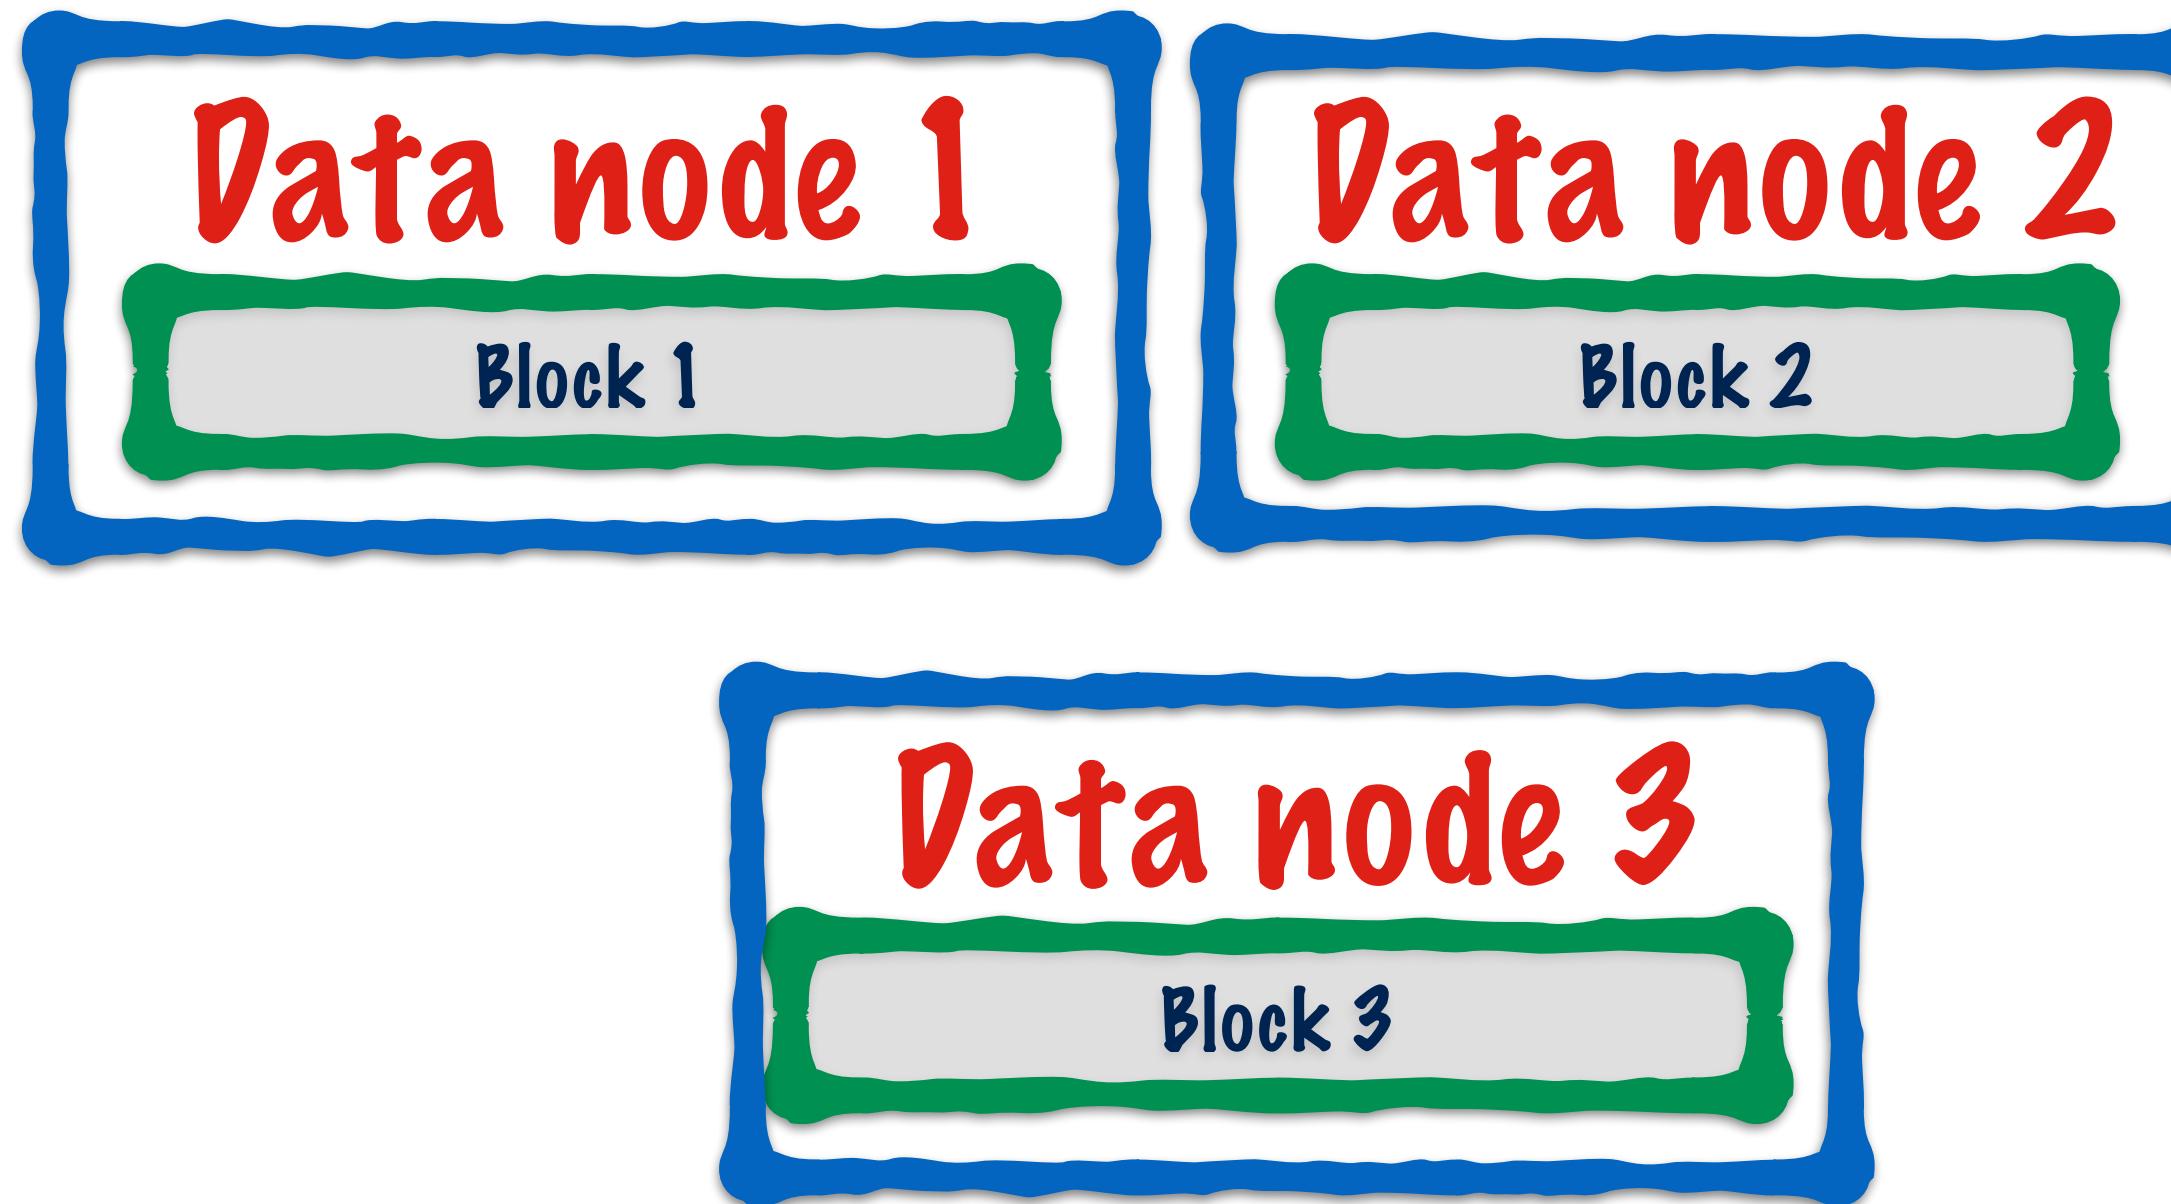
Let's go back to our Word Count example

Objective: Create a
Frequency
Distribution of
words in a text file

Word	Count
because	1
each	4
figure	9
..	..

RECAP

The text file has been divided into blocks and stored in HDFS



Name node

The name node stores metadata

Here is how the data flows in a MapReduce job

RECAP

Data node 1

Hey Diddle Diddle

...

Data node 2

The cat and the fiddle

....

Data node 3

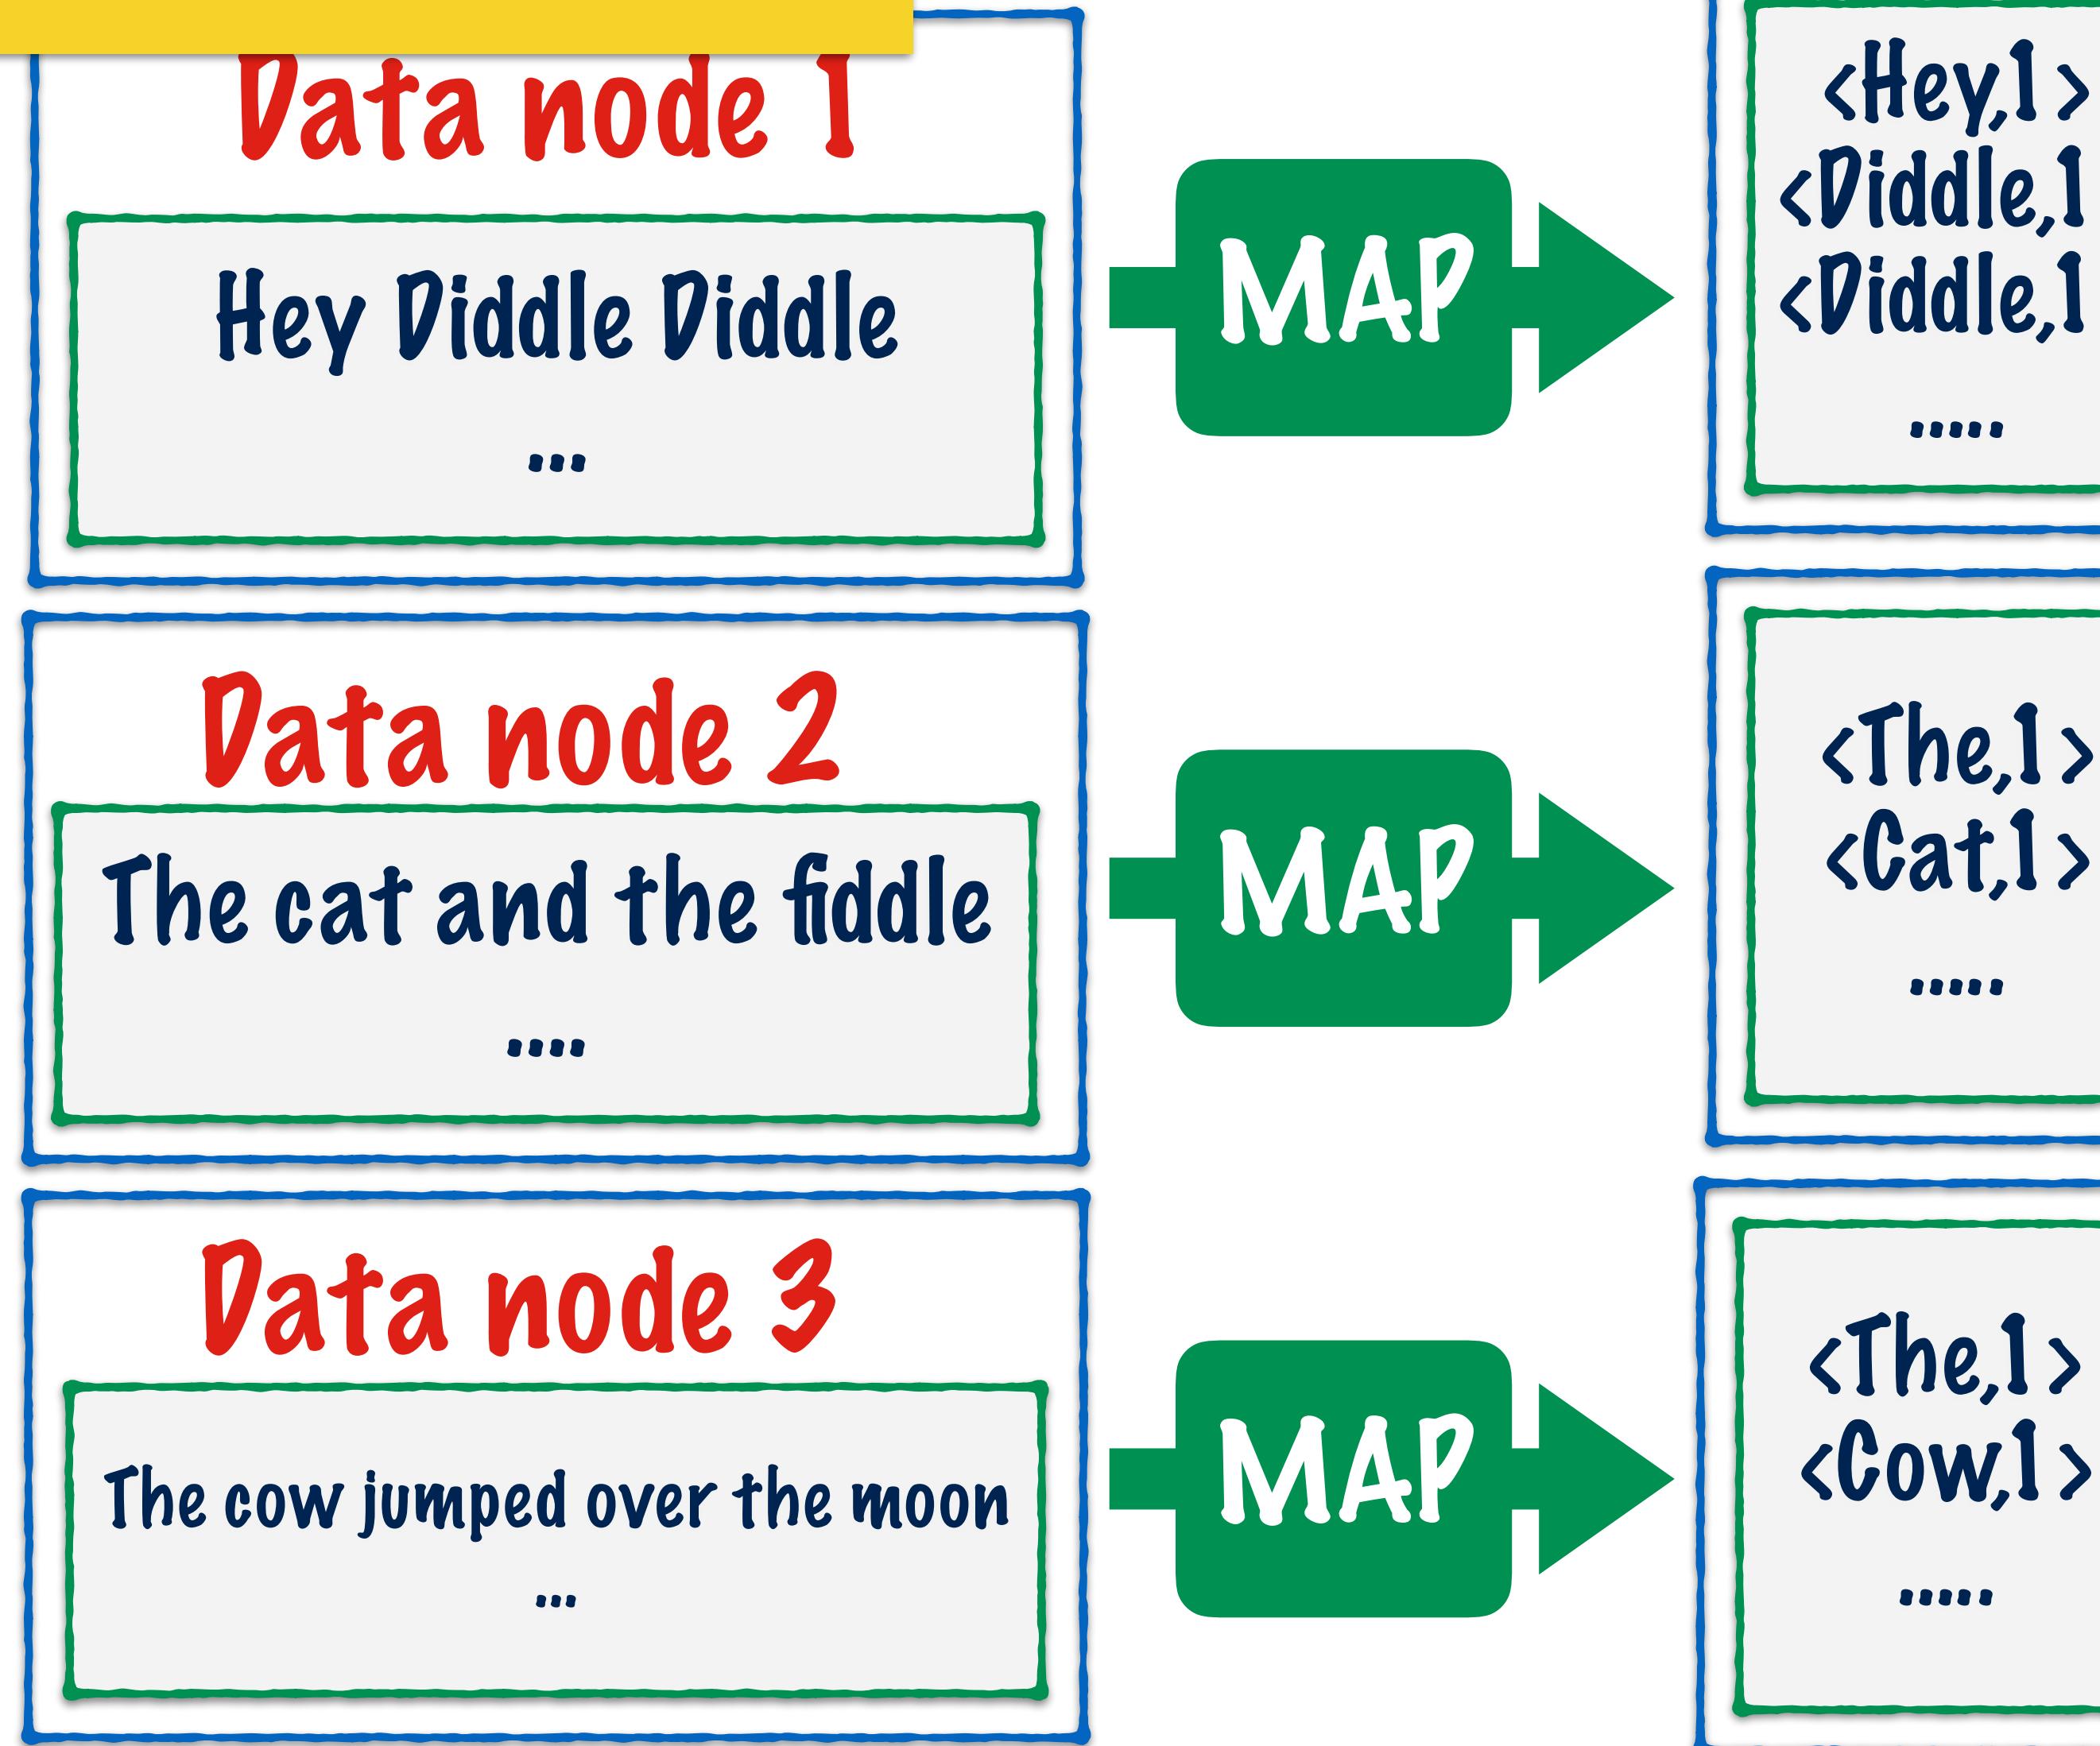
The cow jumped over the moon

...

Each block here would represent a part of the text file

Here is a snapshot of some text from each block

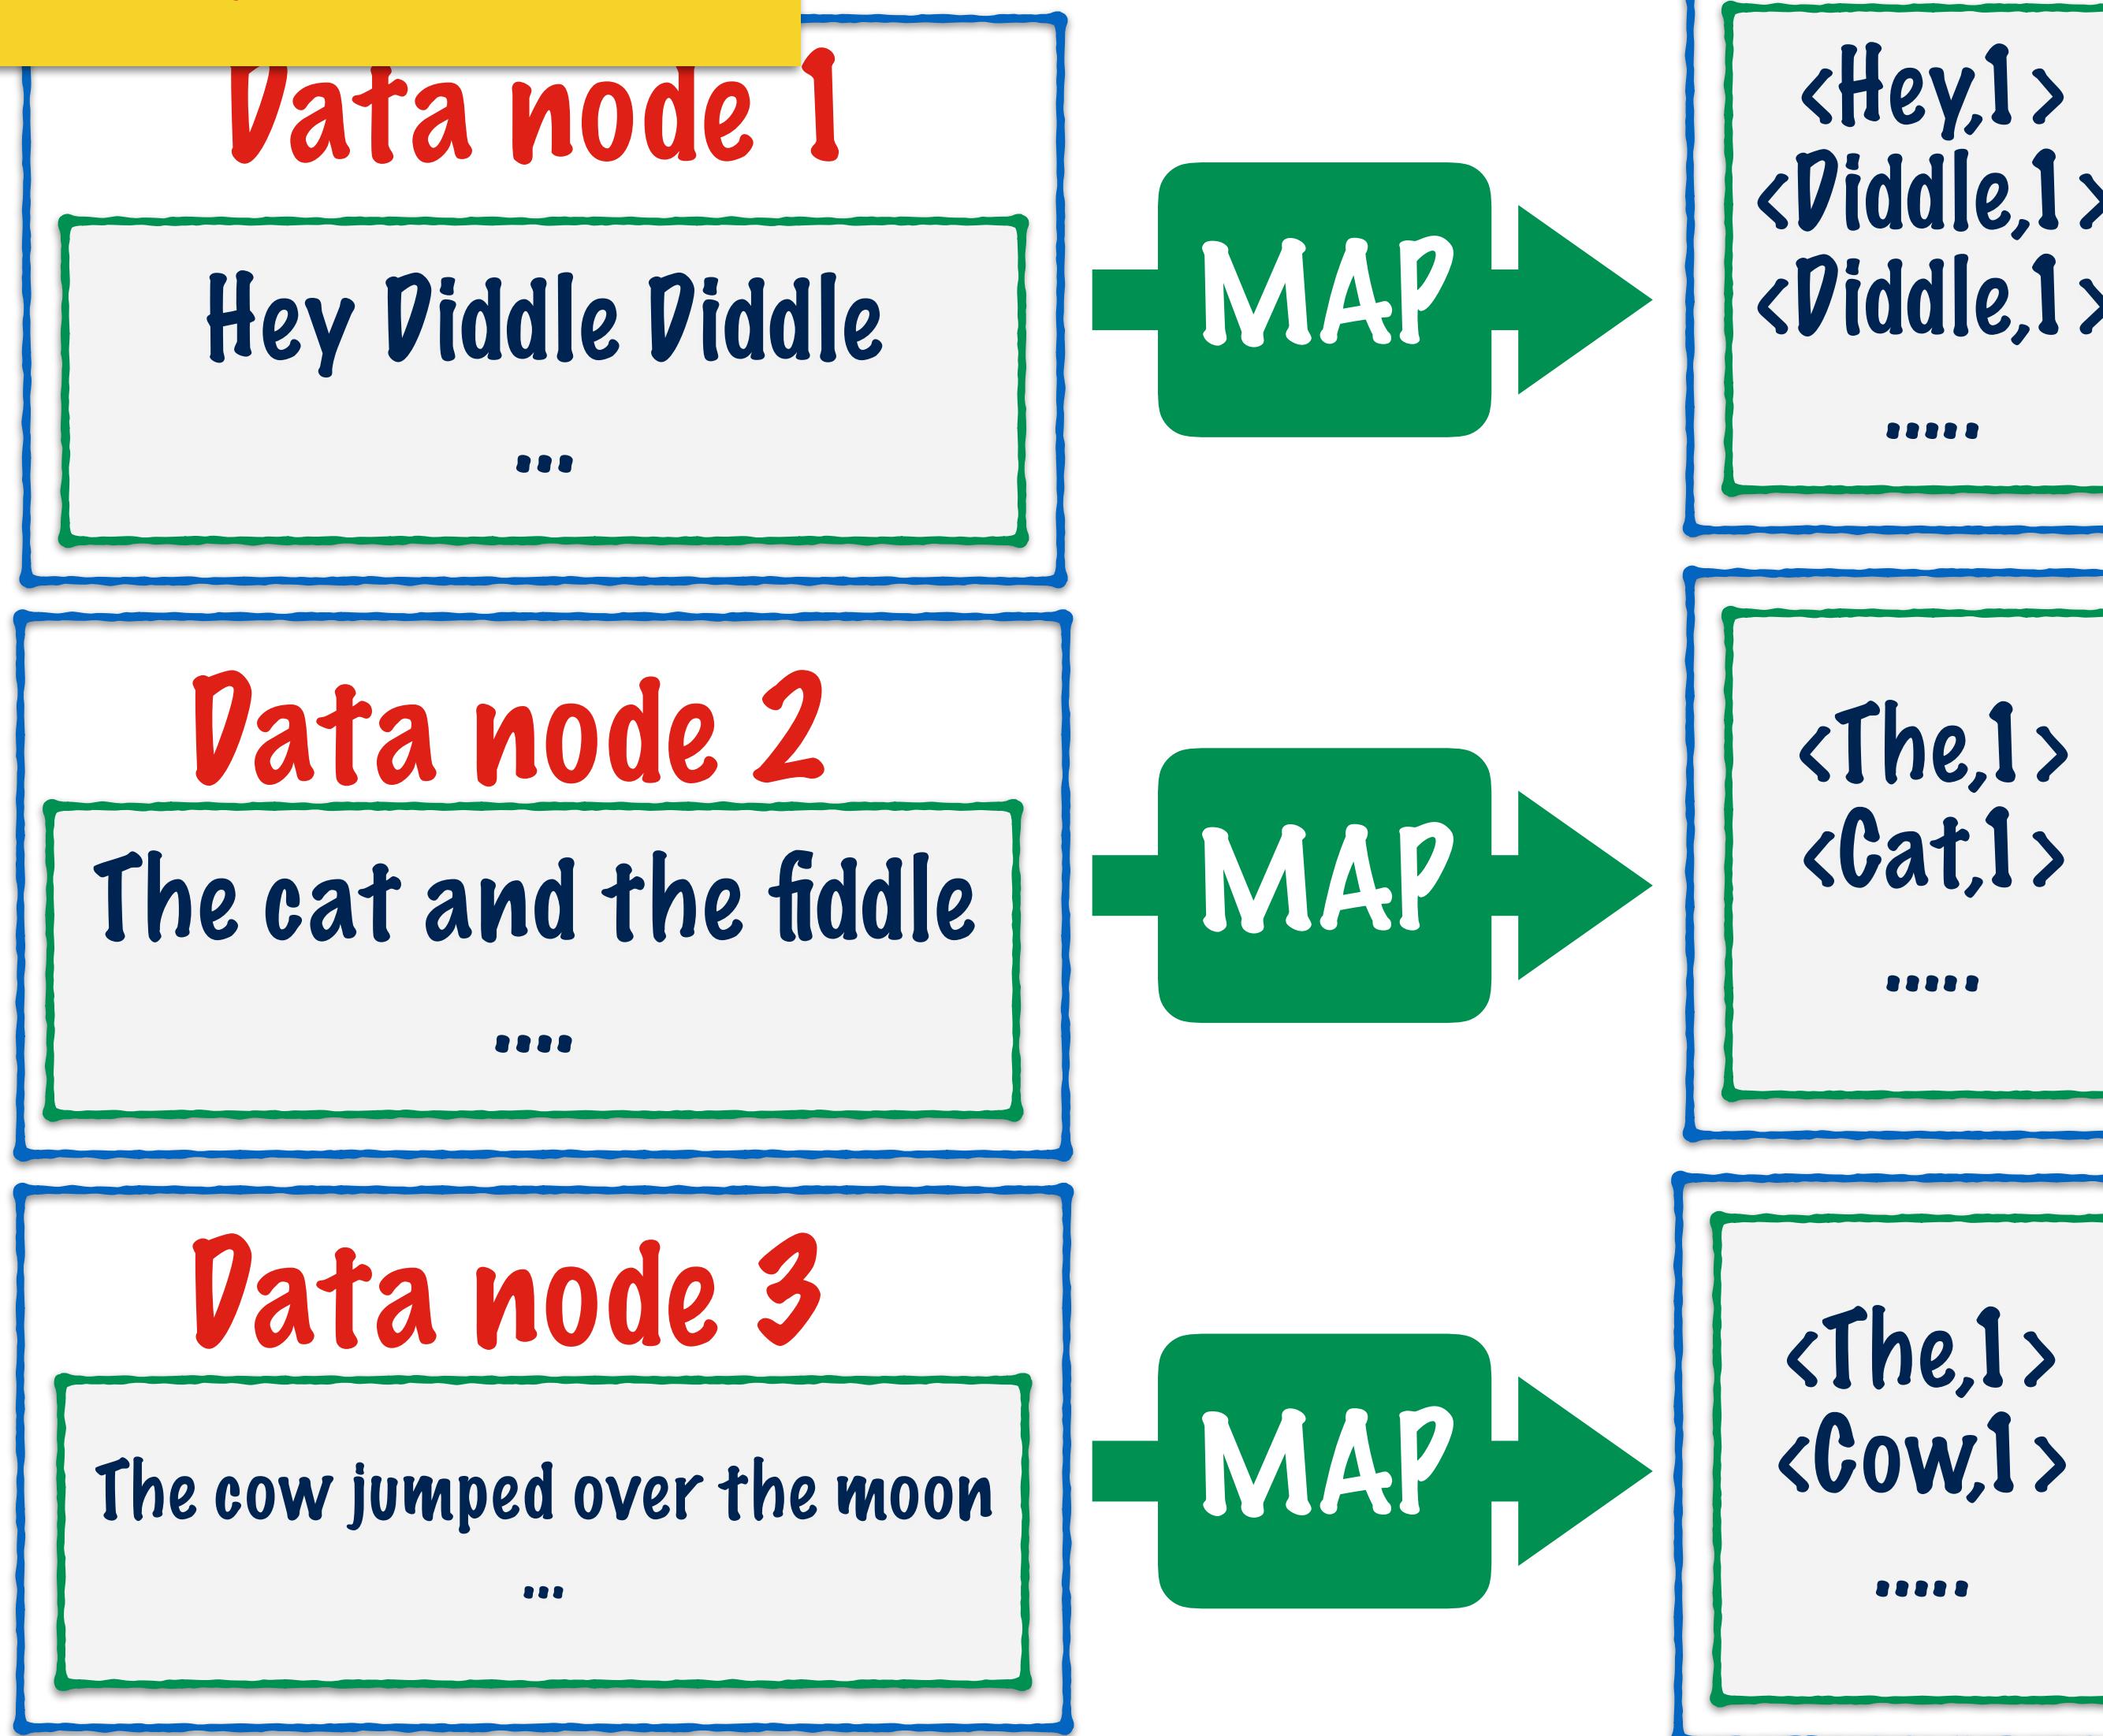
RECAP



The map step will generate a list of key-value pairs on each node

From now on all the inputs and outputs are formatted as <key,value> pairs

RECAP



These are all copied over to one single node

On that node an operation called **Sort/Merge** occurs

RECAP

Data node 1

Hey Diddle Diddle

...

Data node 2

The cat and the fiddle

....

Data node 3

The cow jumped over the moon

...

MAP

MAP

MAP

Sort/Merge

<Hey,1>
<Diddle,1>

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

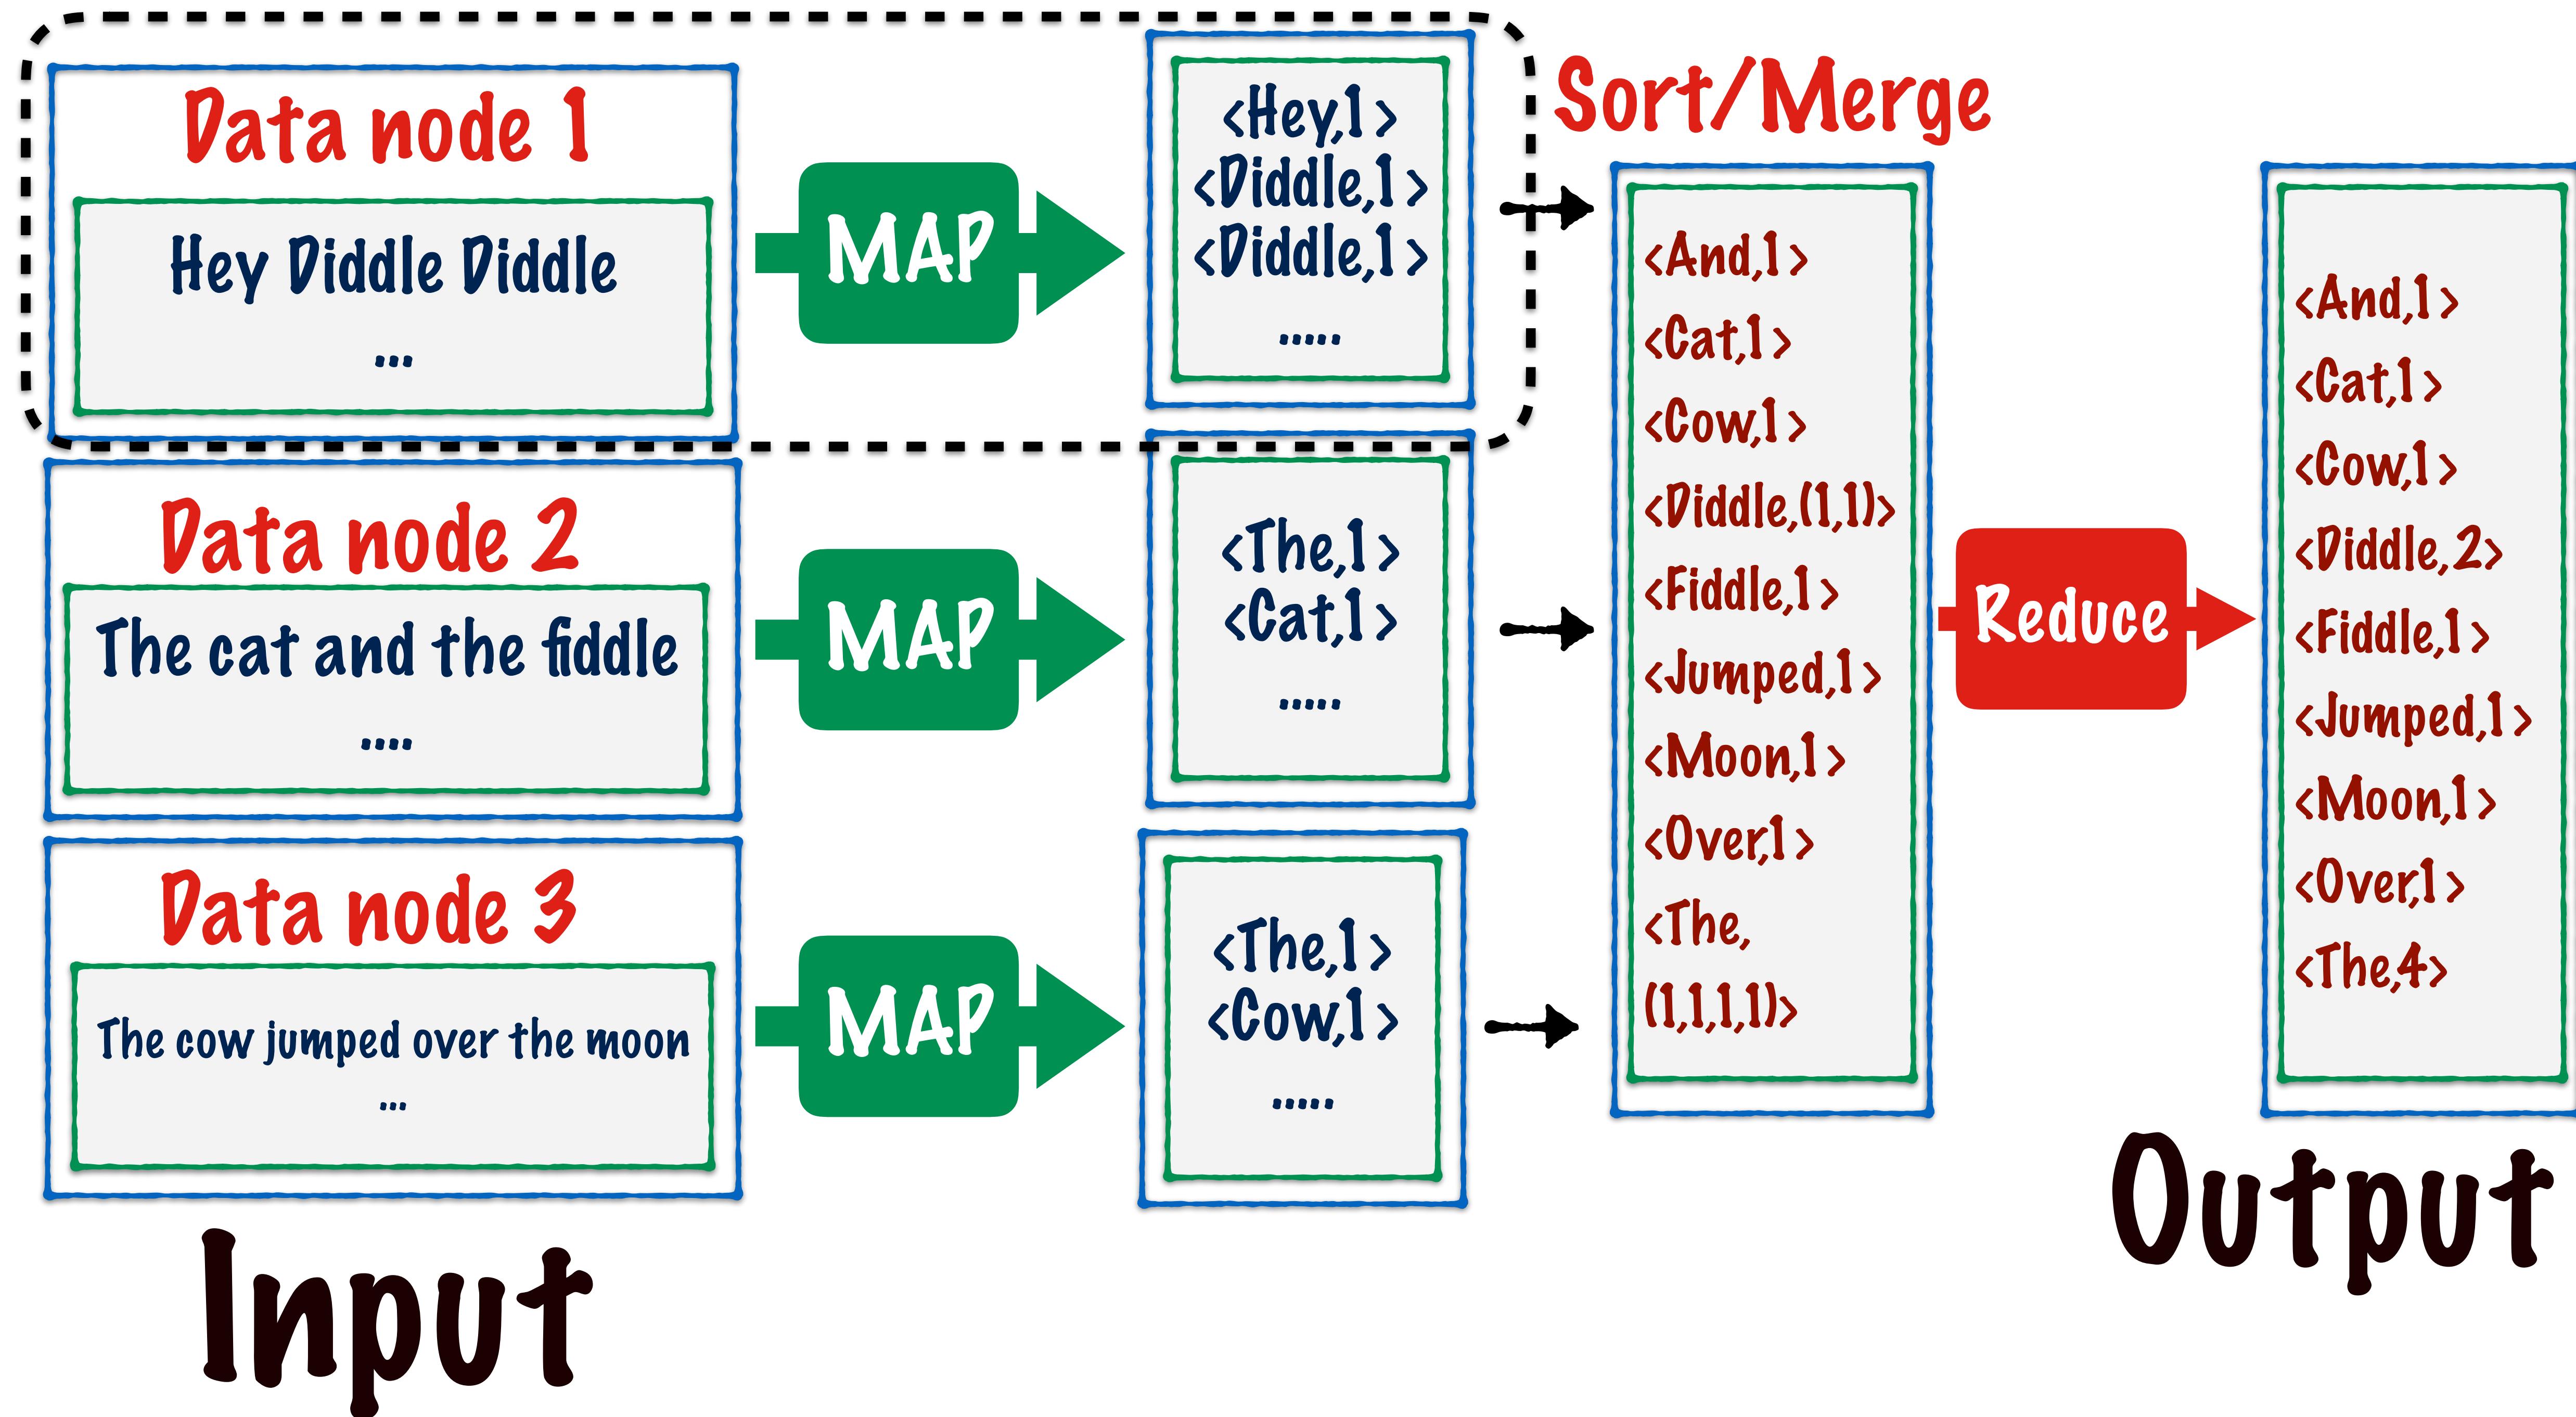
...

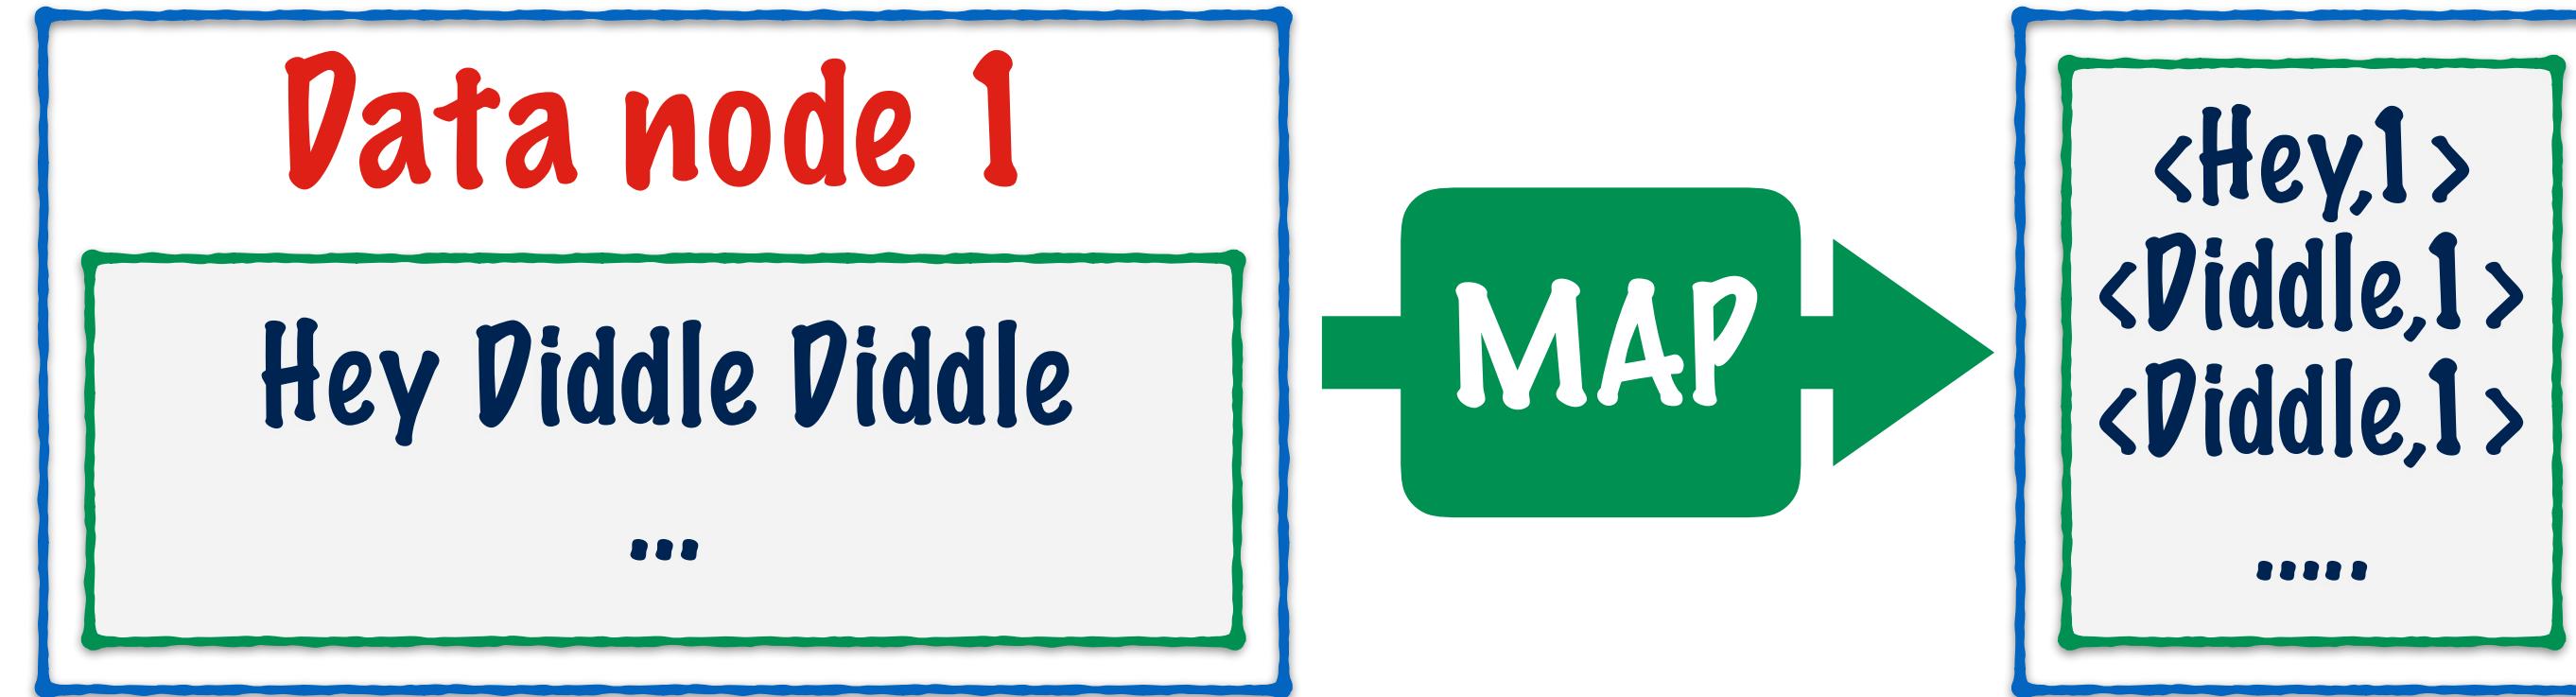
...

...

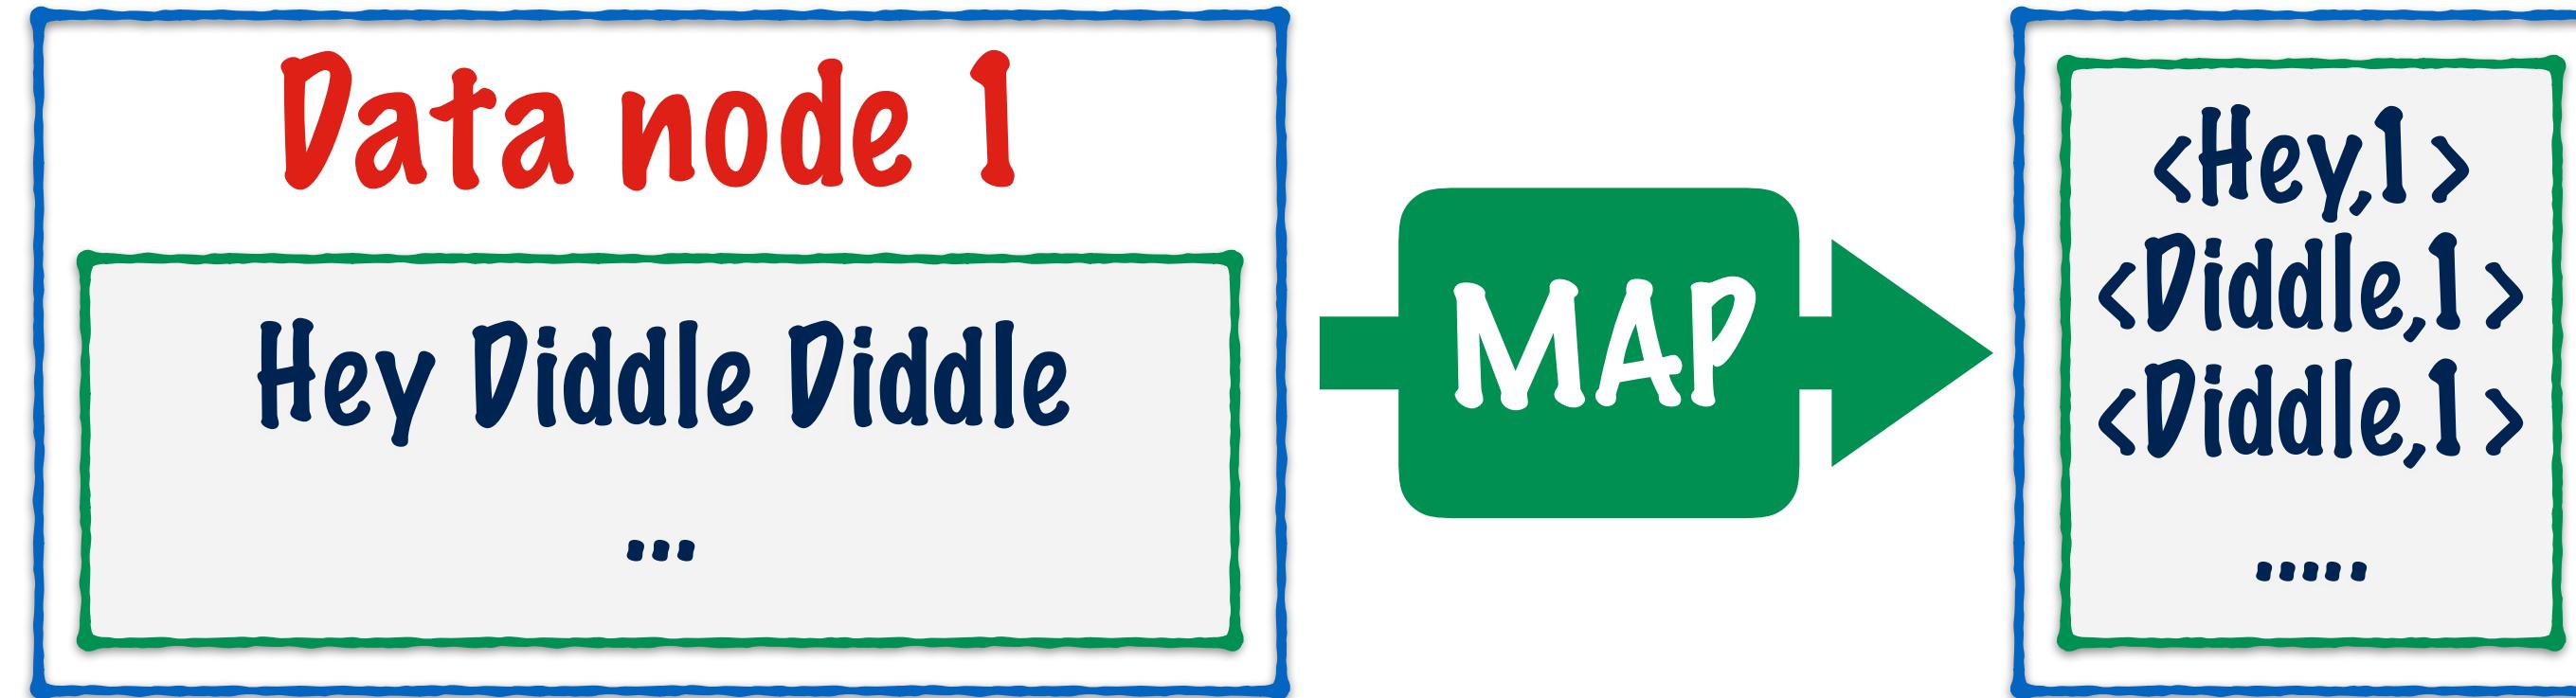
...

...



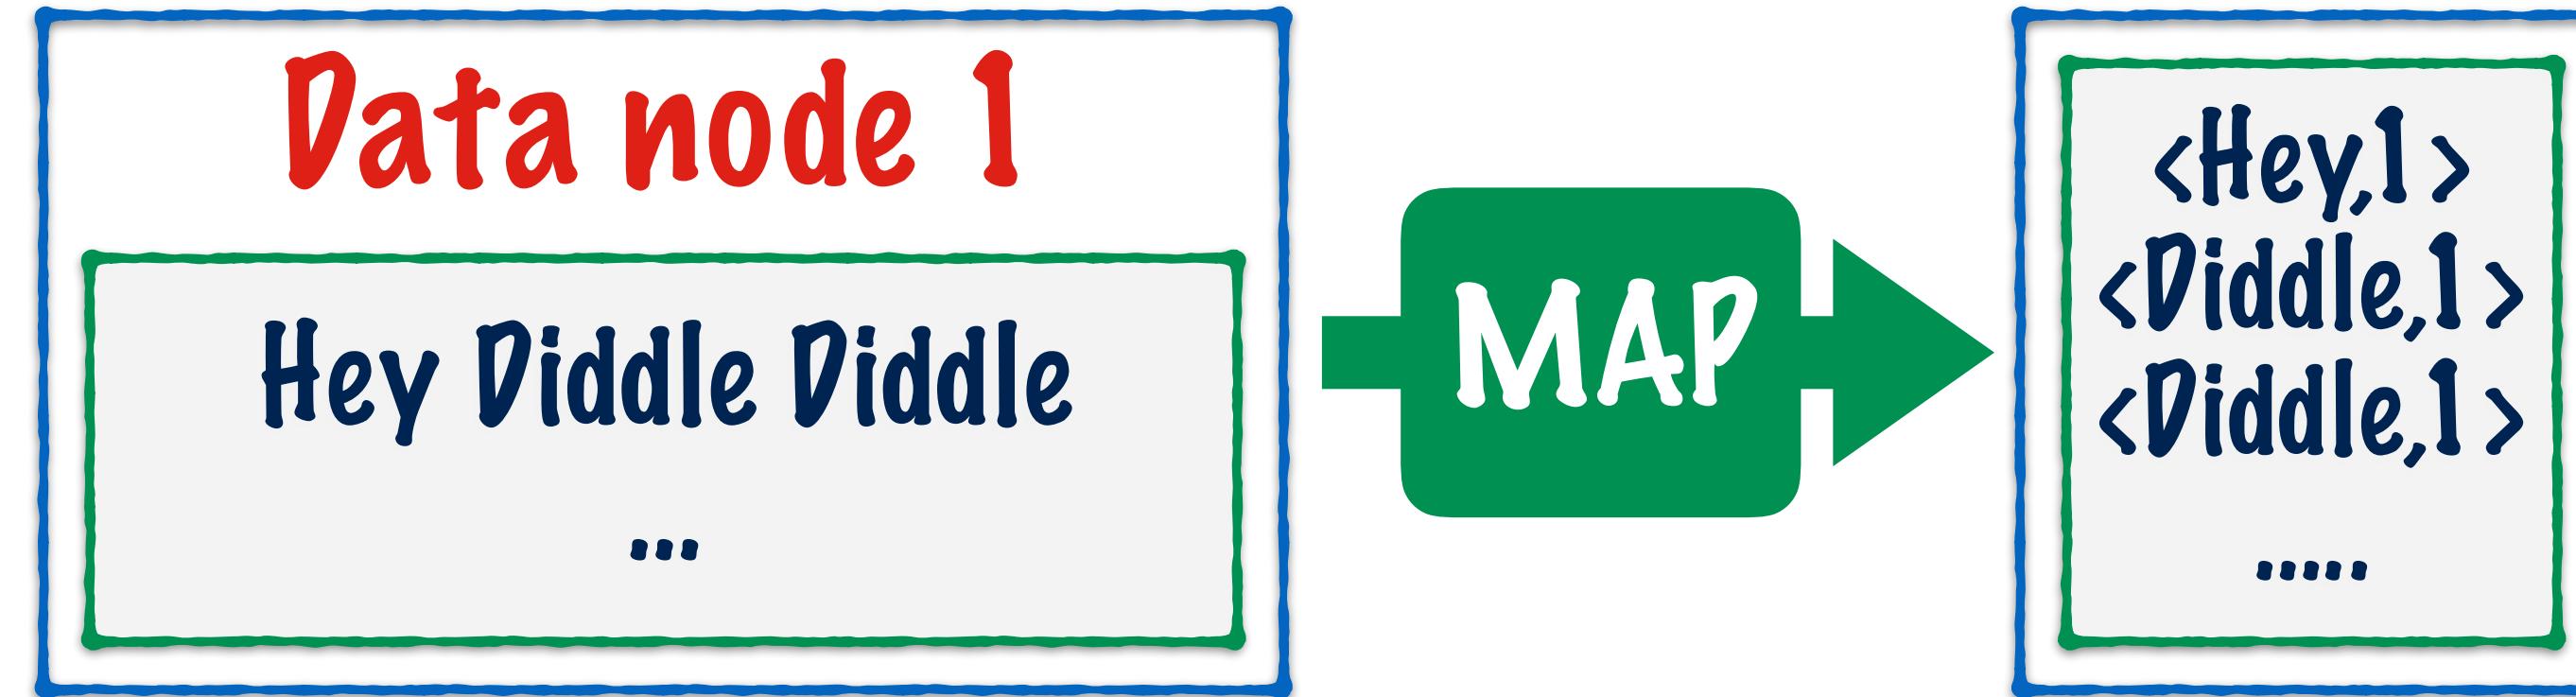


Let's concentrate
on the operation
in each data node



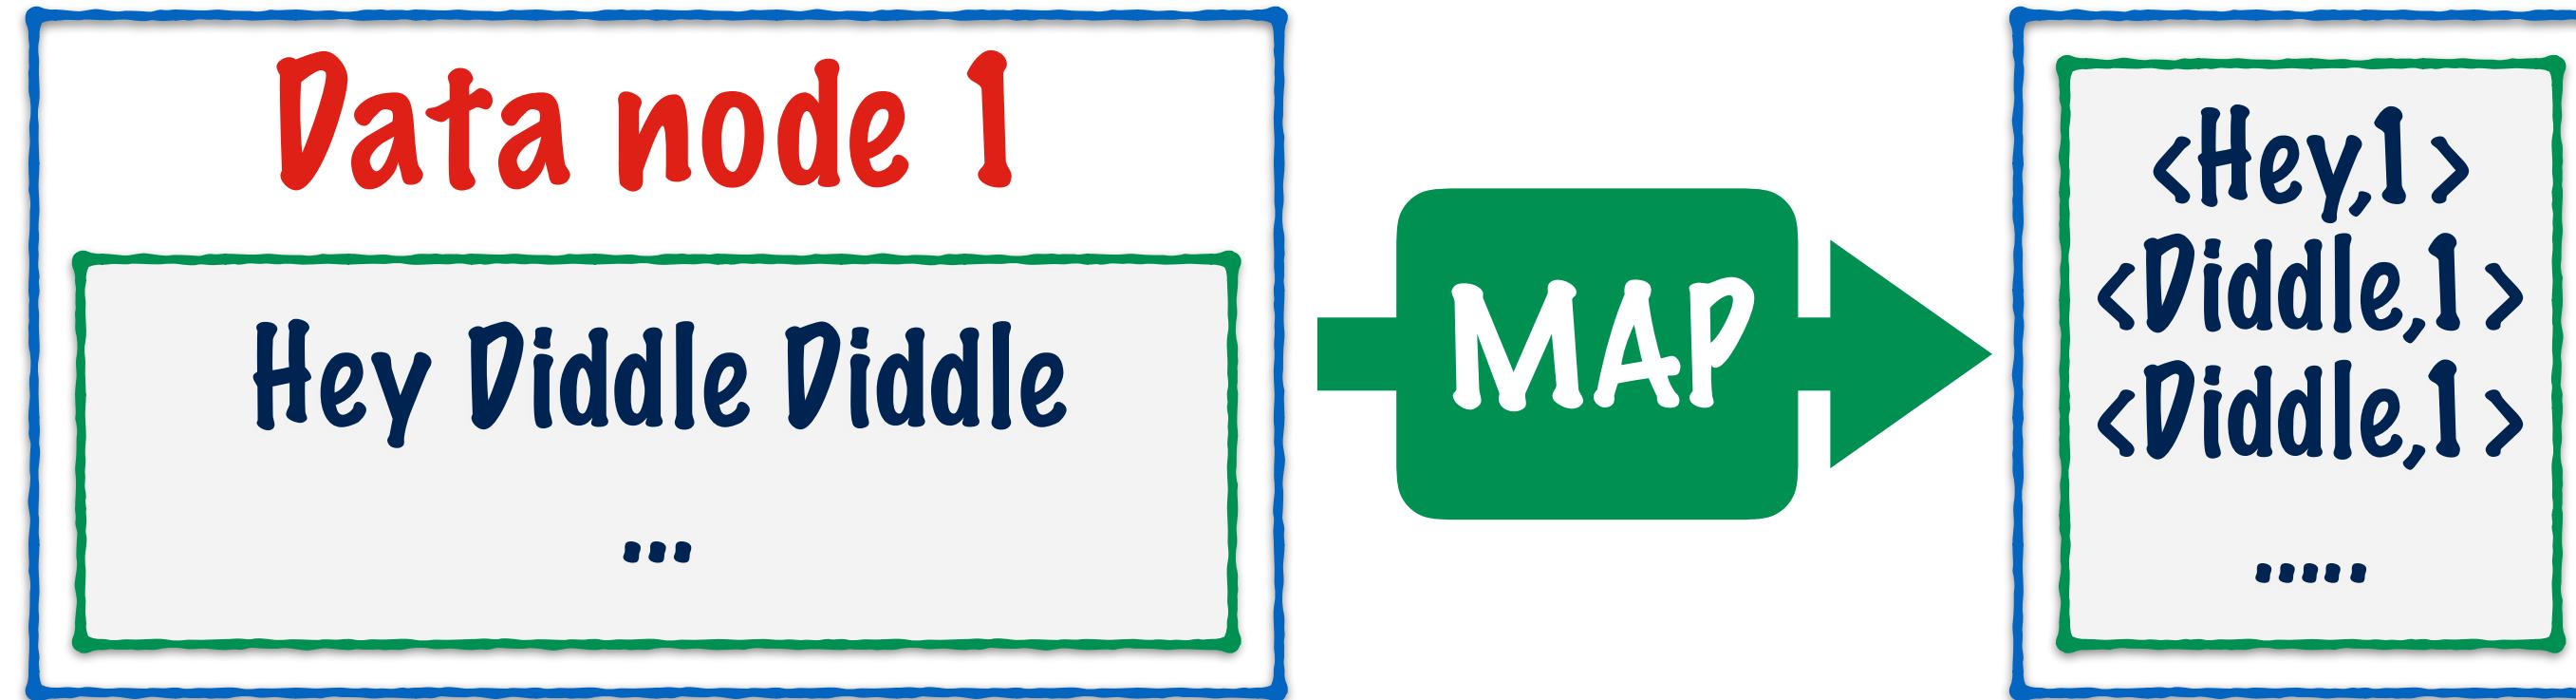
Map operation
produces a list of
key-value pairs

Reduce operation
combines pairs
with the same key

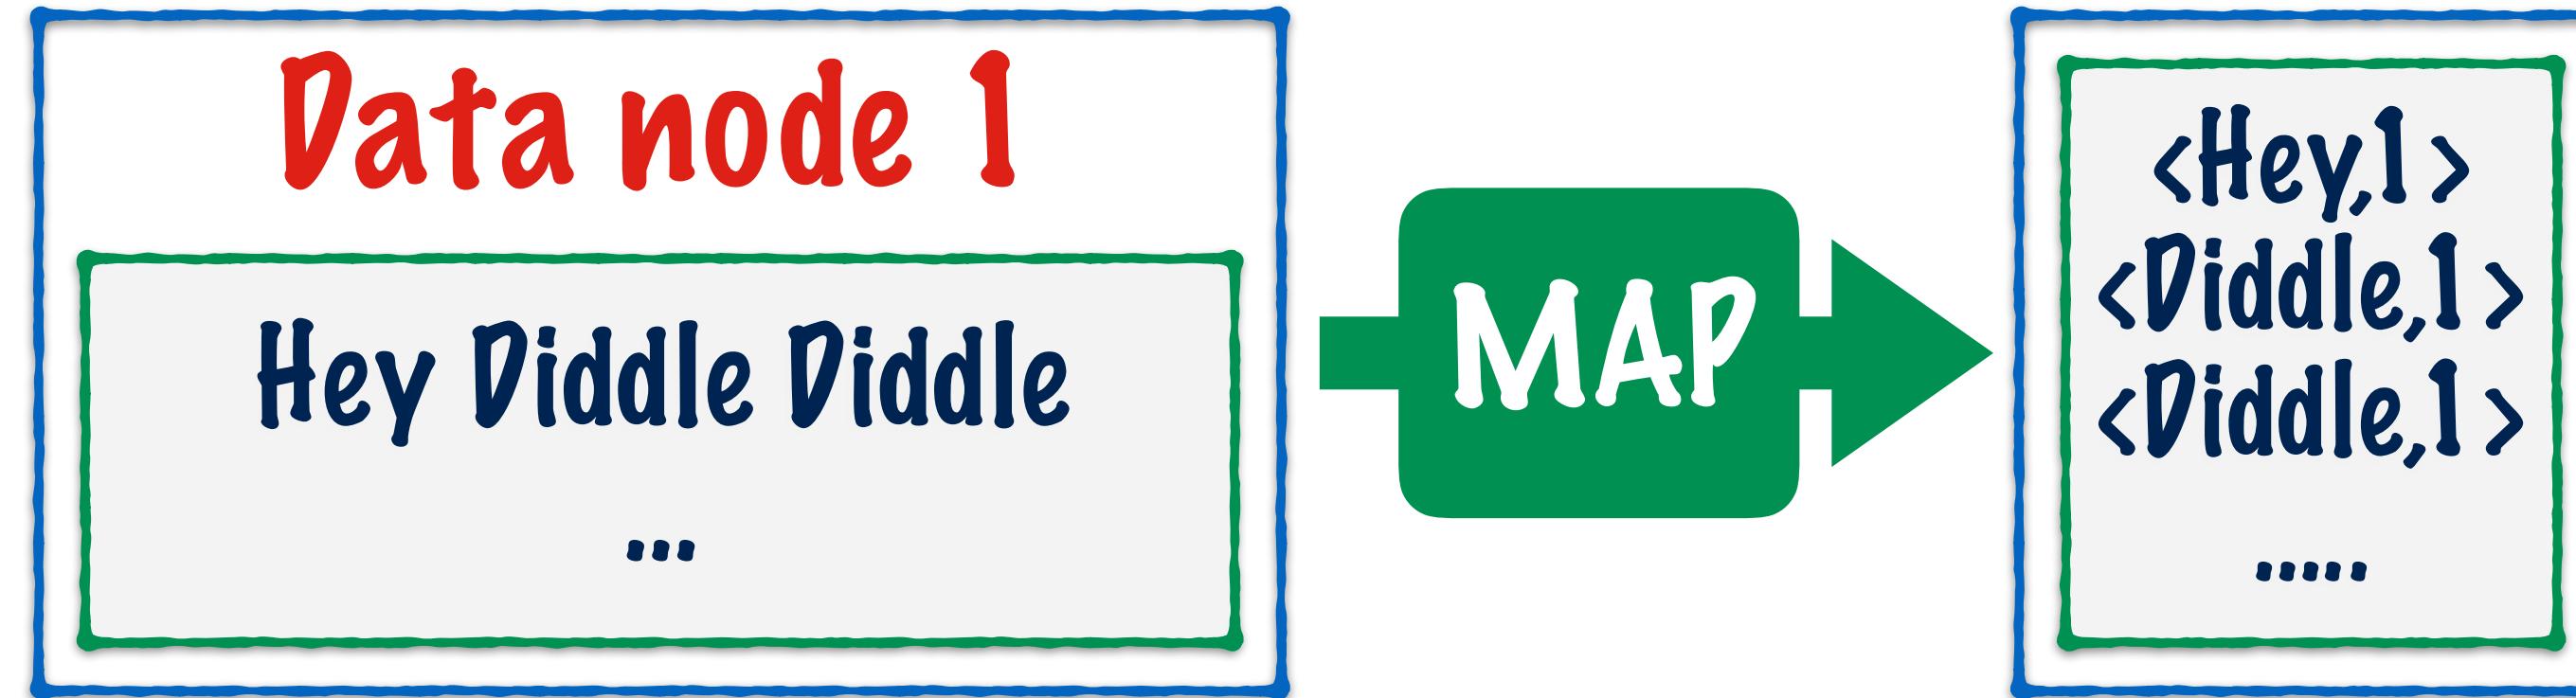


Map operation
is highly
parallelized

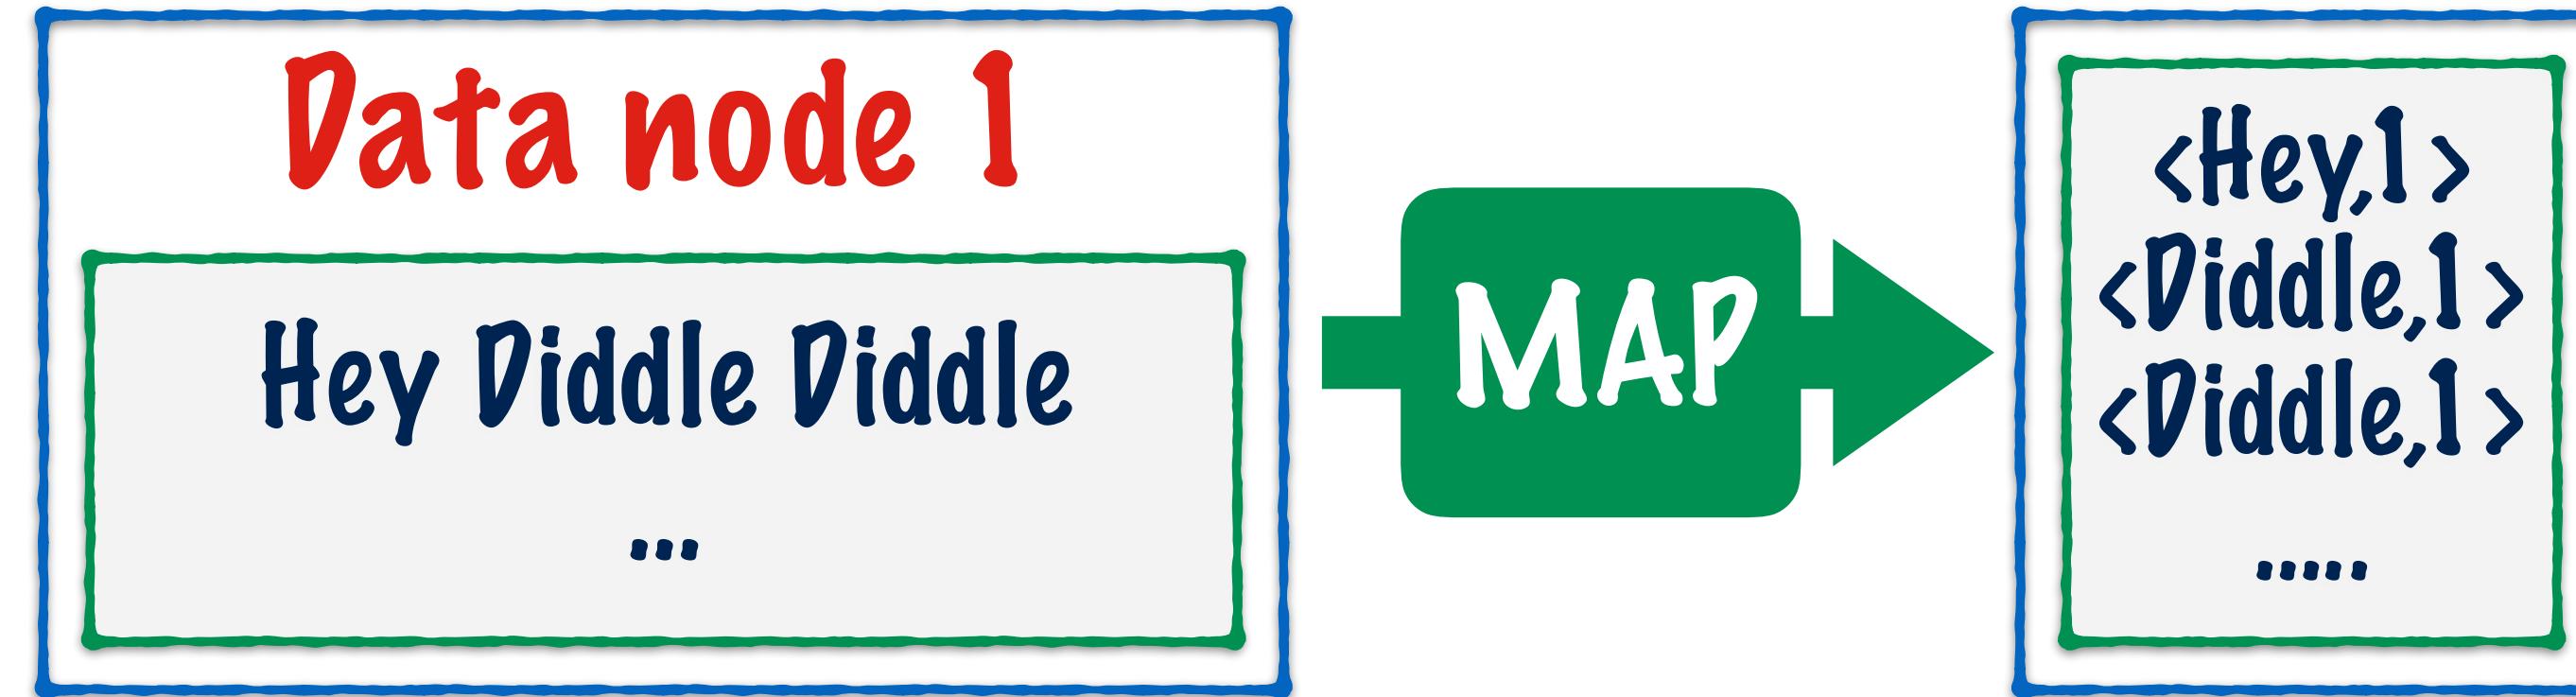
Reduce operation
currently has no
parallelization



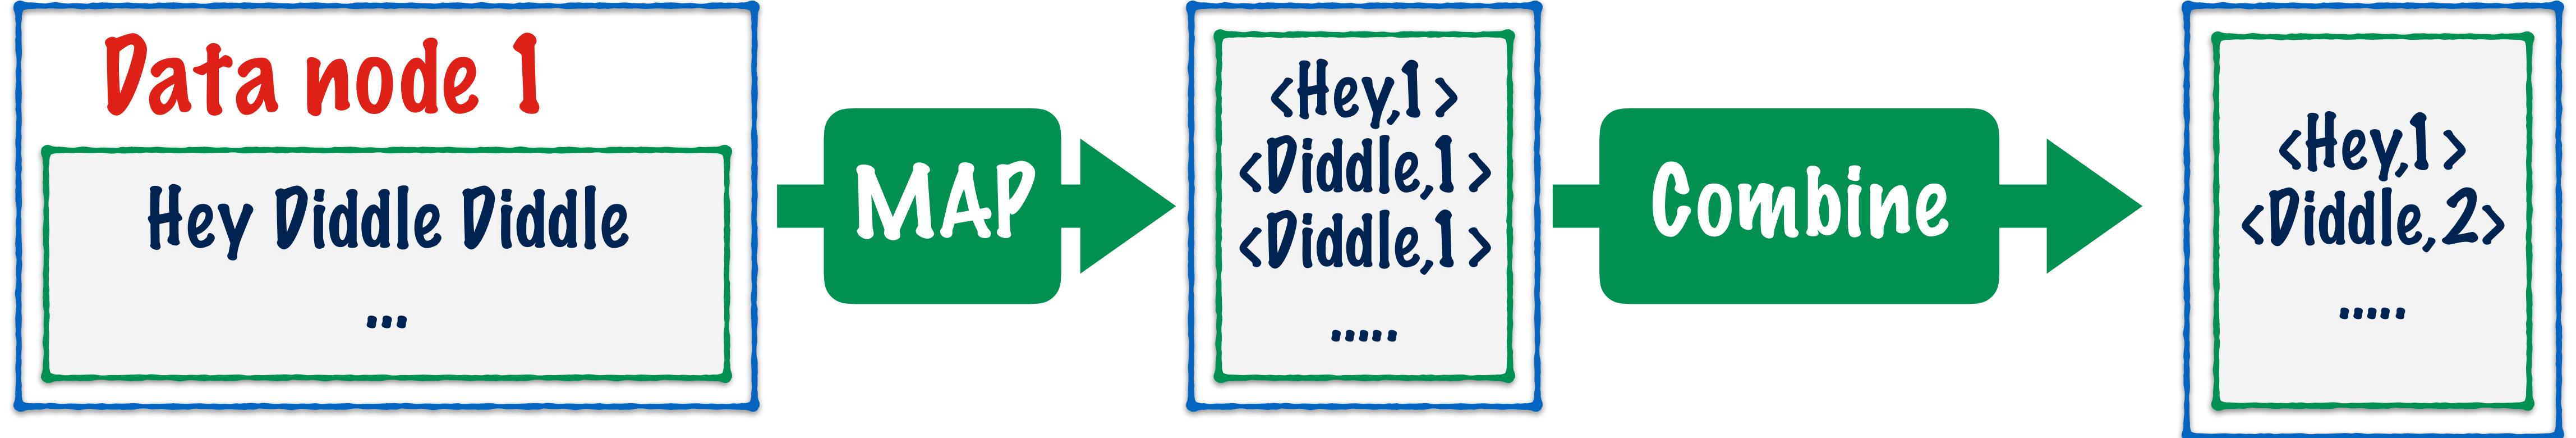
What if some part of
the **Reduce** operation
could happen on each
data node ie. **in parallel**?



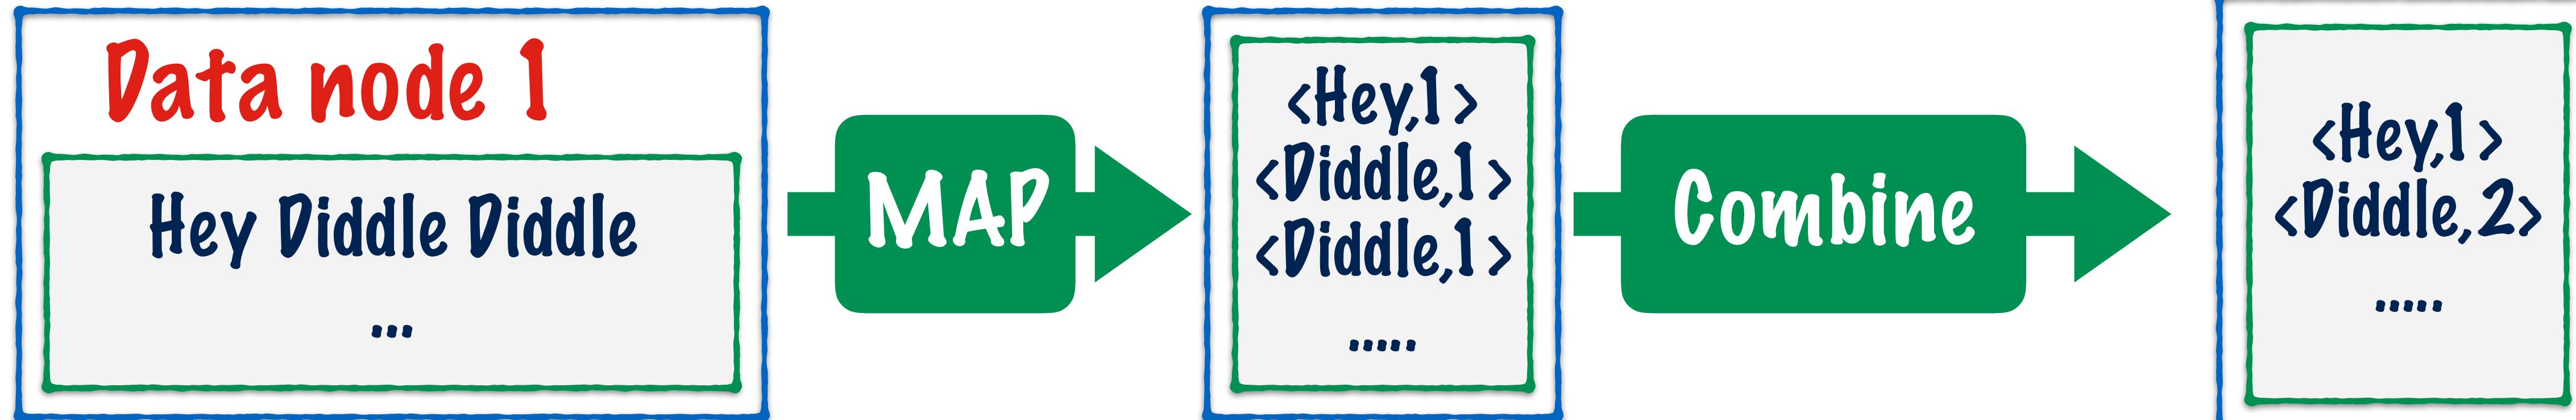
Essentially, we want to combine key-value pairs with same key - before they are copied over for reduce



Use a combiner
function



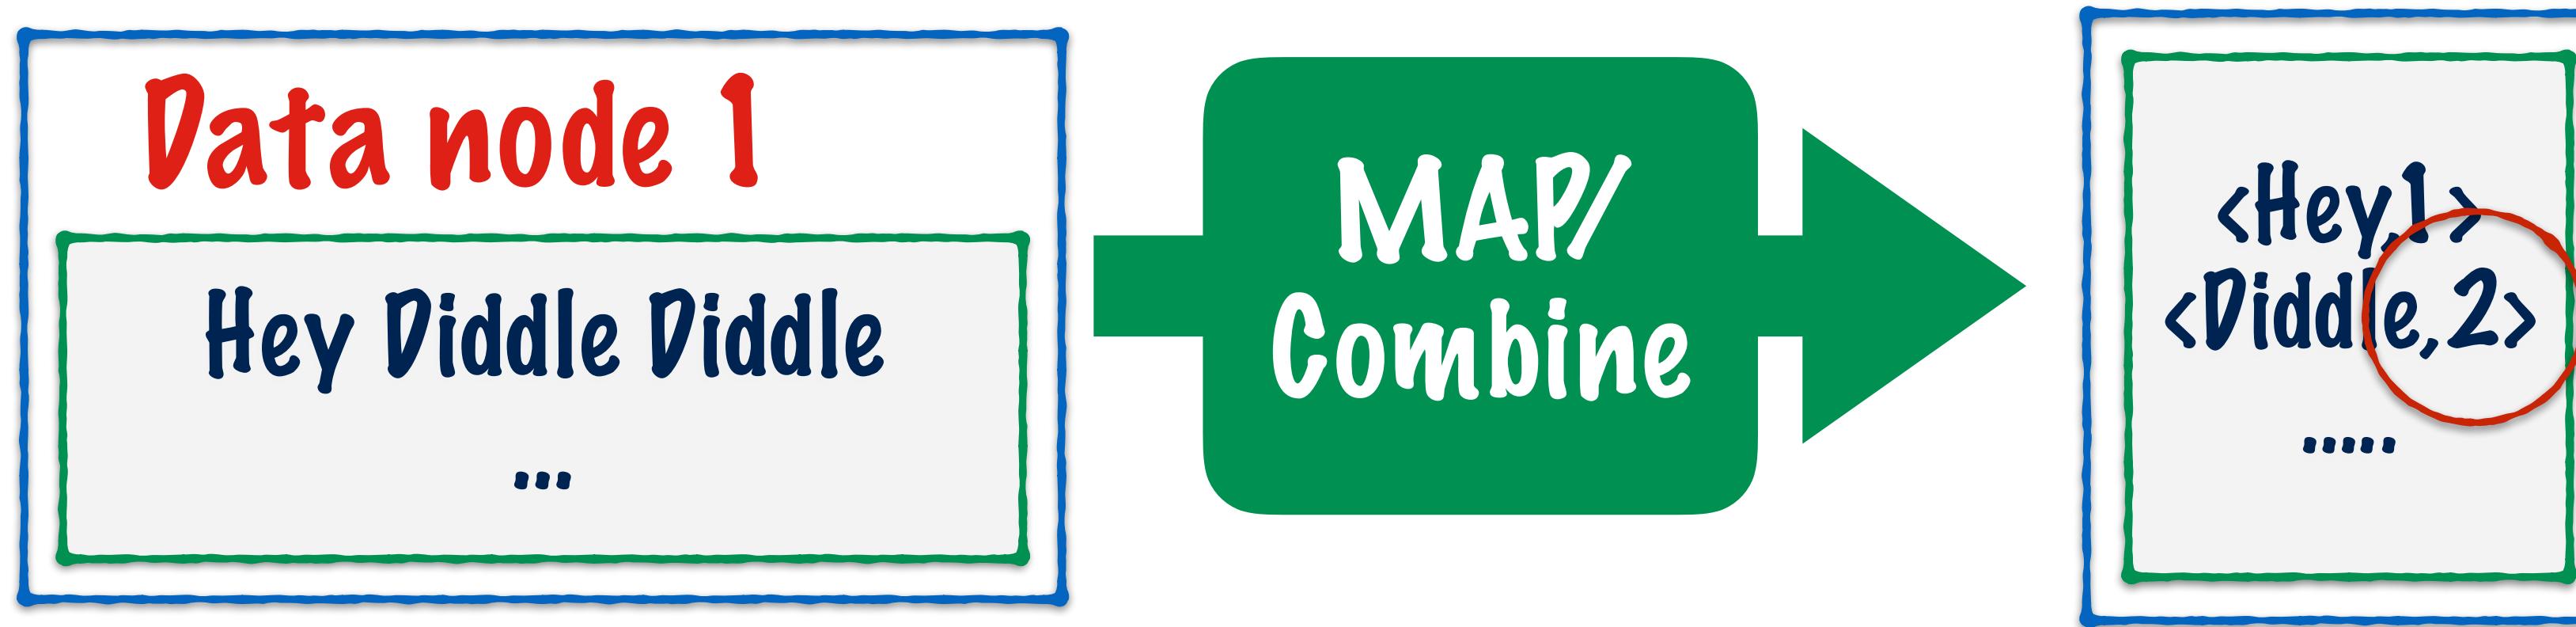
Use a combiner
function



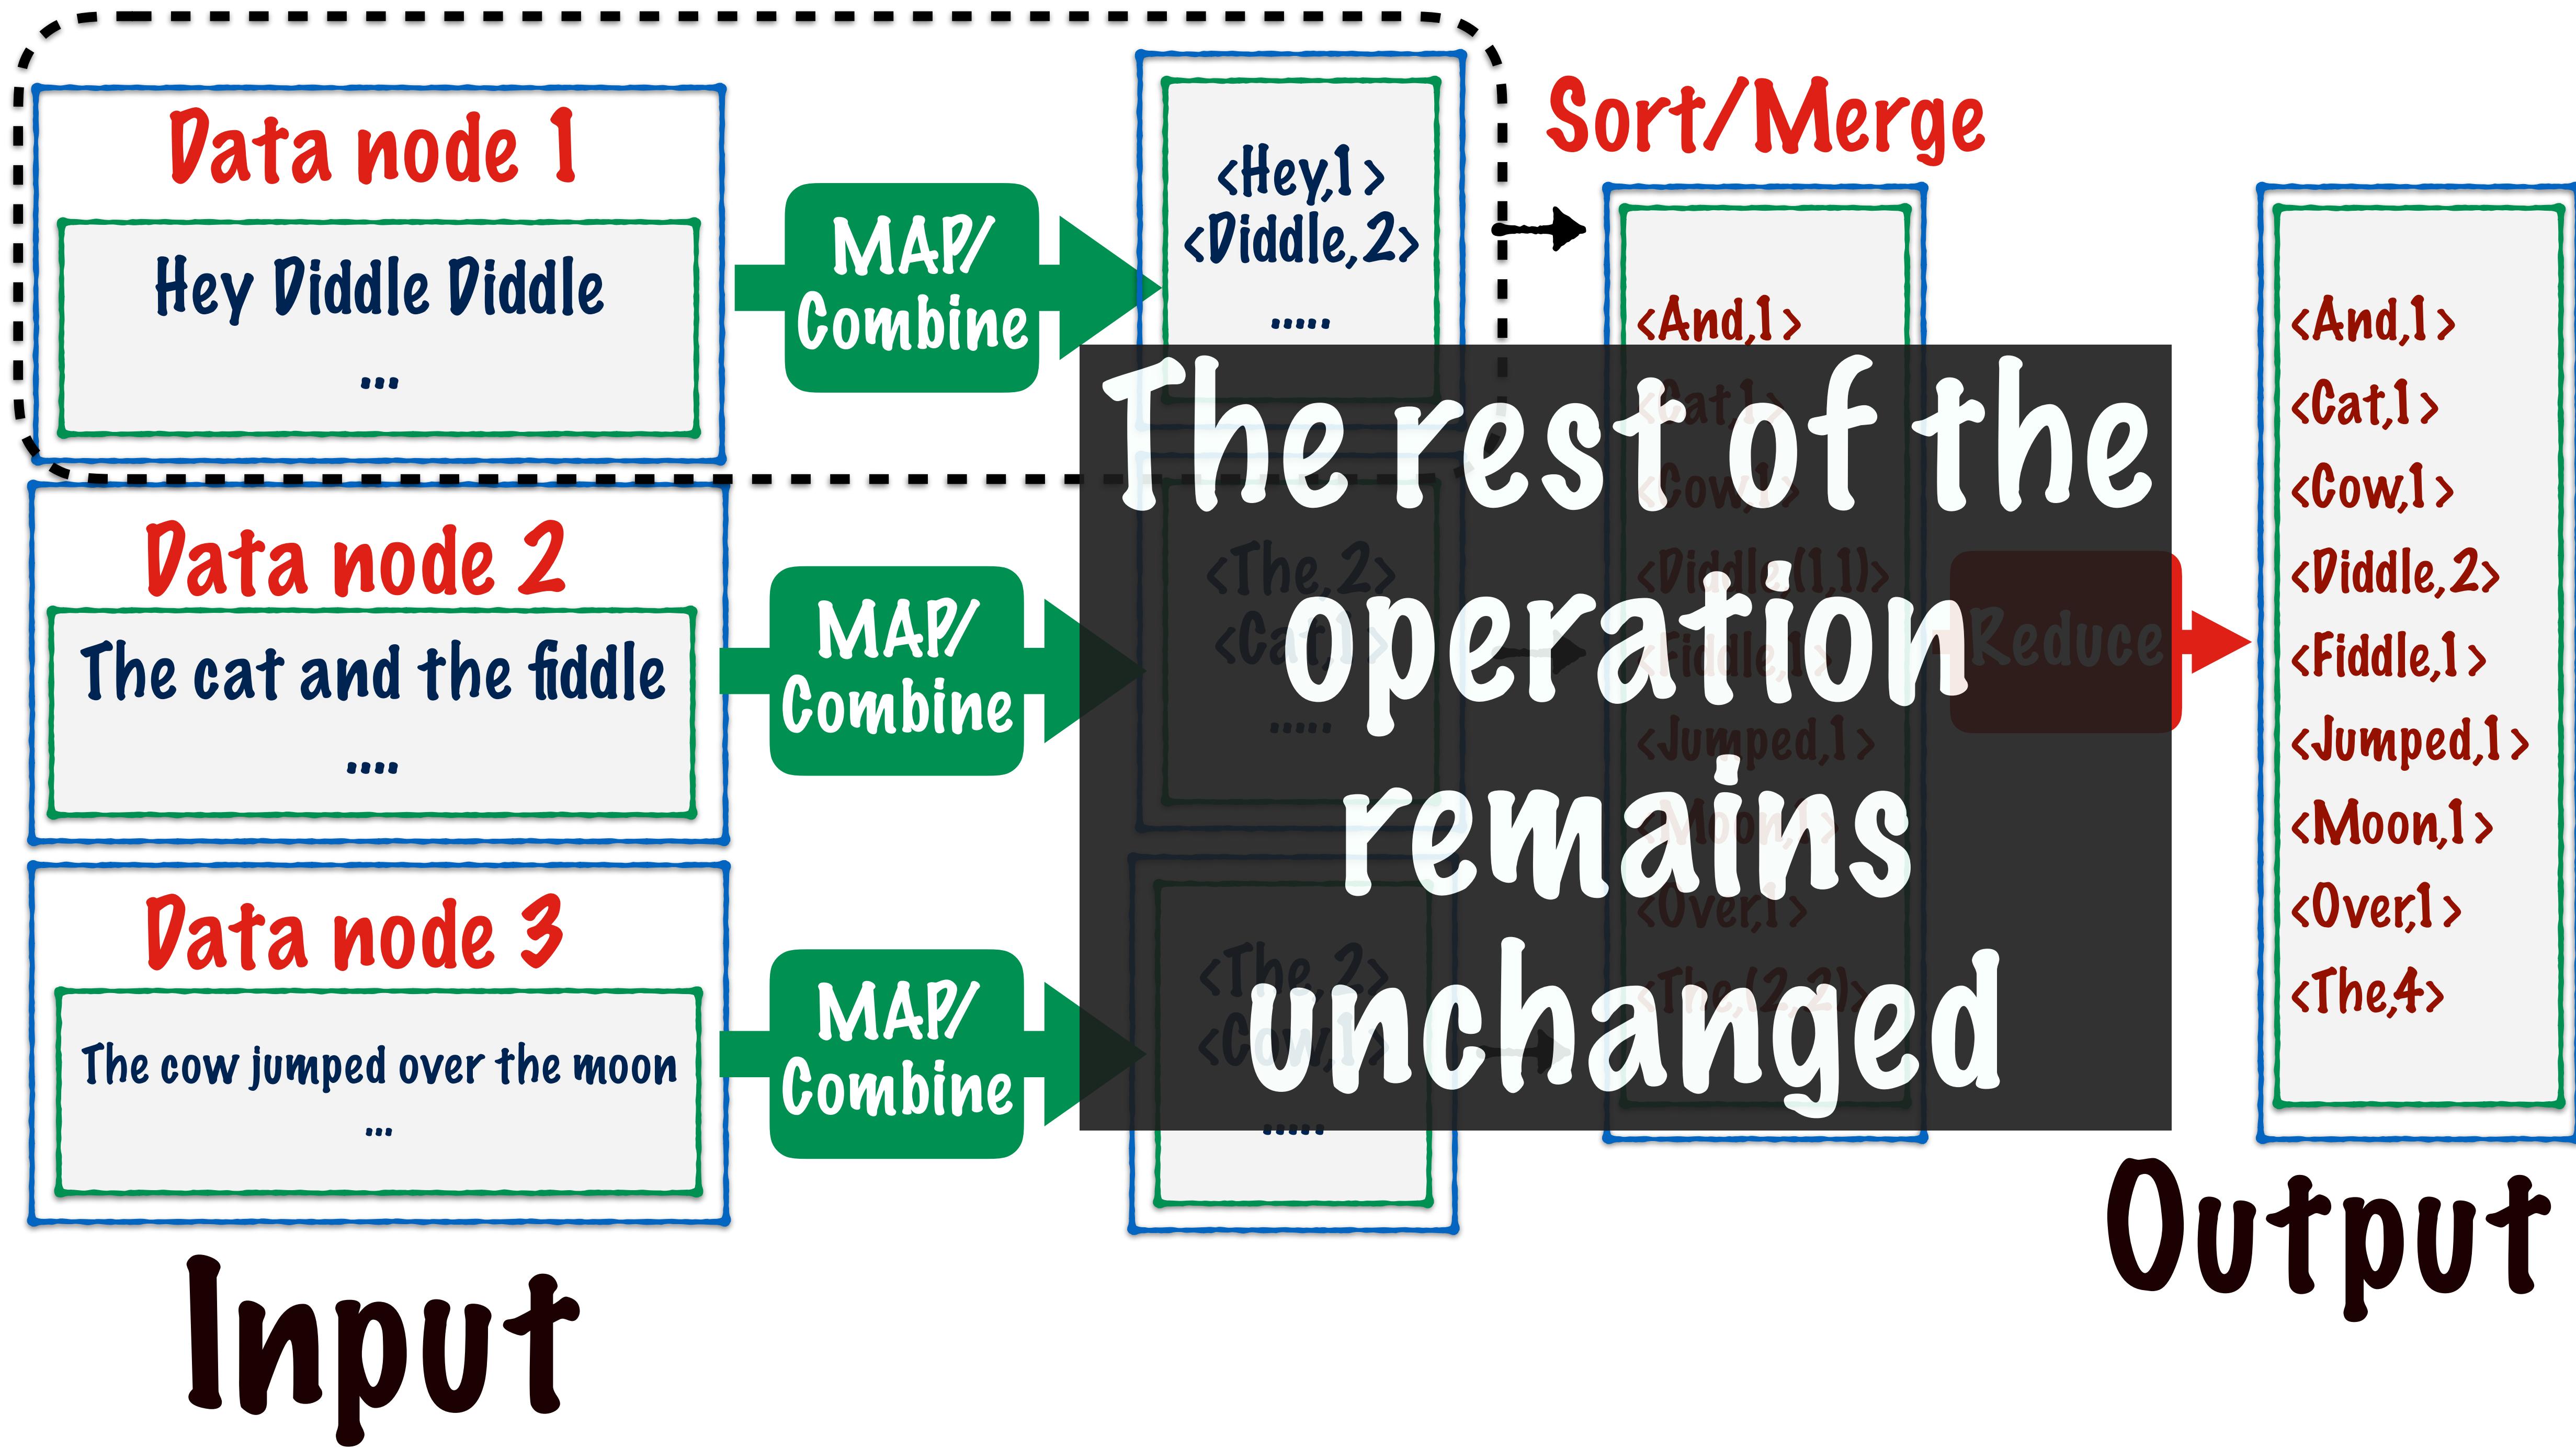
<Word, 1>

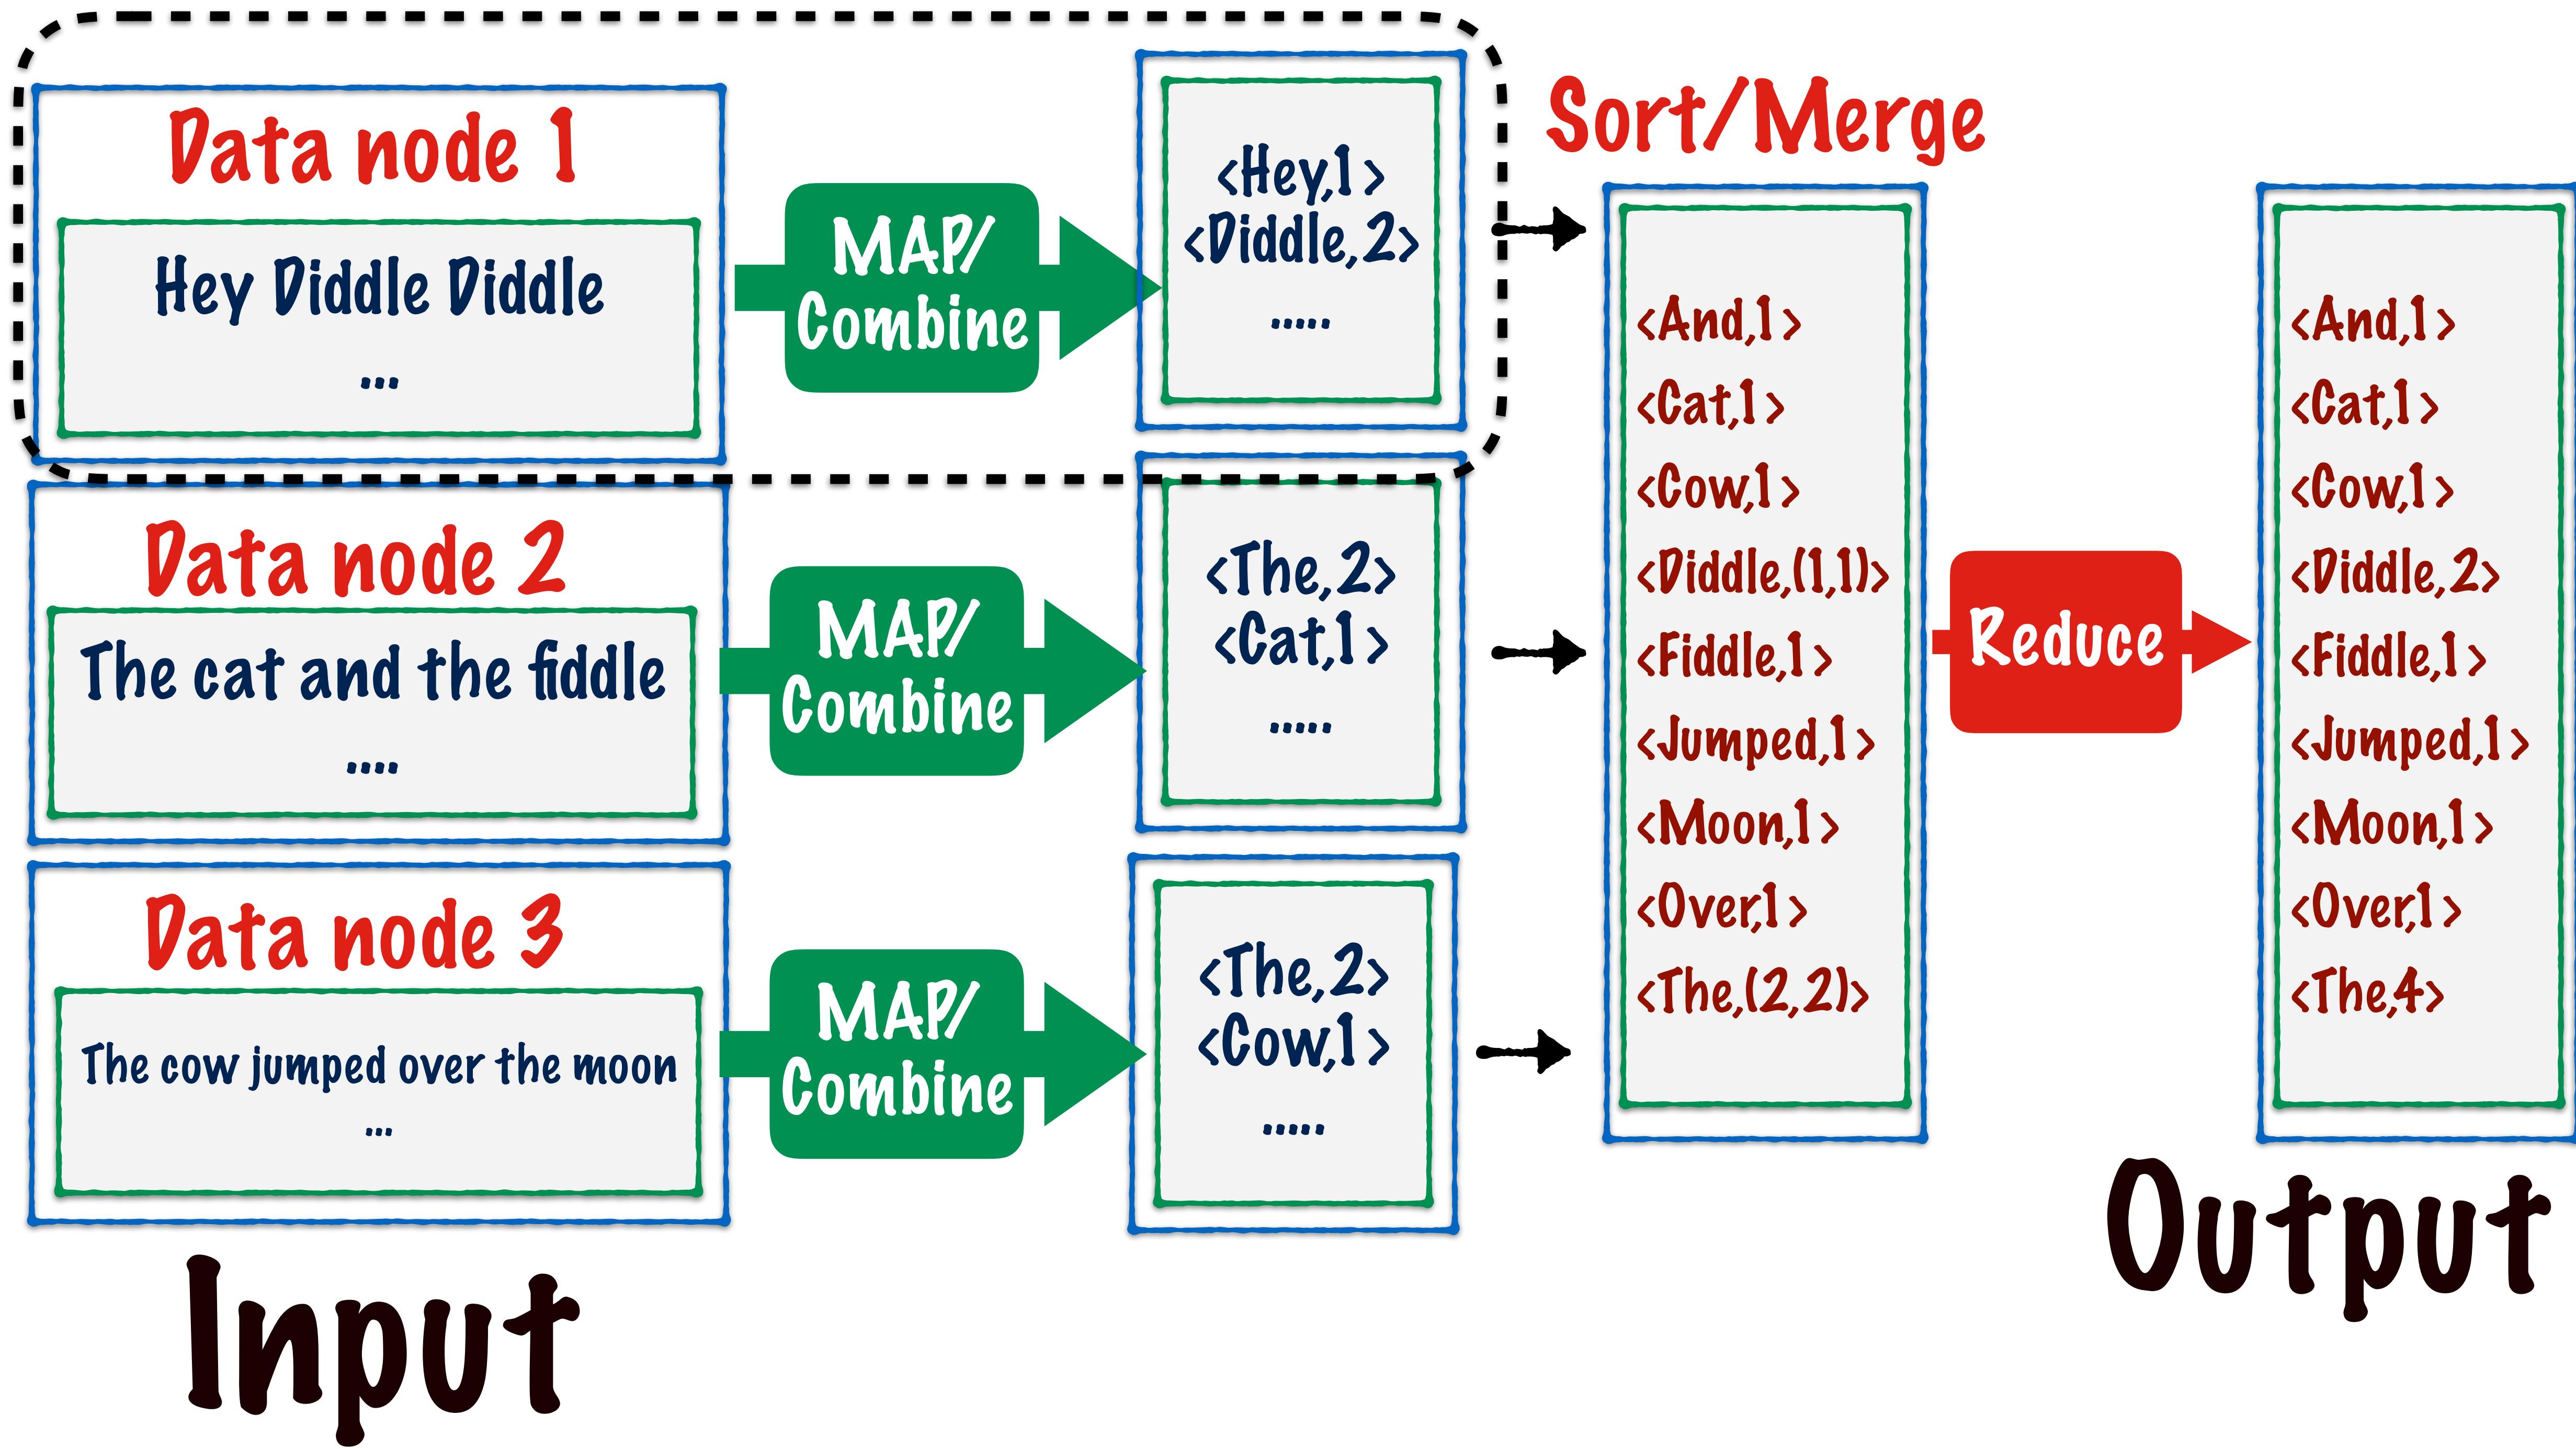
<Word, Count>

This function has
combined key-value
pairs with the same key



This function has
combined key-value
pairs with the same key





Implementing a Combiner in MapReduce

We already wrote code for
the following 3 steps

Step 1: Write a map() function

Step 1a: Write a combiner() function

Step 2: Write a reduce() function

Step 3: Setup a driver that points to our
map and reduce implementations

WordMapper Class

WordReducer Class

We have these 3
classes already

WordCount Class

Job Object

Input filepath

Output filepath

Mapper class

Reducer class

Output data types

WordMapper Class

The Combiner function is implemented by extending the Reducer class

The same class that WordReducer inherits from

WordReducer Class

WordCount Class

Job Object

Input filepath

Output

Mapper class

Reducer class

Output data types

WordMapper Class

We'll simply point
the combiner
class in Job to
WordReducer

WordReducer Class

WordCount Class
Job Object
Mapper class
Reducer class
Combiner class
Output data types

WordMapper Class

In other words
the WordReducer
can act as the
combiner as well!

WordReducer Class

WordCount Class
Job Object
Mapper class
Reducer class
Combiner class
Output data types

We don't need
to change
these 2 classes

WordCount.JAR

WordMapper Class

WordReducer Class

WordCount Class

WordCount Class

```
package com.company;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class WordCount {

    public static void main(String[] args) throws Exception {
        if(args.length !=2){
            System.err.println("Invalid Command");
            System.err.println("Usage: WordCount <input path> <output path>");
            System.exit(0);
        }
        Configuration conf = new Configuration();
        Job job = new Job(conf, "wordcount");
        job.setJarByClass(WordCount.class);

        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));
        job.setMapperClass(WordCountMapper.class);
        job.setReducerClass(WordCountReducer.class);

        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);

        System.exit(job.waitForCompletion(true)?0:1);
    }
}
```

Here we just
need to make
1 change

WordCount Class

```
package com.company;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class WordCount {

    public static void main(String[] args) throws Exception {
        if(args.length !=2){
            System.err.println("Invalid Command");
            System.err.println("Usage: WordCount <input path> <output path>");
            System.exit(0);
        }
        Configuration conf = new Configuration();

        Job job = new Job(conf, "wordcount");

        job.setJarByClass(WordCount.class);
        job.setJobName("Word Count");
        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));

        job.setMapperClass(WordCountMapper.class);
        job.setReducerClass(WordCountReducer.class);

        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);

        System.exit(job.waitForCompletion(true)?0:1);
    }
}
```

We set the Mapper class
and Reducer class to the
classes we've written

WordCount Class

```
package com.company;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class WordCount {

    public static void main(String[] args) throws Exception {
        if(args.length !=2){
            System.err.println("Invalid Command");
            System.err.println("Usage: WordCount <input path> <output path>");
            System.exit(0);
        }
        Configuration conf = new Configuration();
        Job job = new Job(conf, "wordcount");
        job.setJarByClass(WordCount.class);
        job.setJobName("Word Count");
        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));

        job.setMapperClass(WordCountMapper.class);
        job.setReducerClass(WordCountReducer.class);
        job.setCombinerClass(WordCountReducer.class);
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);

        System.exit(job.waitForCompletion(true)?0:1);
    }
}
```

Set the combiner class also

Constraints on the combiner function

The combiner function makes
the MapReduce task faster

1. The reduce operation is partially parallelized
2. Less data needs to be transferred
across the network

The combiner function should have no impact on the final result itself

Whether it is used or not, the results should remain the same

In fact, the decision of whether to use the combiner or not ultimately rests with Hadoop

Even if you specify a combiner function,
Hadoop **might decide not to call it**

So, the outcome of MapReduce should
be independent of the combiner

Let's see an example of a combiner
that does not satisfy this constraint

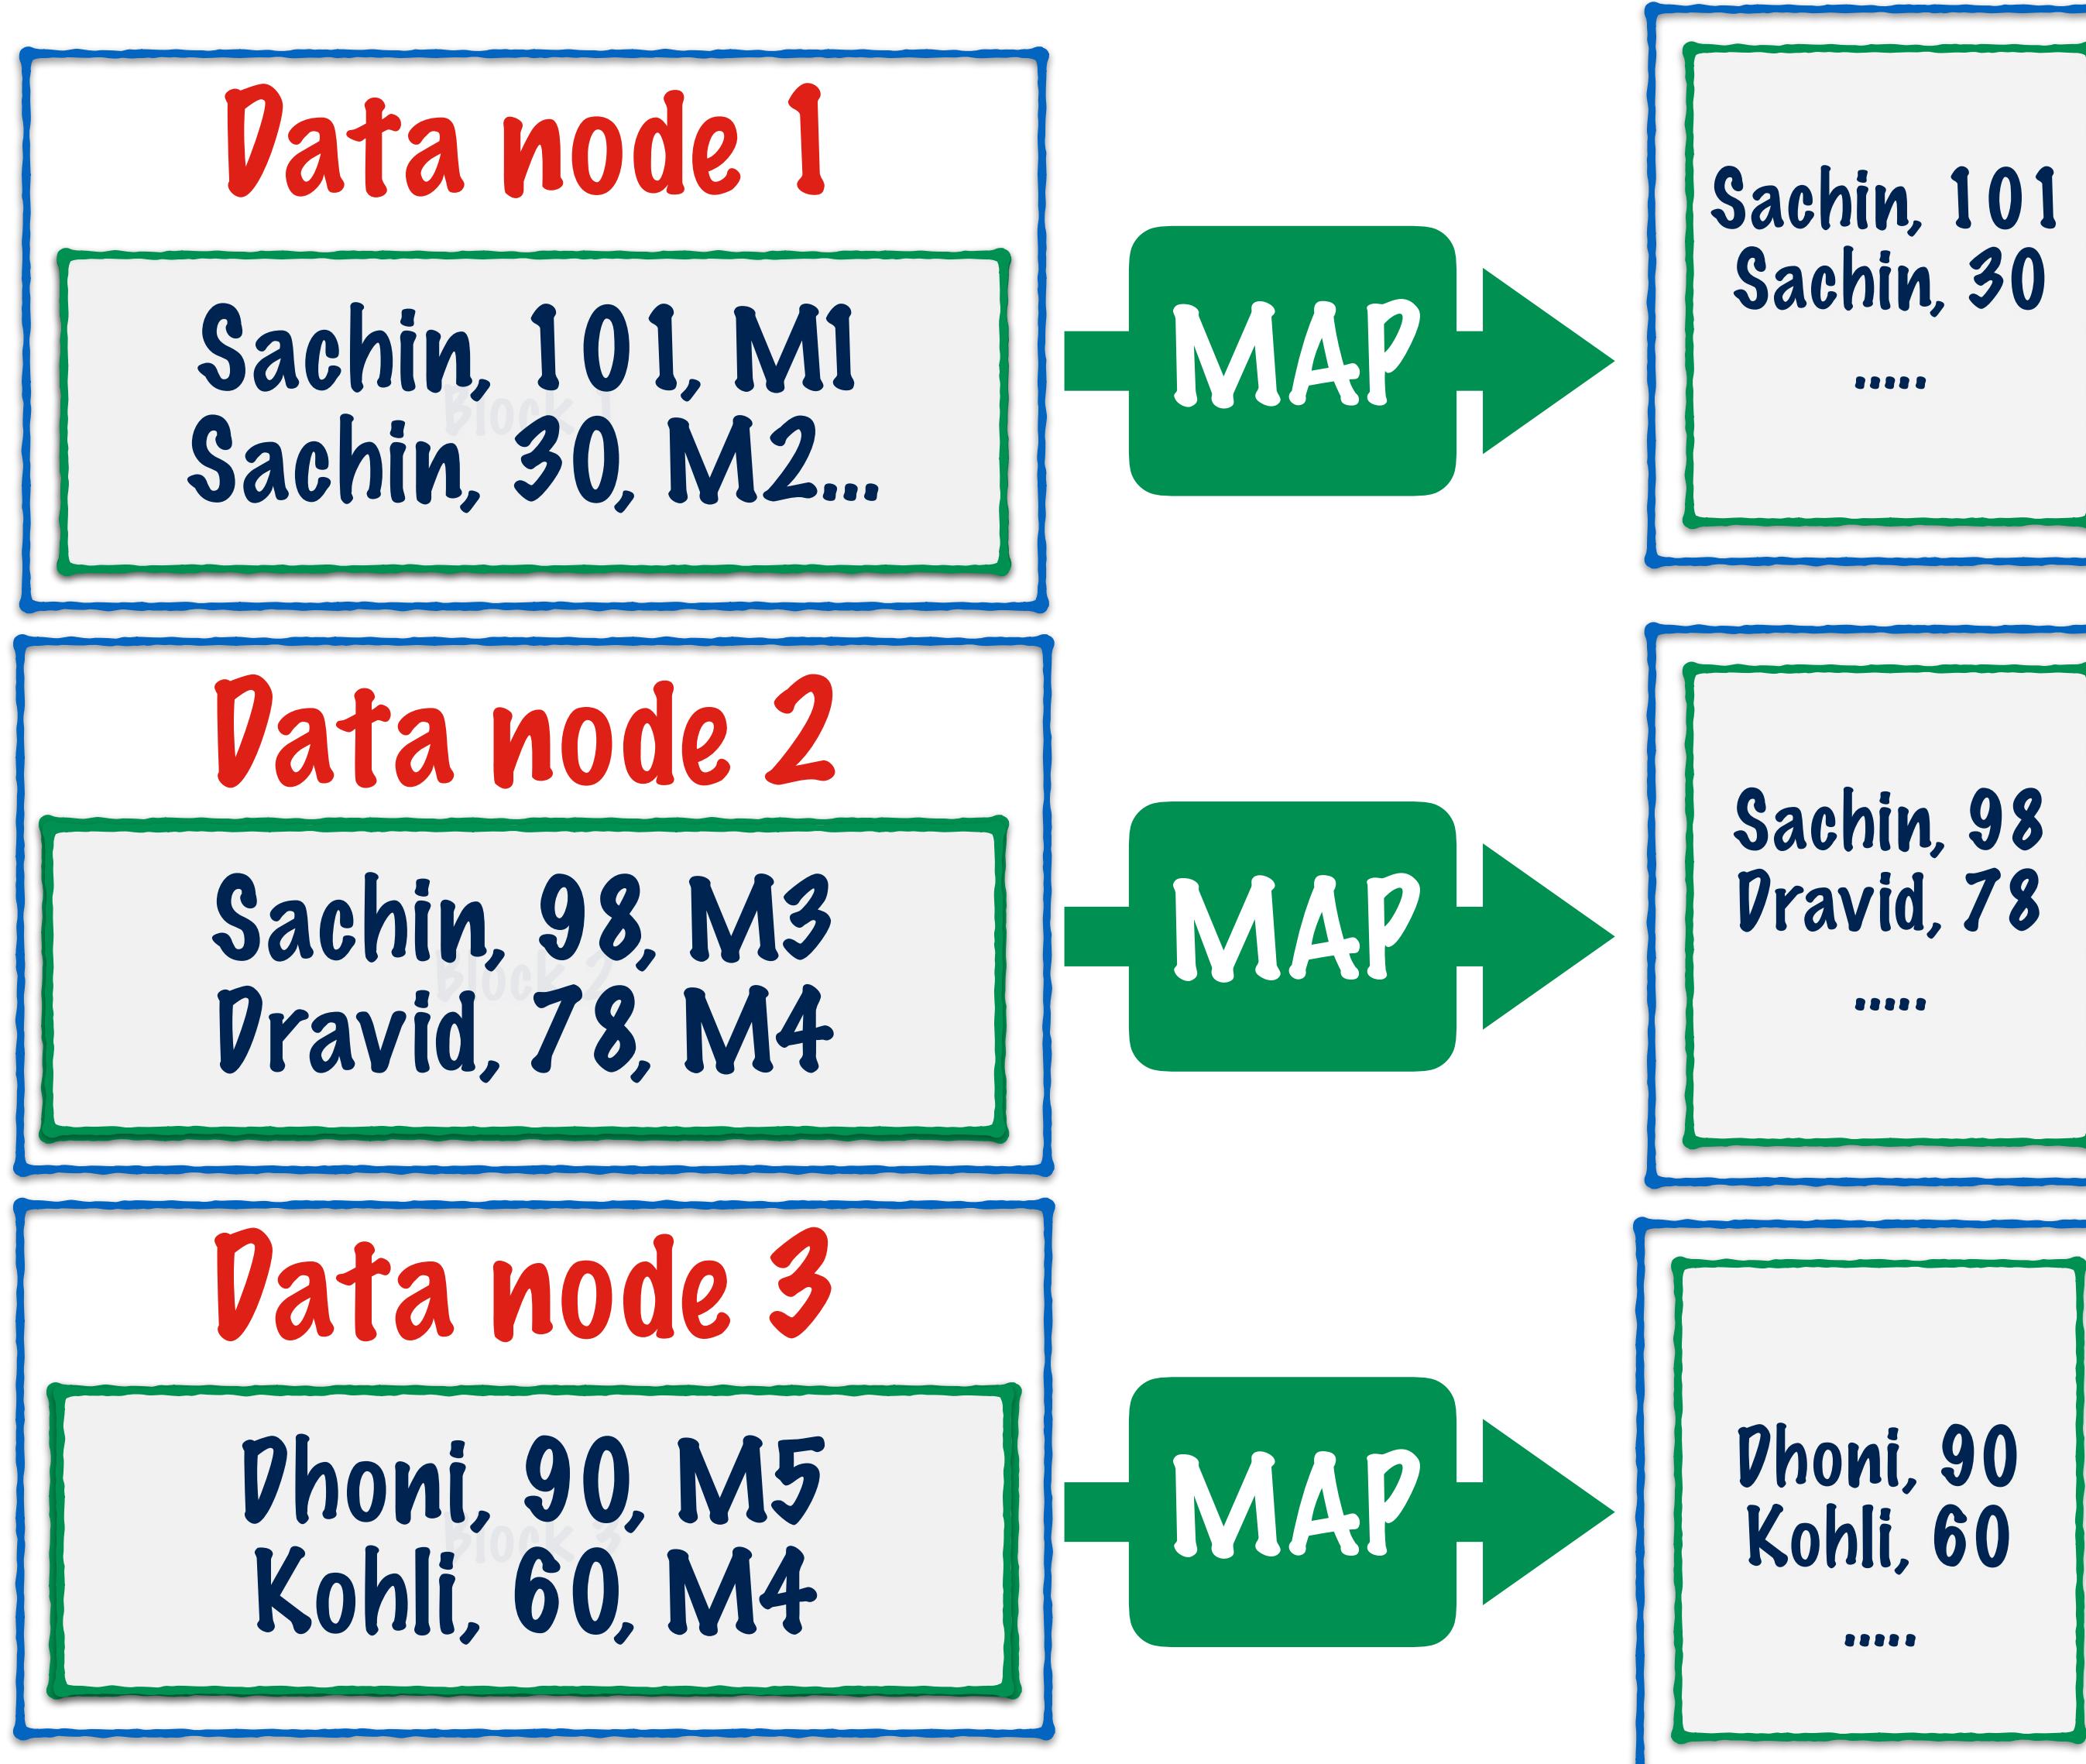
Example: A file contains 3 columns

Player, Runs scored, Match

Objective: Compute the batting average
for each player

Can we use `averagel()` as the combiner
function?

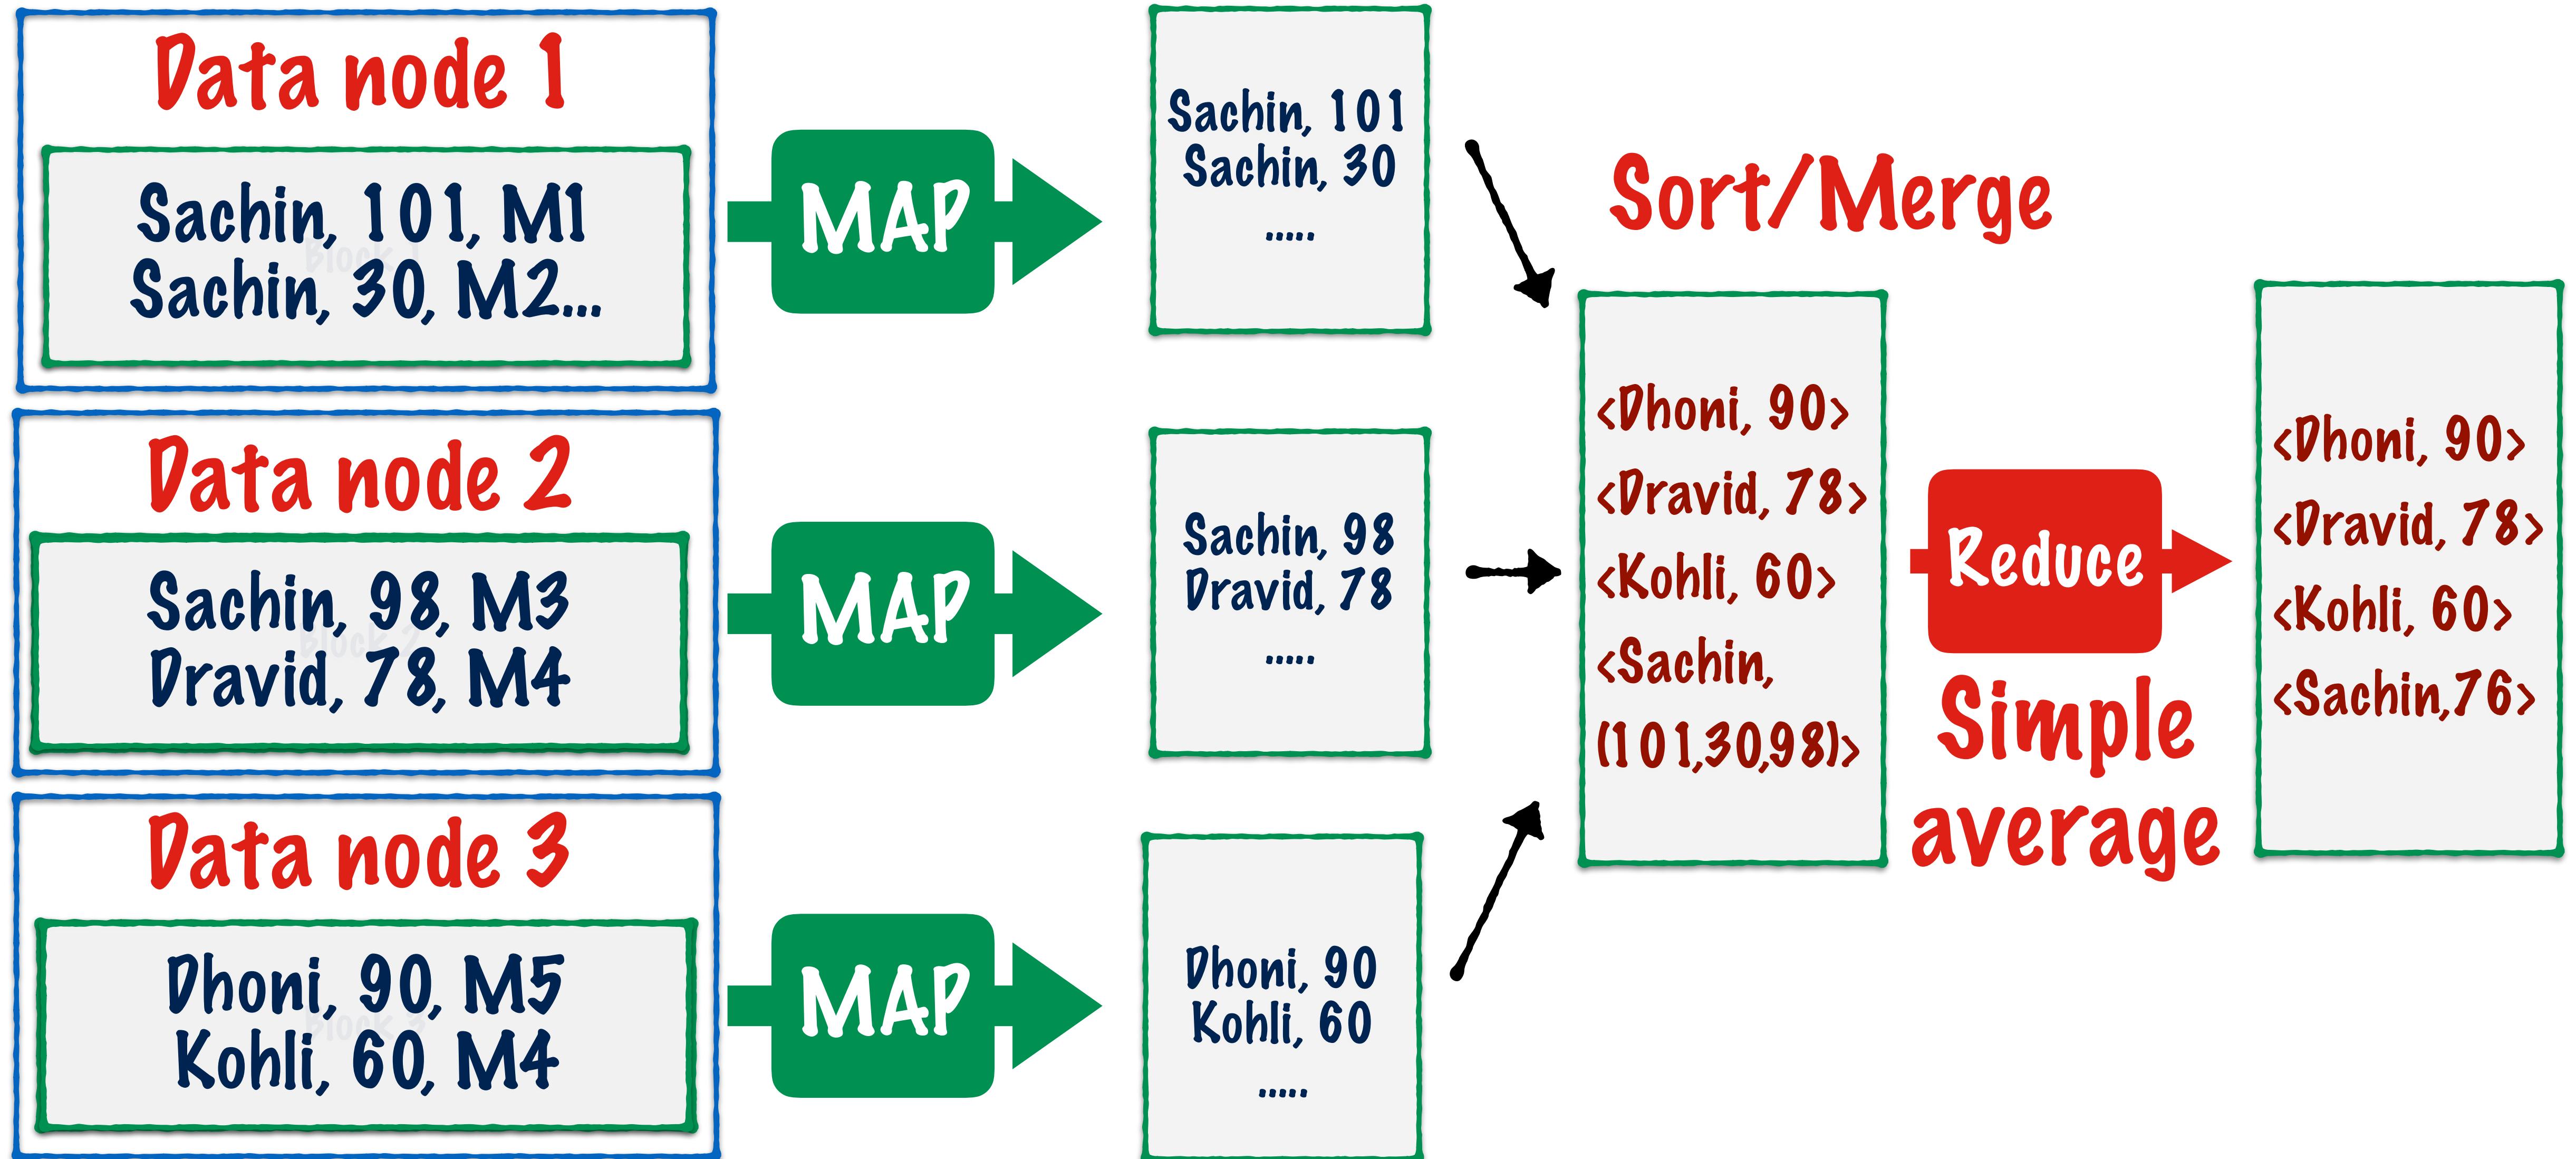
Objective: Compute the batting average for each player



Let's first see
this without a
combiner

Map step just
picks the player
name and score

Objective: Compute the batting average for each player

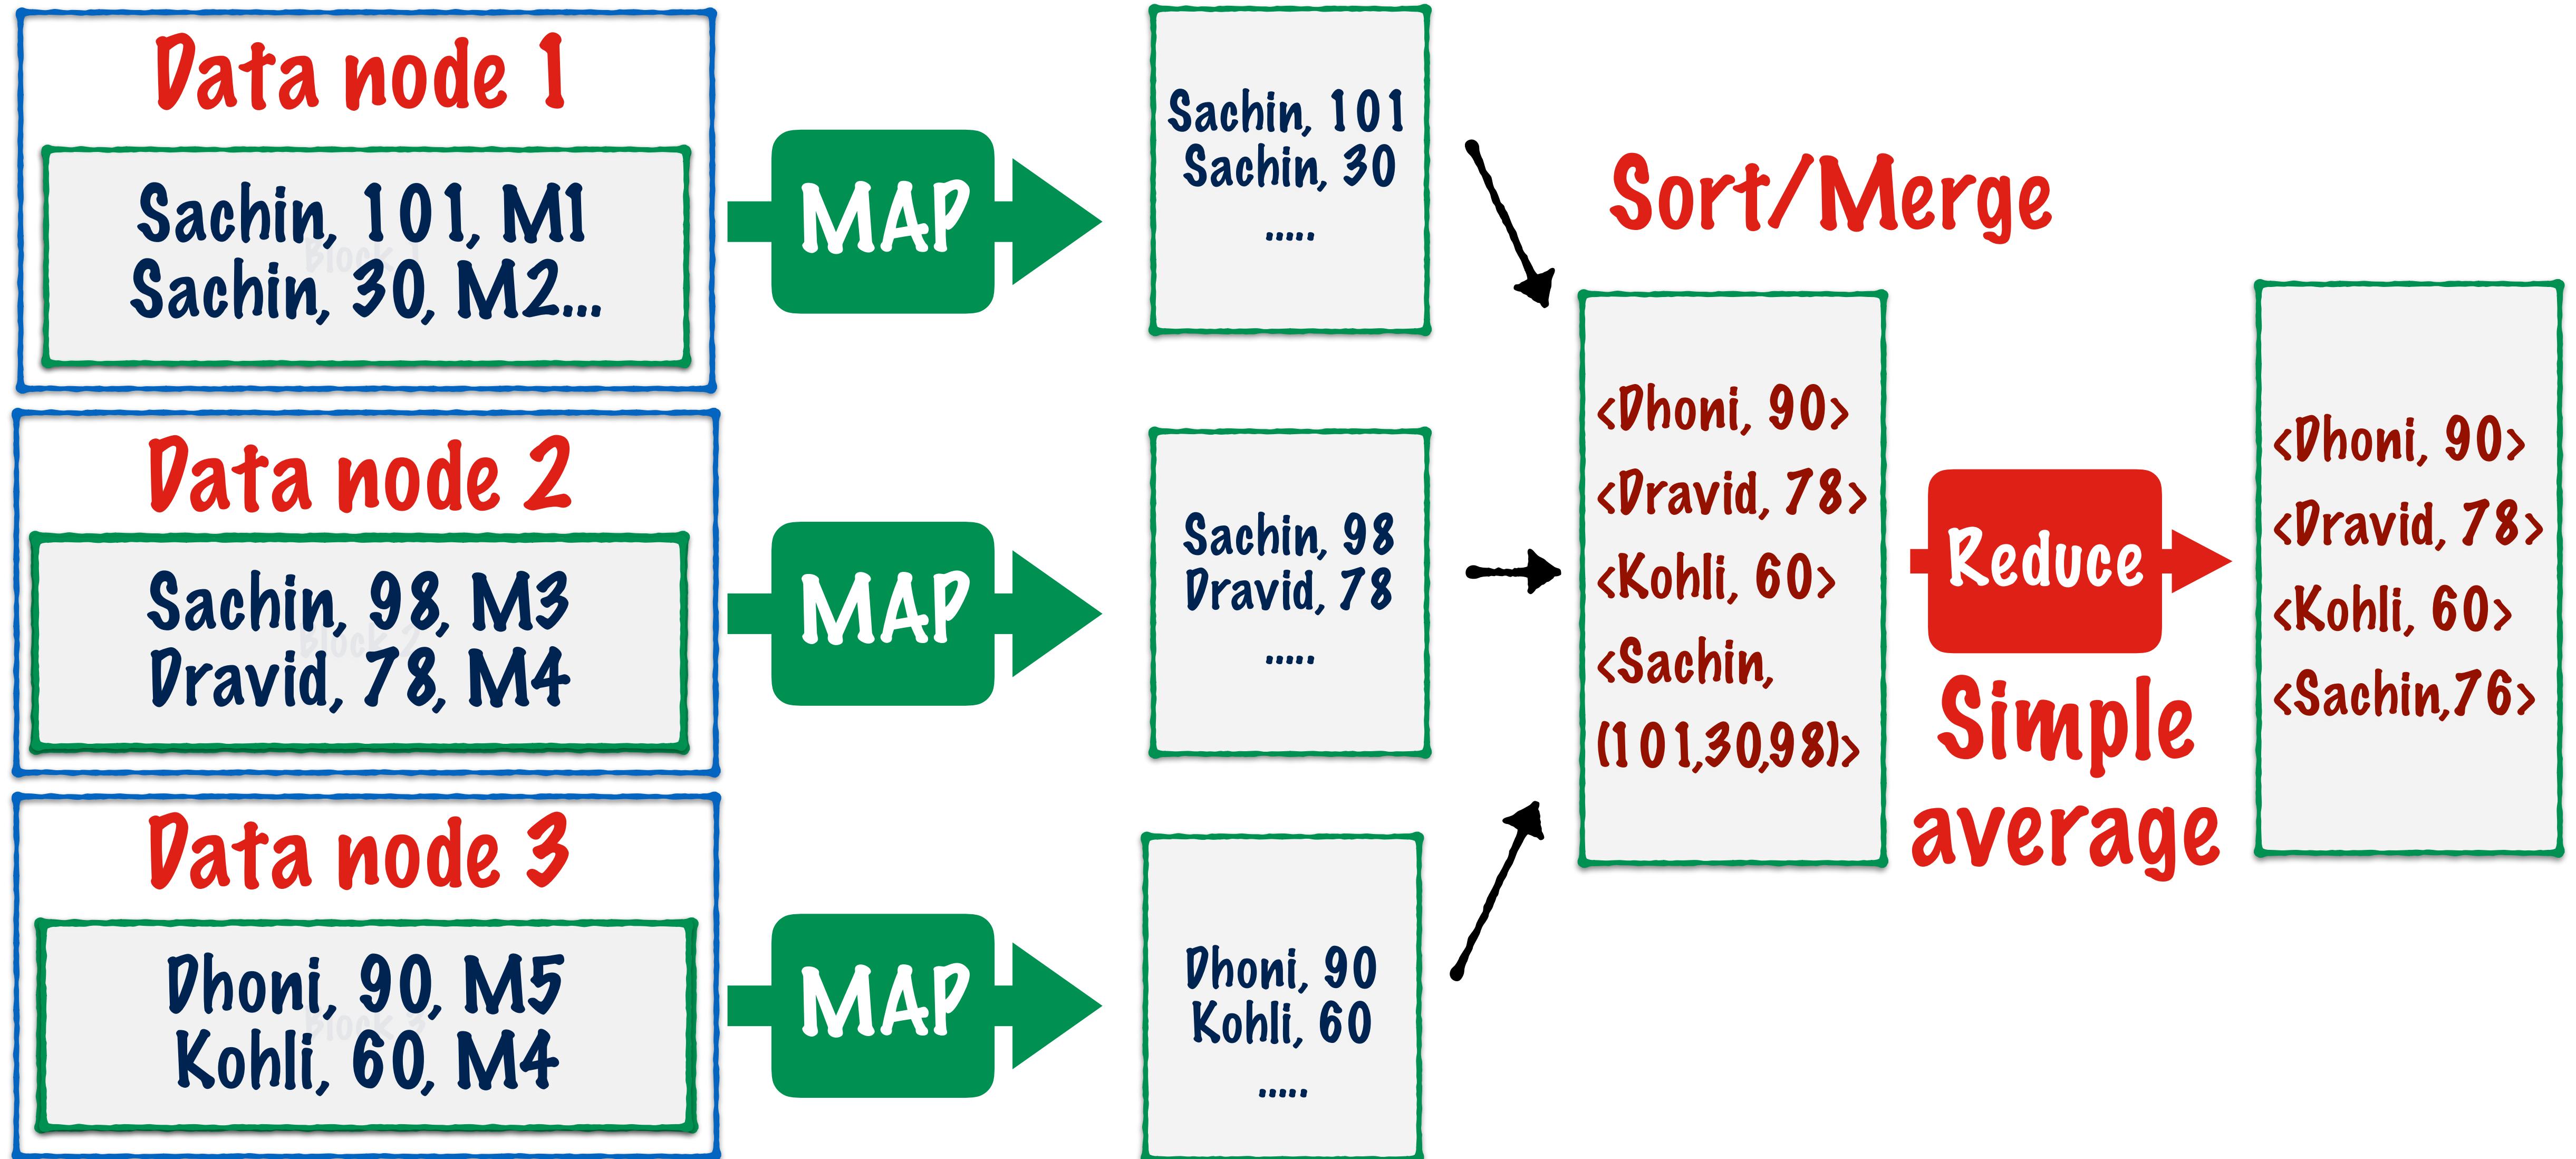


Objective: Compute the batting average for each player

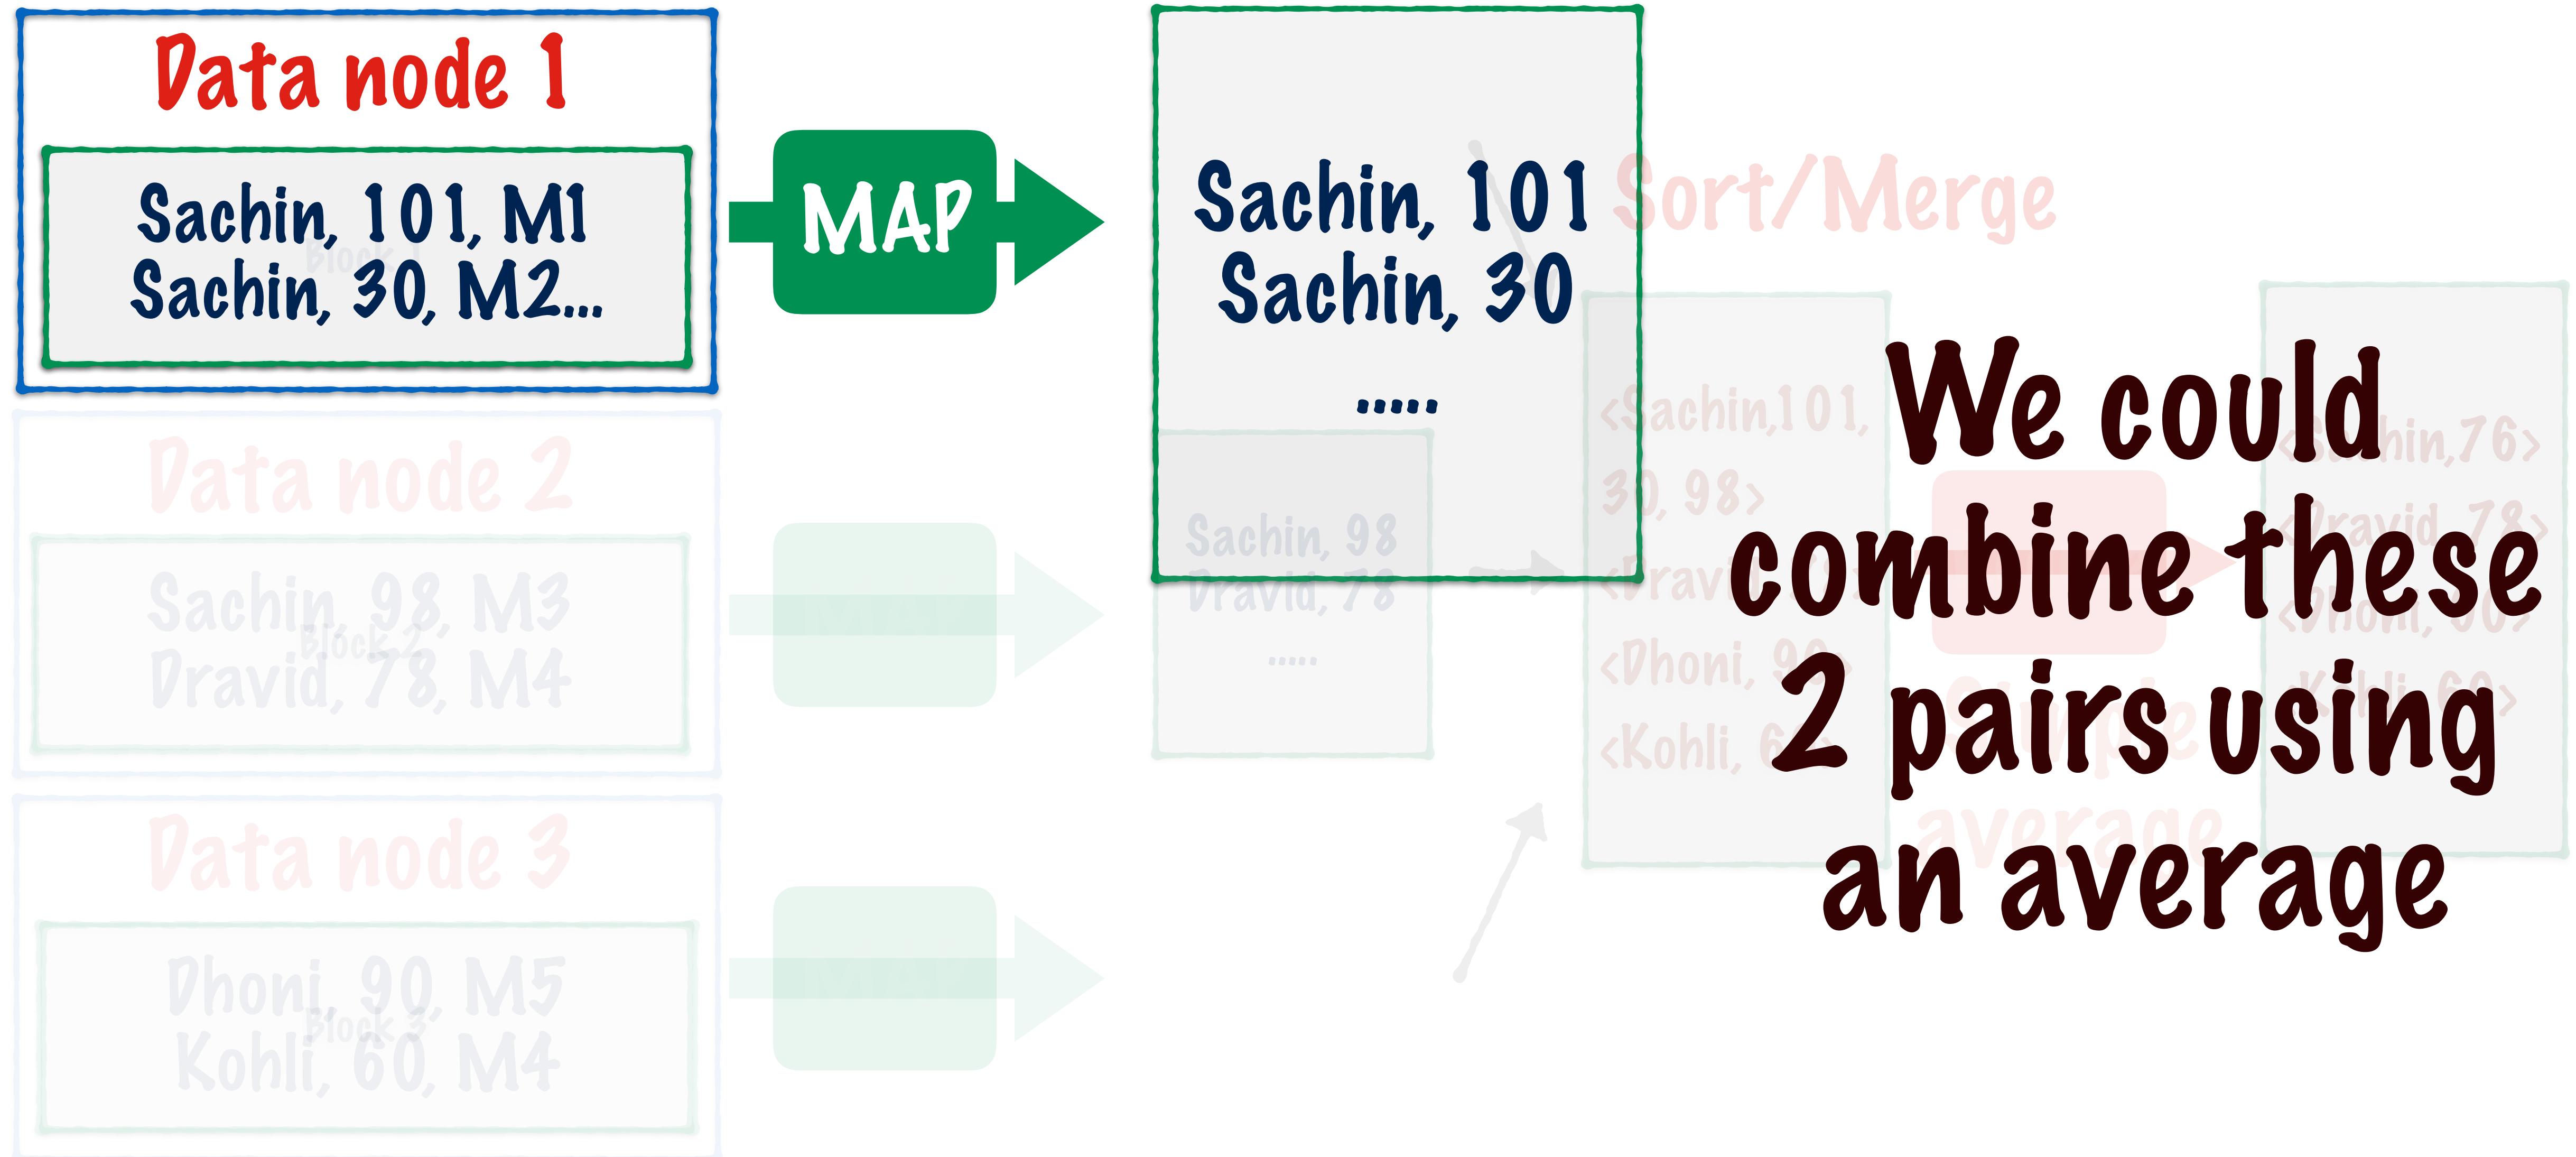
Let's see what's different with `averagel` as the combiner function

<Dhoni, 90>
<Dravid, 78>
<Kohli, 60>
<Sachin,76>

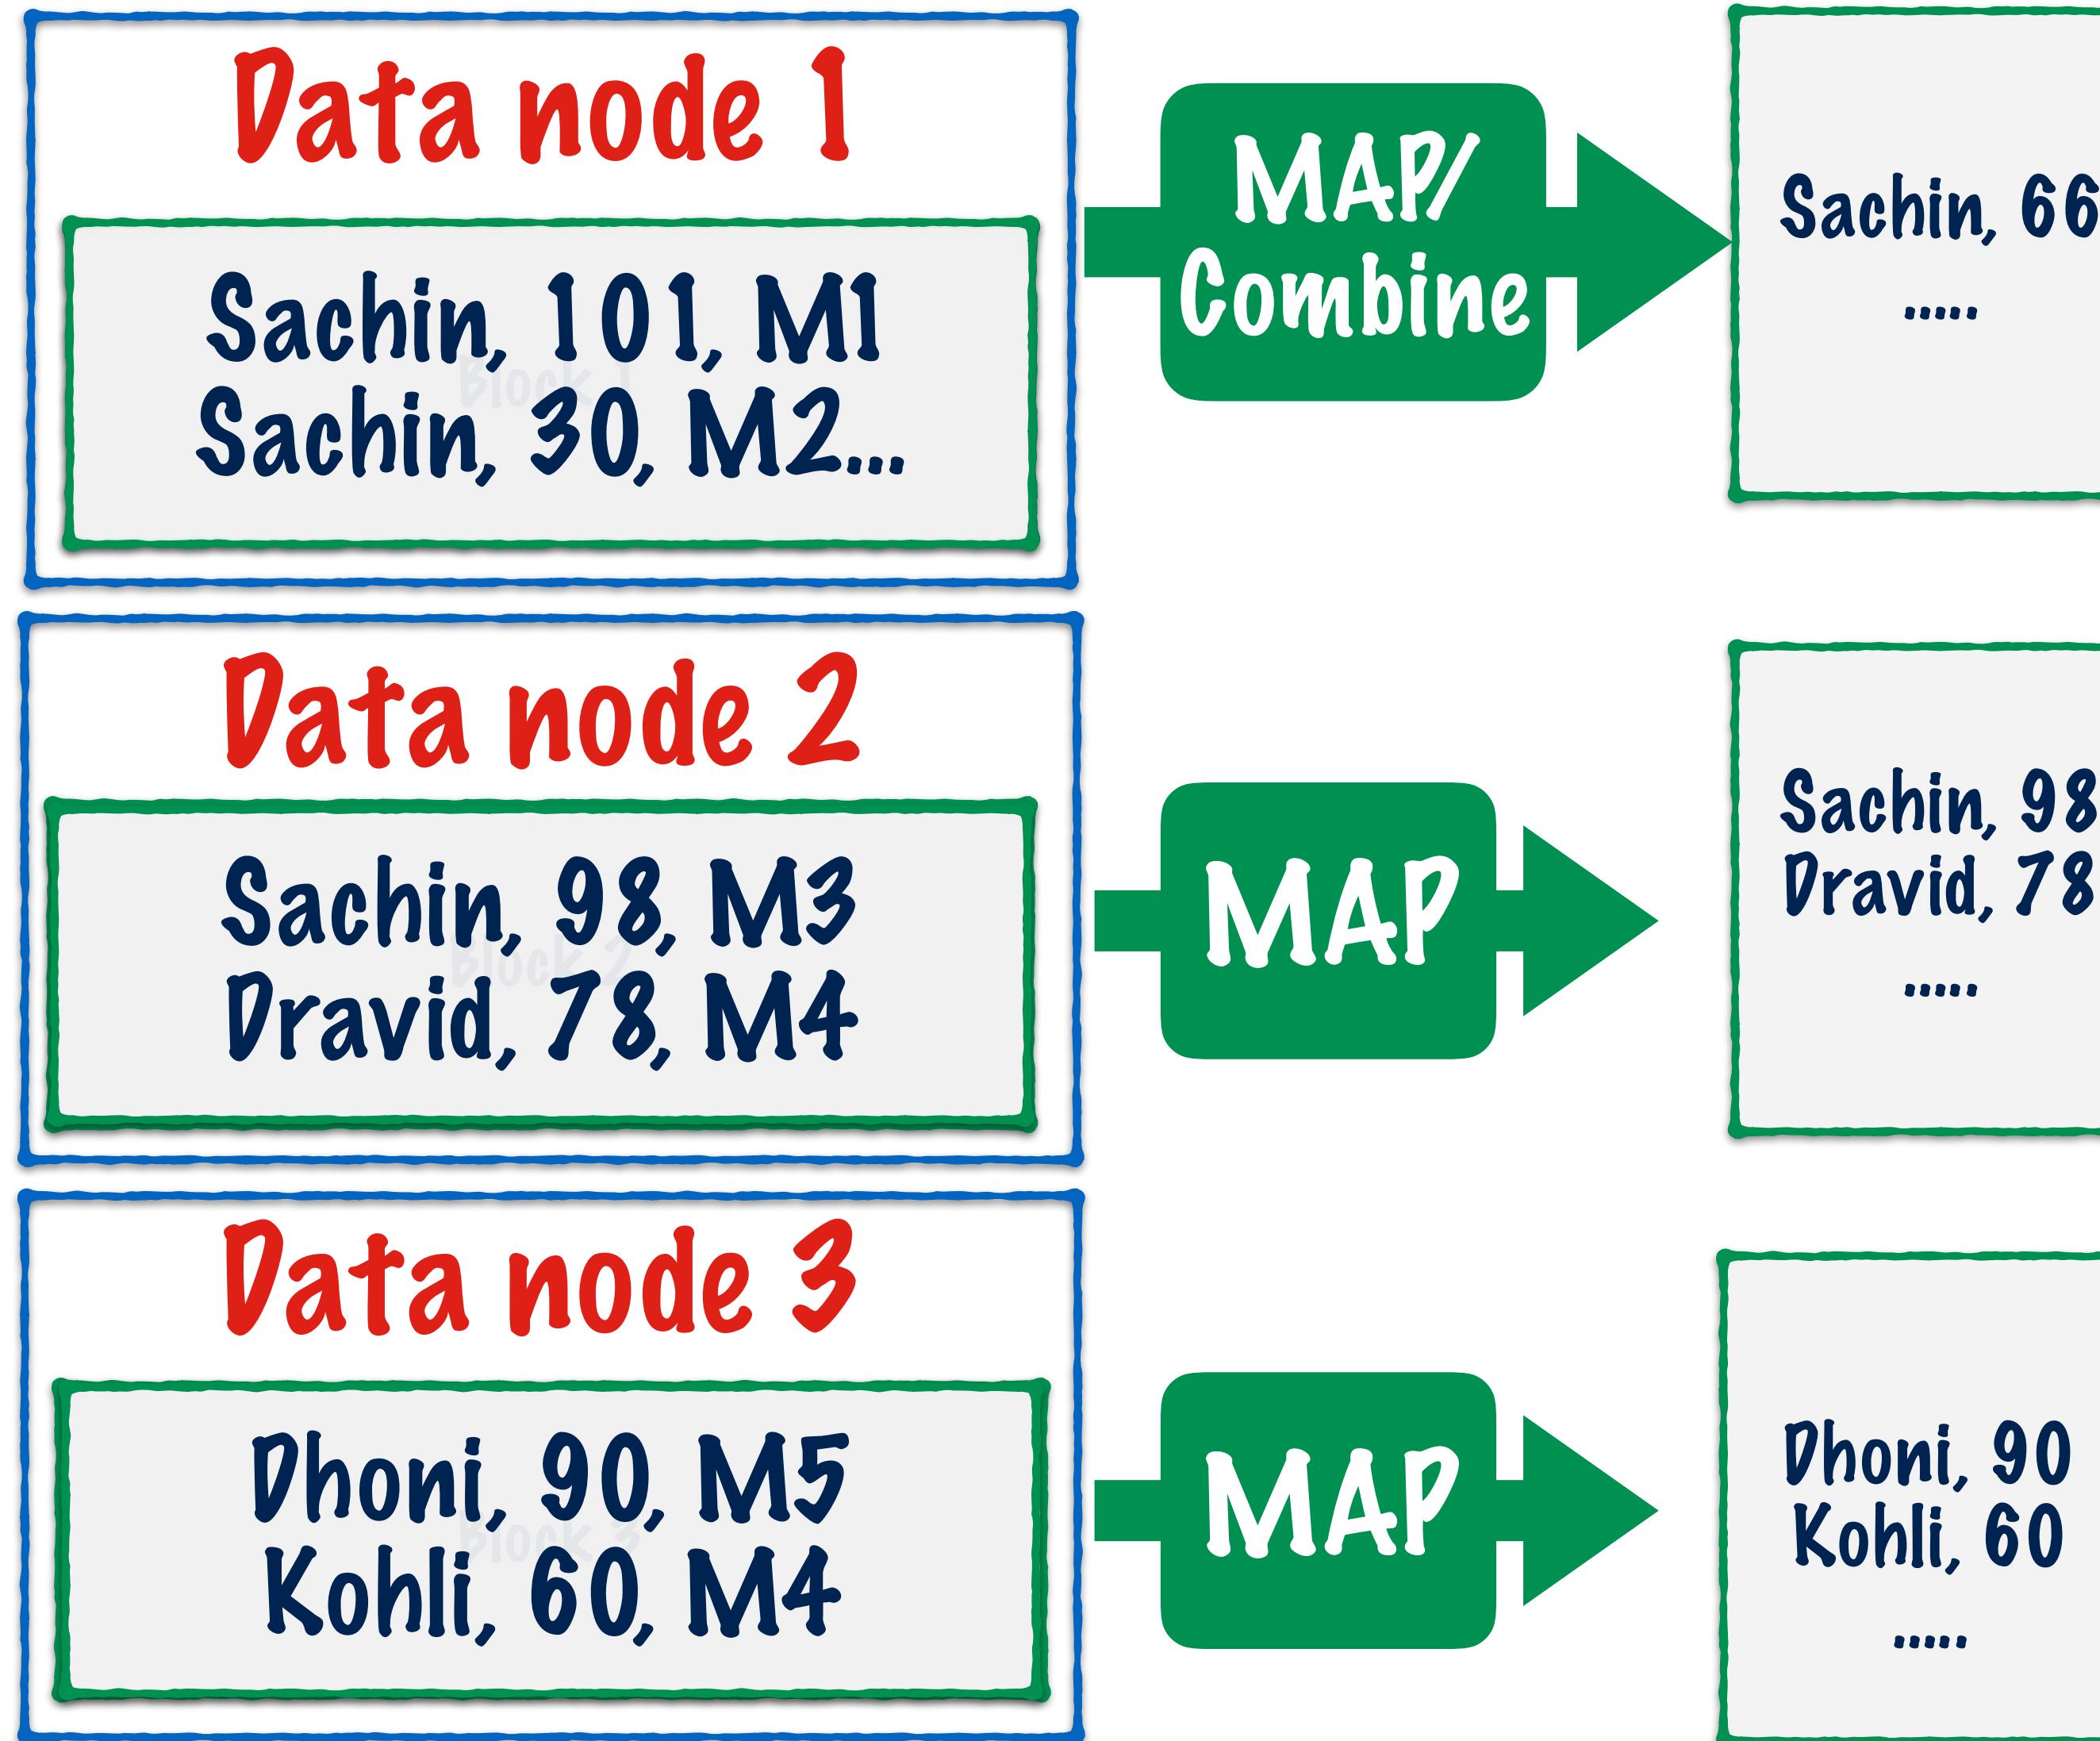
Objective: Compute the batting average for each player



Objective: Compute the batting average for each player

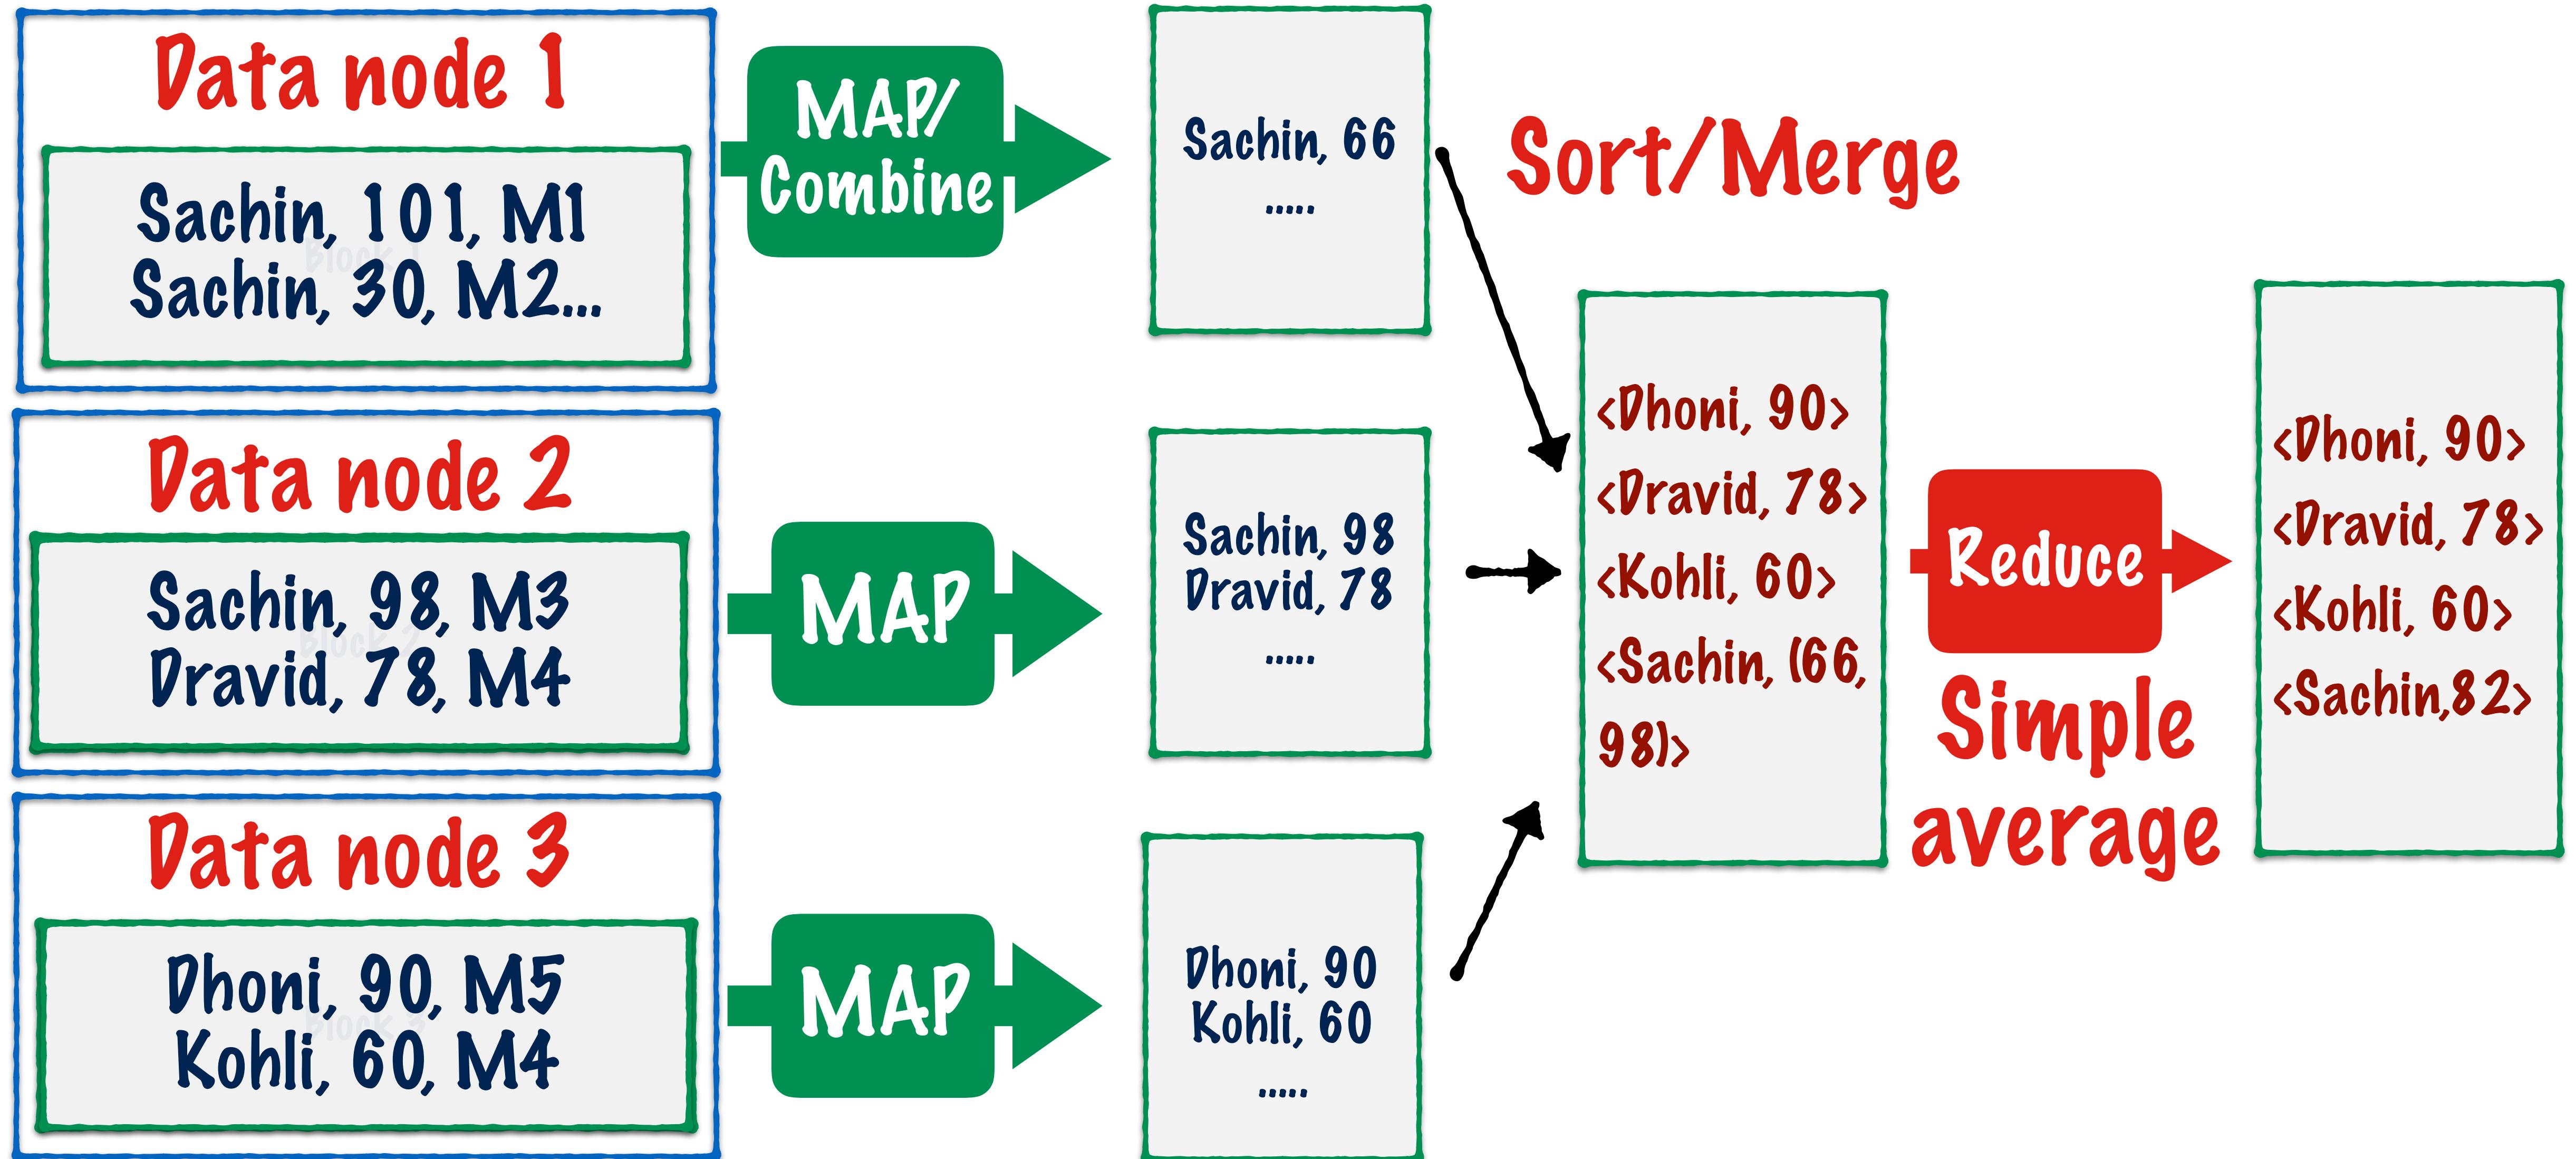


Objective: Compute the batting average for each player



Combiner
won't have any
impact on the
other 2 nodes

Objective: Compute the batting average for each player



Objective: Compute the batting average for each player

Without
Combiner

<Dhoni, 90>
<Dravid, 78>
<Kohli, 60>
<Sachin,76>

With
Combiner

<Dhoni, 90>
<Dravid, 78>
<Kohli, 60>
<Sachin,82>

The result is
different with and
without the
combiner

Objective: Compute the batting average for each player

This is because

Average of a
set of numbers <>

Average (Averages
of subsets)

Objective: Compute the batting average for each player

In our example

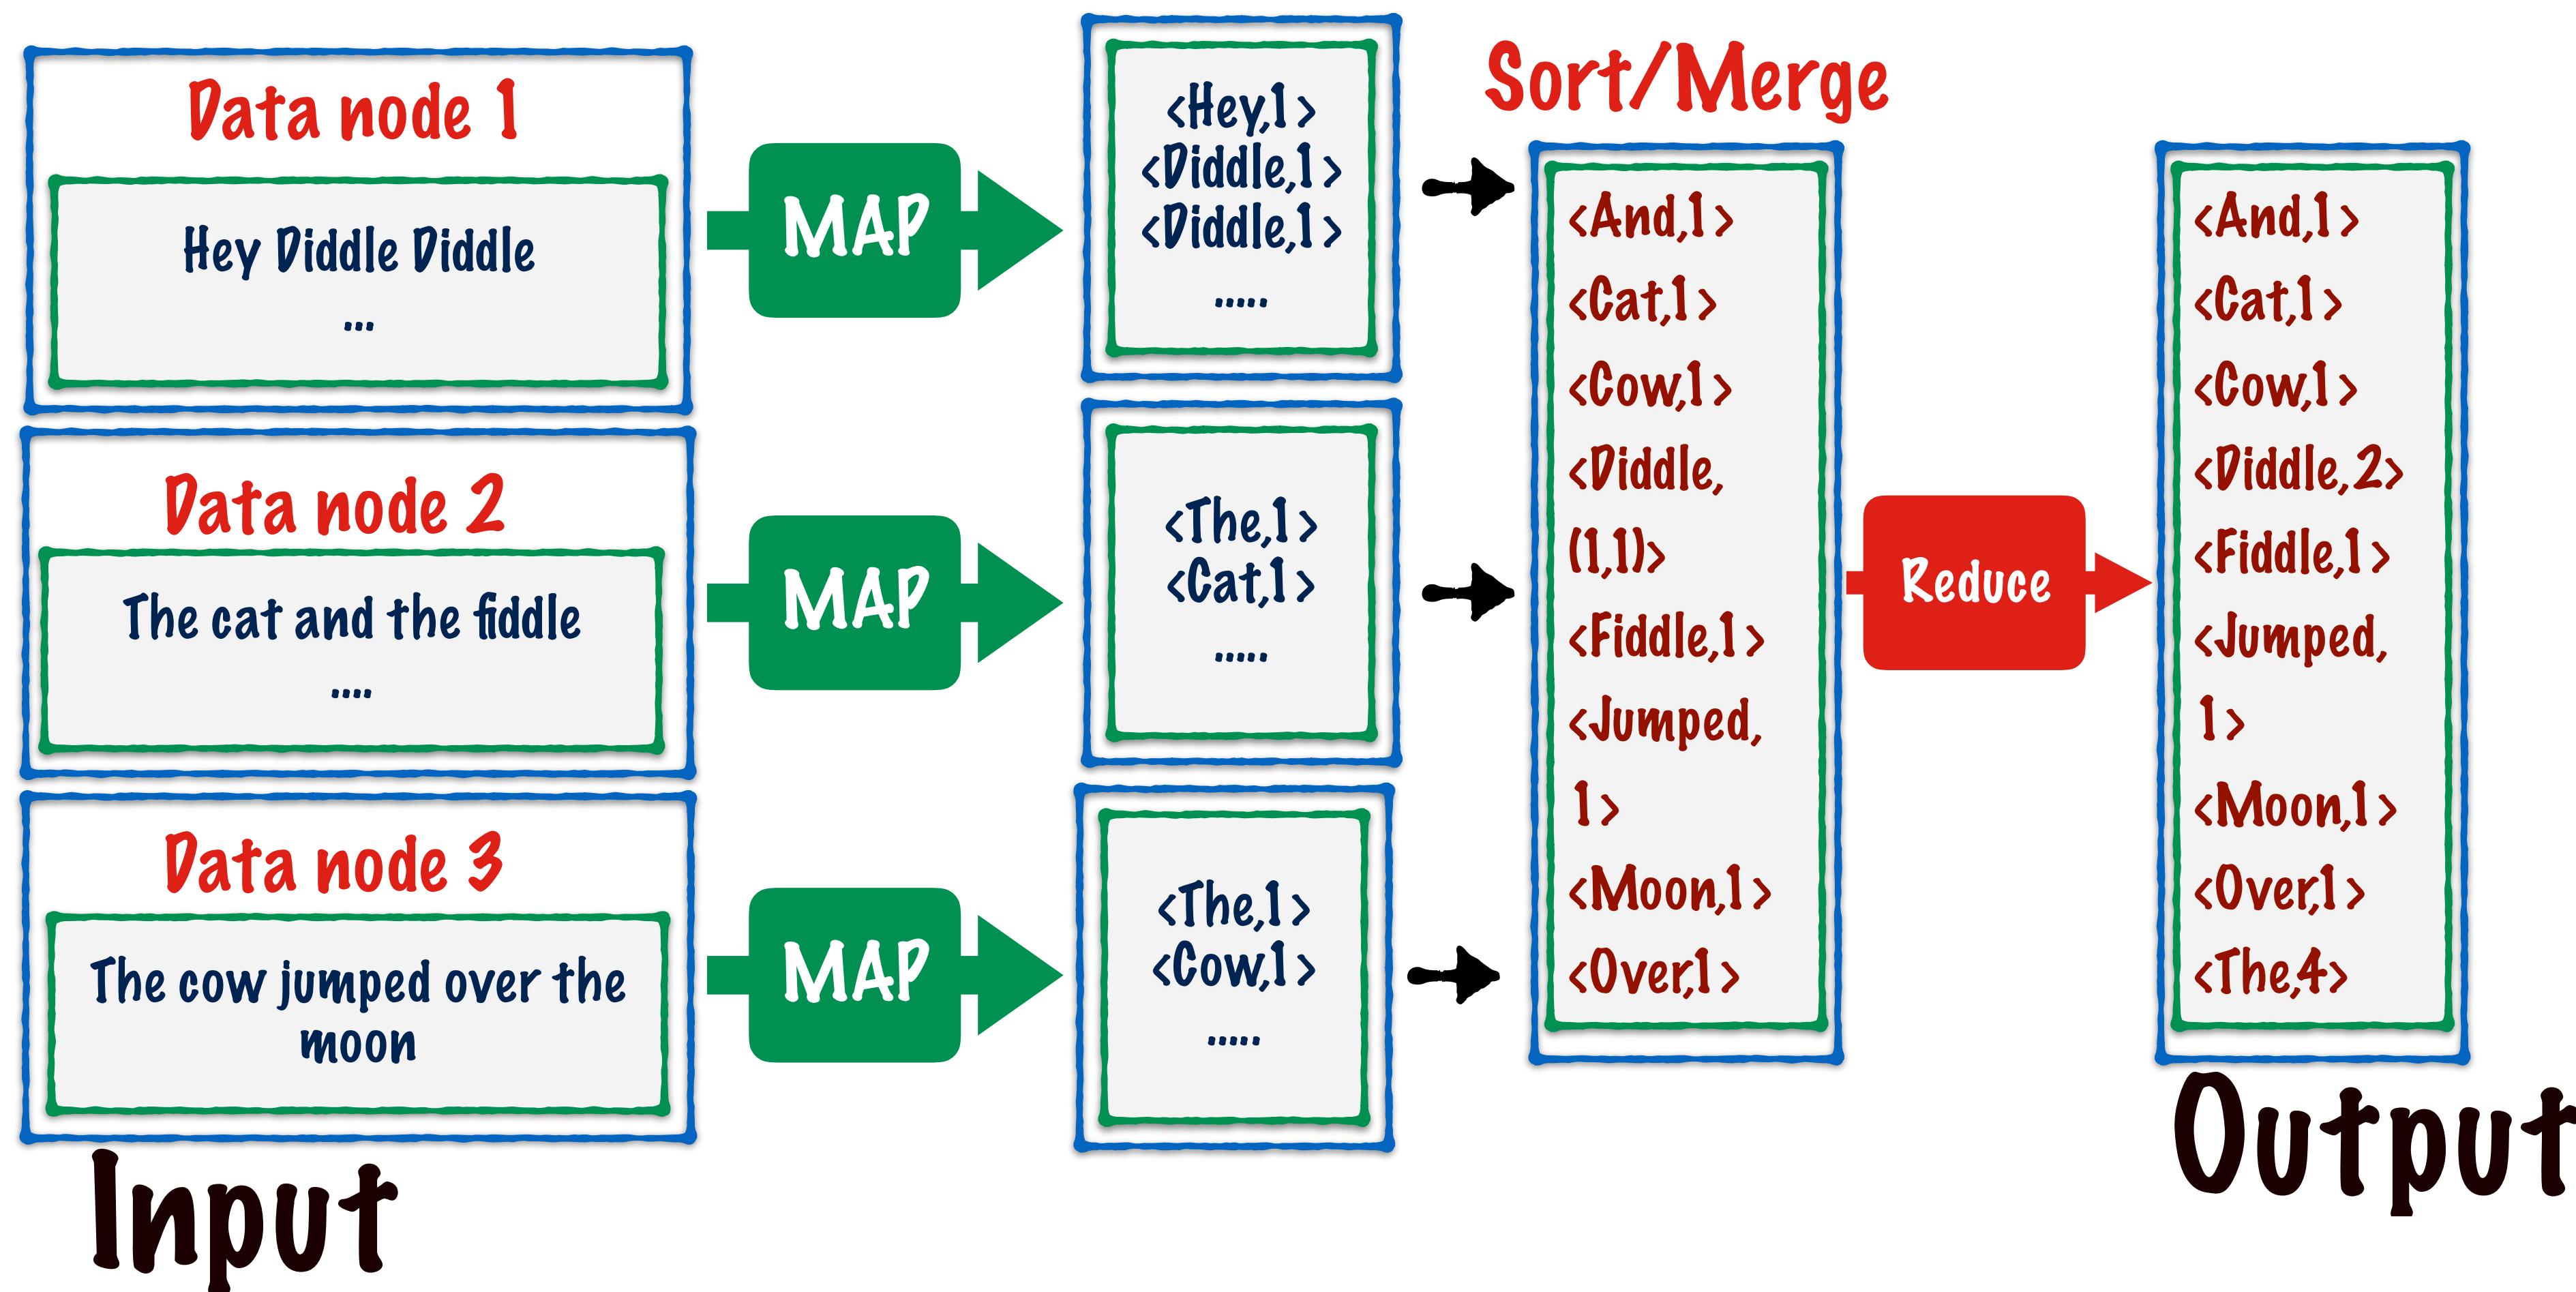
Average(101, 30, 98) < > Average(Avg(101, 30), 98)

We can't use `average()` as
the combiner function

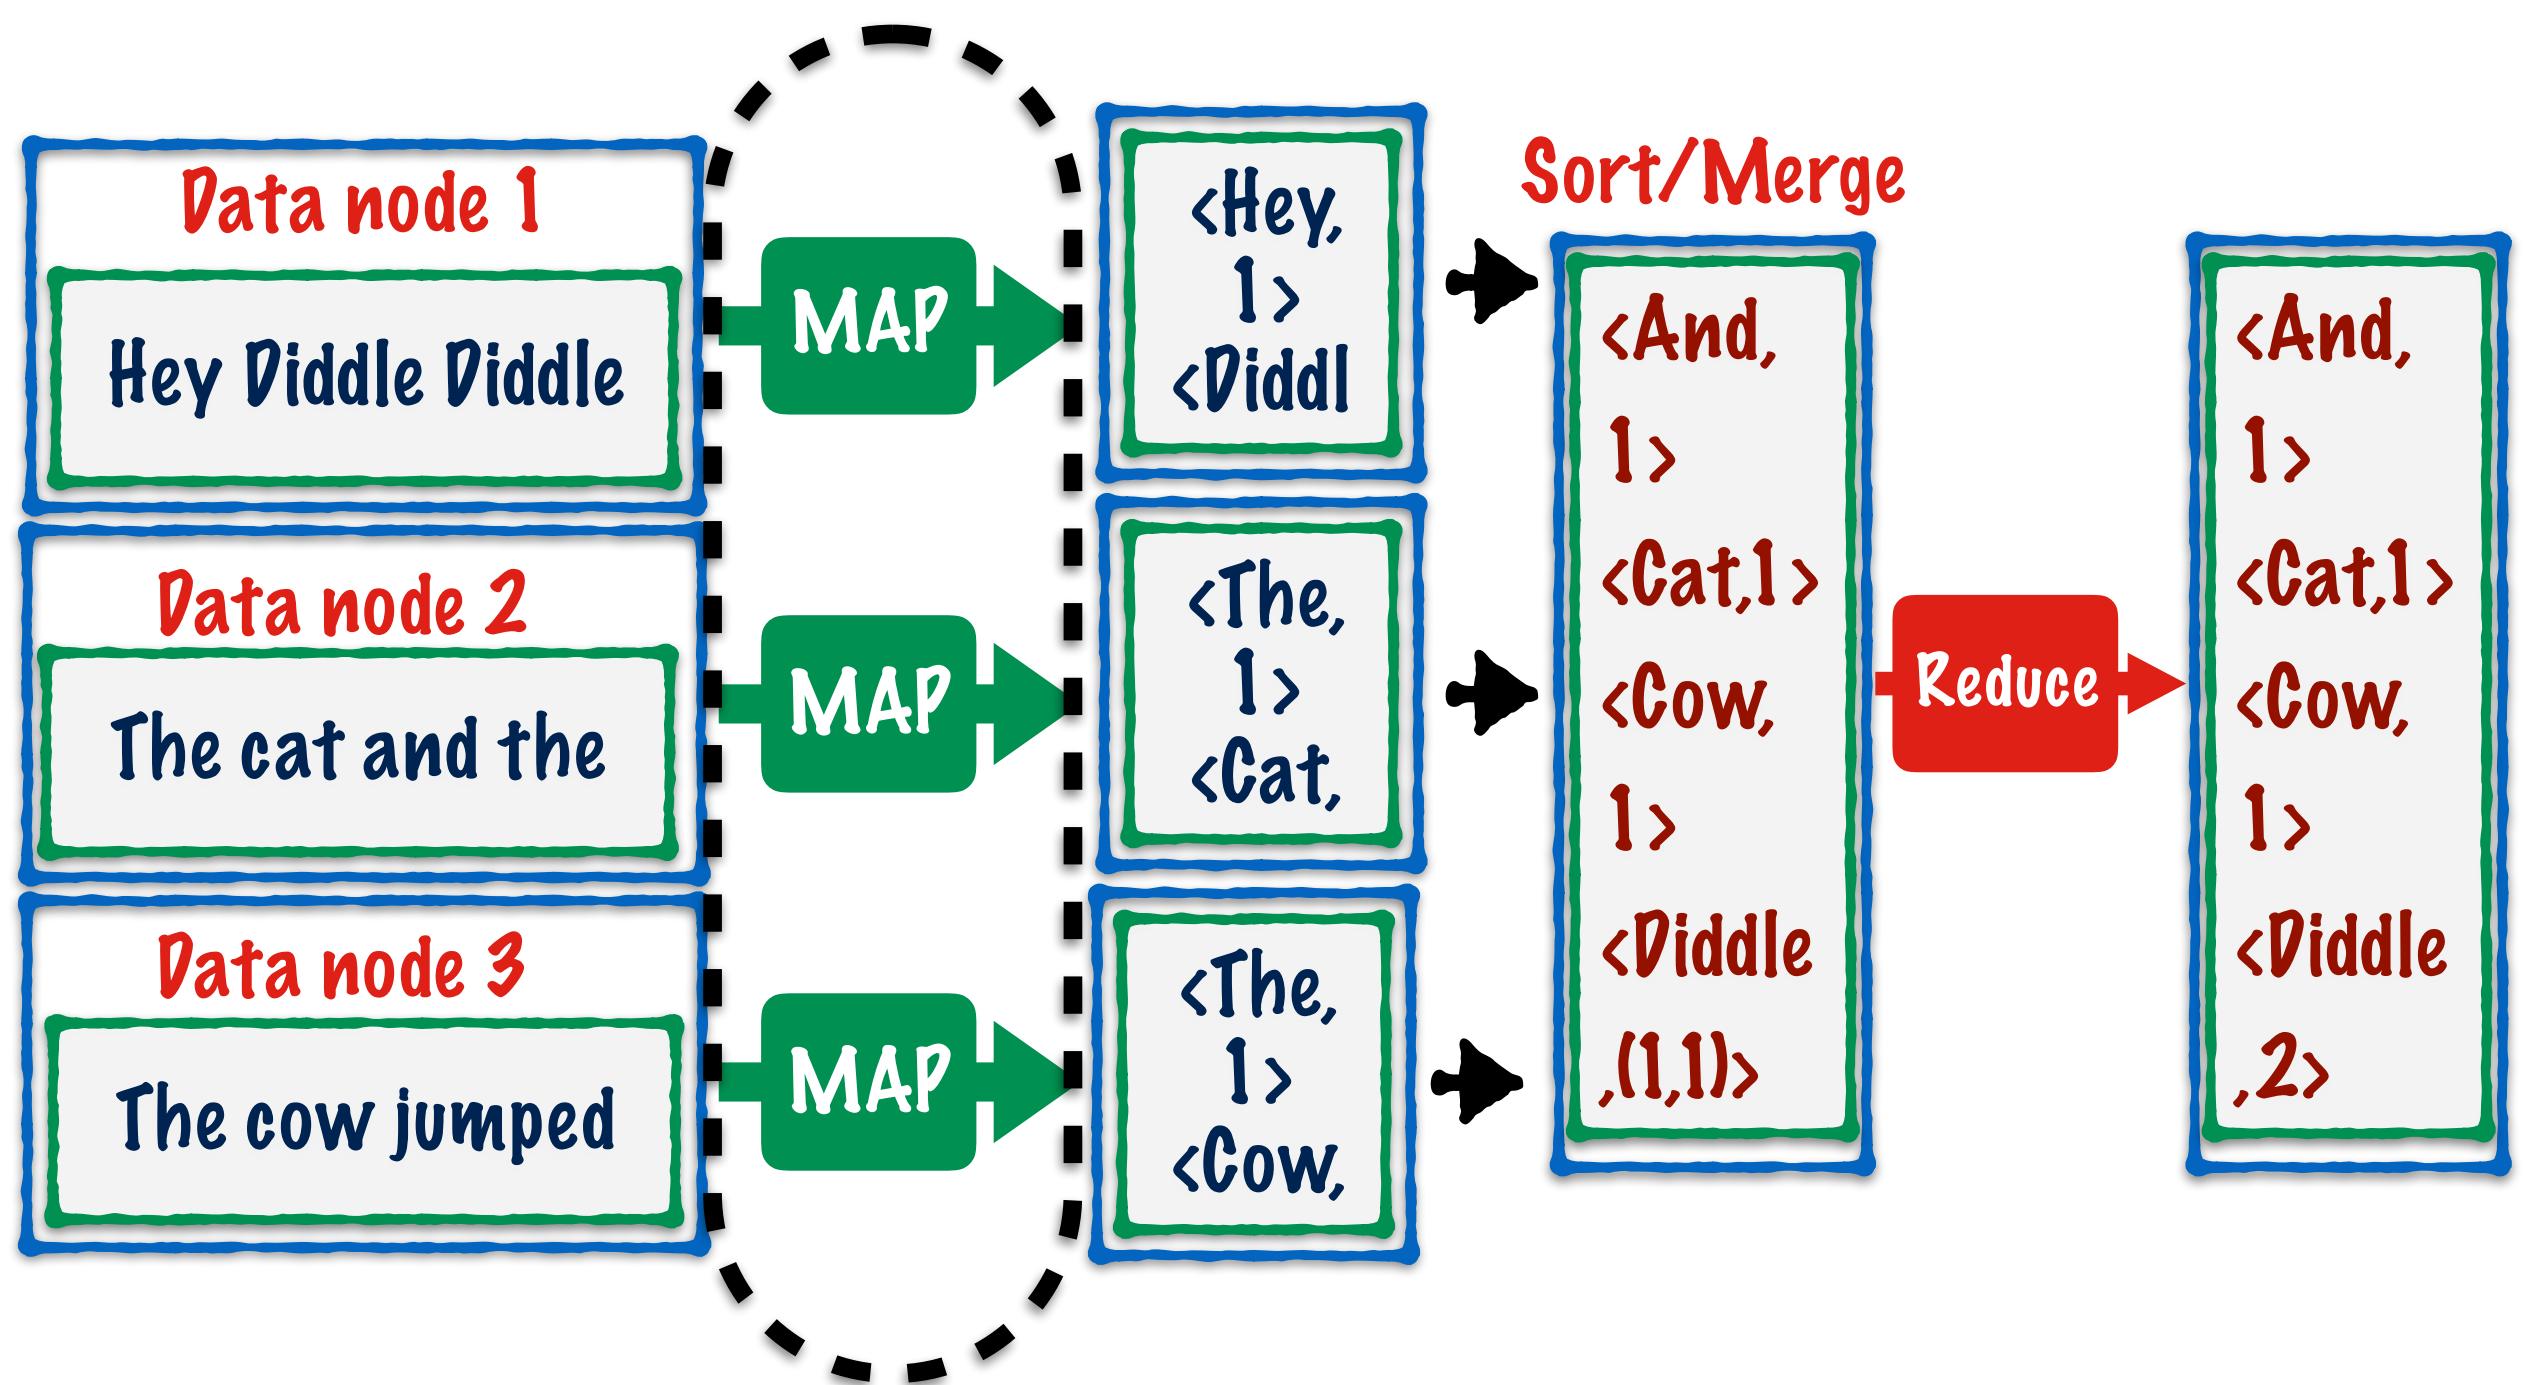
Shuffle and Sort Using Multiple Reducers

Shuffle and Sort

Here is a MapReduce representation that we are familiar with



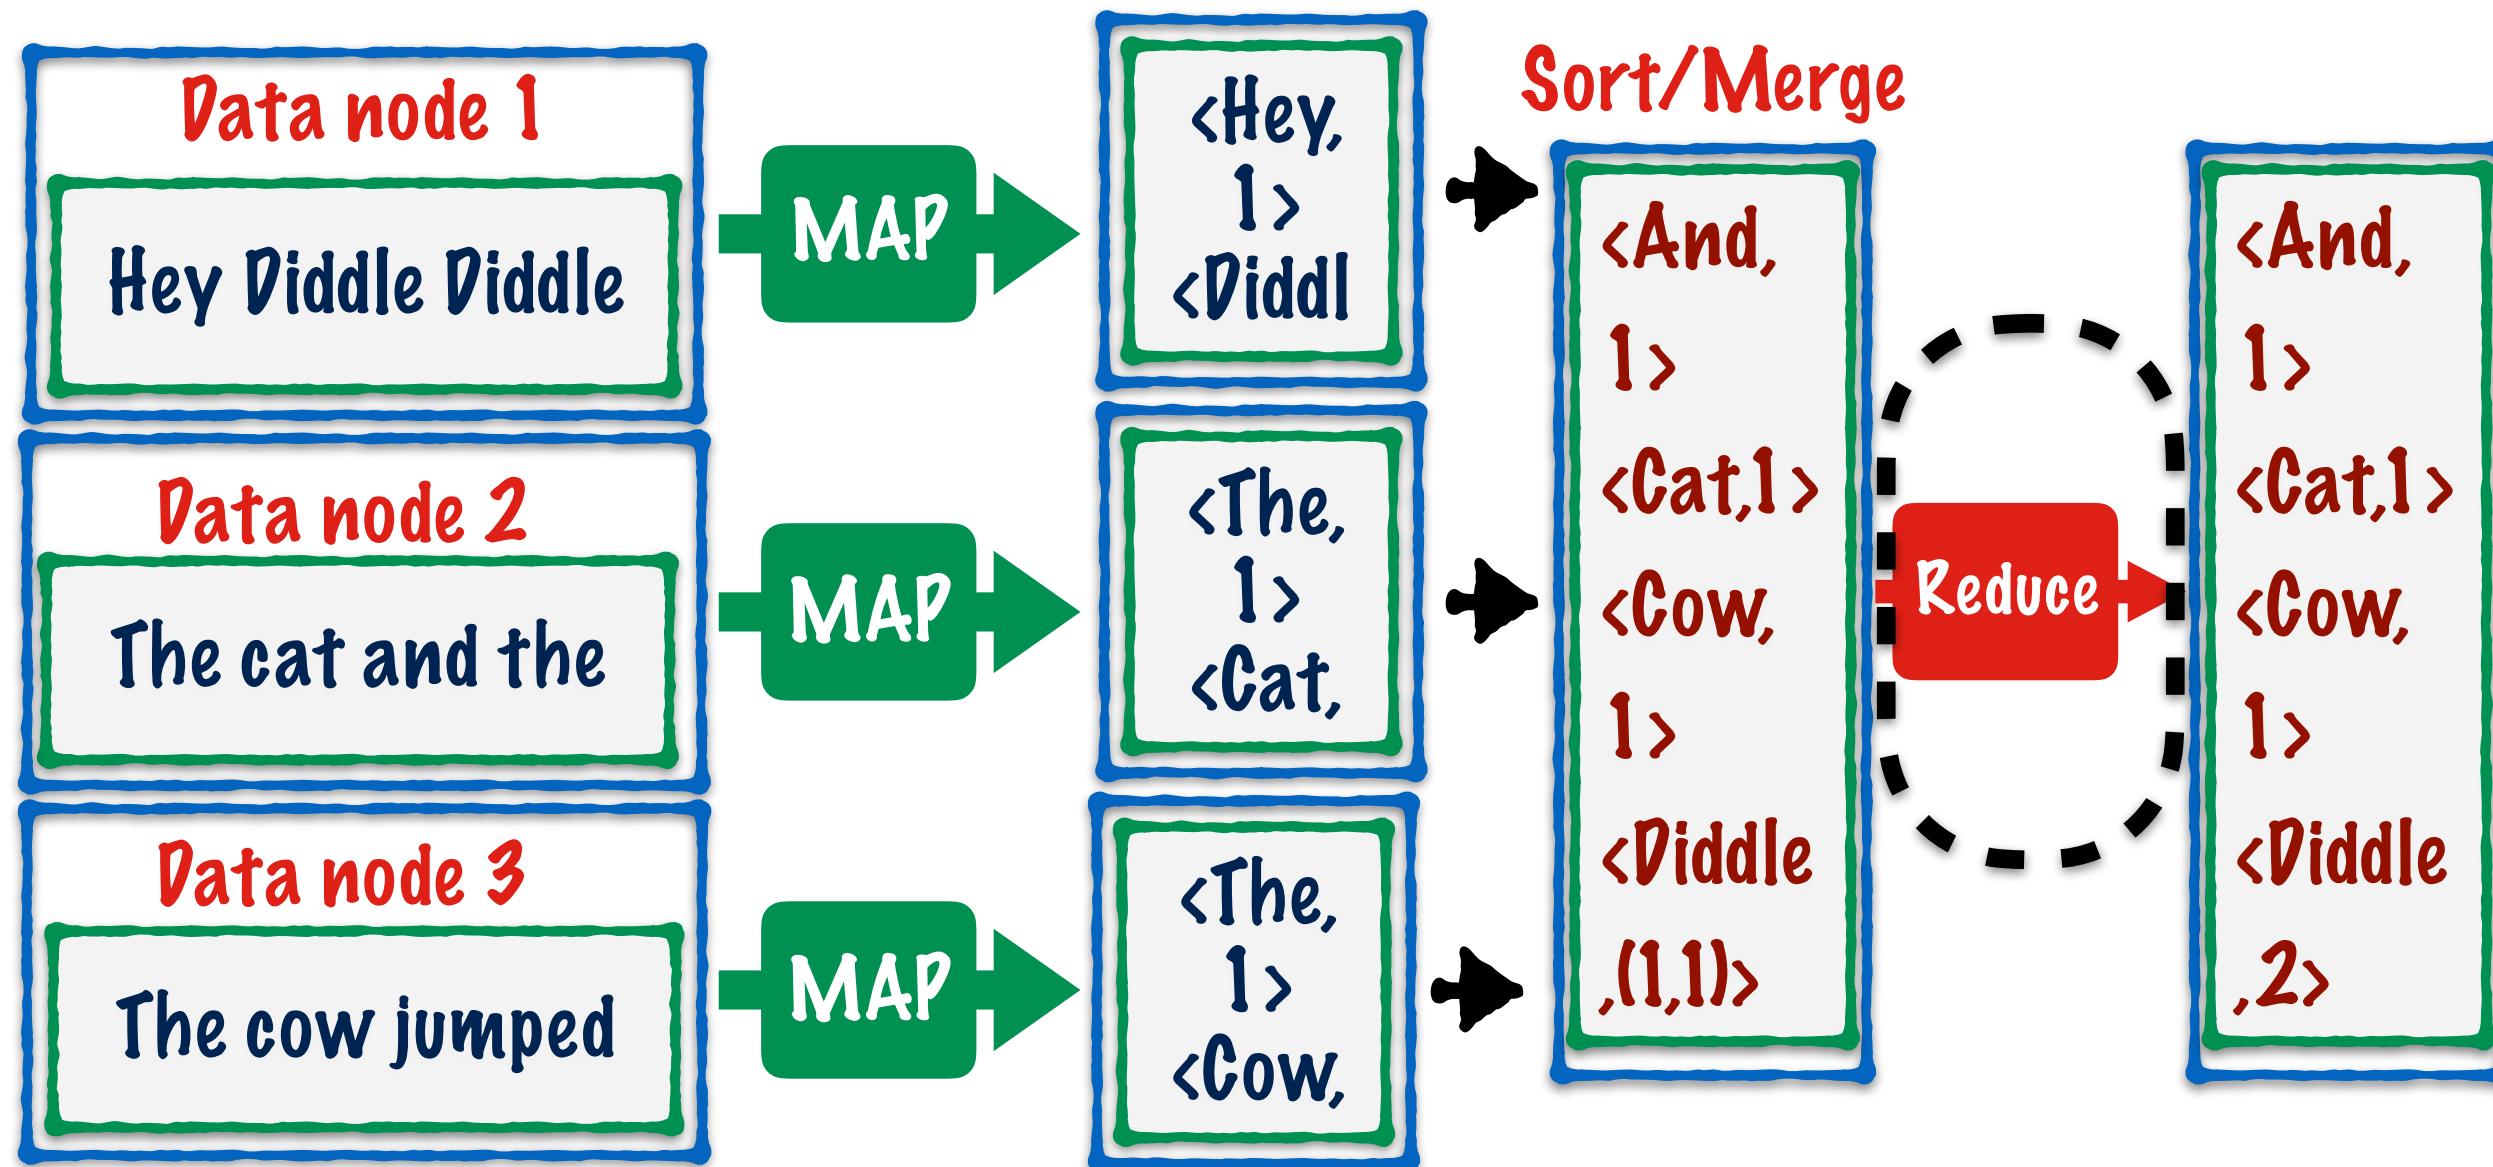
Shuffle and Sort



These represent a bunch of **map tasks** that run **in parallel**

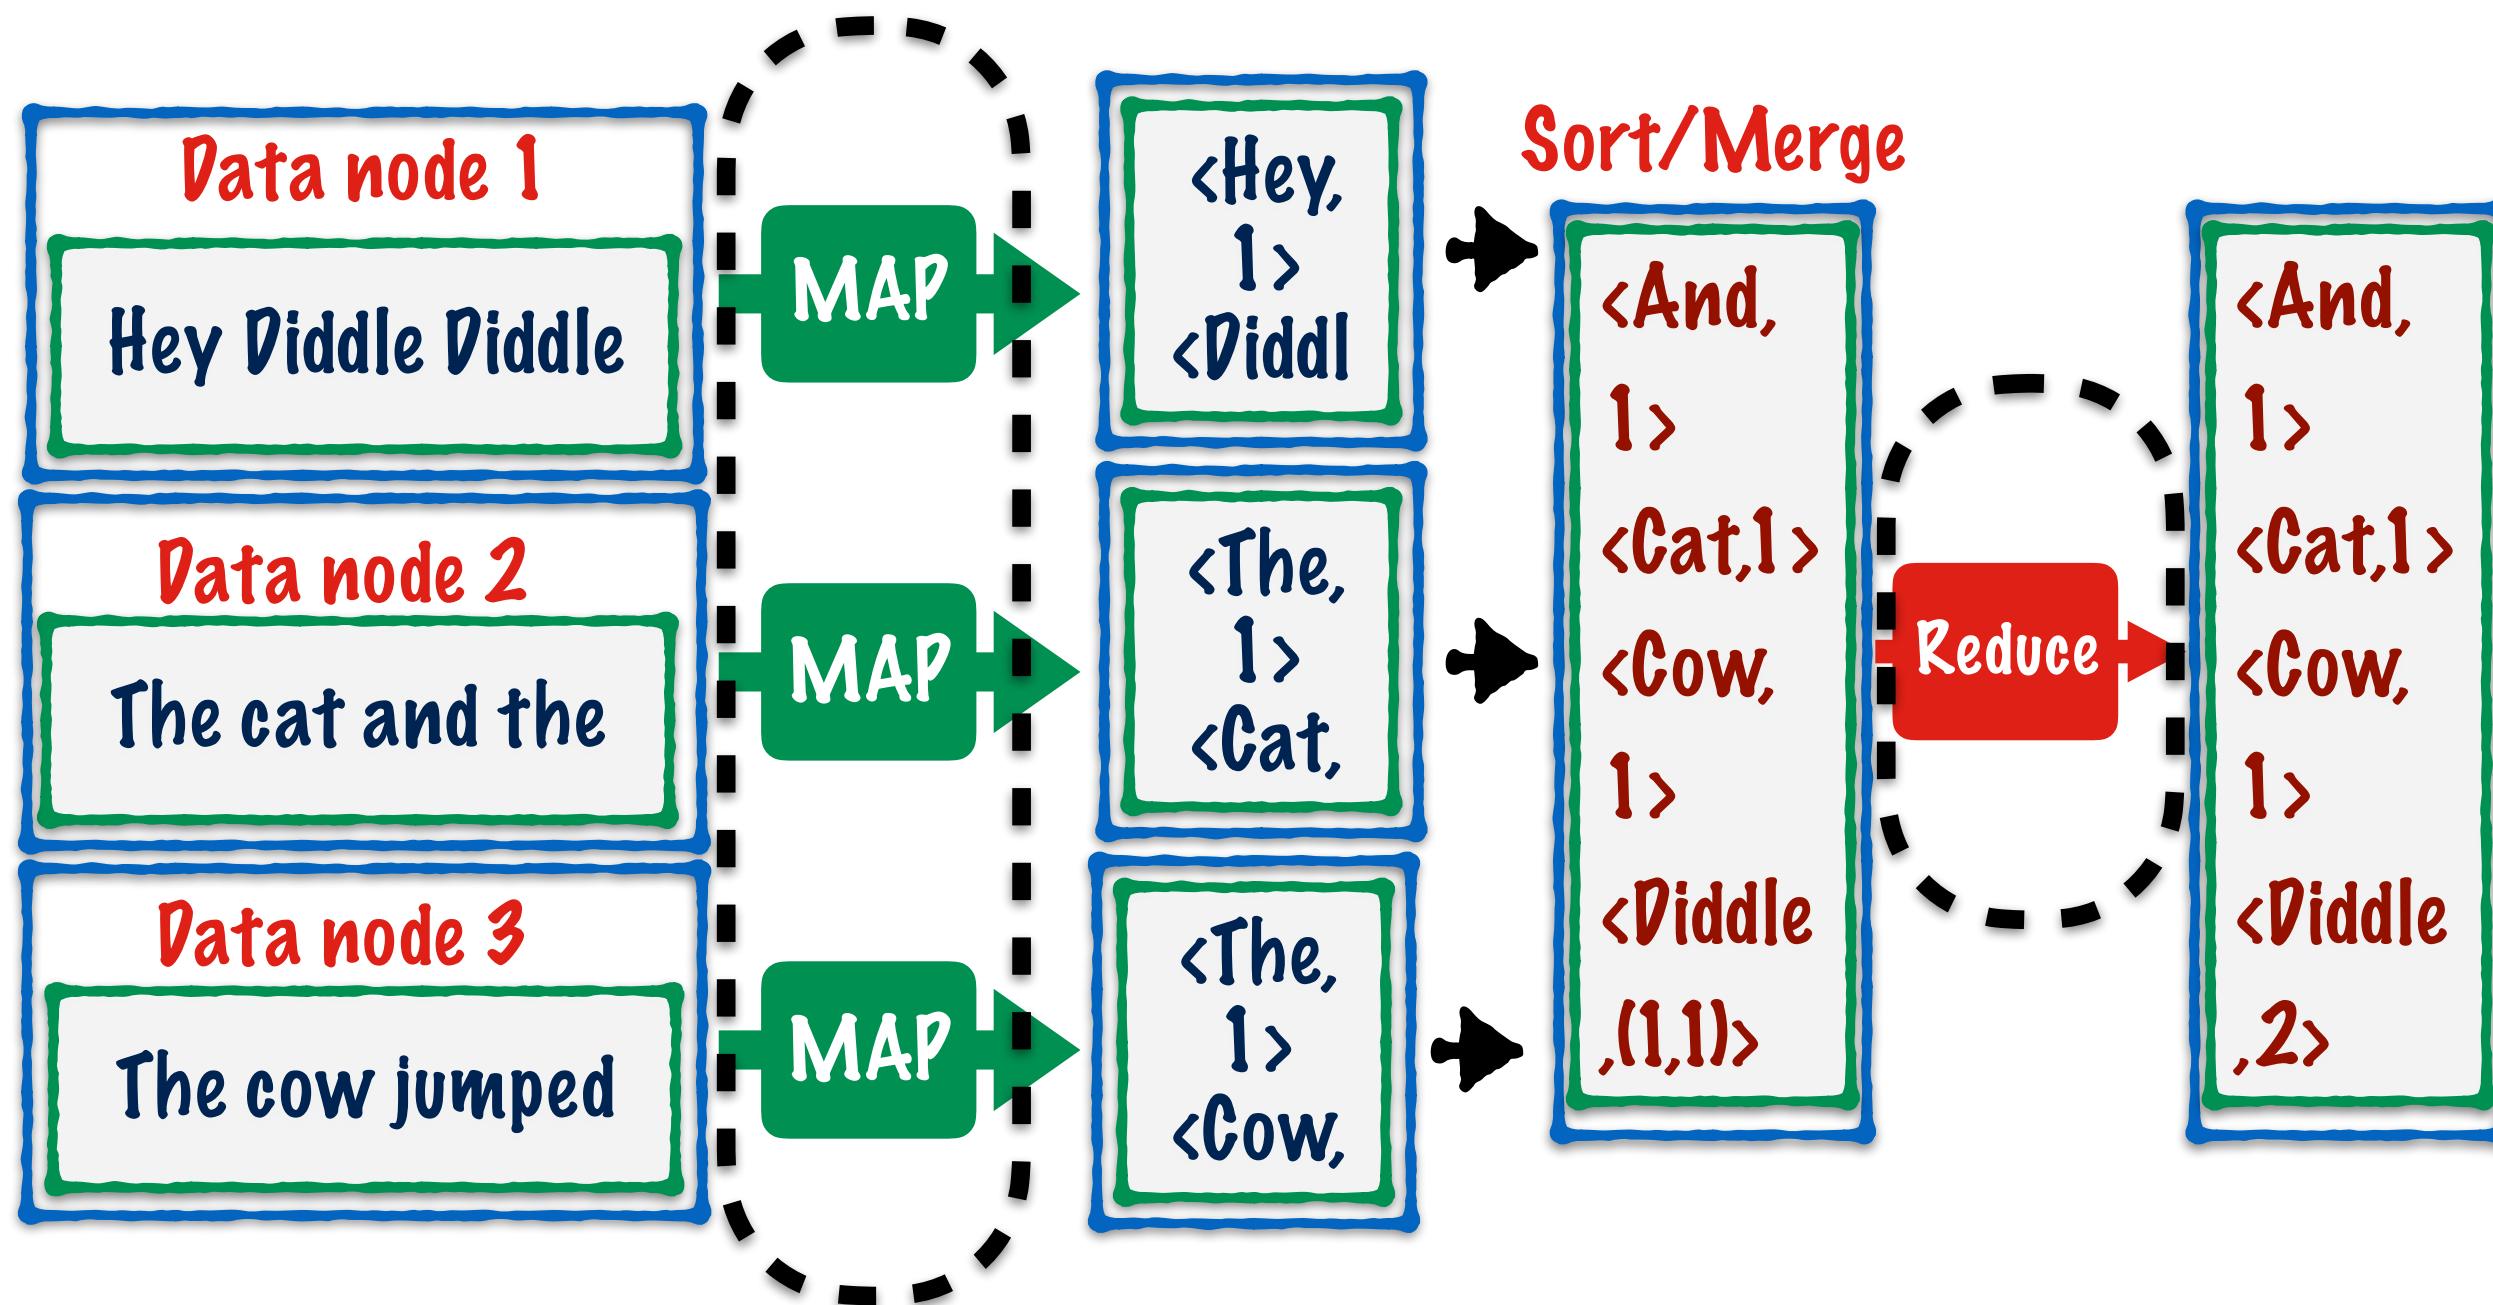
Each of these map tasks is executed by a “process” on a node in the cluster

Shuffle and Sort



Similarly, the reduce task is also executed by a “process”

Shuffle and Sort



The map processes
are called
mappers

All these processes
are launched and
managed by **YARN**

The reduce
processes are called
reducers

Shuffle and Sort

mappers

reducers

How many mappers and reducers
does a MapReduce job require?

Shuffle and Sort

mappers

reducers

The number of
mappers is usually
controlled by
YARN

Shuffle and Sort

mappers

reducers

Each mapper
operates on
one HDFS block

Shuffle and Sort

mappers

reducers

It depends on the
number of “splits” i.e.
the number of blocks
your input file is
broken into

Shuffle and Sort

mappers

= number of “splits”

reducers

This effectively means
multiple mappers can be
assigned to 1 Data node

Shuffle and Sort

mappers

= number of “splits”

The number of
mappers can only be
controlled to a limited
extent by the user

reducers

Shuffle and Sort

mappers

= number of “splits”

reducers

The number of
reducers can be
controlled by the
user

Shuffle and Sort

mappers

= number of “splits”

reducers

By default, there
is only 1 reducer
which is what we
have seen so far

Shuffle and Sort

mappers

= number of “splits”

reducers

With multiple
reducers, the
reduce task can
also be parallelized

Shuffle and Sort

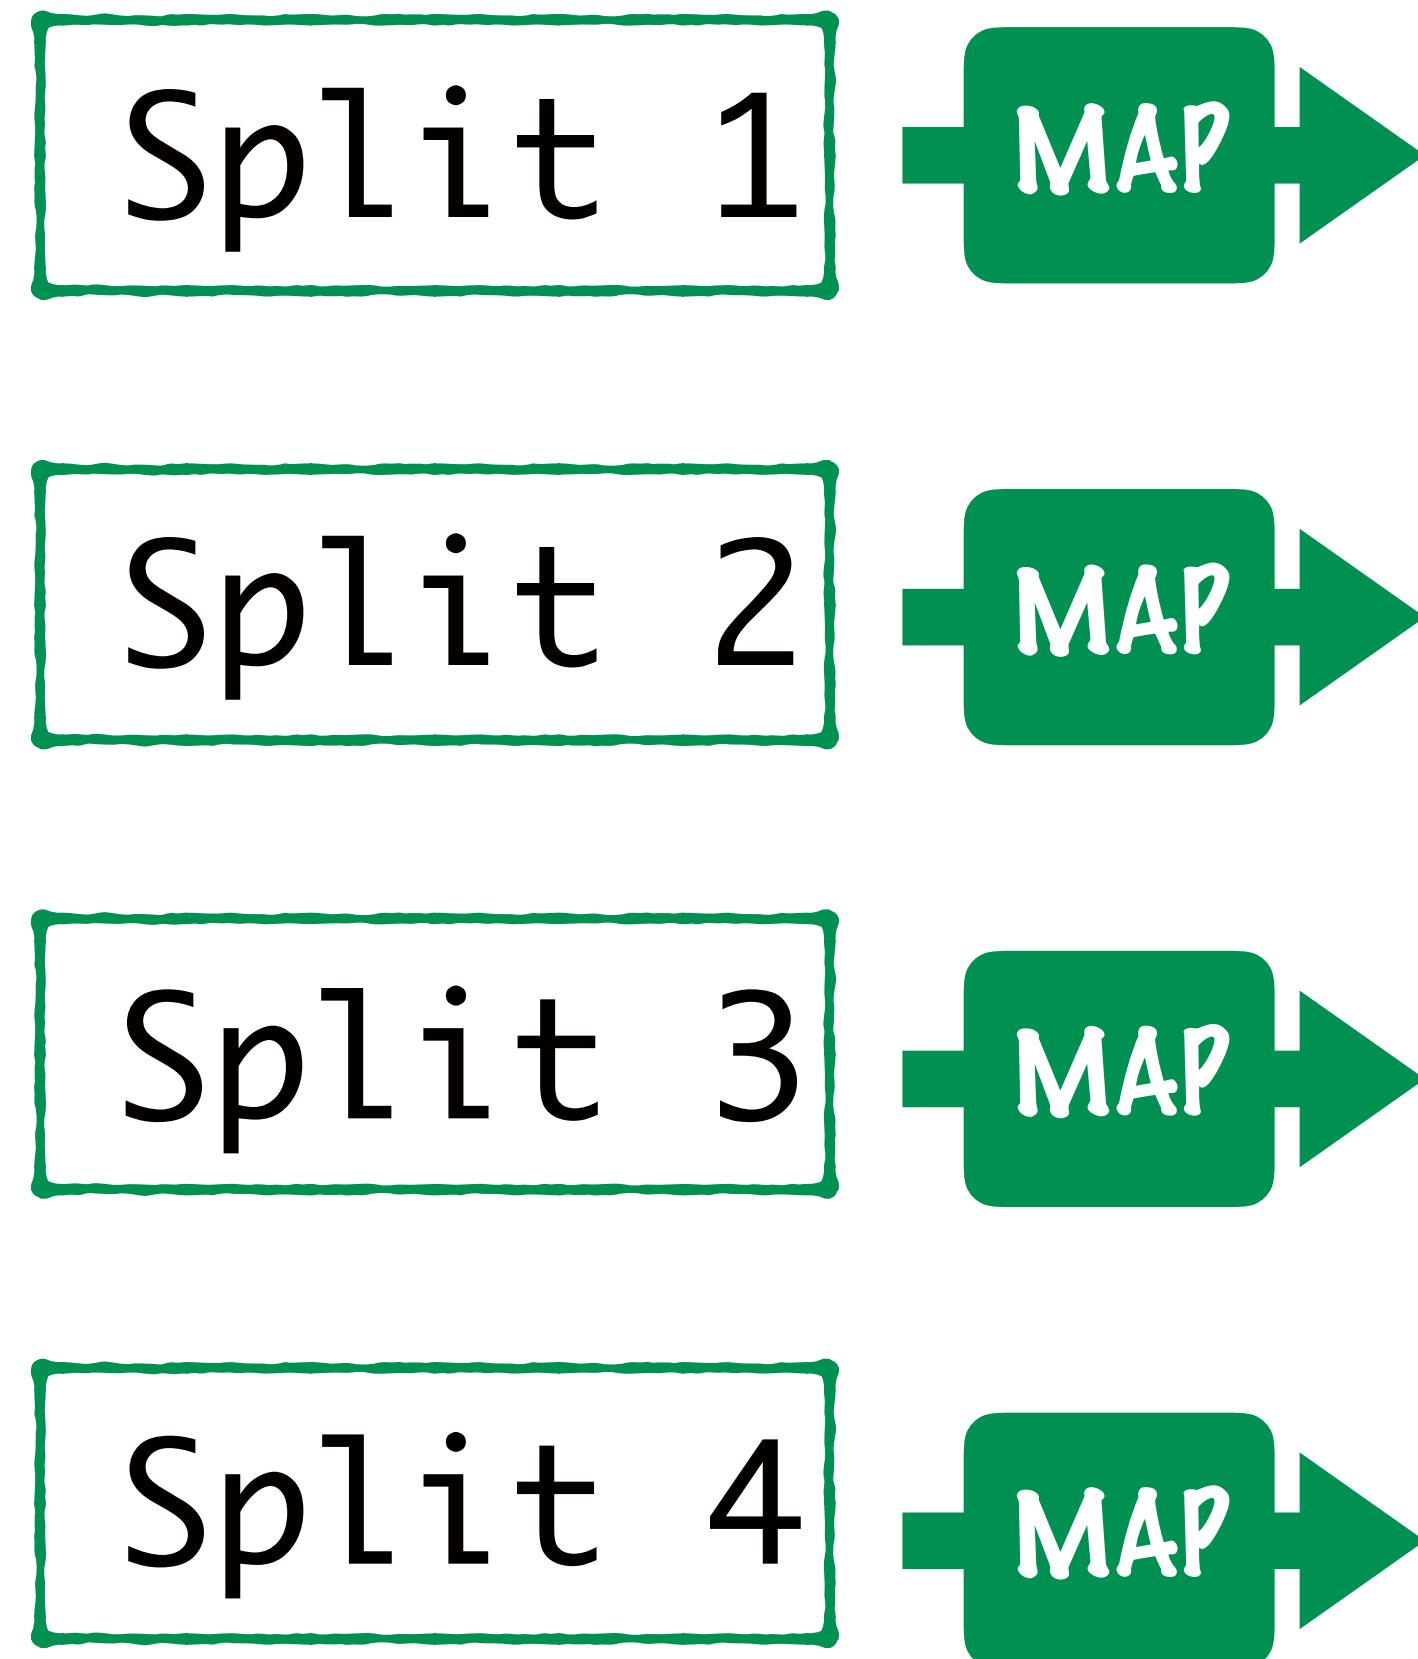
reducers

This is how you normally
run a MapReduce job

```
$ hadoop jar <JARpath> <mainclass> <inputFile> <outputDir>  
      -D mapred.reduce.tasks=2
```

The number of reducers can be
specified when you run the job

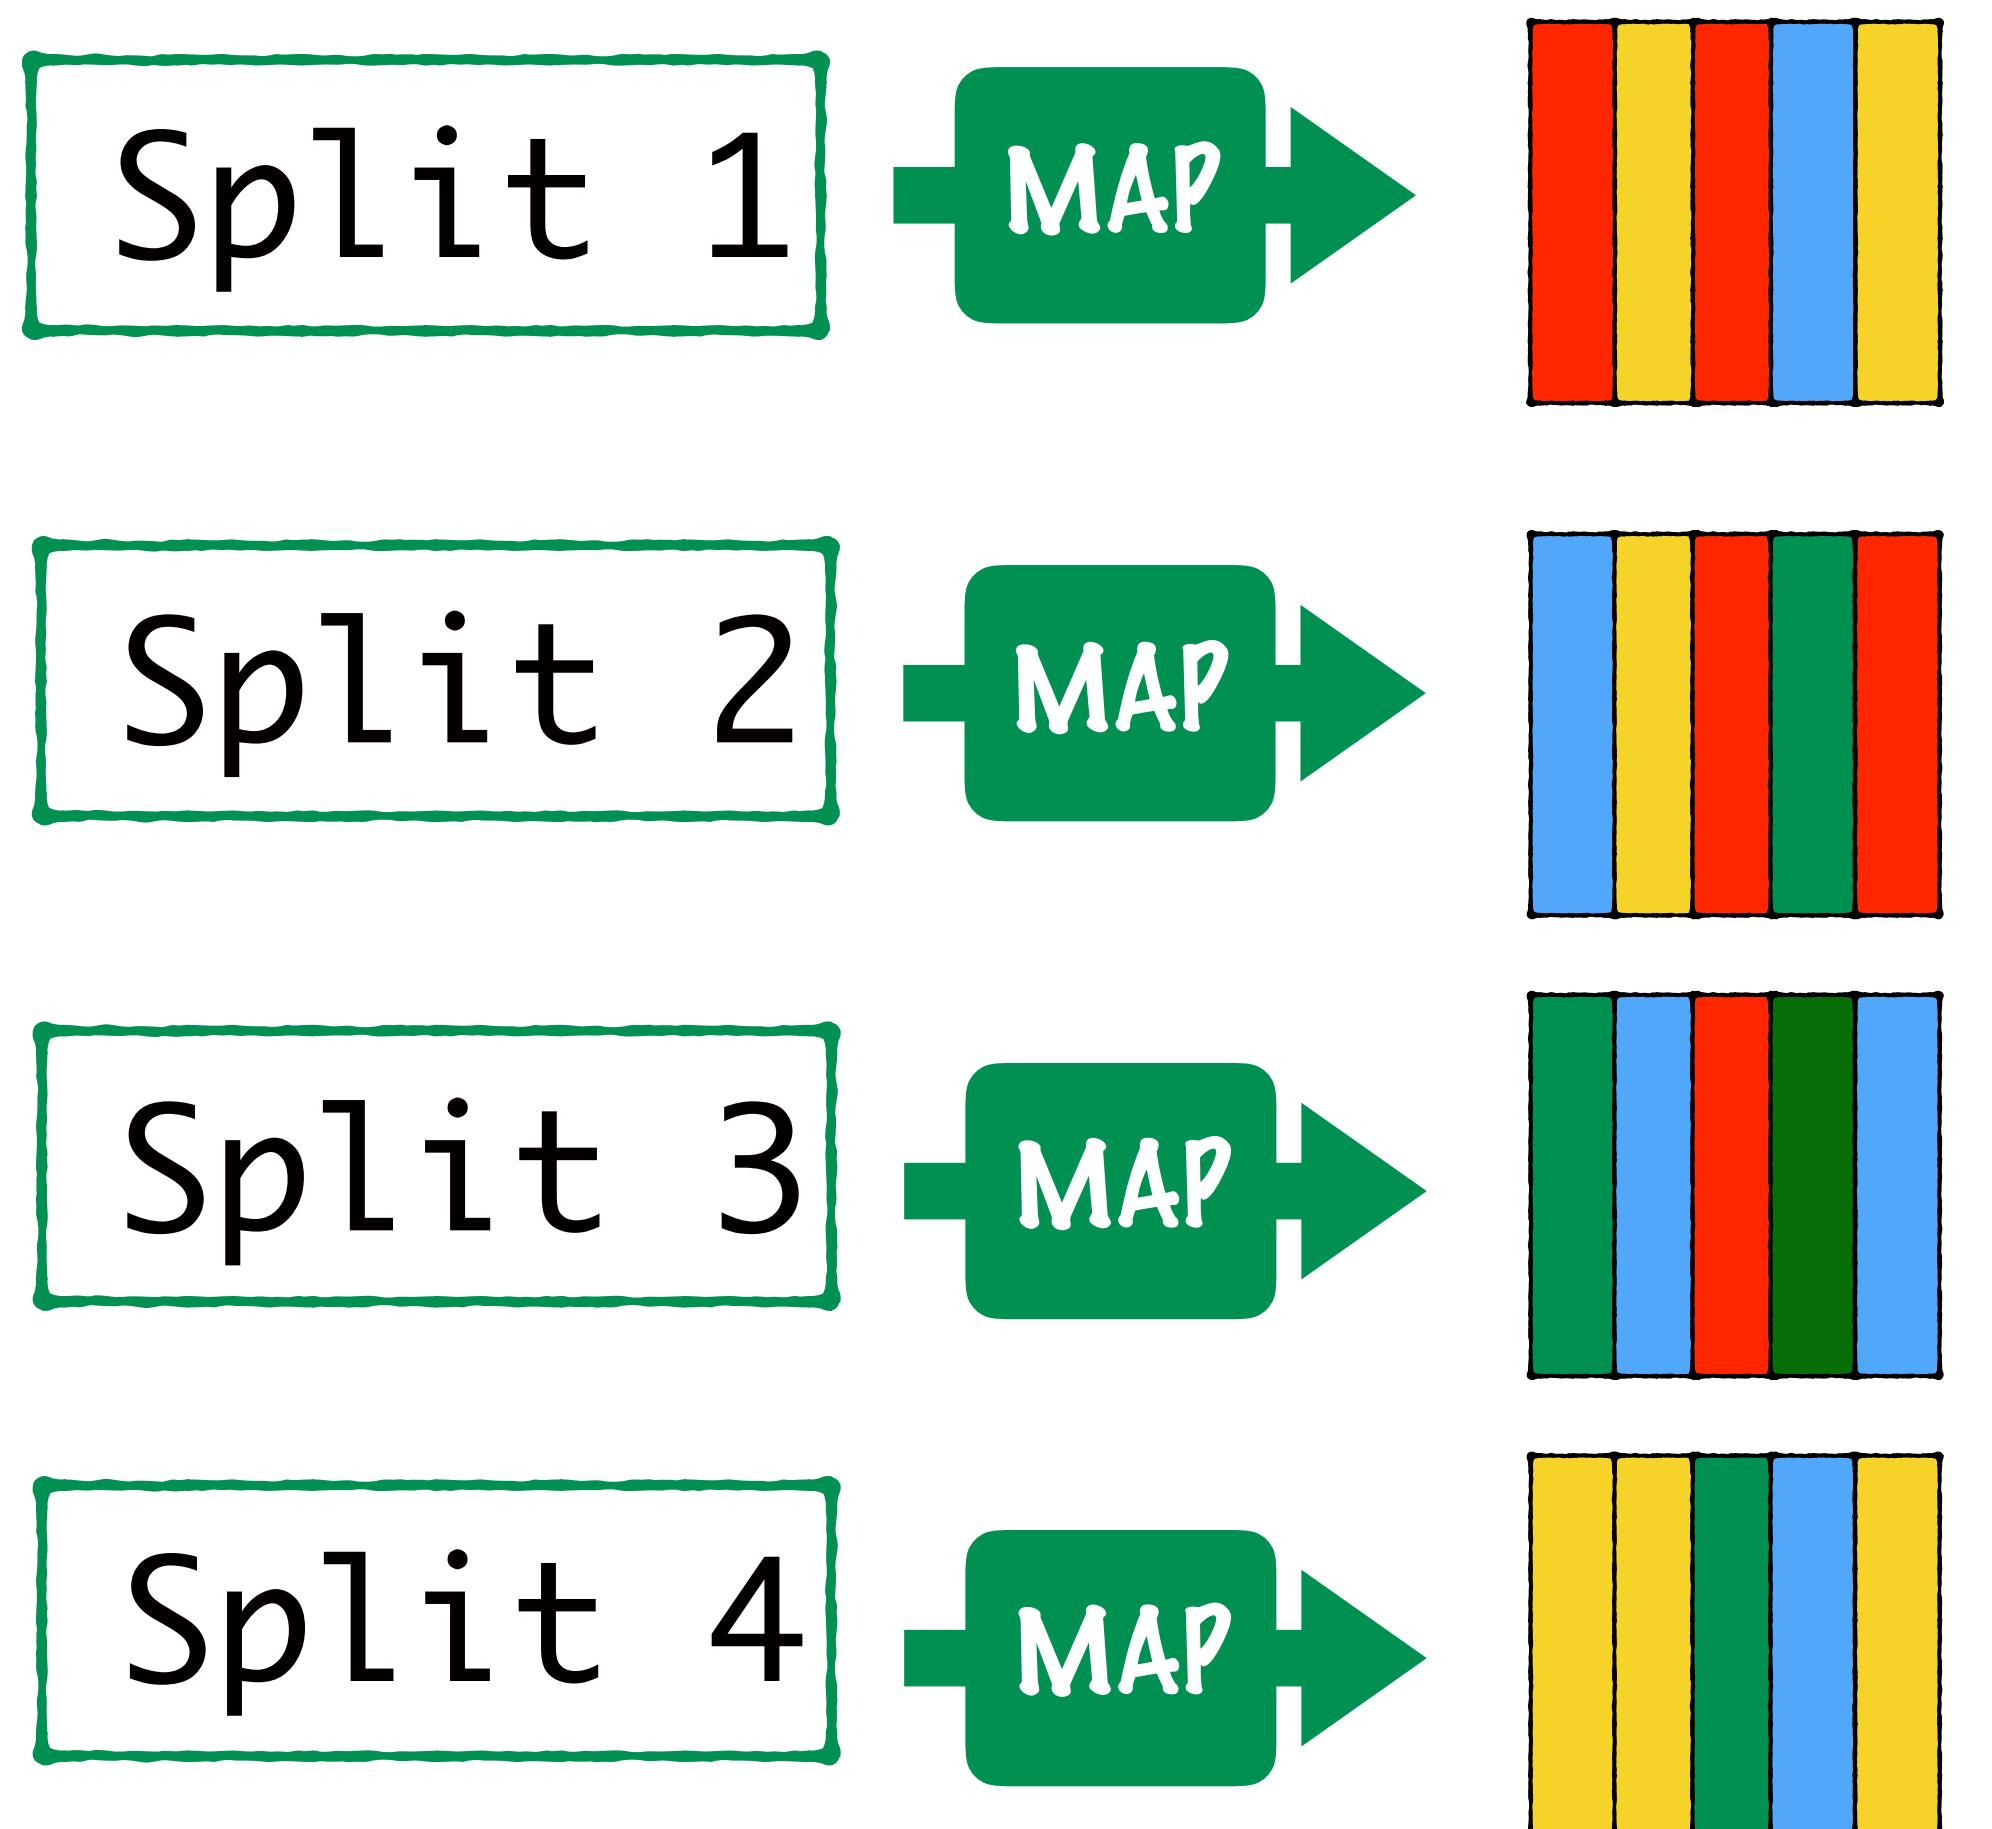
Shuffle and Sort



Let's say our input file
had been **split into 4**
blocks by HDFS

There will 4 mappers

Shuffle and Sort

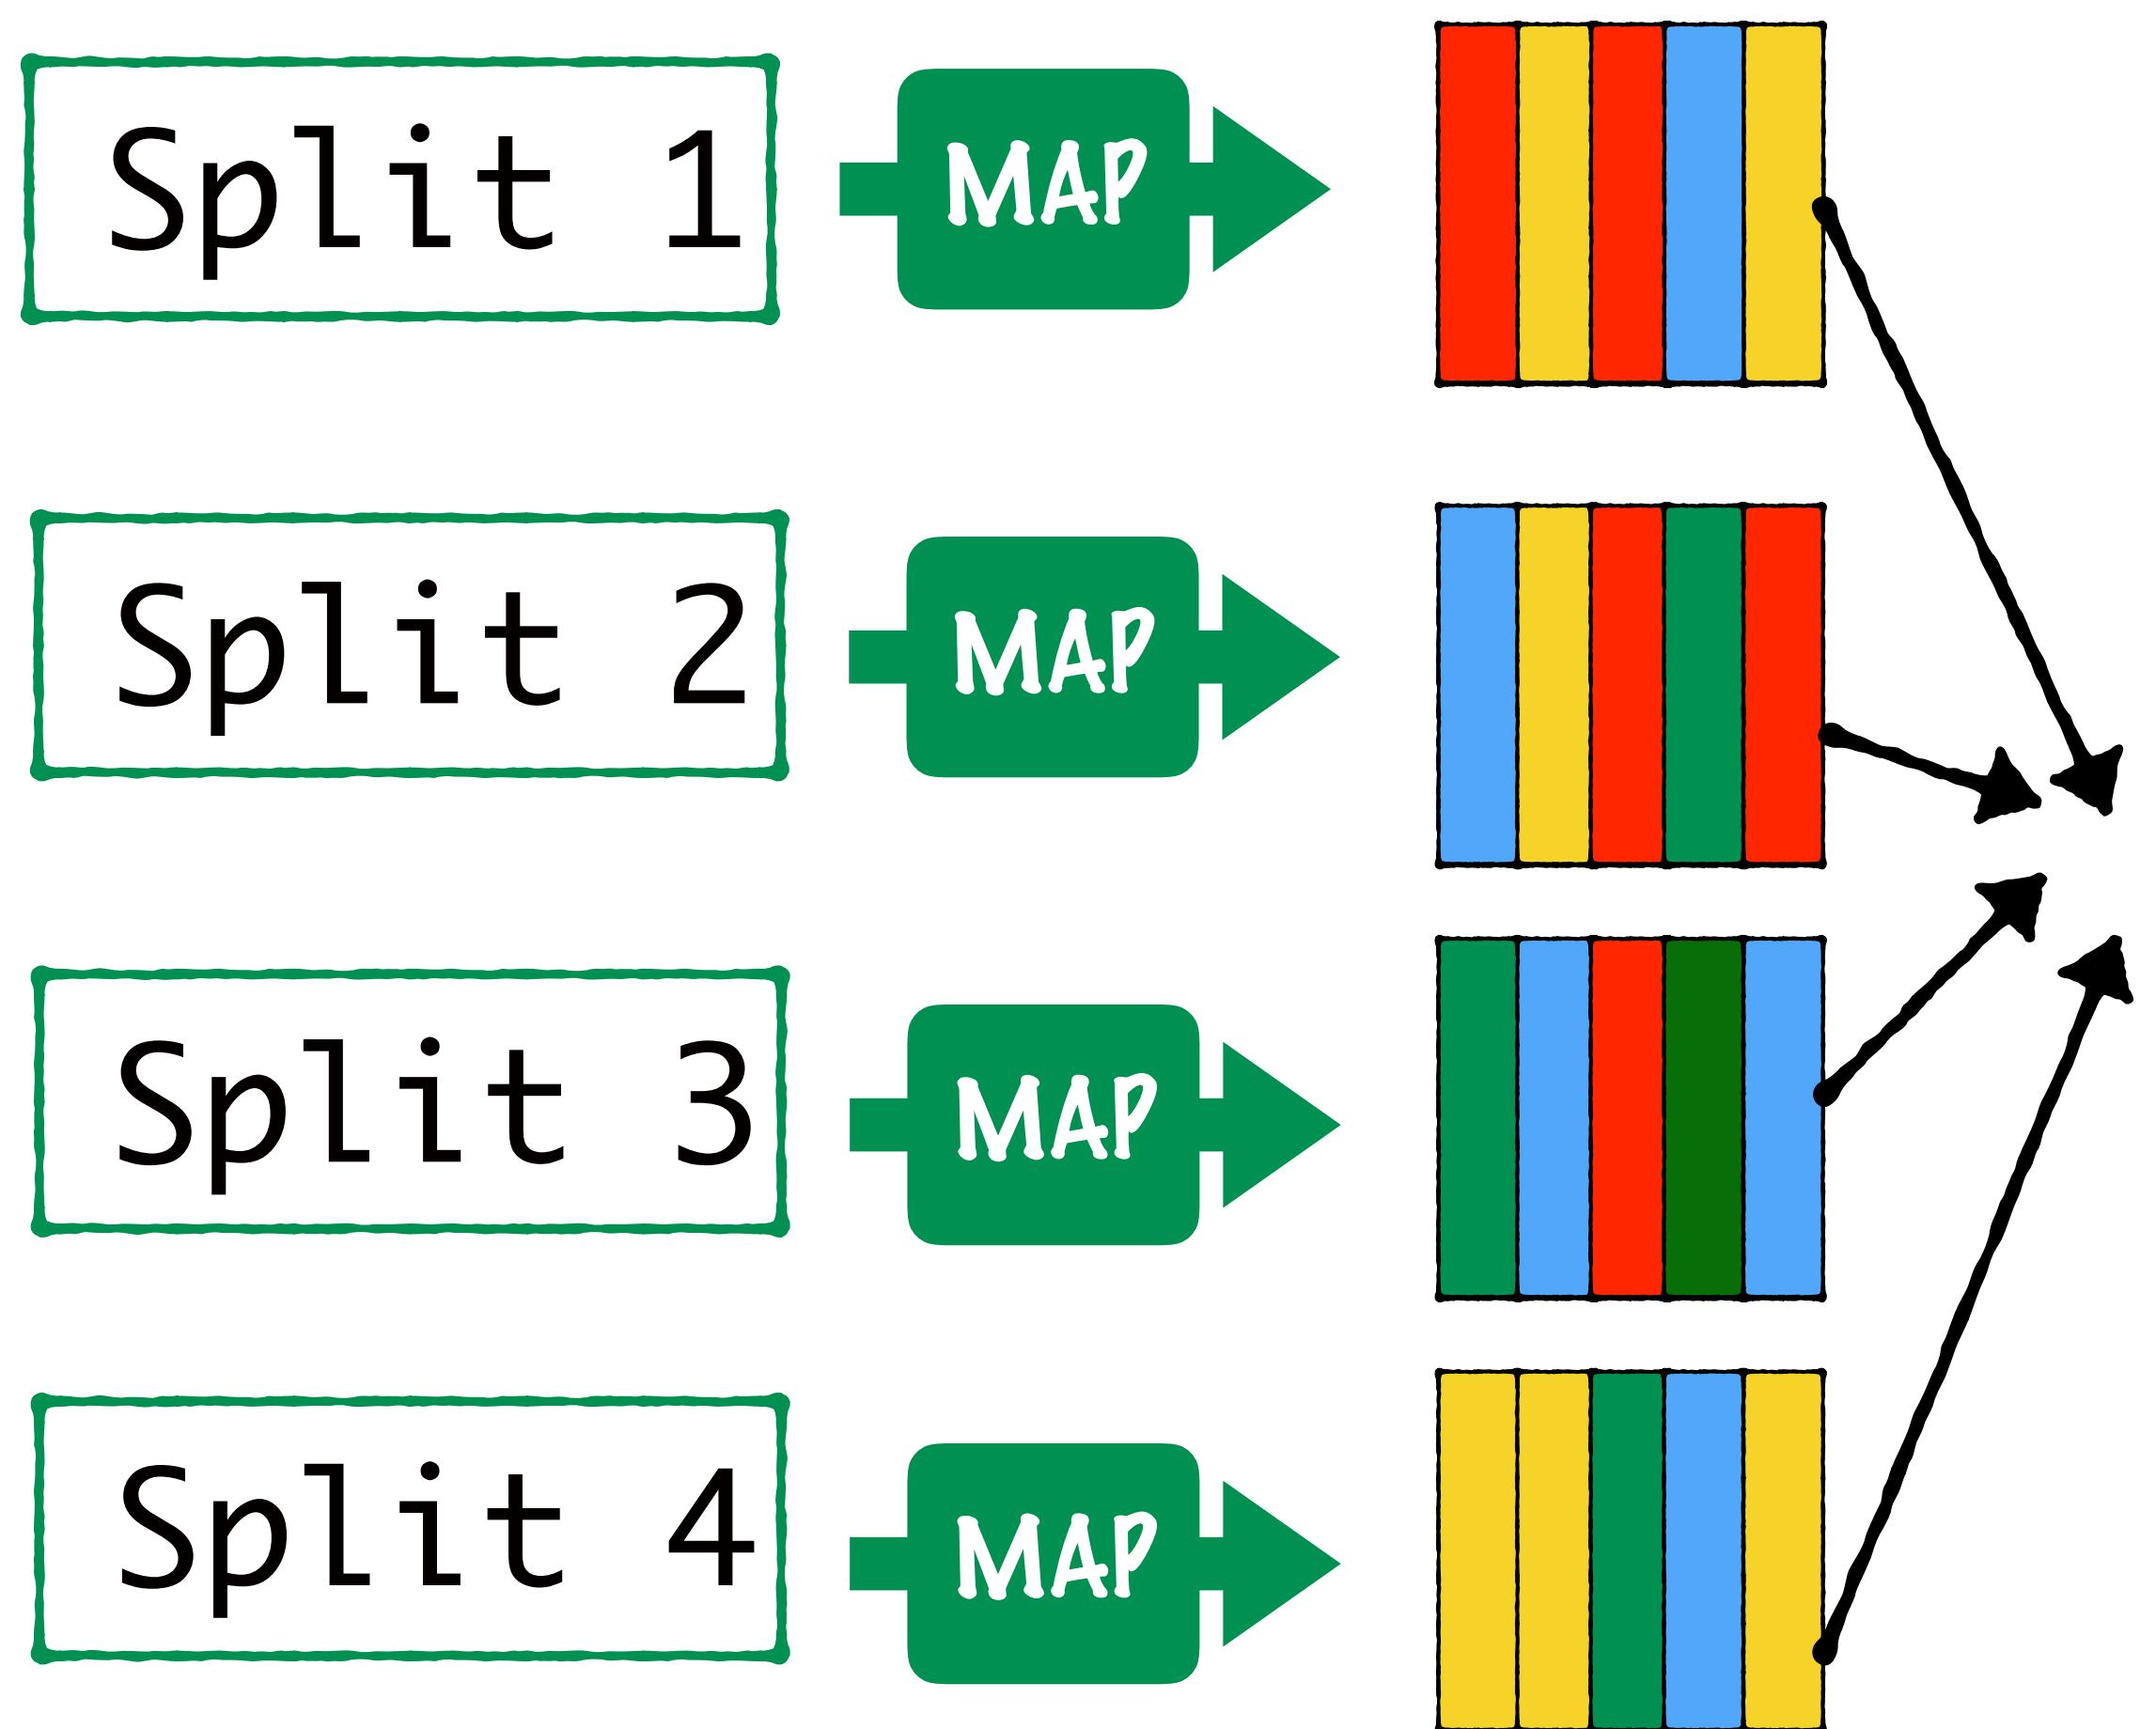


The mappers output
<Key, Value> pairs

Each block here represents
a **<Key, Value>** pair

The color
represents the key

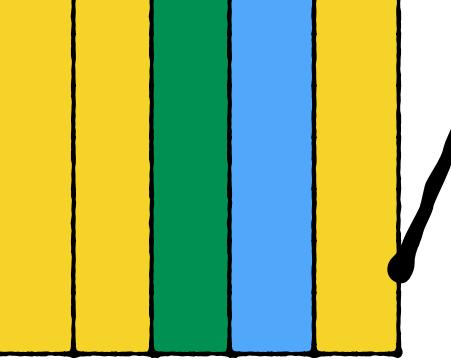
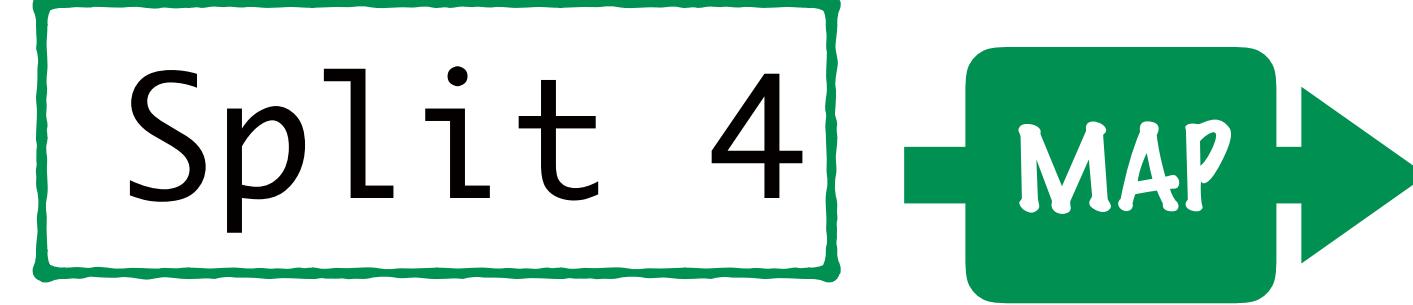
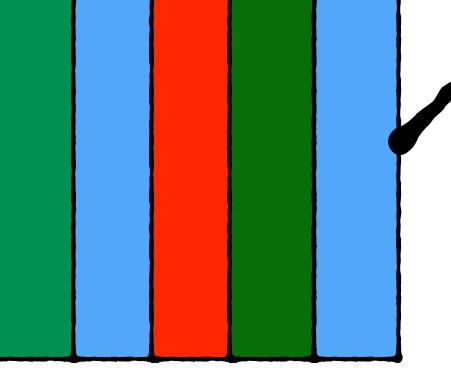
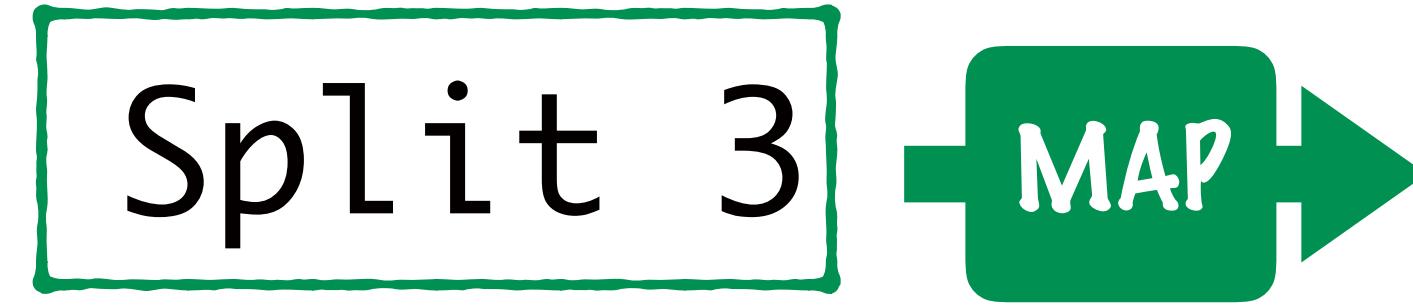
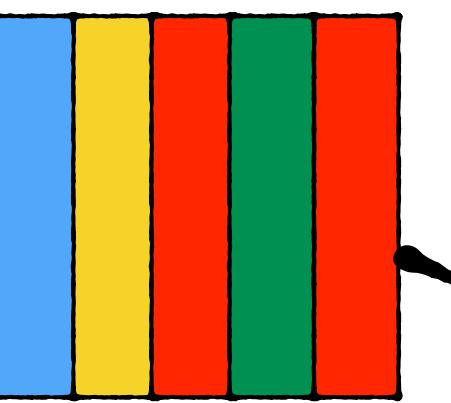
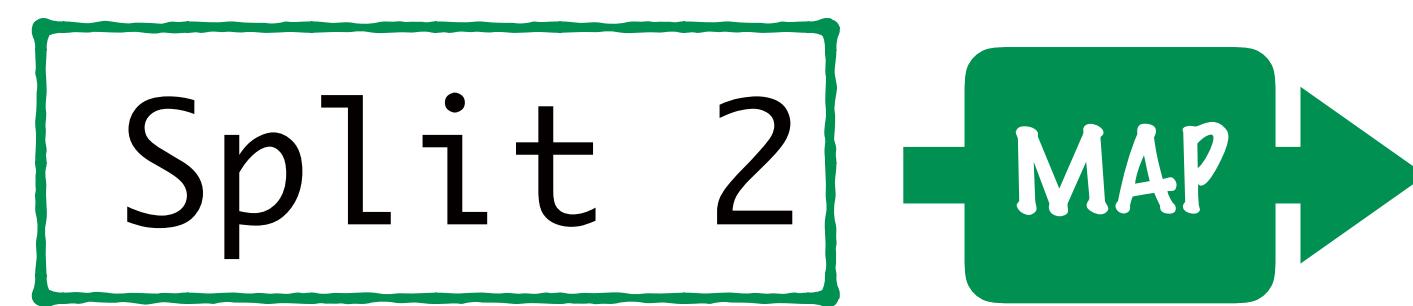
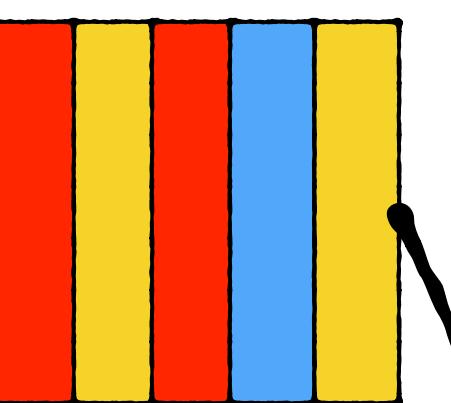
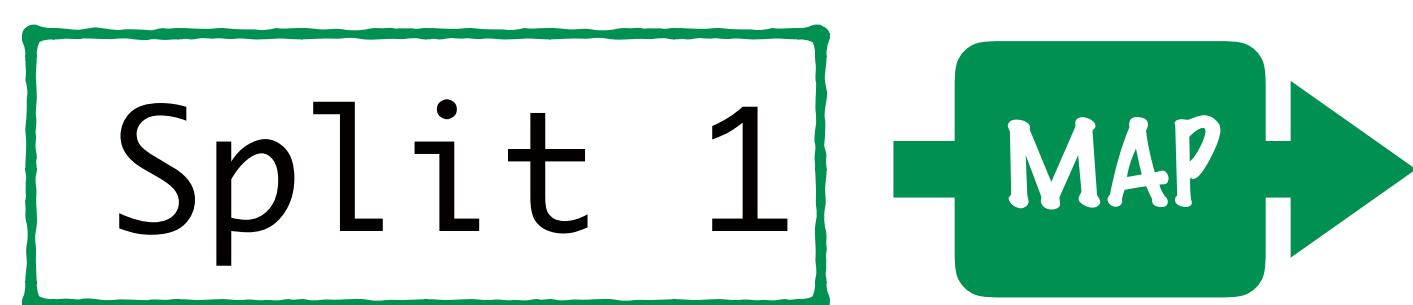
Shuffle and Sort



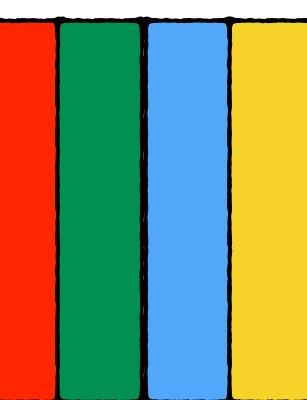
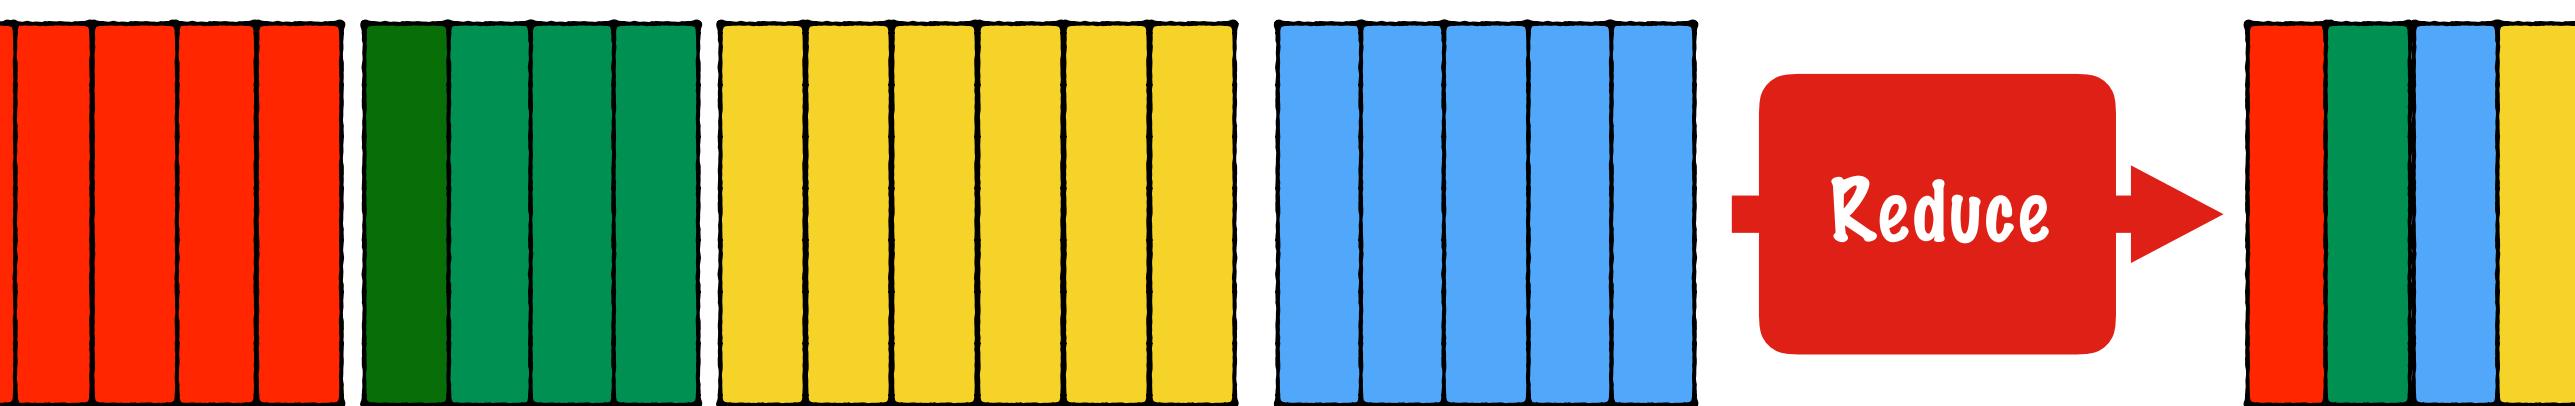
When there is only
1 reducer, these are
simply sorted by key
and then reduced

Shuffle and Sort

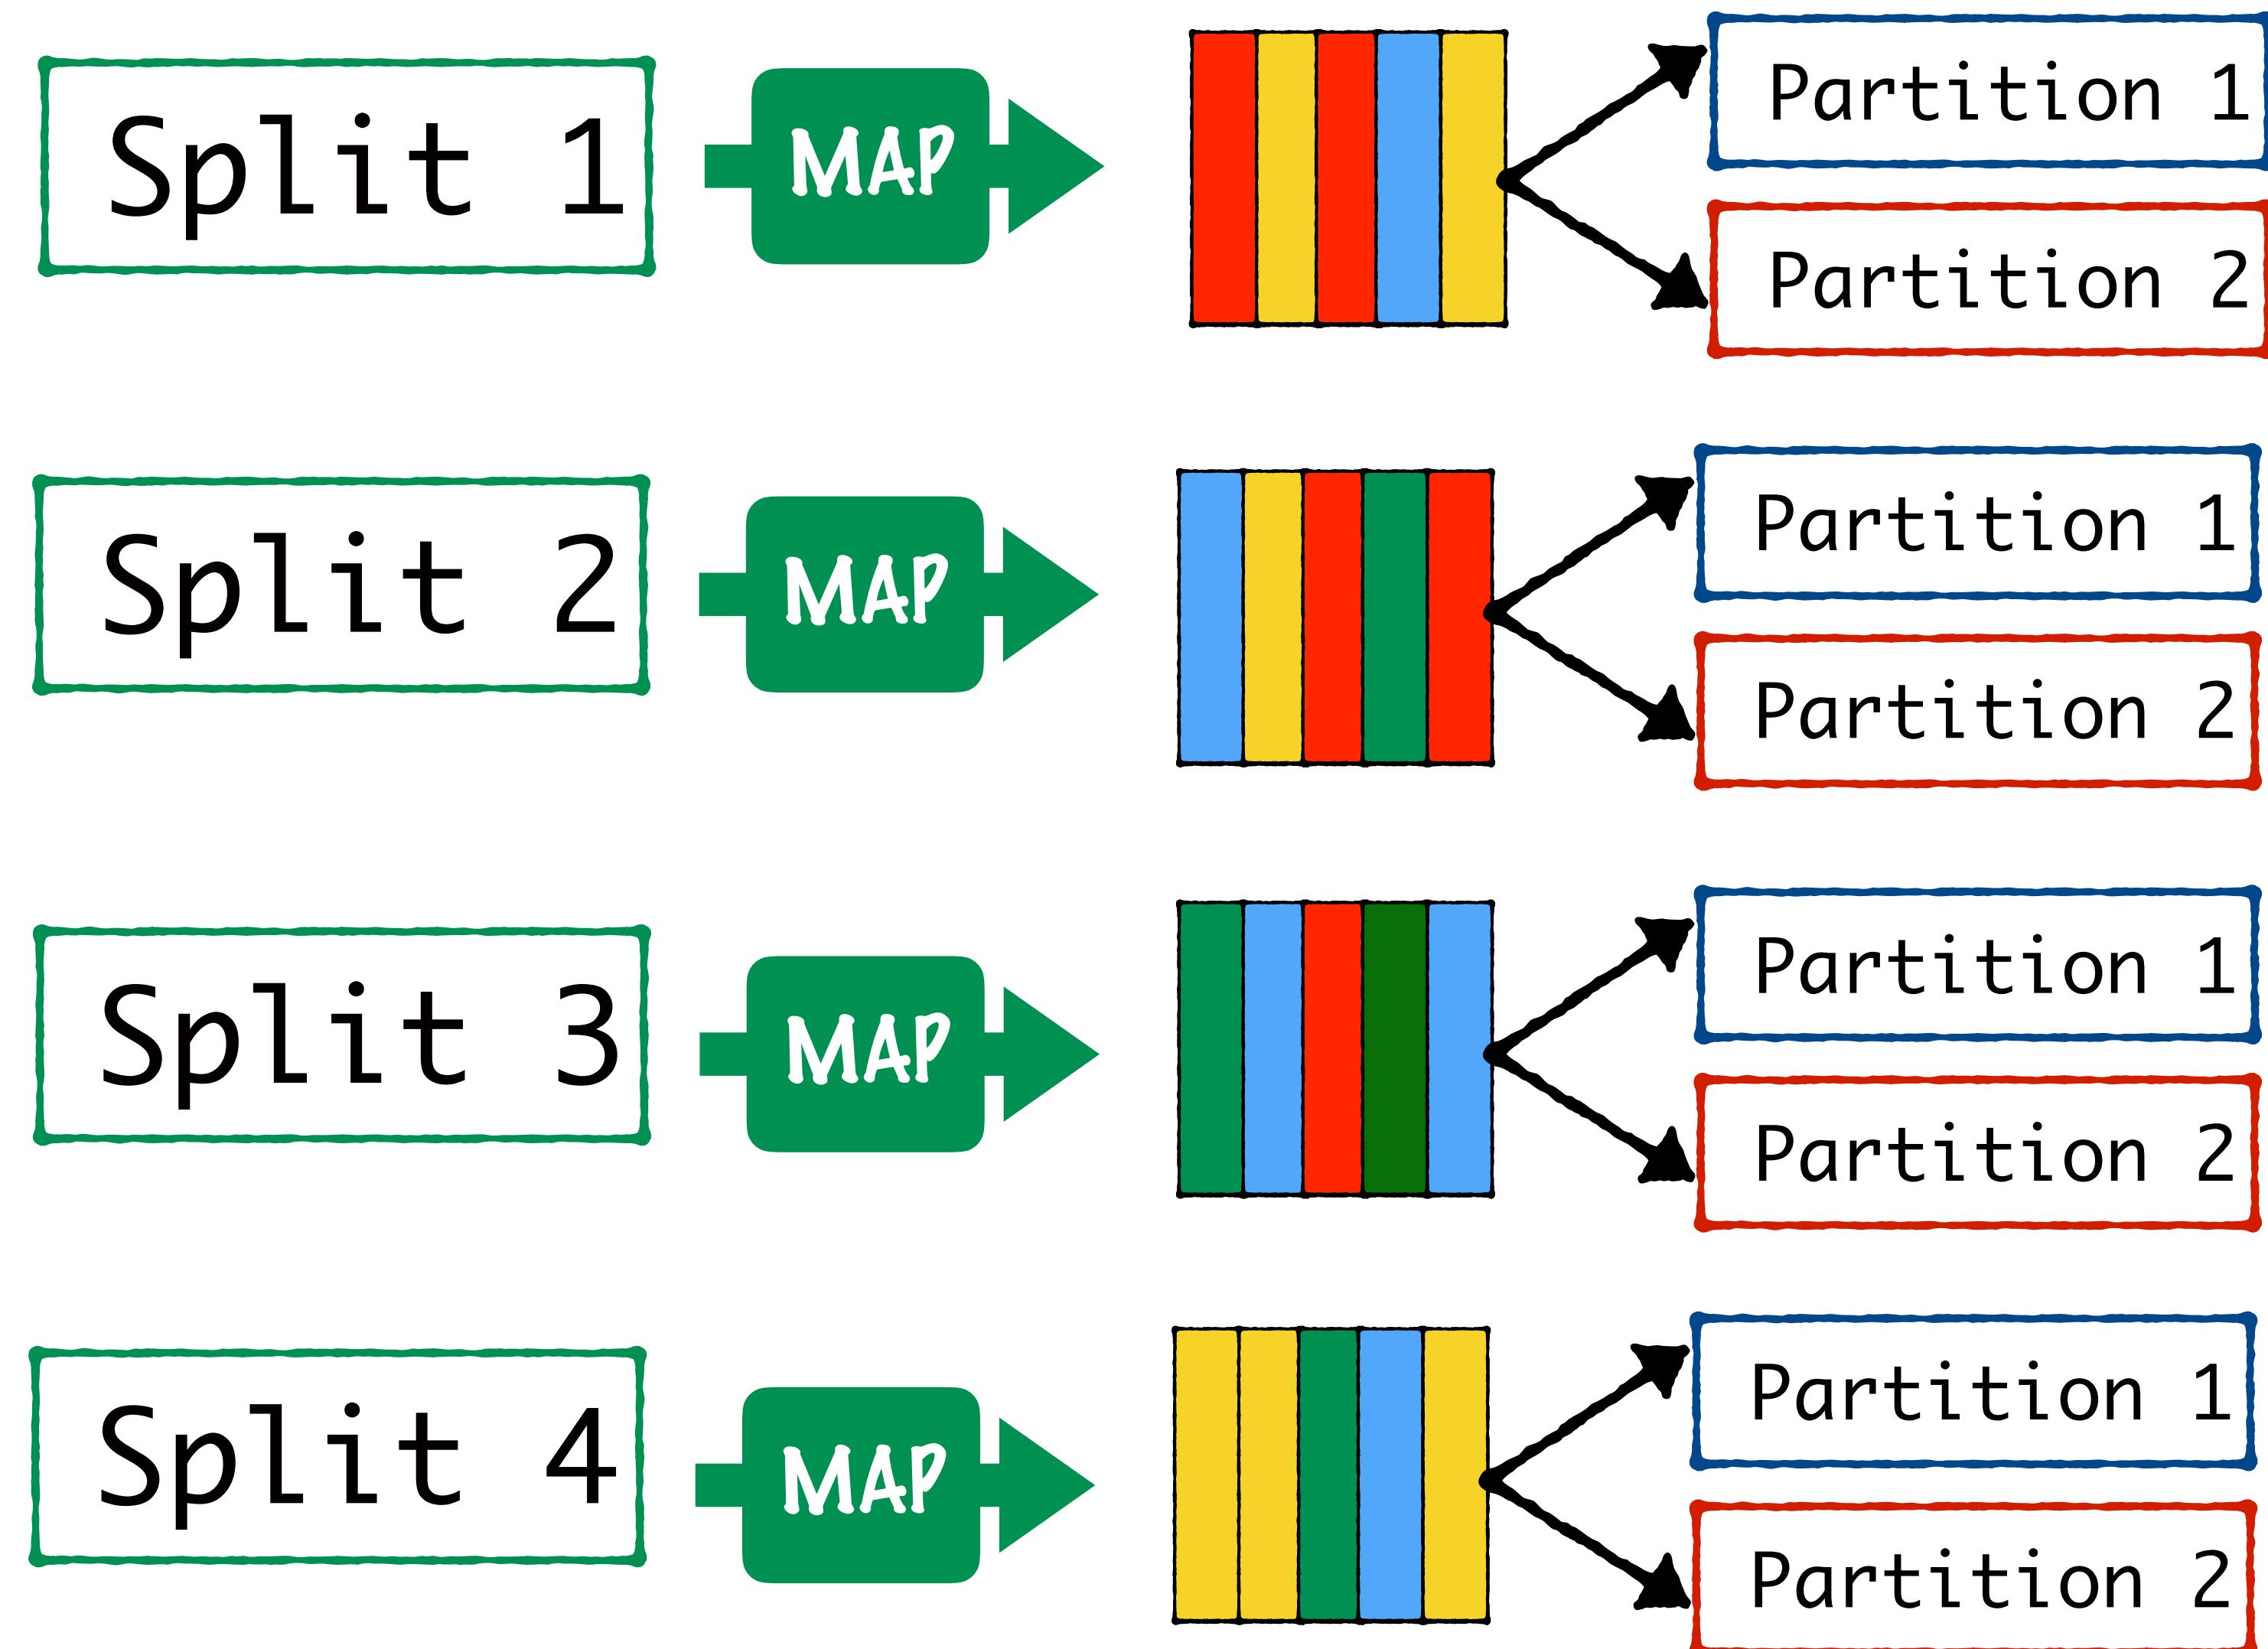
1 reducer



Sort/Merge

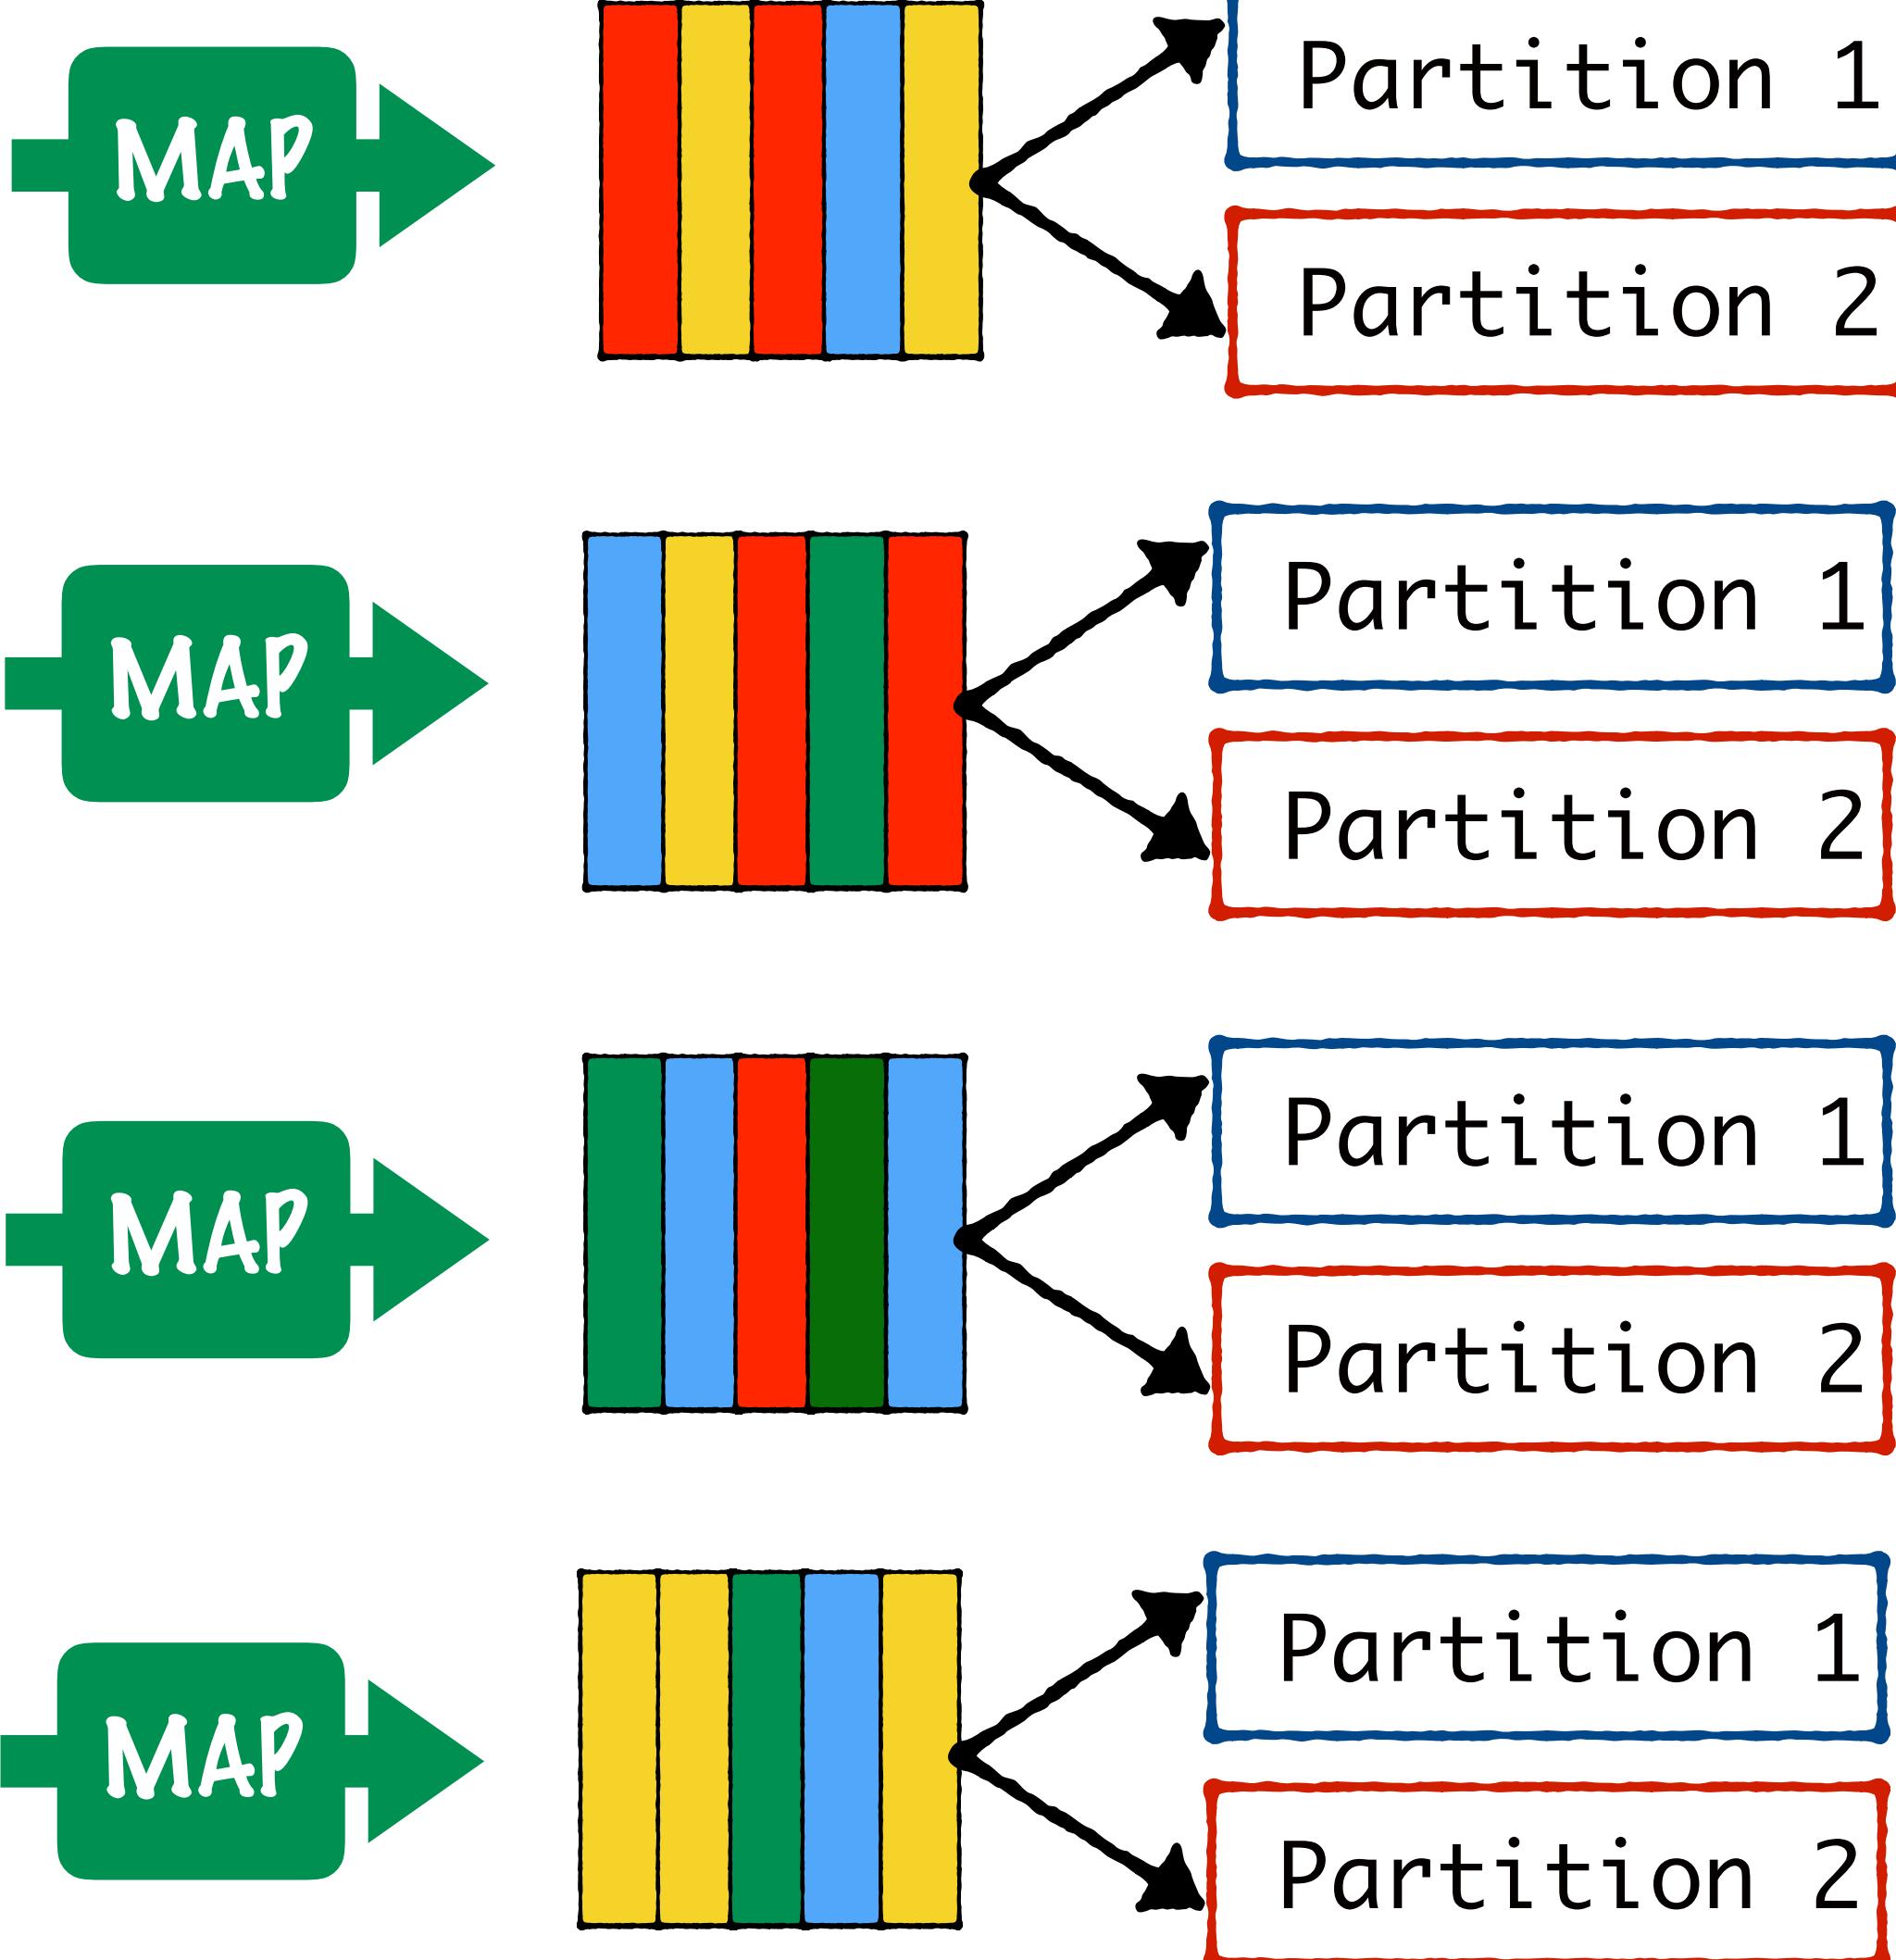


Shuffle and Sort



When there are 2 reducers, the output of each map is first partitioned into 2 partitions

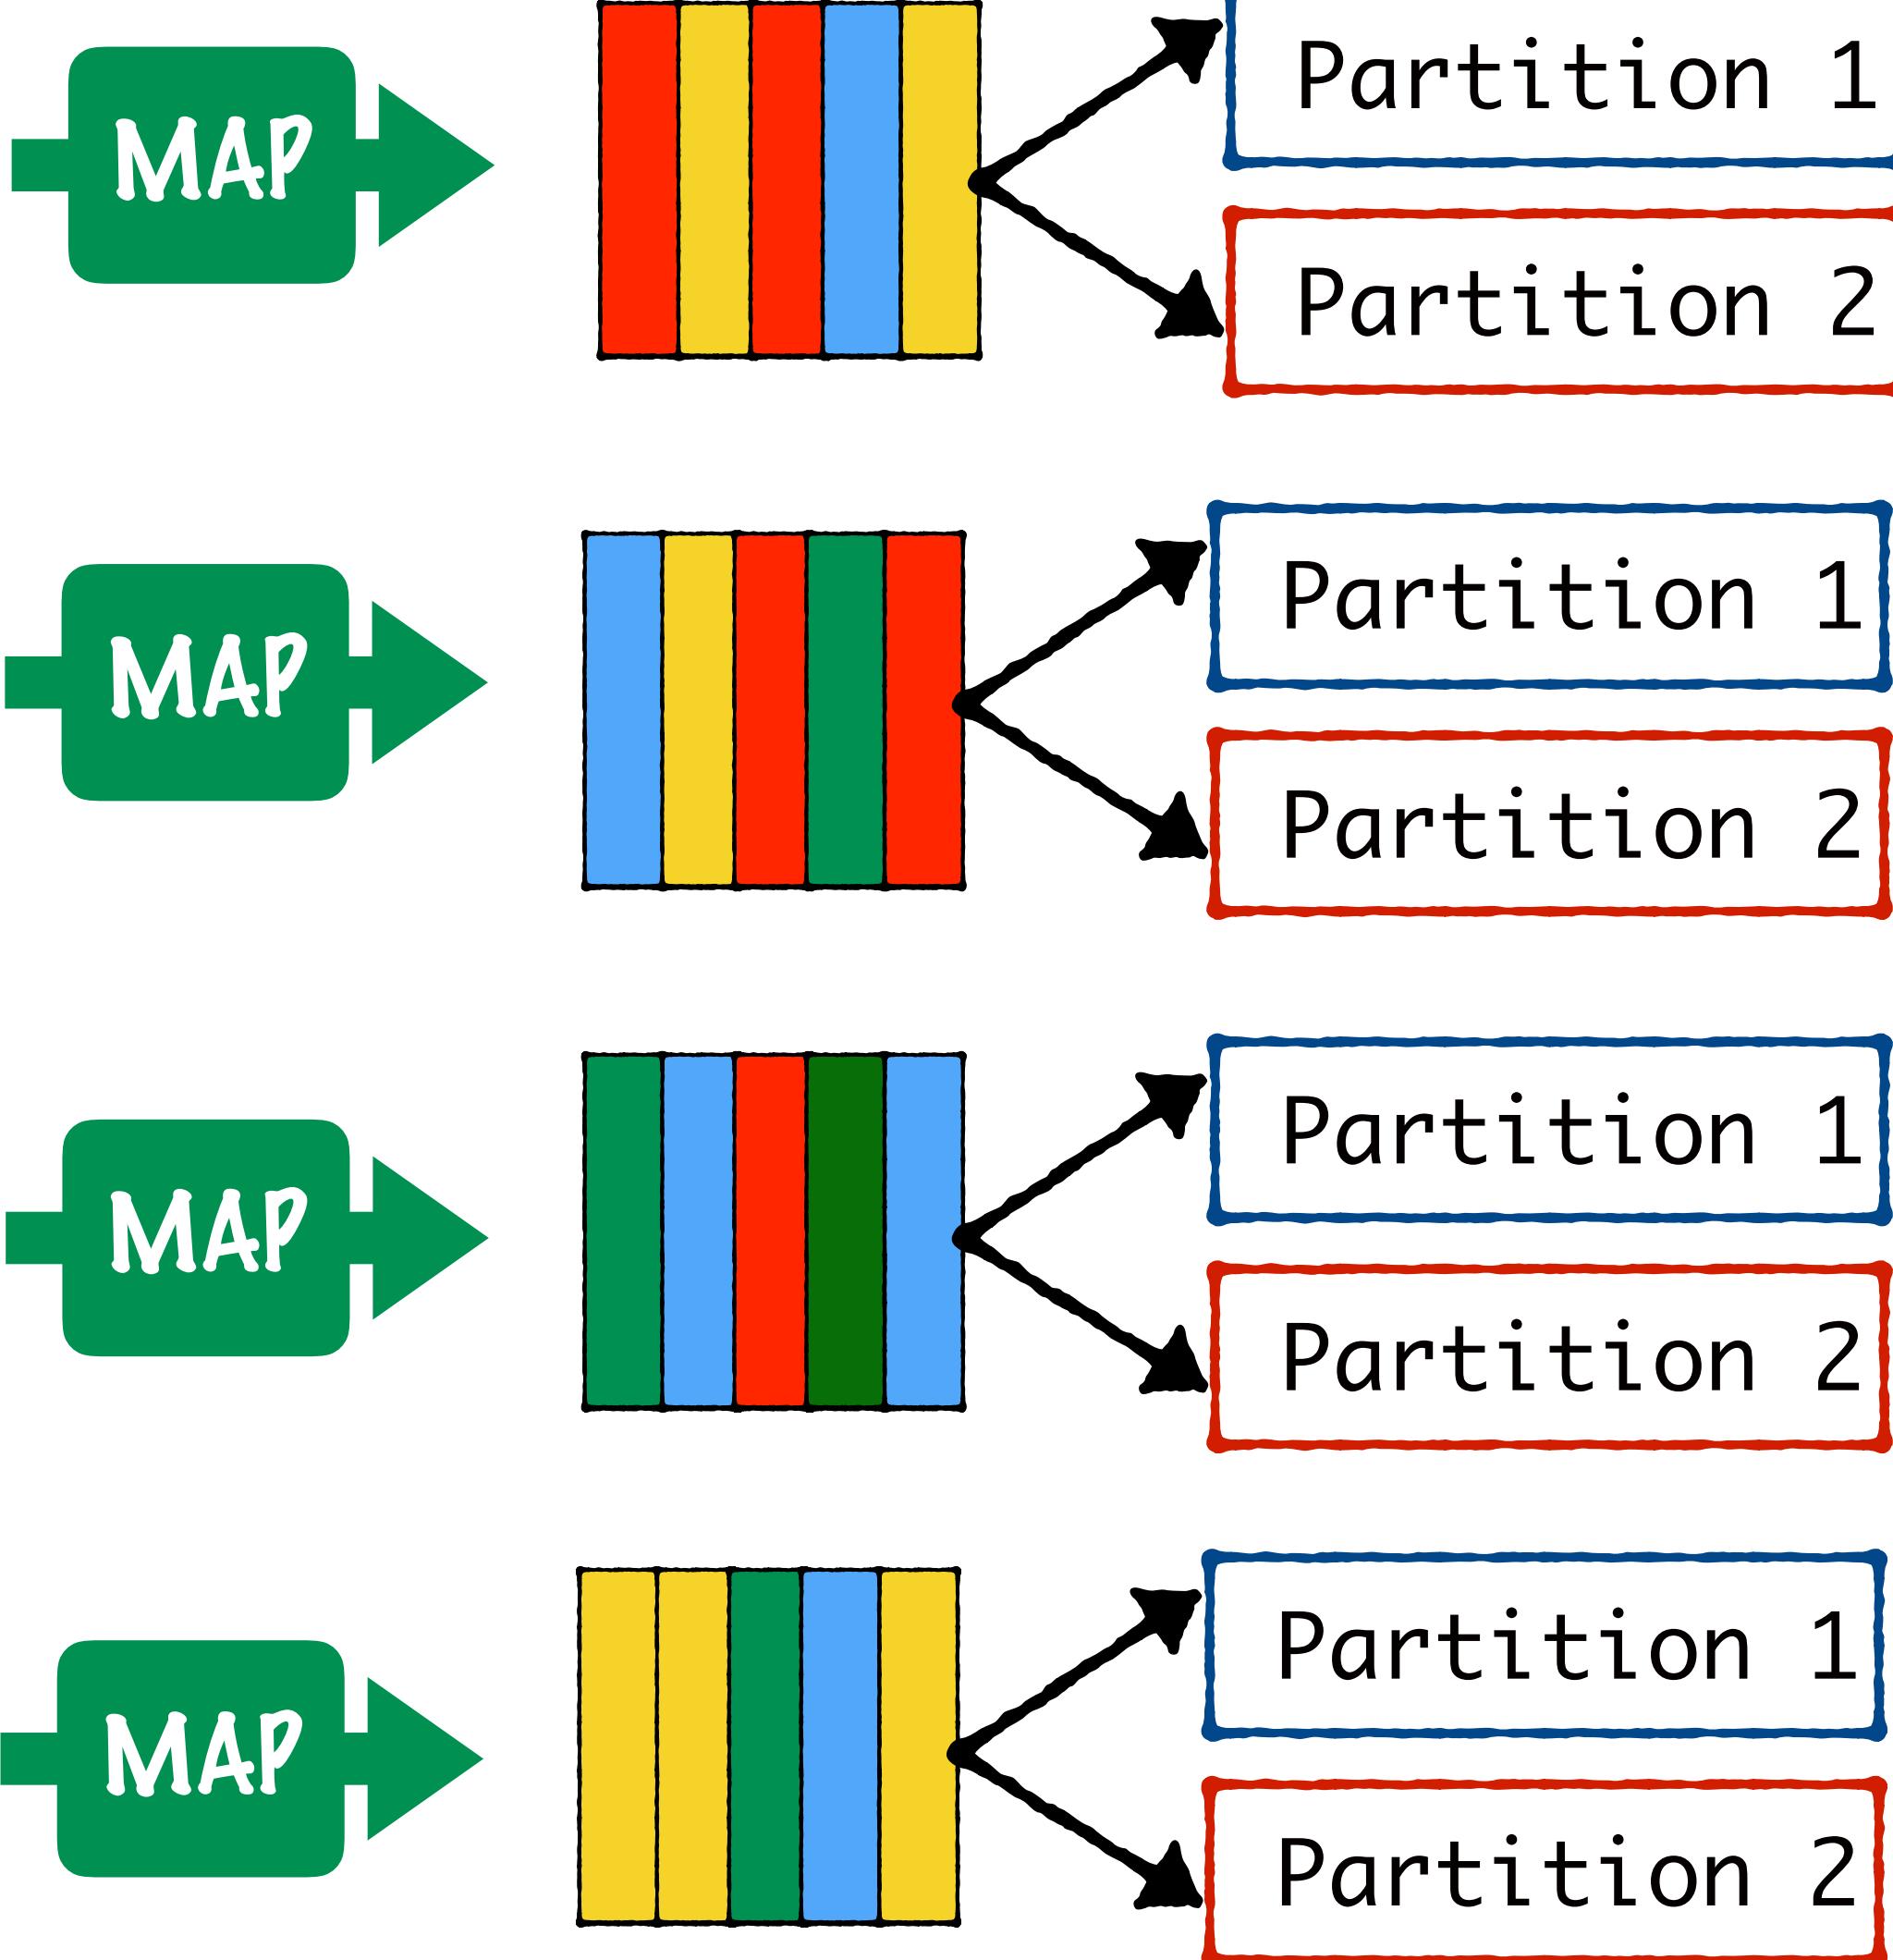
Shuffle and Sort



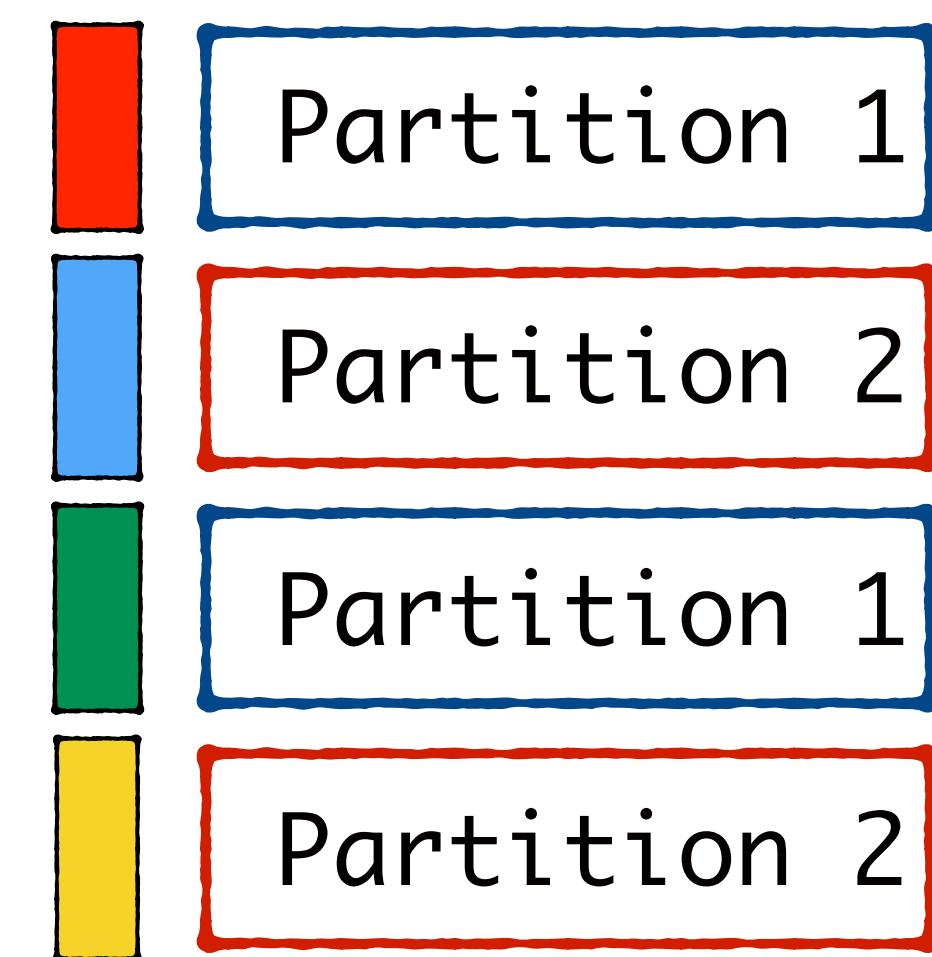
The map output is
assigned to a partition
based on the key

All values with
the same key must
be assigned to the
same partition

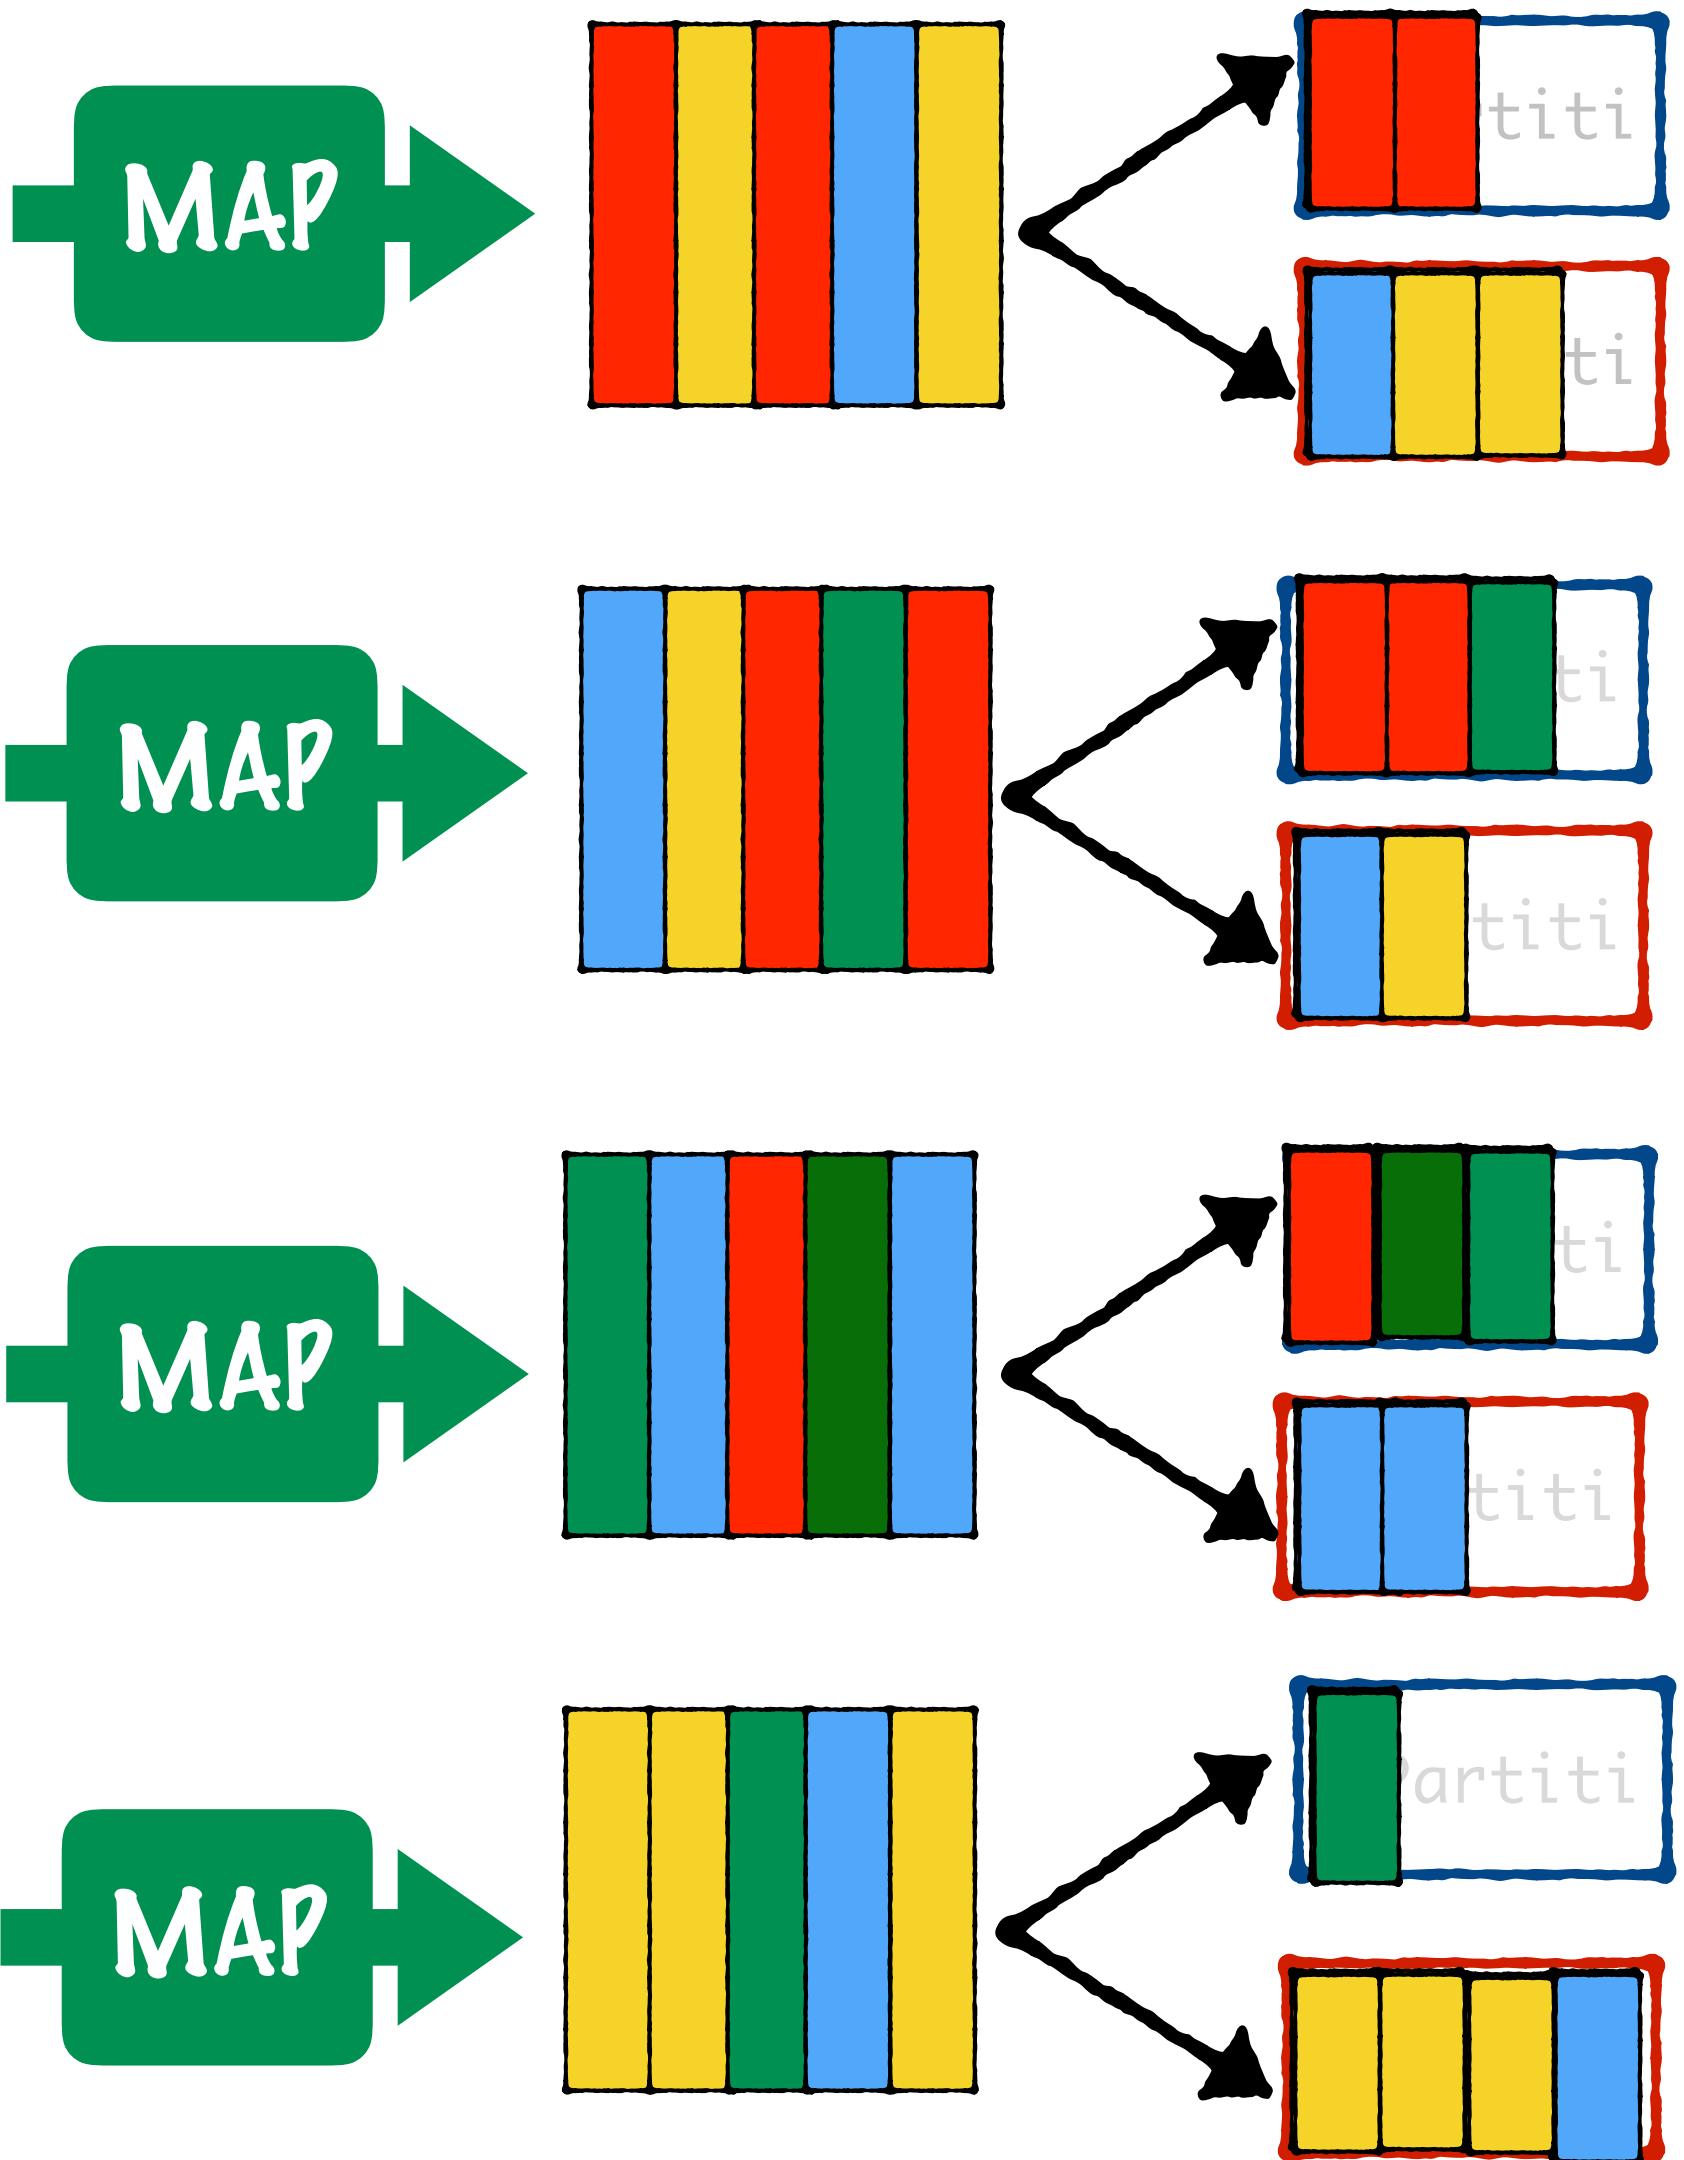
Shuffle and Sort



All values with the
same key must be
assigned to the same
partition



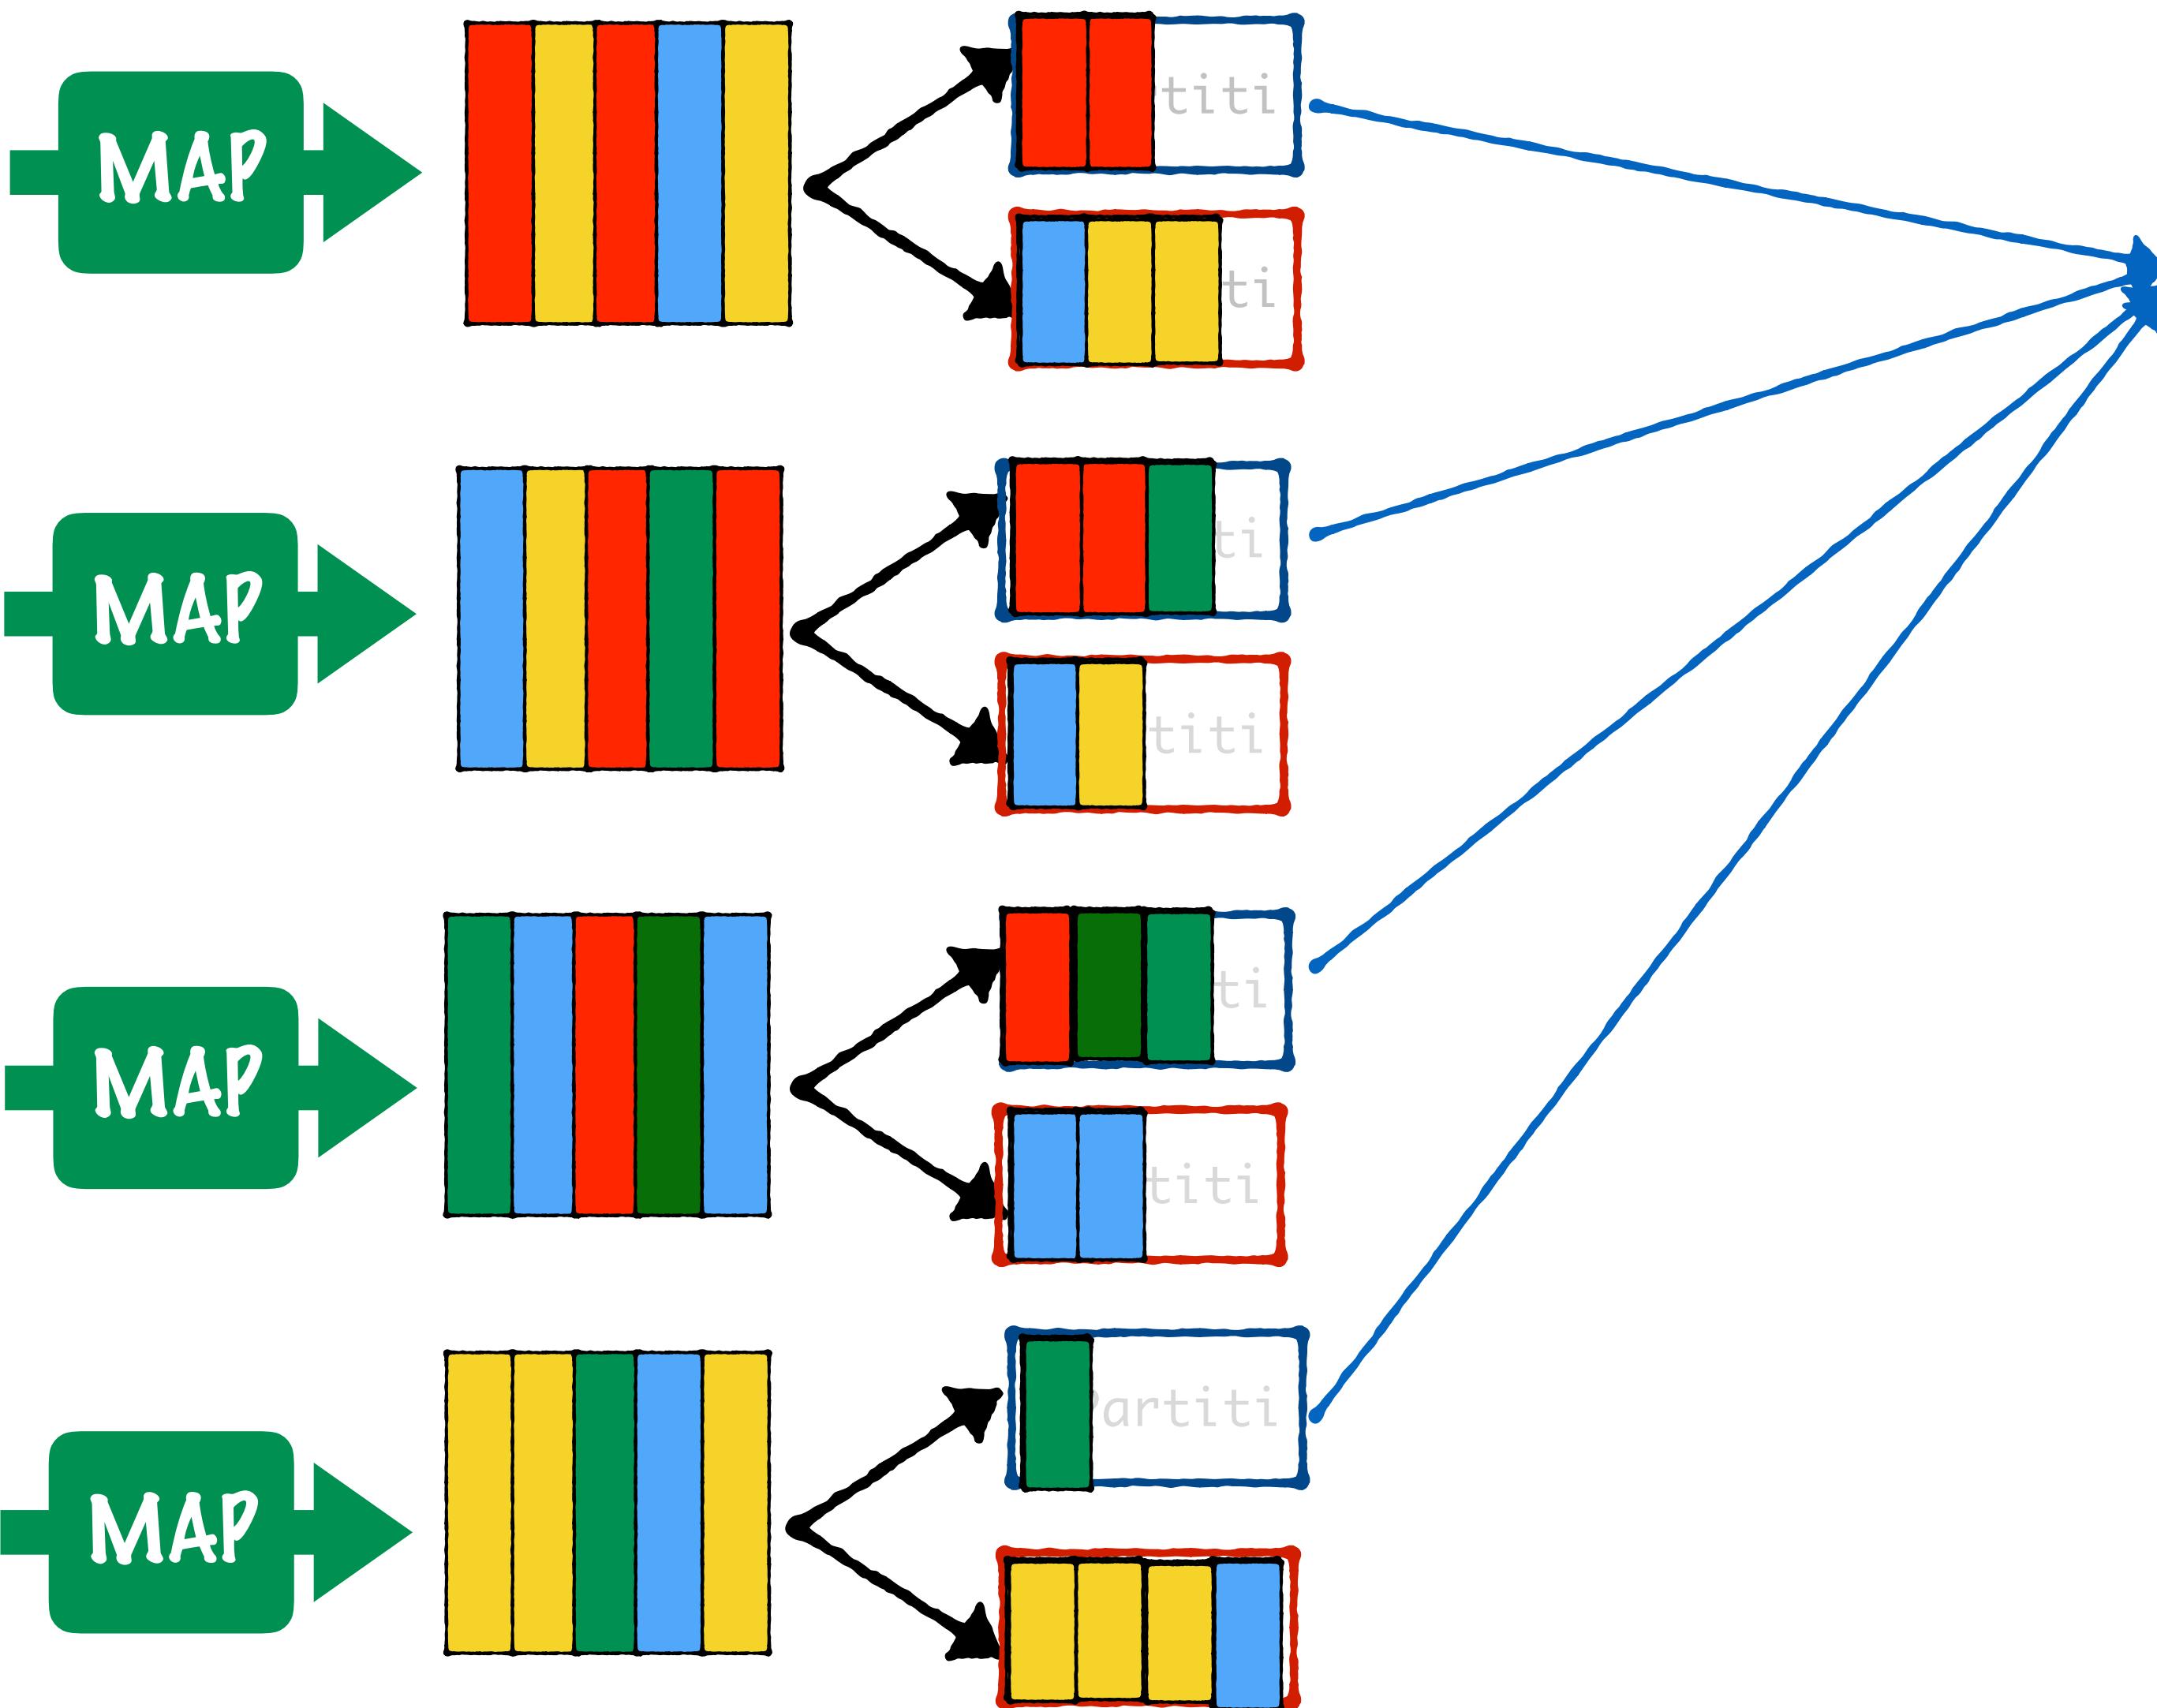
Shuffle and Sort



All values with the
same key must be
assigned to the same
partition

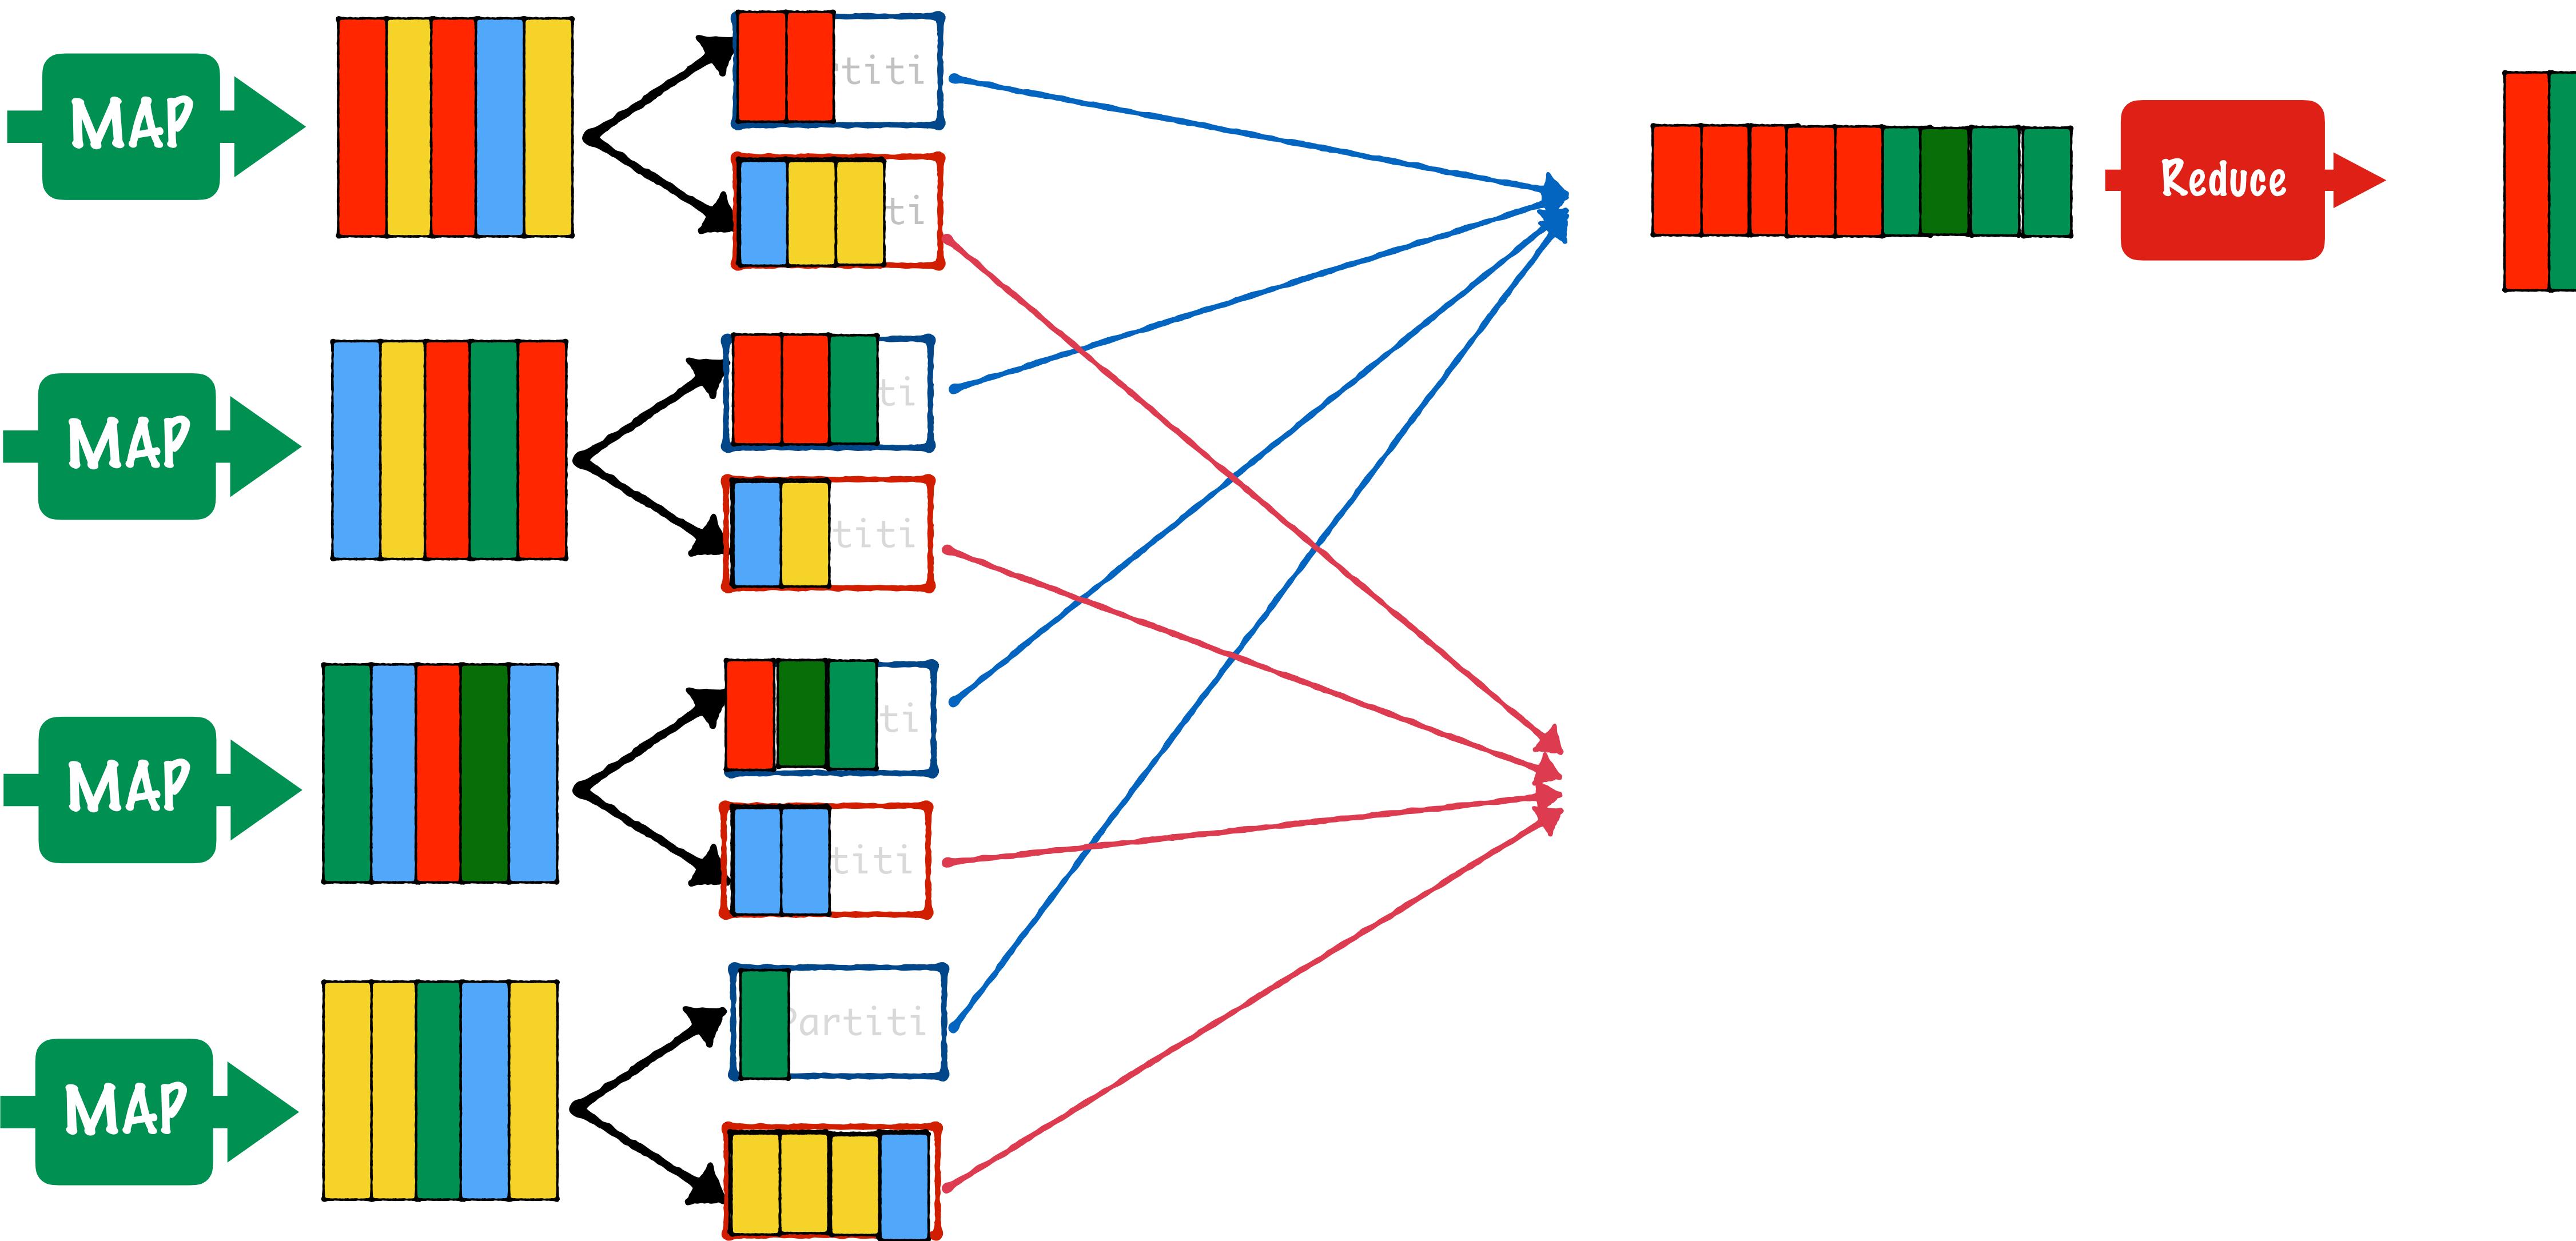


Shuffle and Sort Partitioning



The data in each partition is sent to 1 reducer

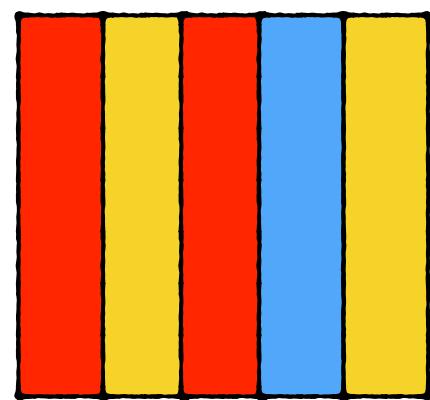
Shuffle and Sort Partitioning



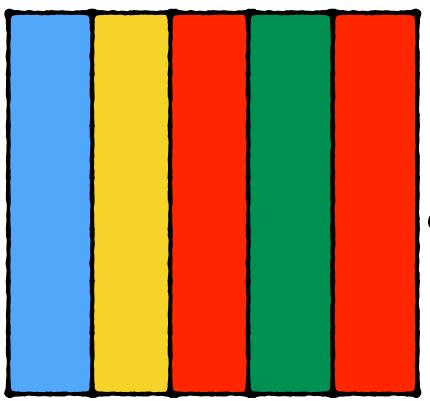
Shuffle and Sort

Partitioning

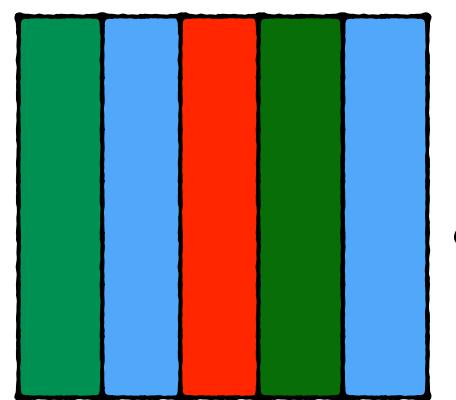
MAP



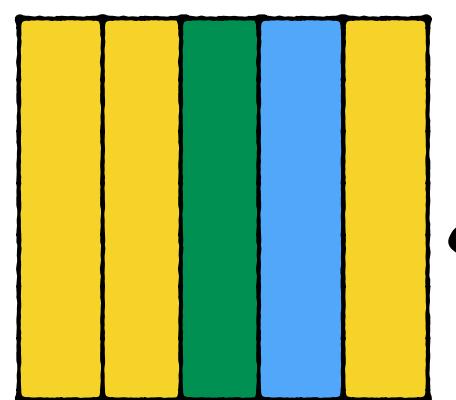
MAP



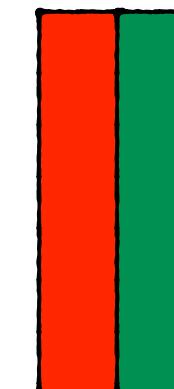
MAP



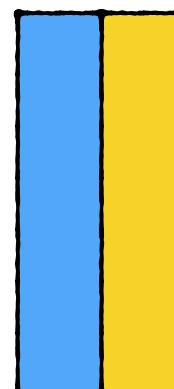
MAP



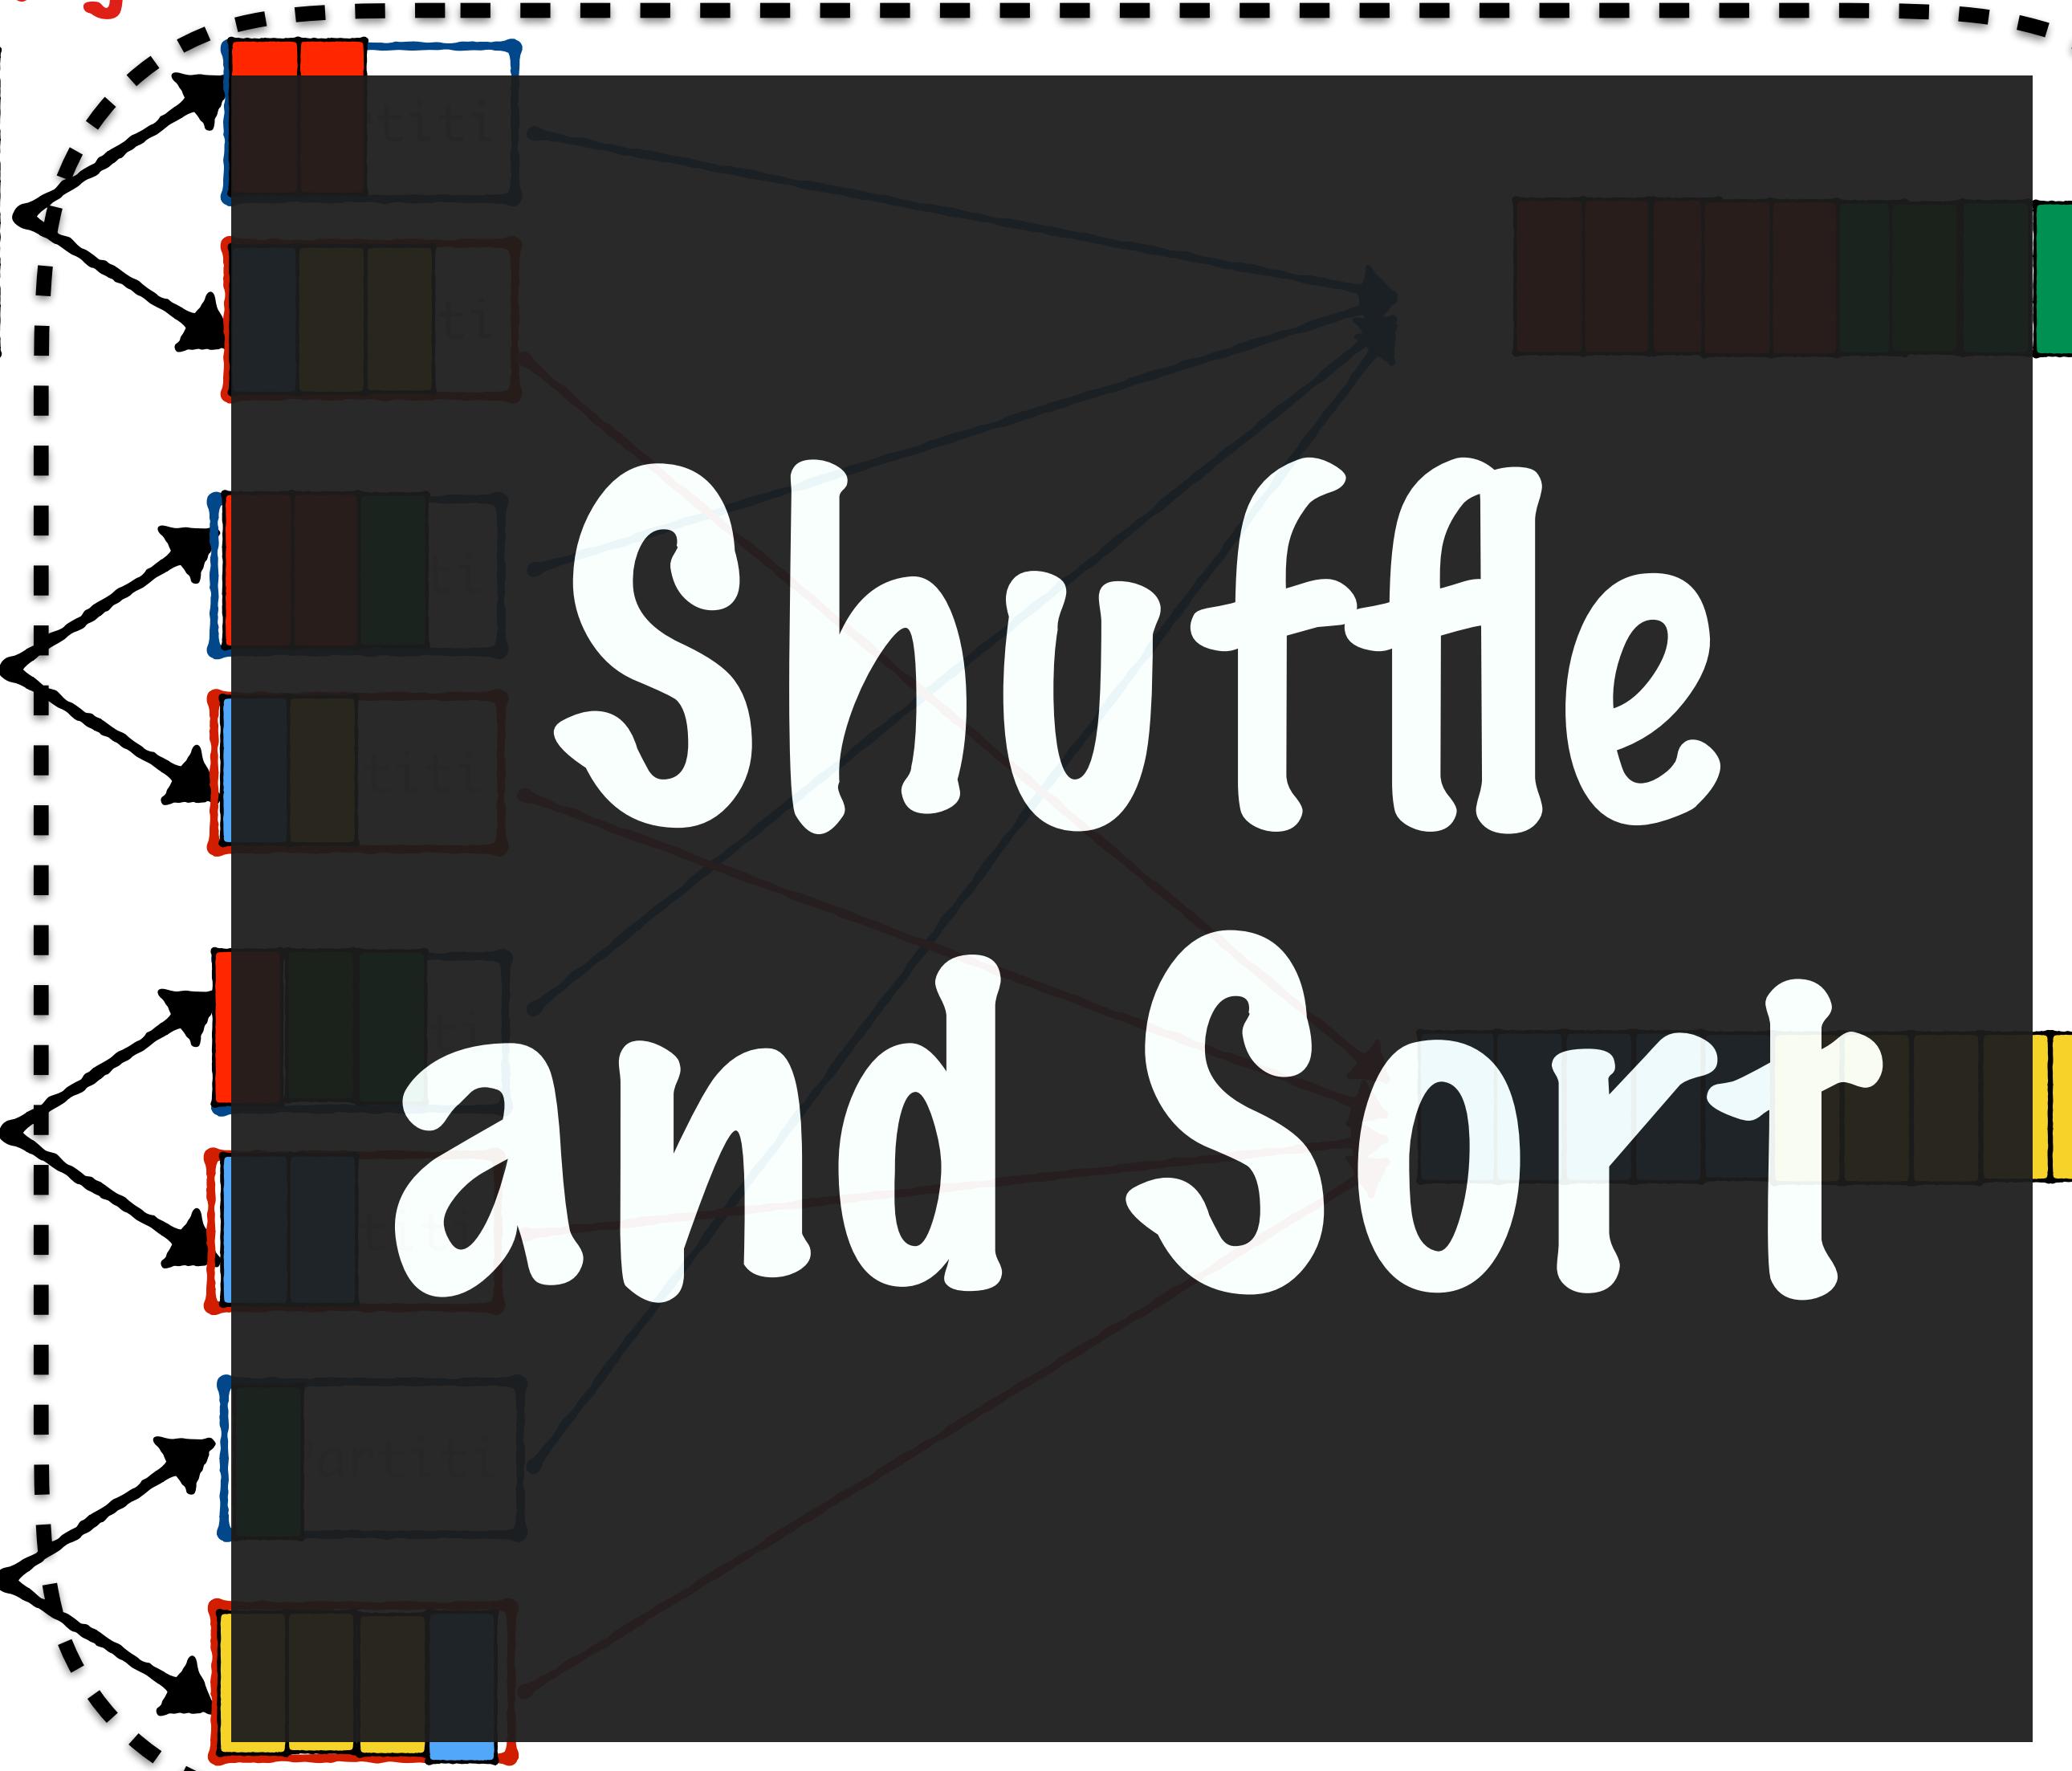
Reduce



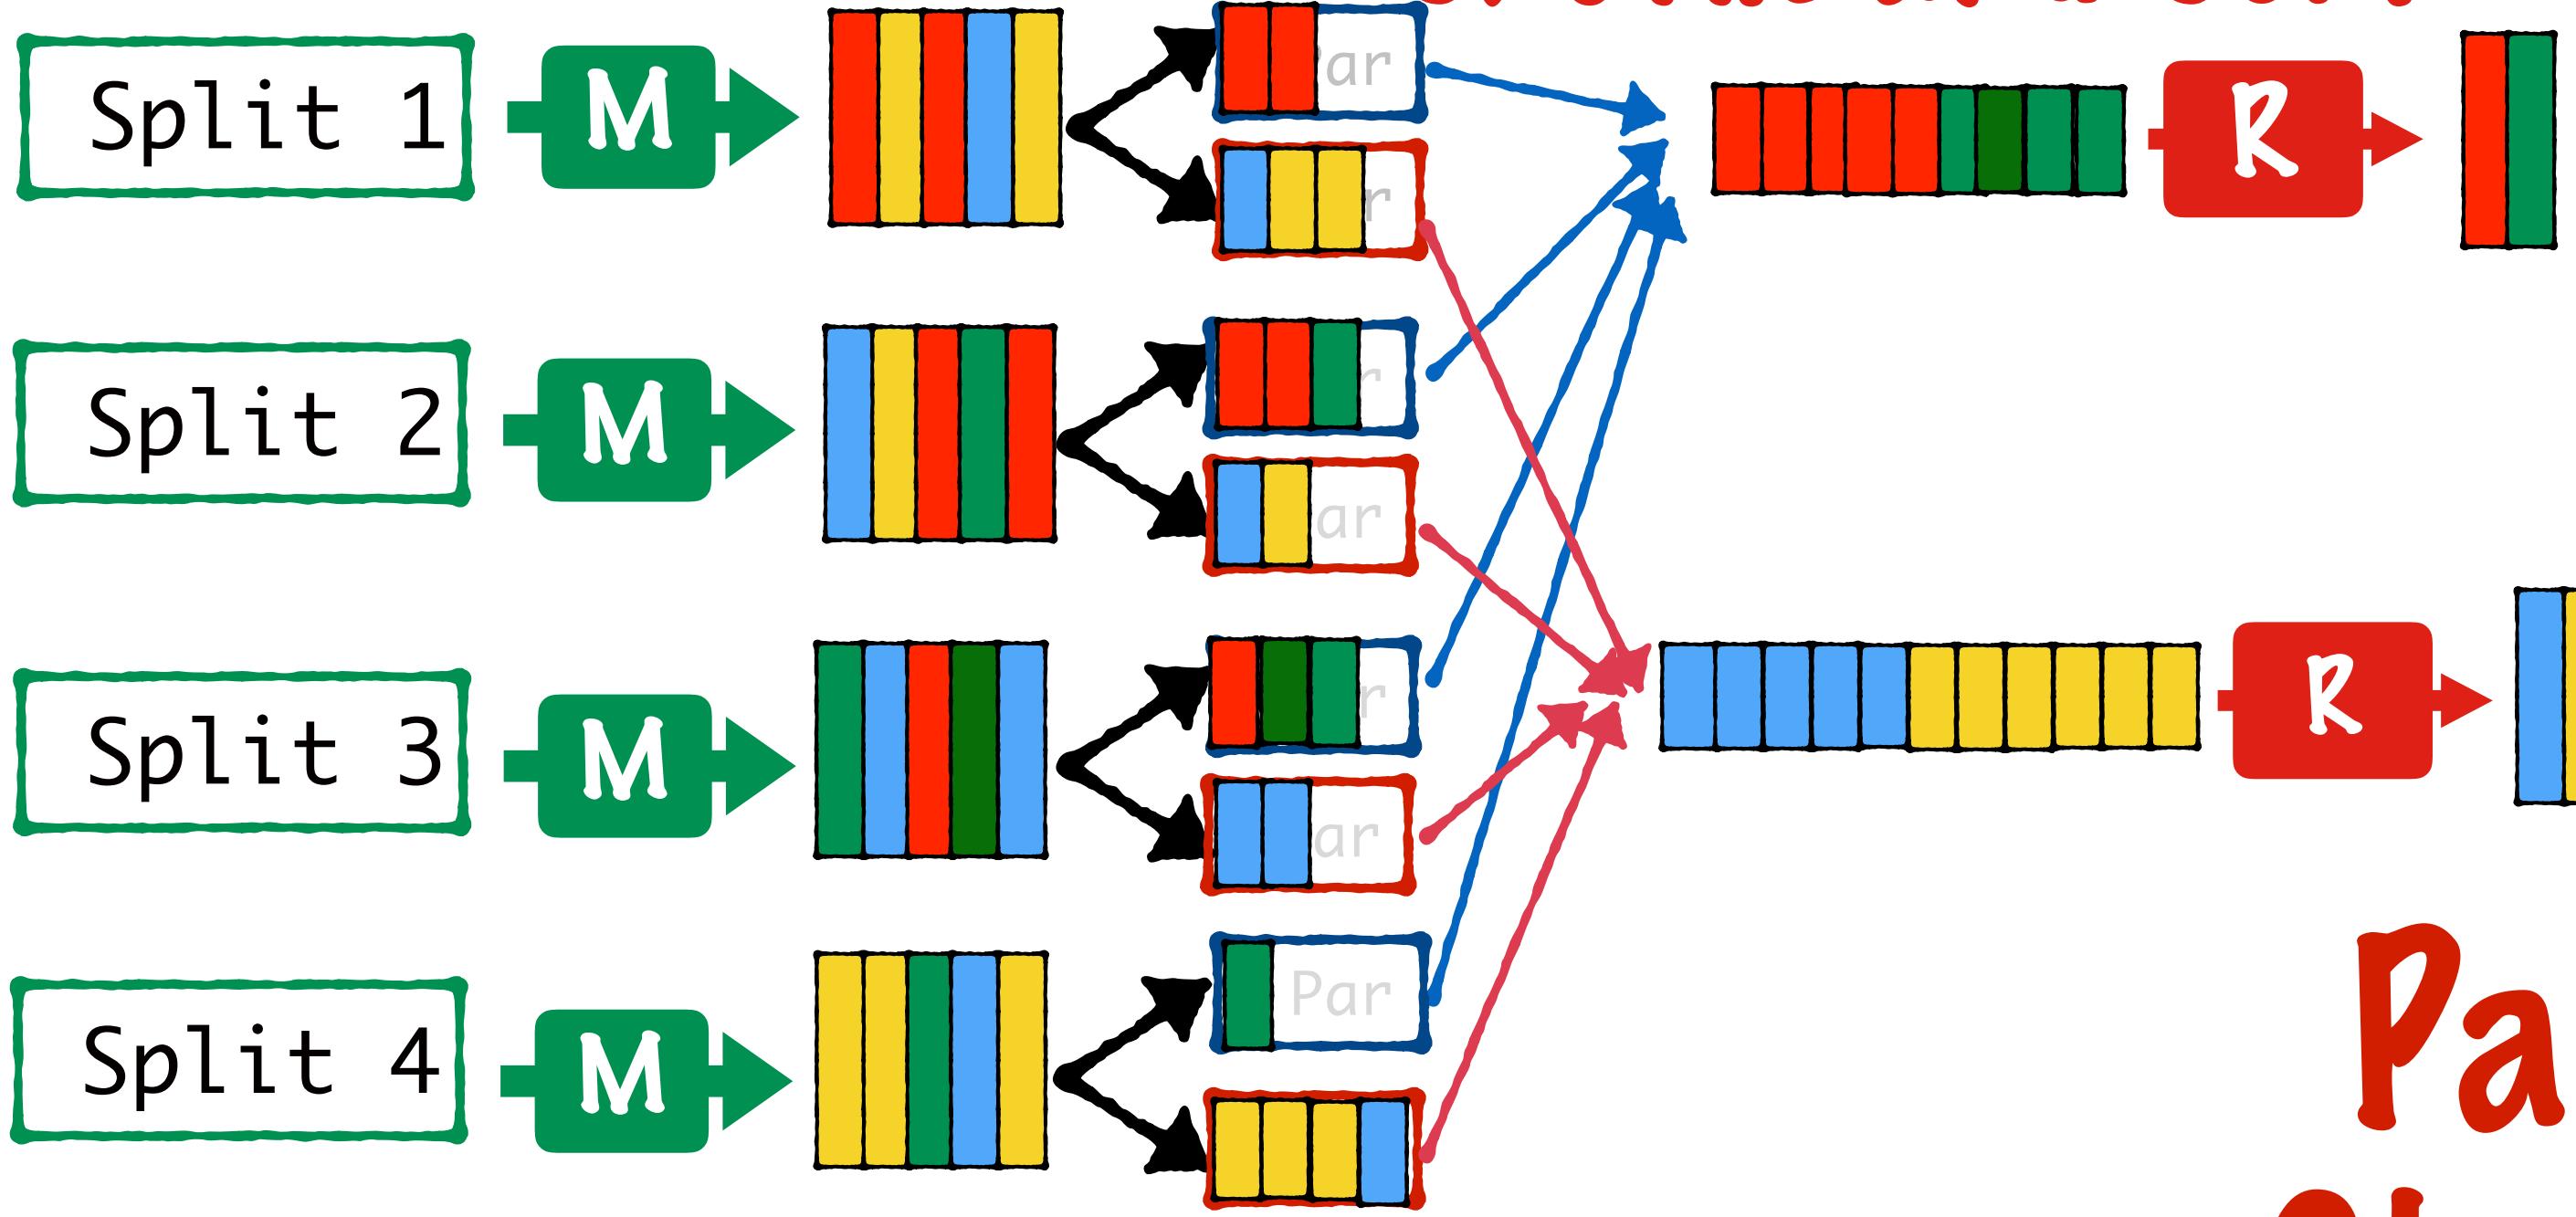
Reduce



Shuffle
and Sort

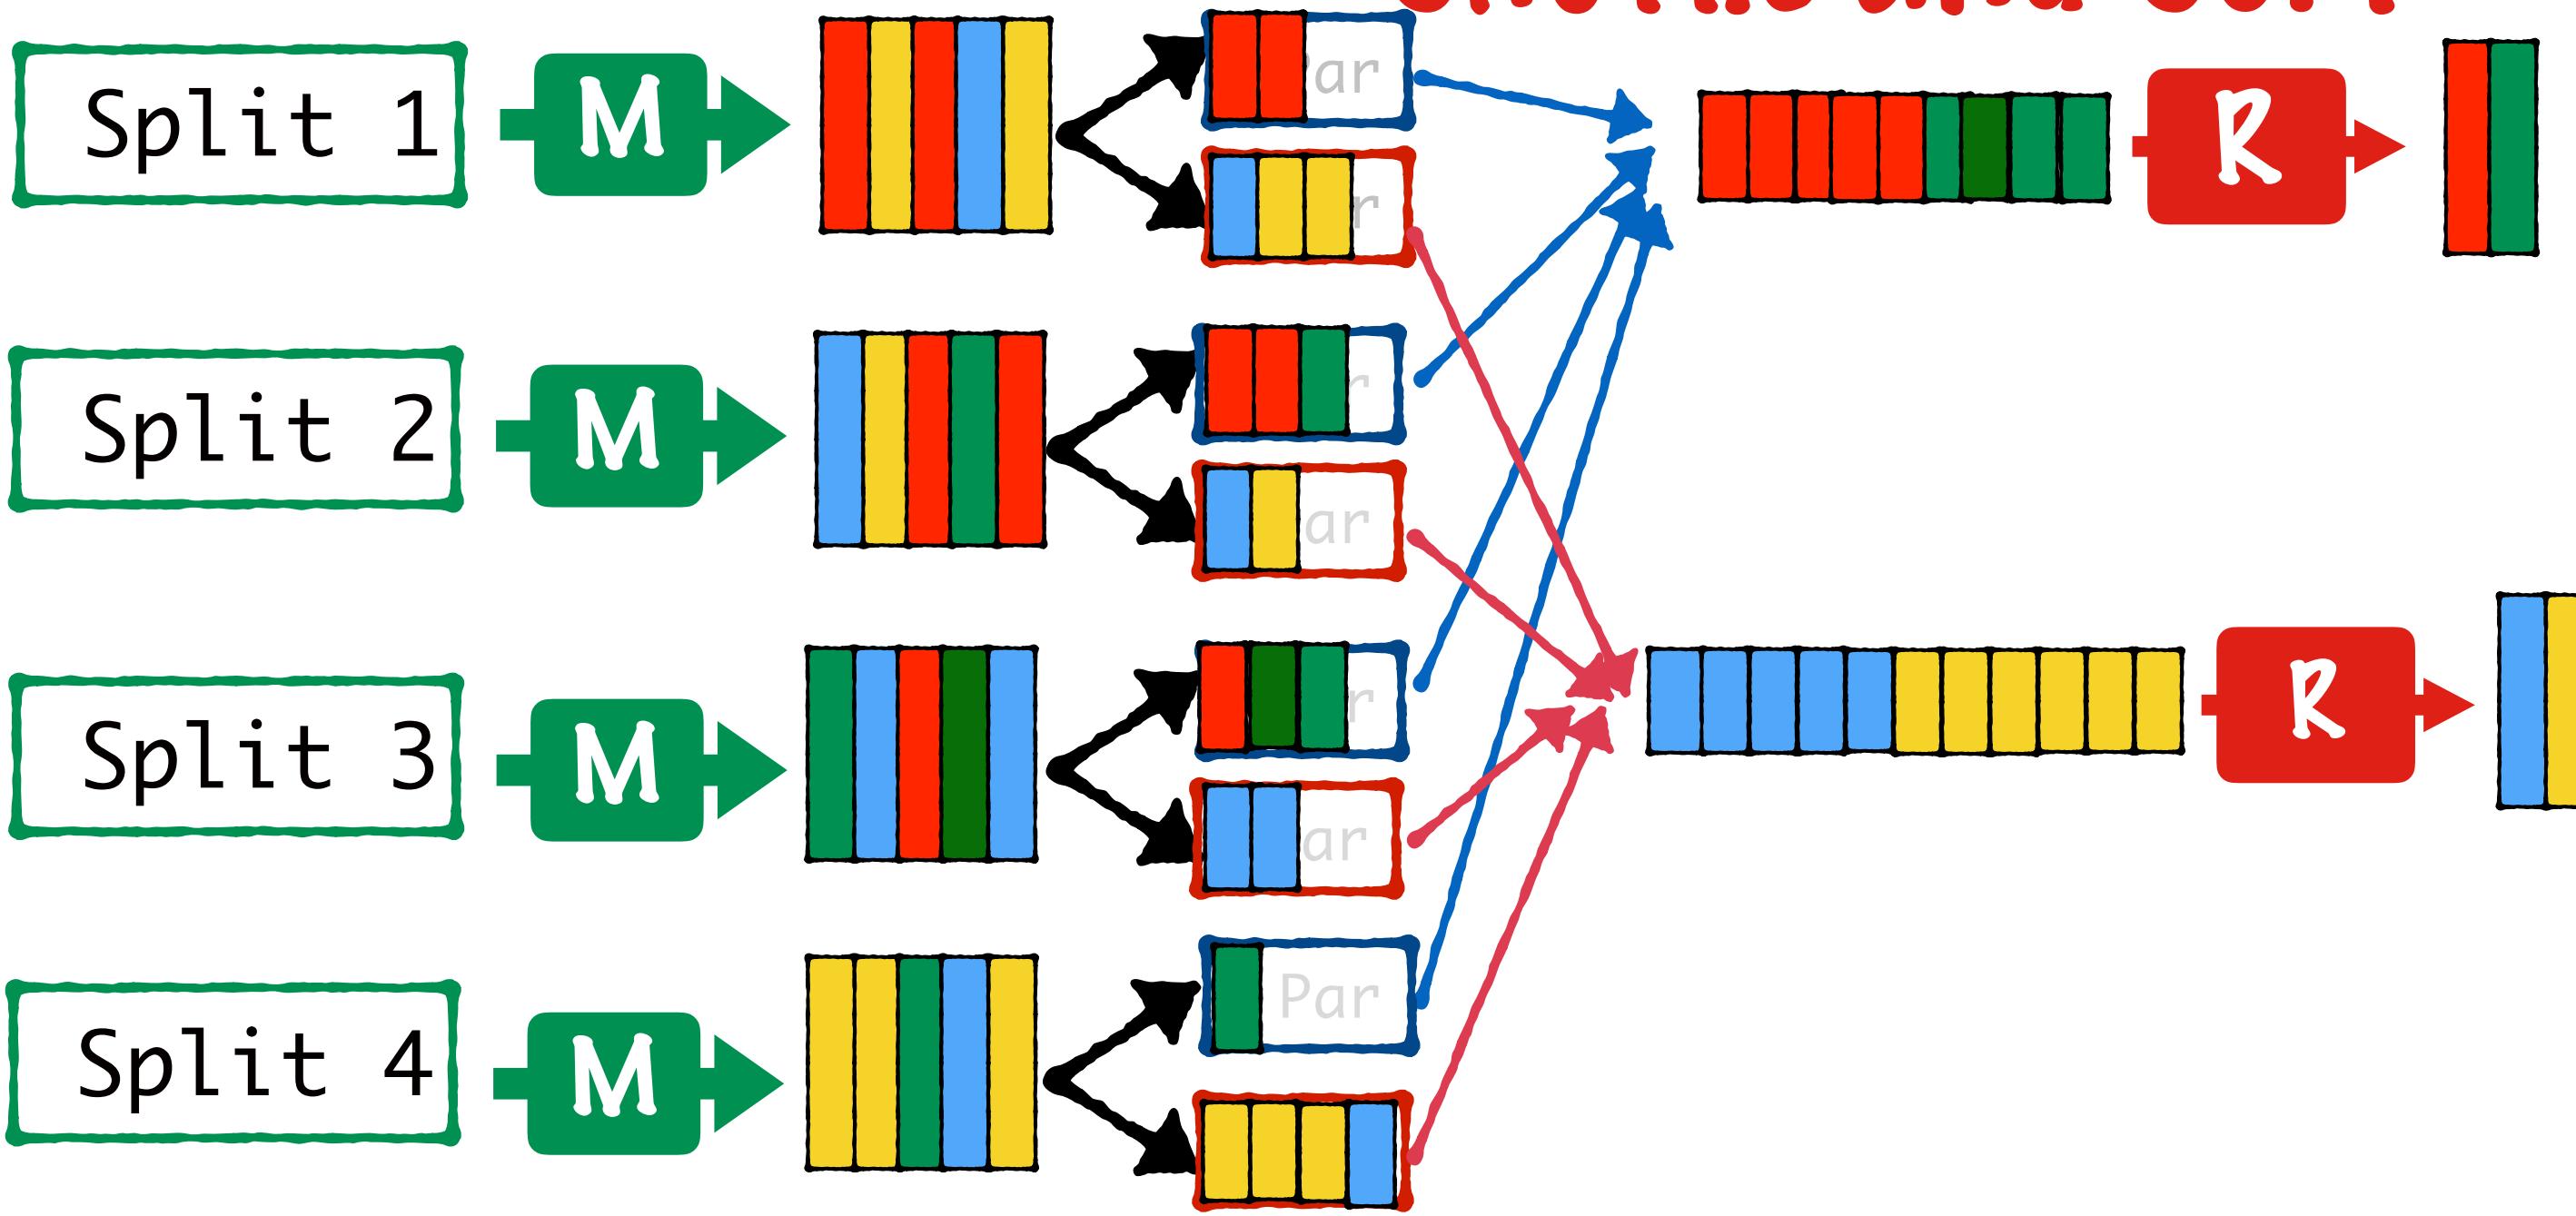


Partitioning Shuffle and Sort



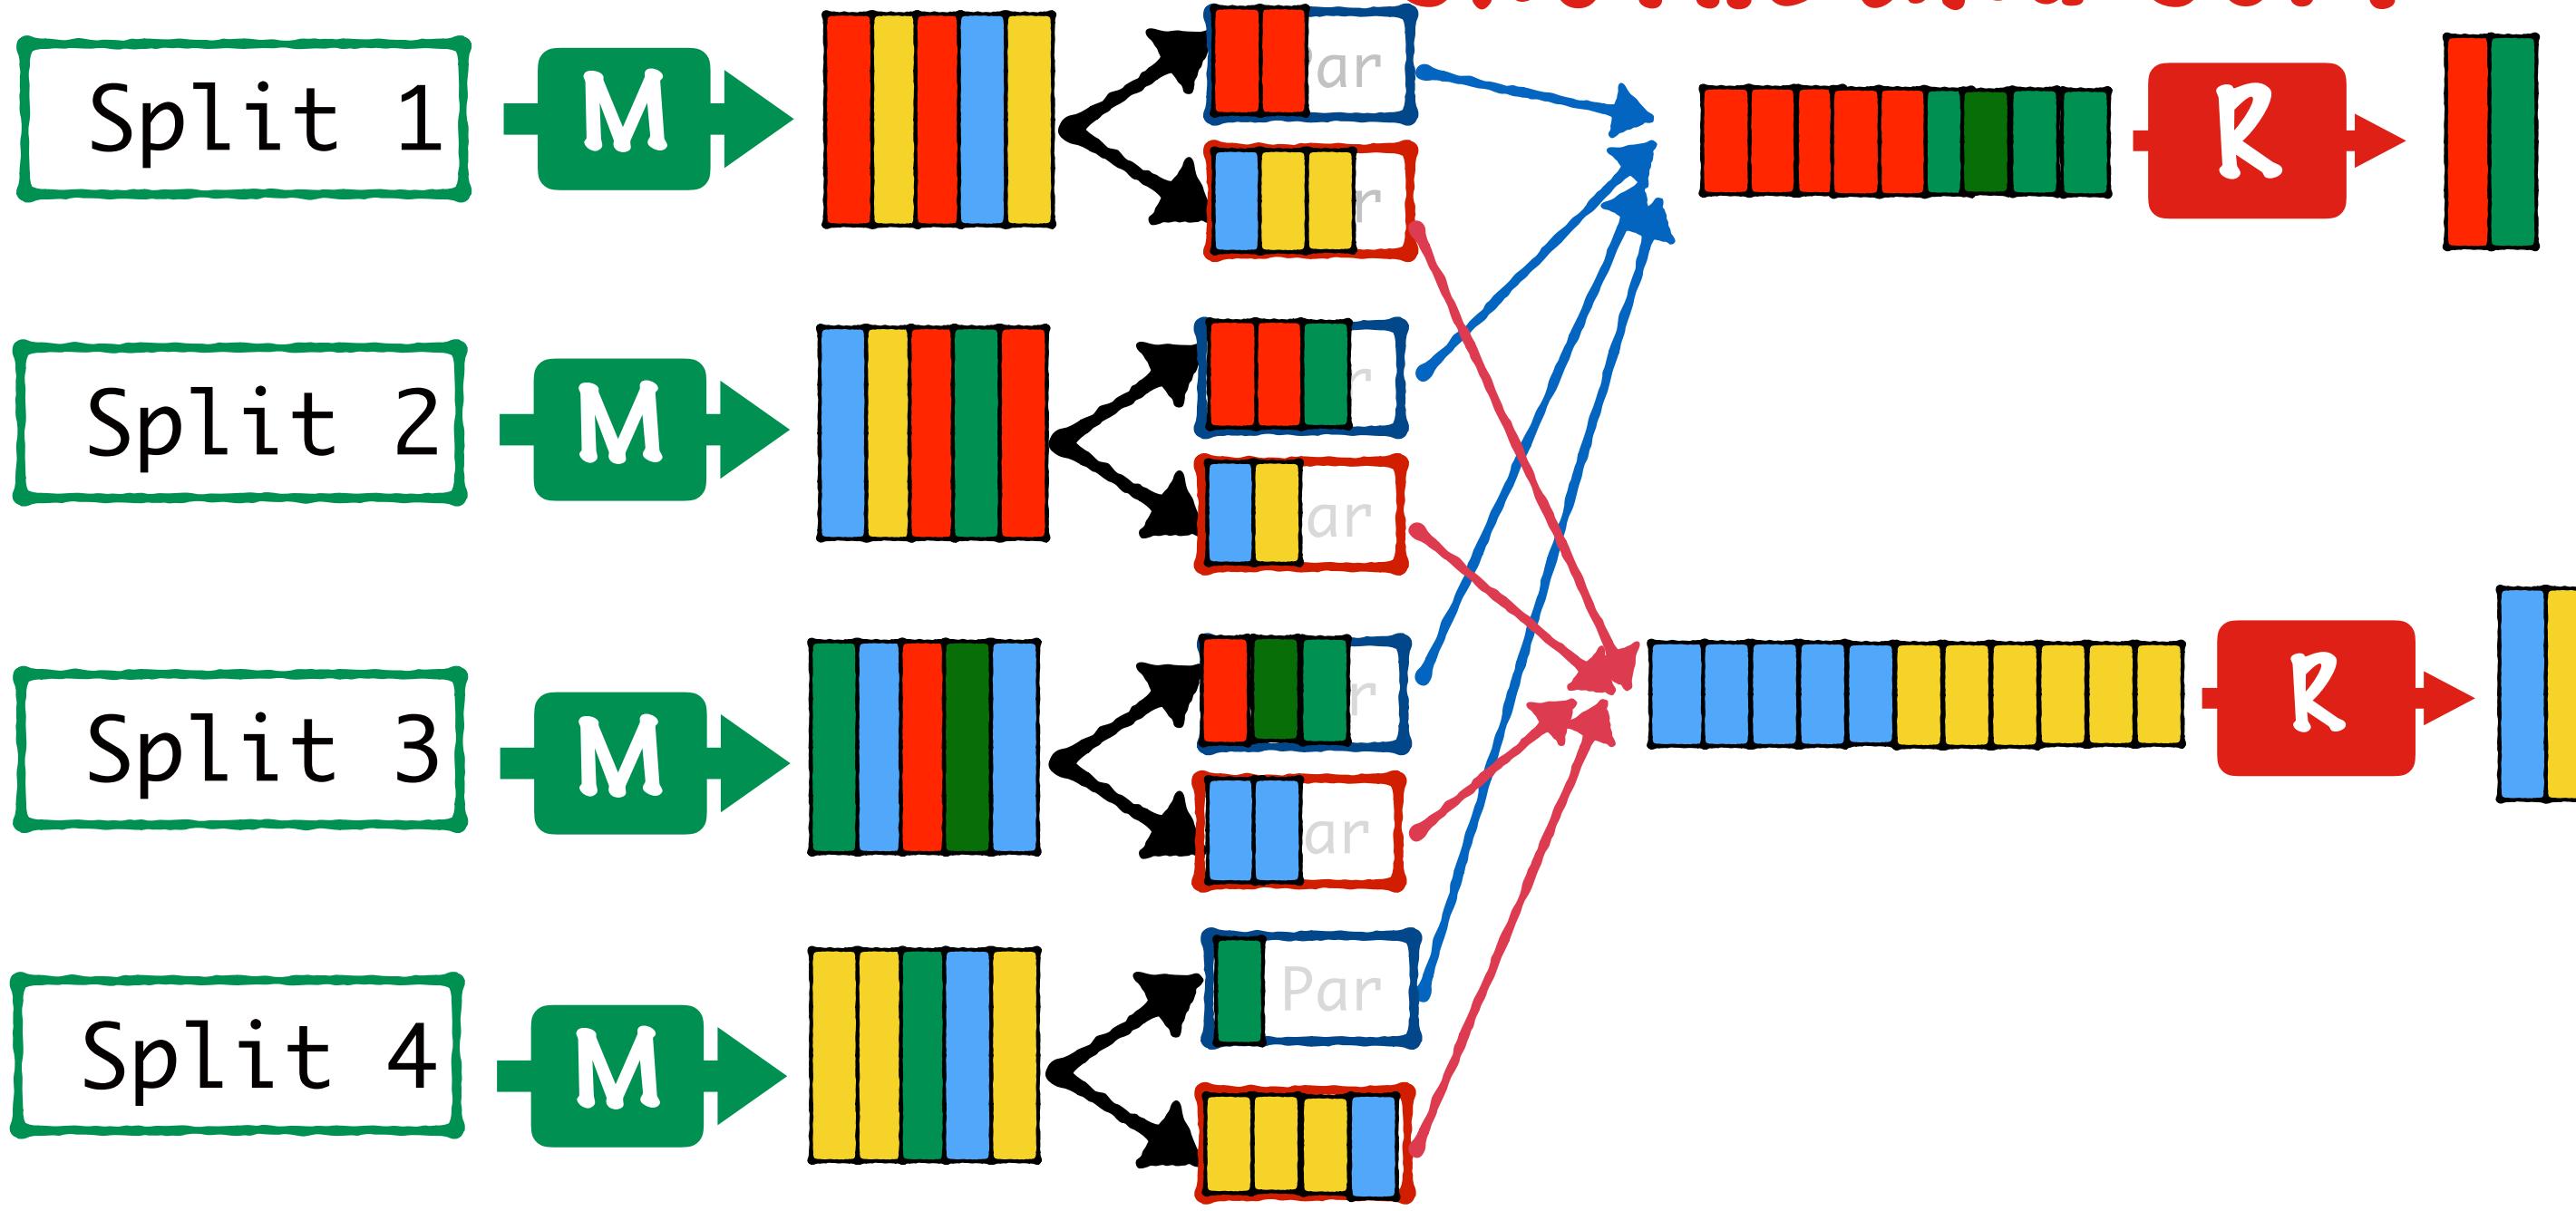
Partitioning and
Shuffle/Sort are 2
extra operations that
happen for MR jobs
with multiple reducers

Partitioning Shuffle and Sort



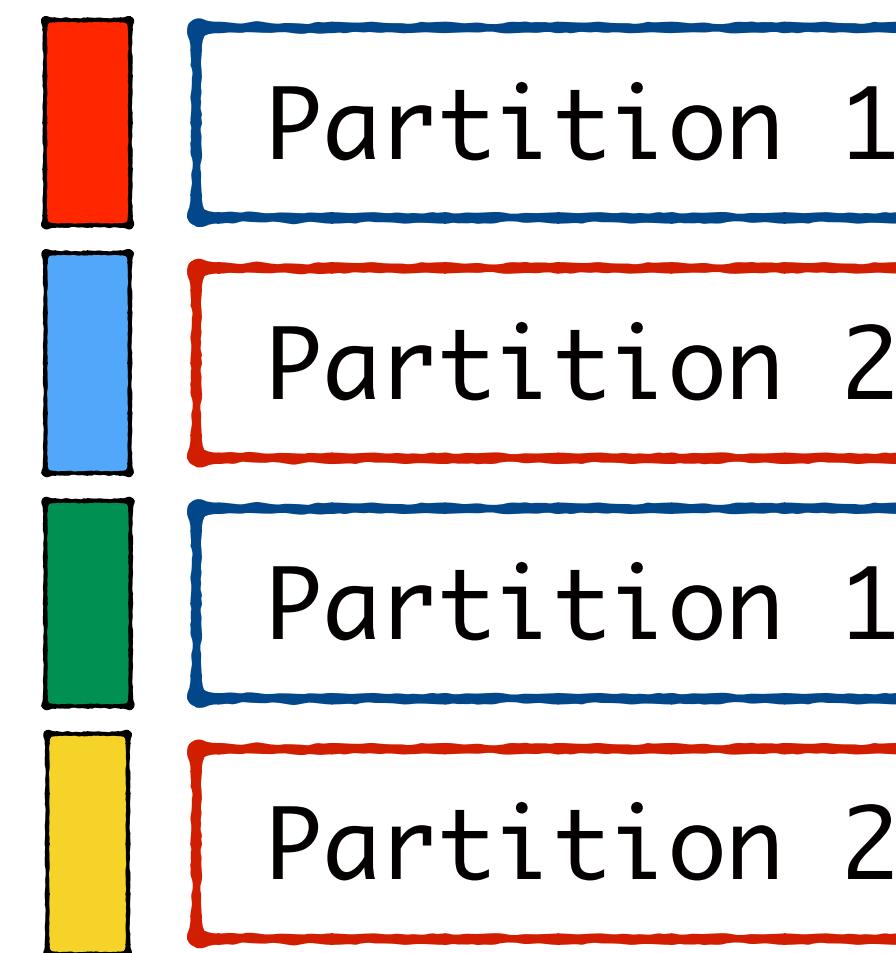
Number of partitions =
Number of reducers

Partitioning Shuffle and Sort



Hadoop takes care of
mapping keys to partitions

Hadoop takes care of mapping keys to partitions



The partition number of a key depends
on its hash value

The partition number of a key depends on its hash value

The default partitioning logic is
implemented in the `getPartition()`
method of **HashPartitioner**

```
public class HashPartitioner<K, V> extends Partitioner<K, V> {  
    public HashPartitioner() {  
    }  
  
    public int getPartition(K key, V value, int numReduceTasks)  
    {  
        return (key.hashCode() & 2147483647) % numReduceTasks;  
    }  
}
```

The default partitioning logic is implemented in the `getPartition()` method of **HashPartitioner**

```
public class HashPartitioner<K, V> extends Partitioner<K, V>
    public HashPartitioner() {
    }

    public int getPartition(K key, V value, int numReduceTasks) {
        return (key.hashCode() & 2147483647) % numReduceTasks;
    }
}
```

Each key is assigned a unique Partition Id based on this logic

```
public class HashPartitioner<K, V>
  extends Partitioner<K, V> {

    public HashPartitioner() {
    }

    public int getPartition(K key, V value, int numReduceTasks) {
      return (key.hashCode() & 2147483647) % numReduceTasks;
    }
}
```

You can write your own Partitioning logic
by **extending the Partitioner Class** and
implementing the `getPartition()` method

Partitioning logic should be such that keys
are evenly distributed among the partitions

HashPartitioner does a pretty
good job of this

MapReduce in Python with the Hadoop Streaming API

Hadoop Streaming API

Since Hadoop is written in Java, it's natural to write MapReduce programs in Java

Hadoop Streaming API

But, you can implement the
map() and reduce() functions
in ANY LANGUAGE

Hadoop Streaming API

Hadoop provides the
Streaming API for
exactly this purpose

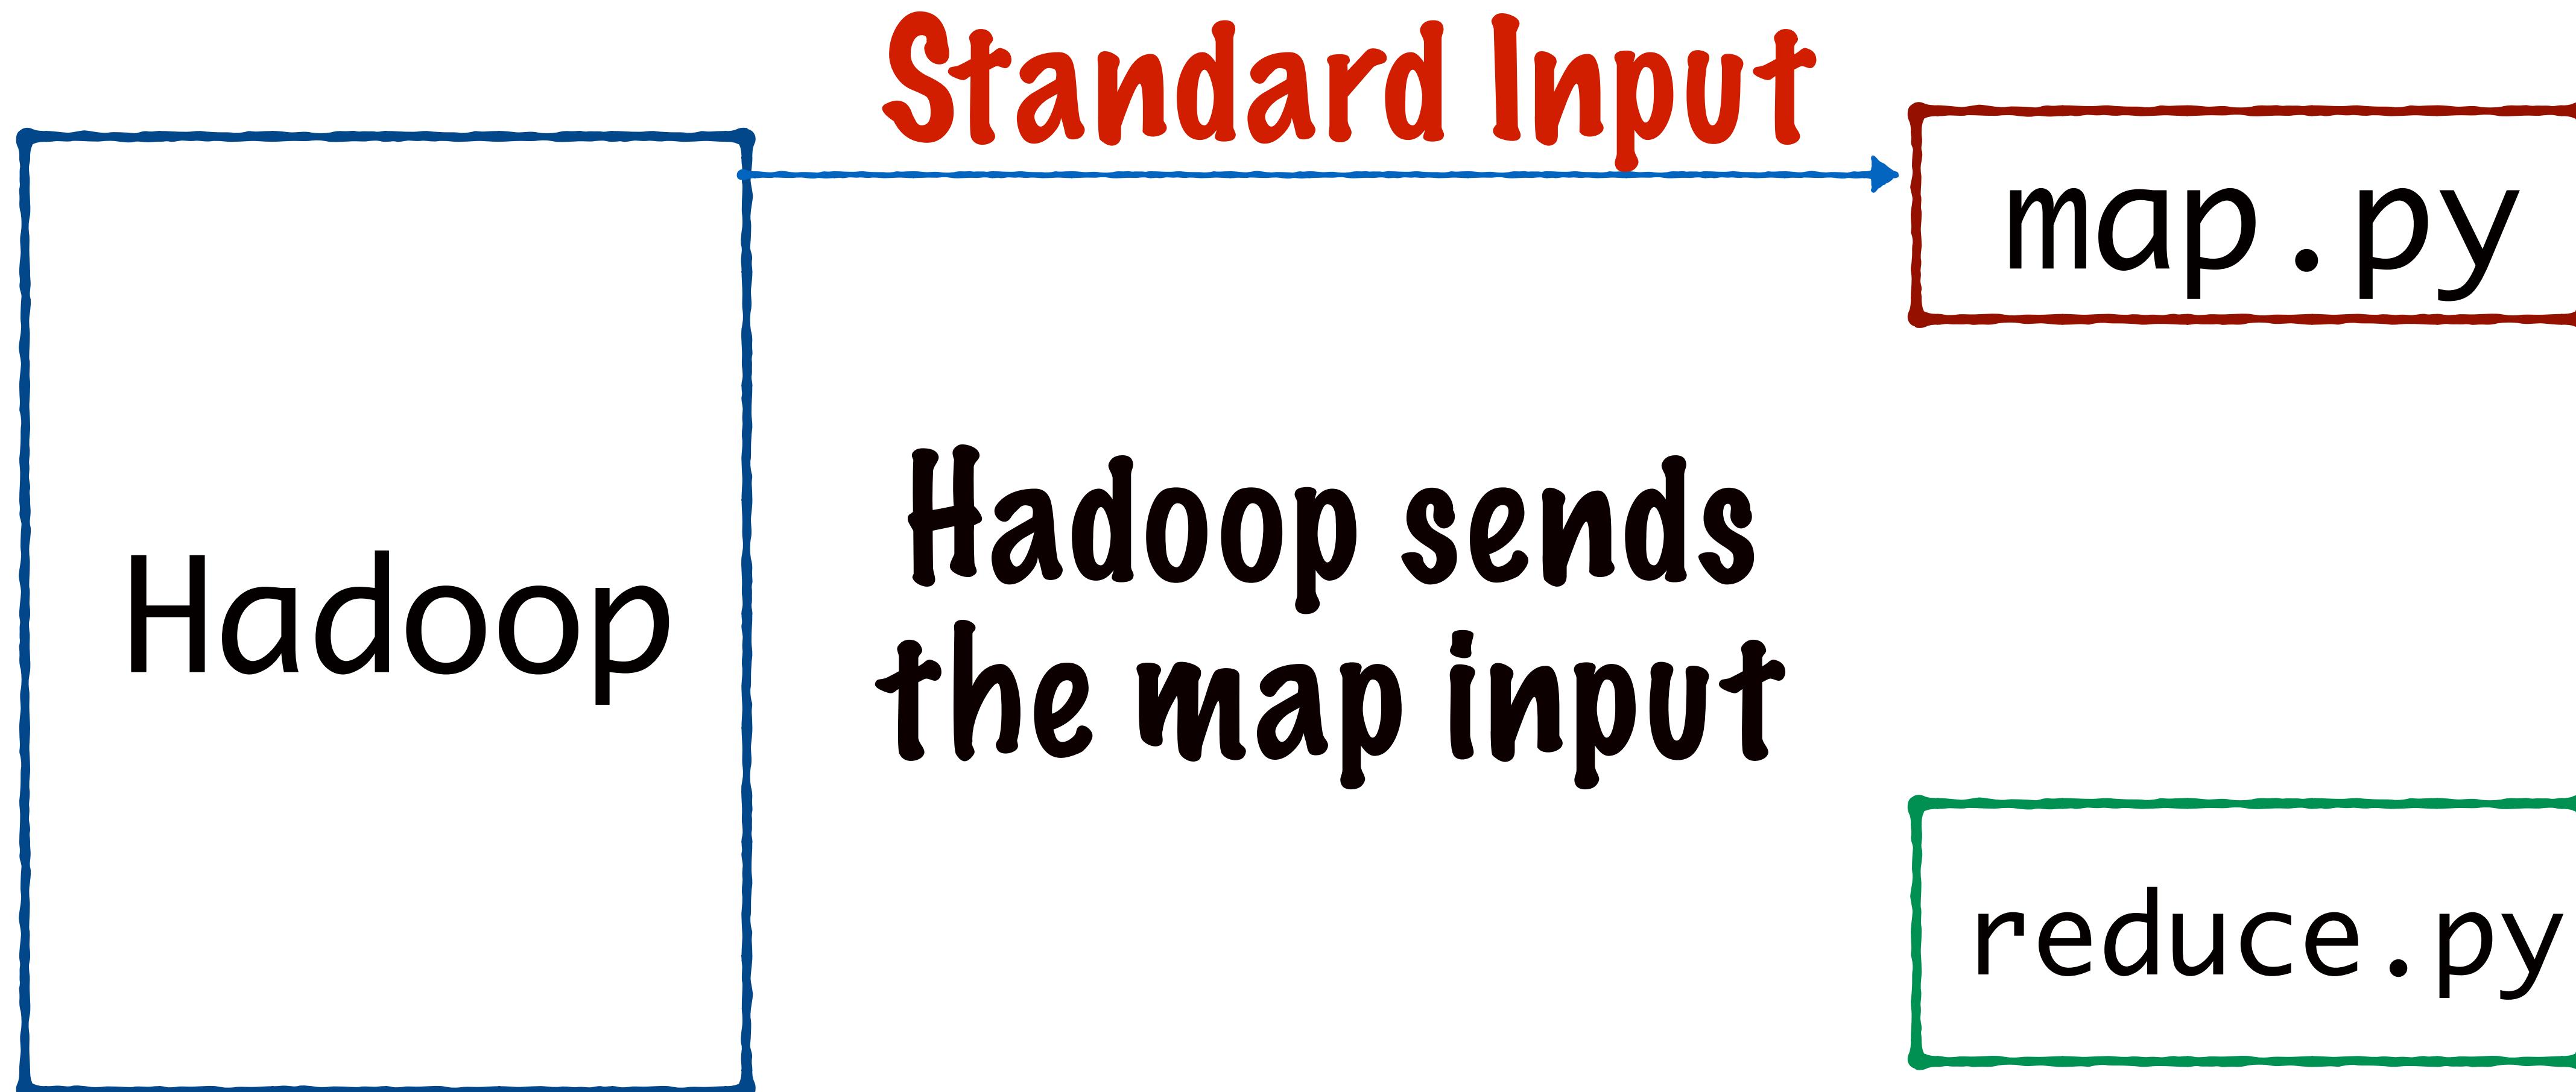
Hadoop Streaming API

The Streaming API uses
Standard Input/Output to
communicate with your
program

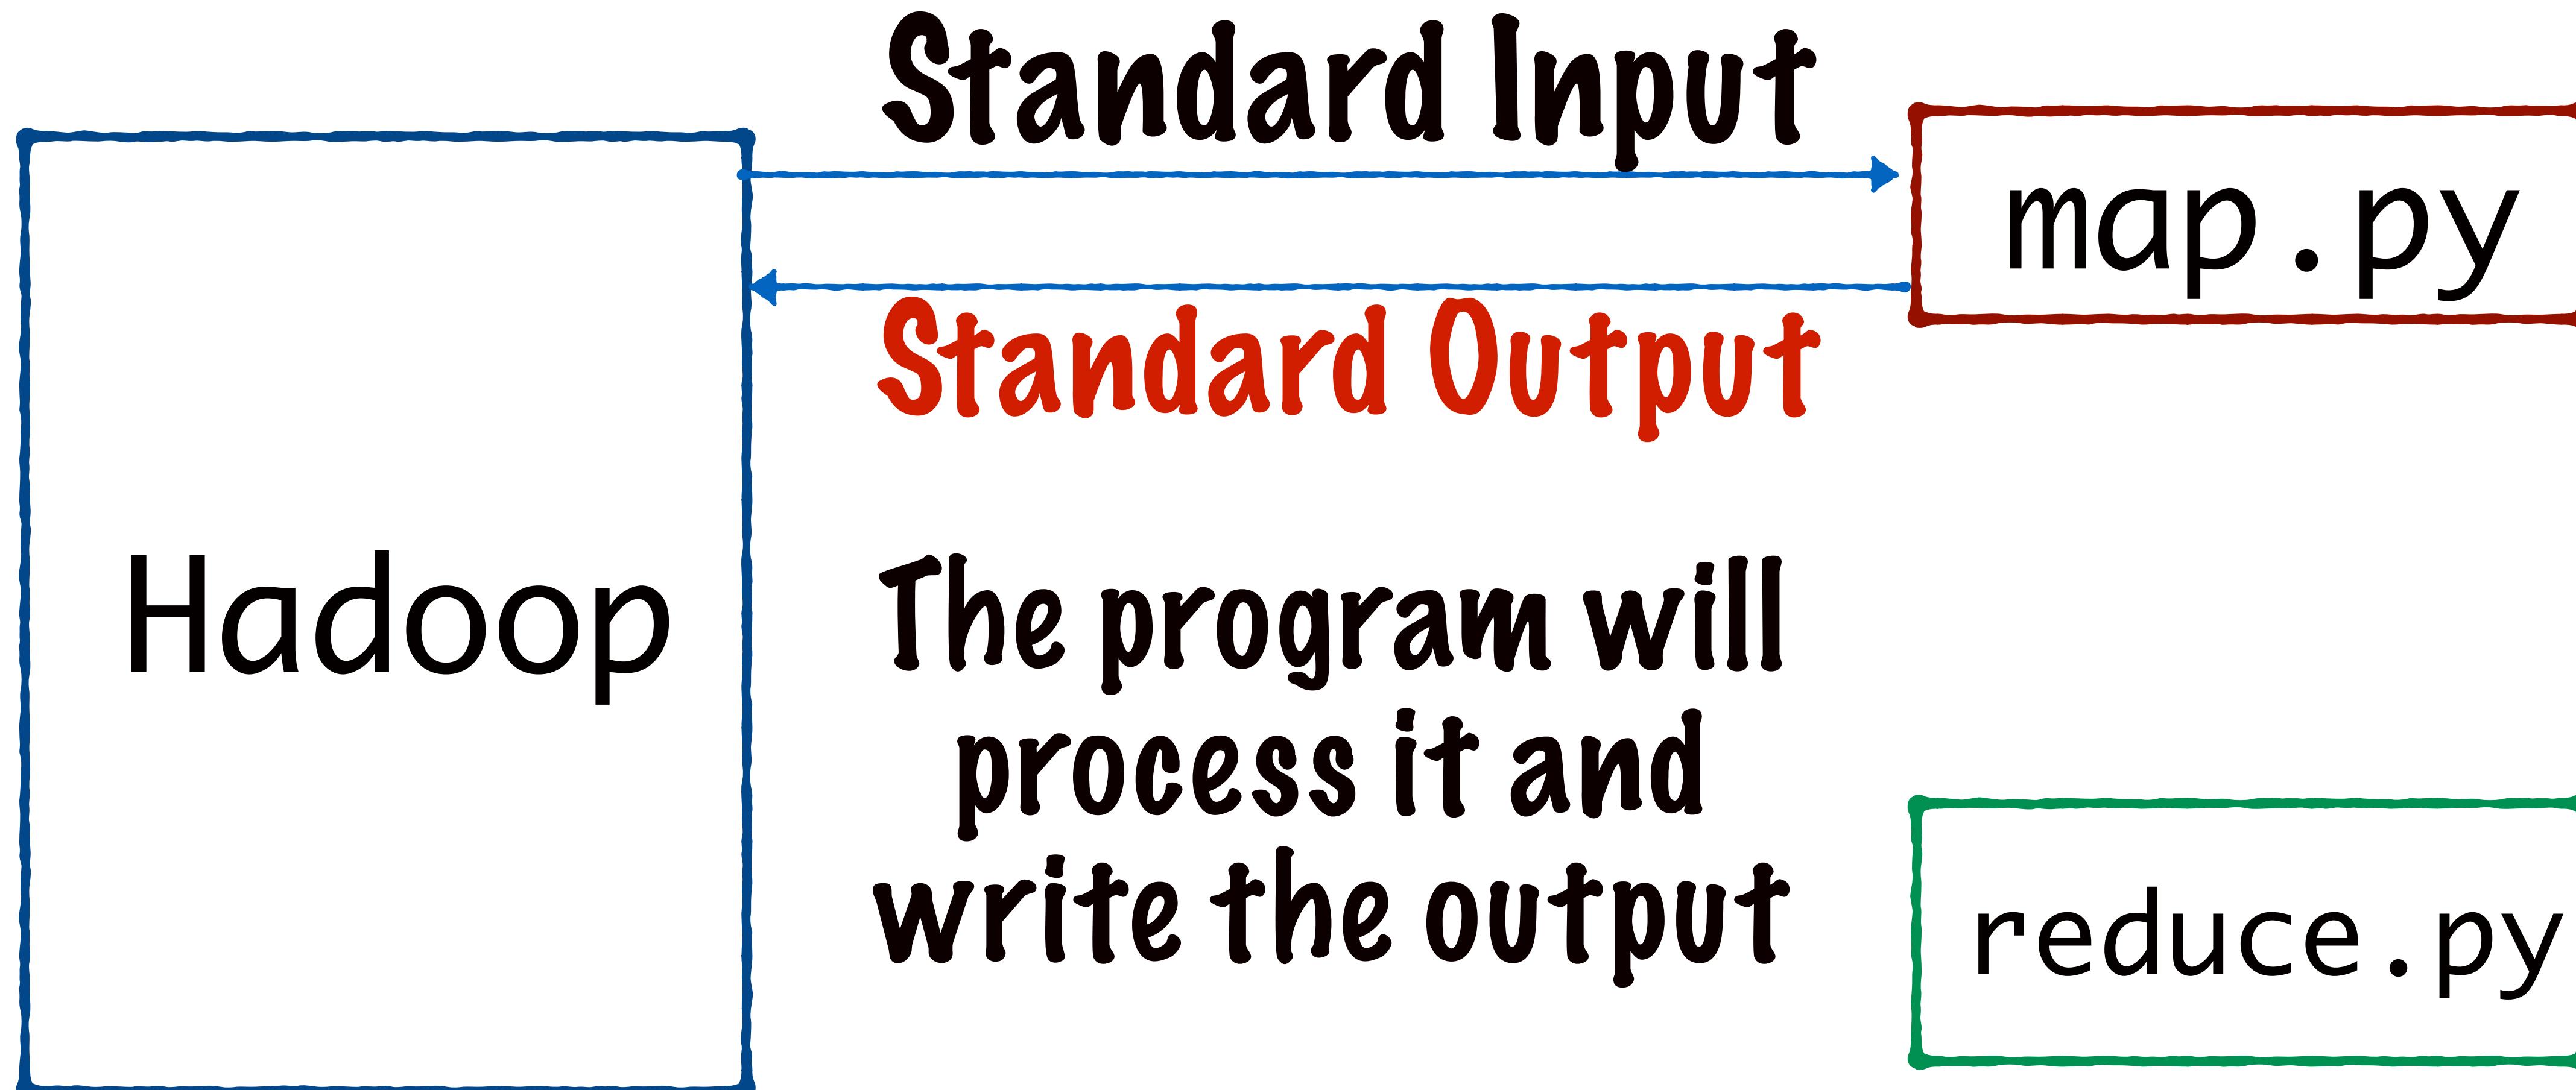
Hadoop Streaming API

Let's say we implemented
the map and reduce
functions in Python

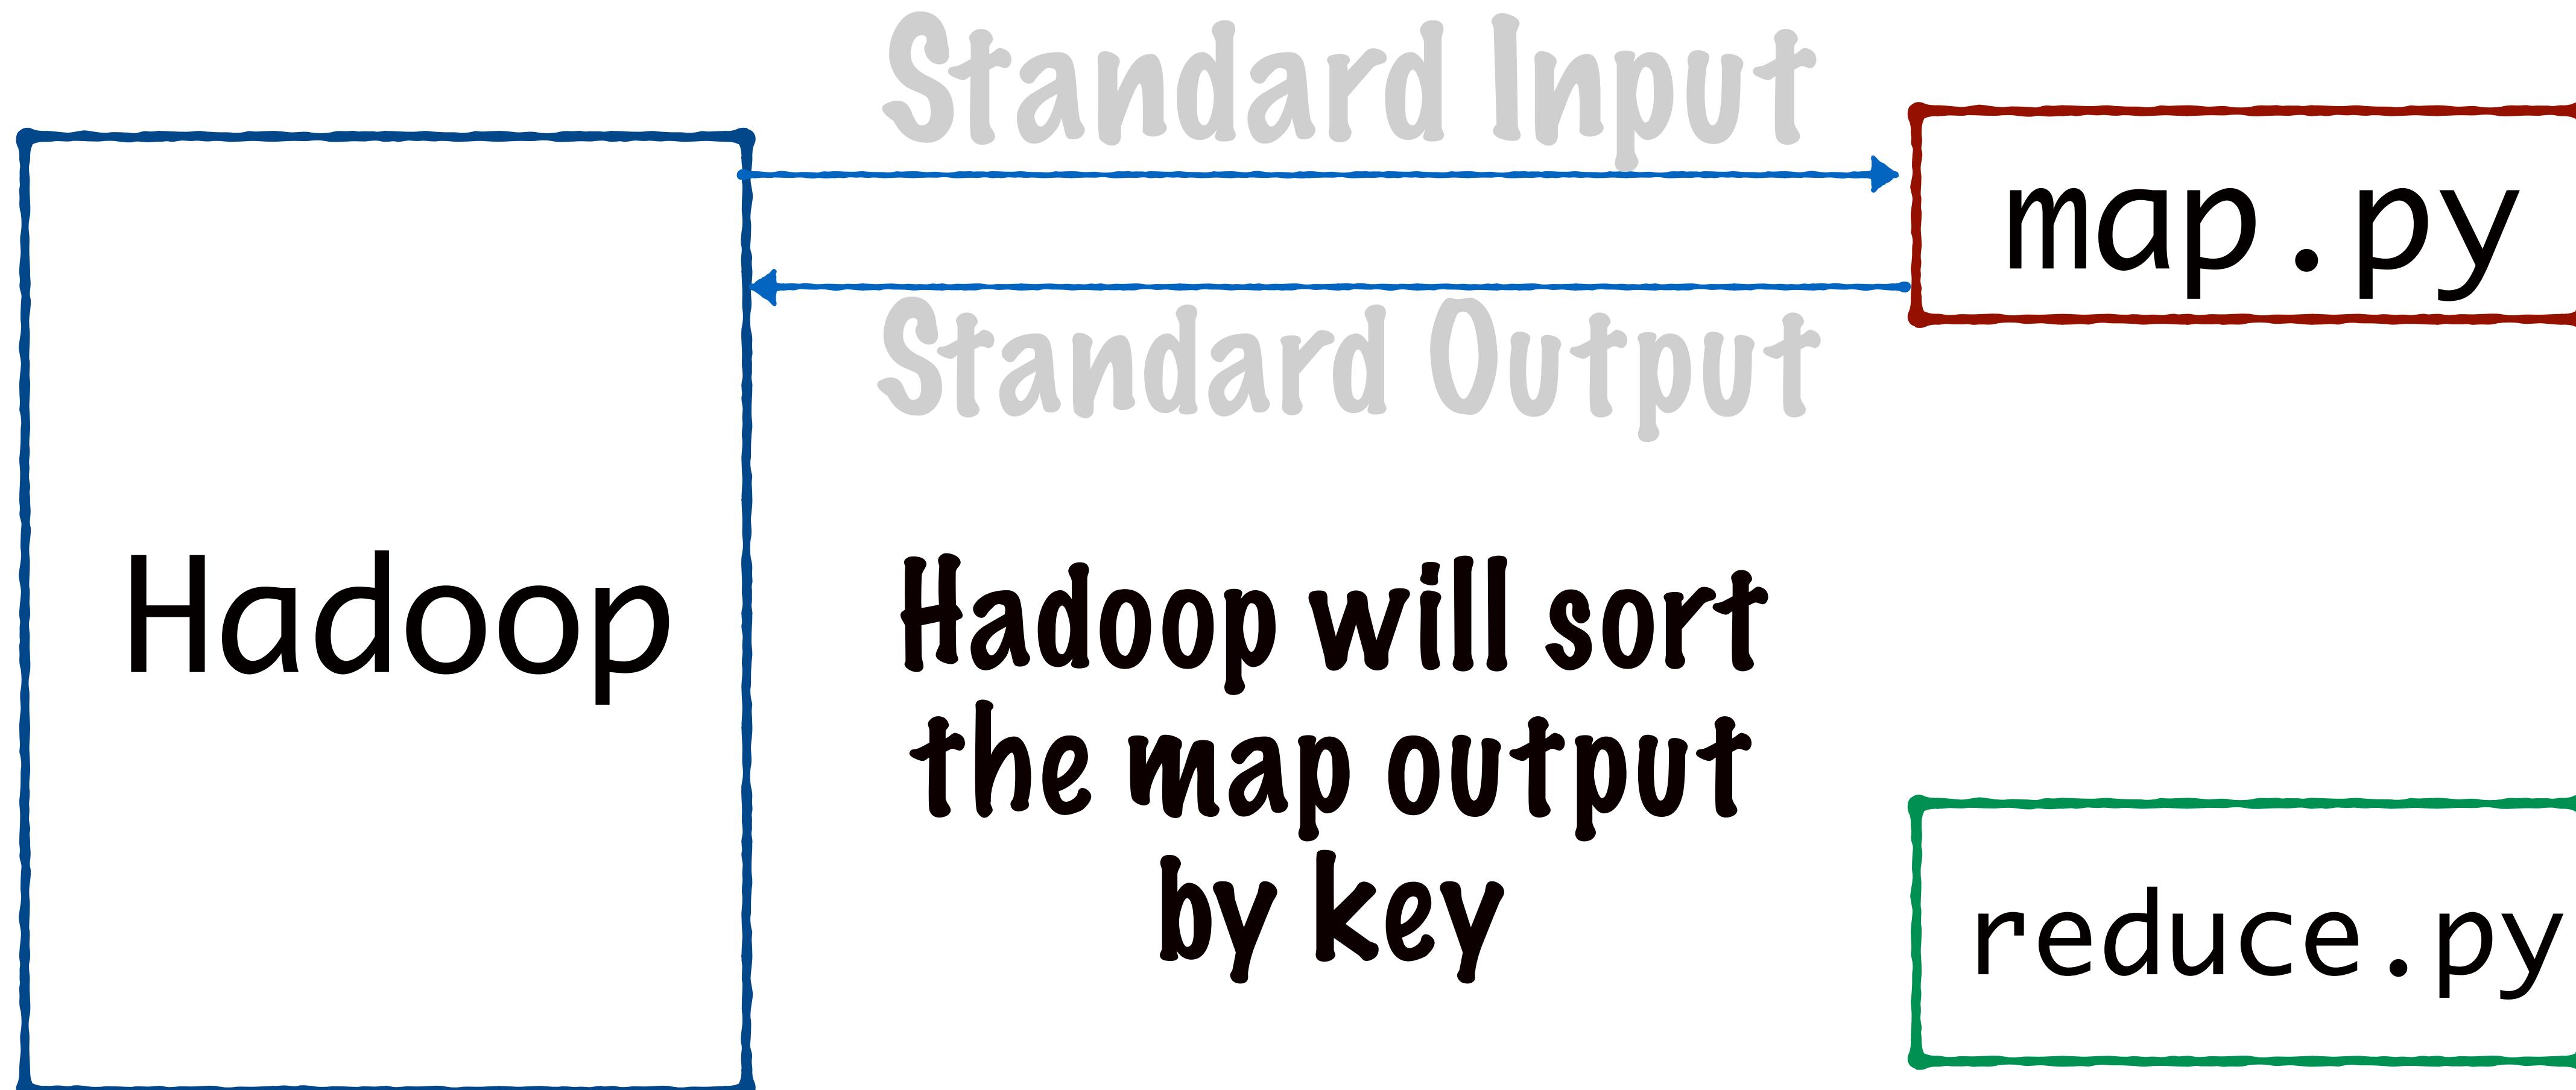
Hadoop Streaming API



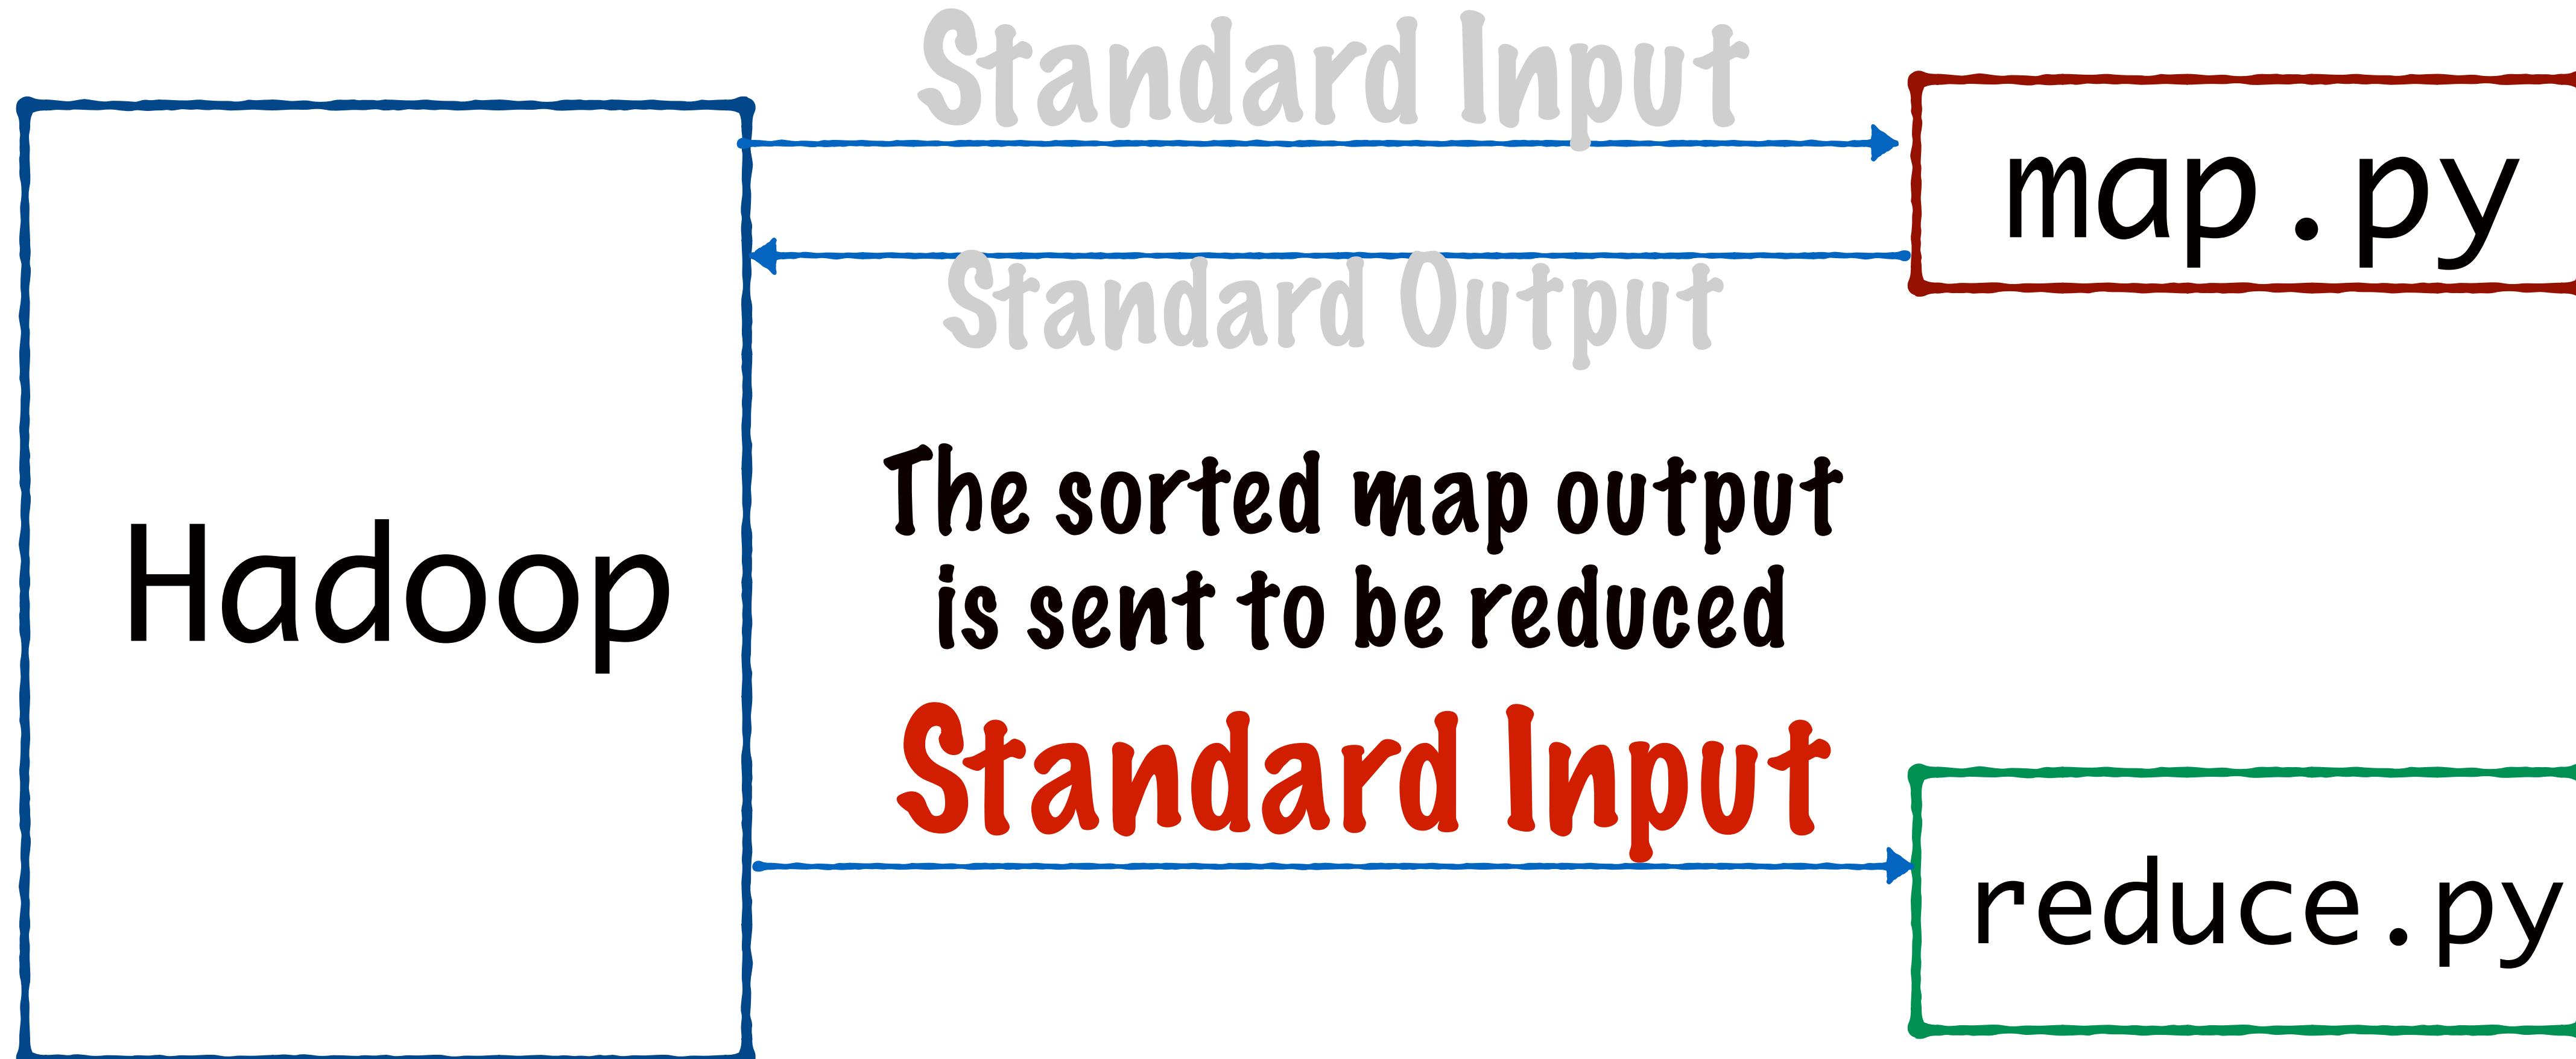
Hadoop Streaming API



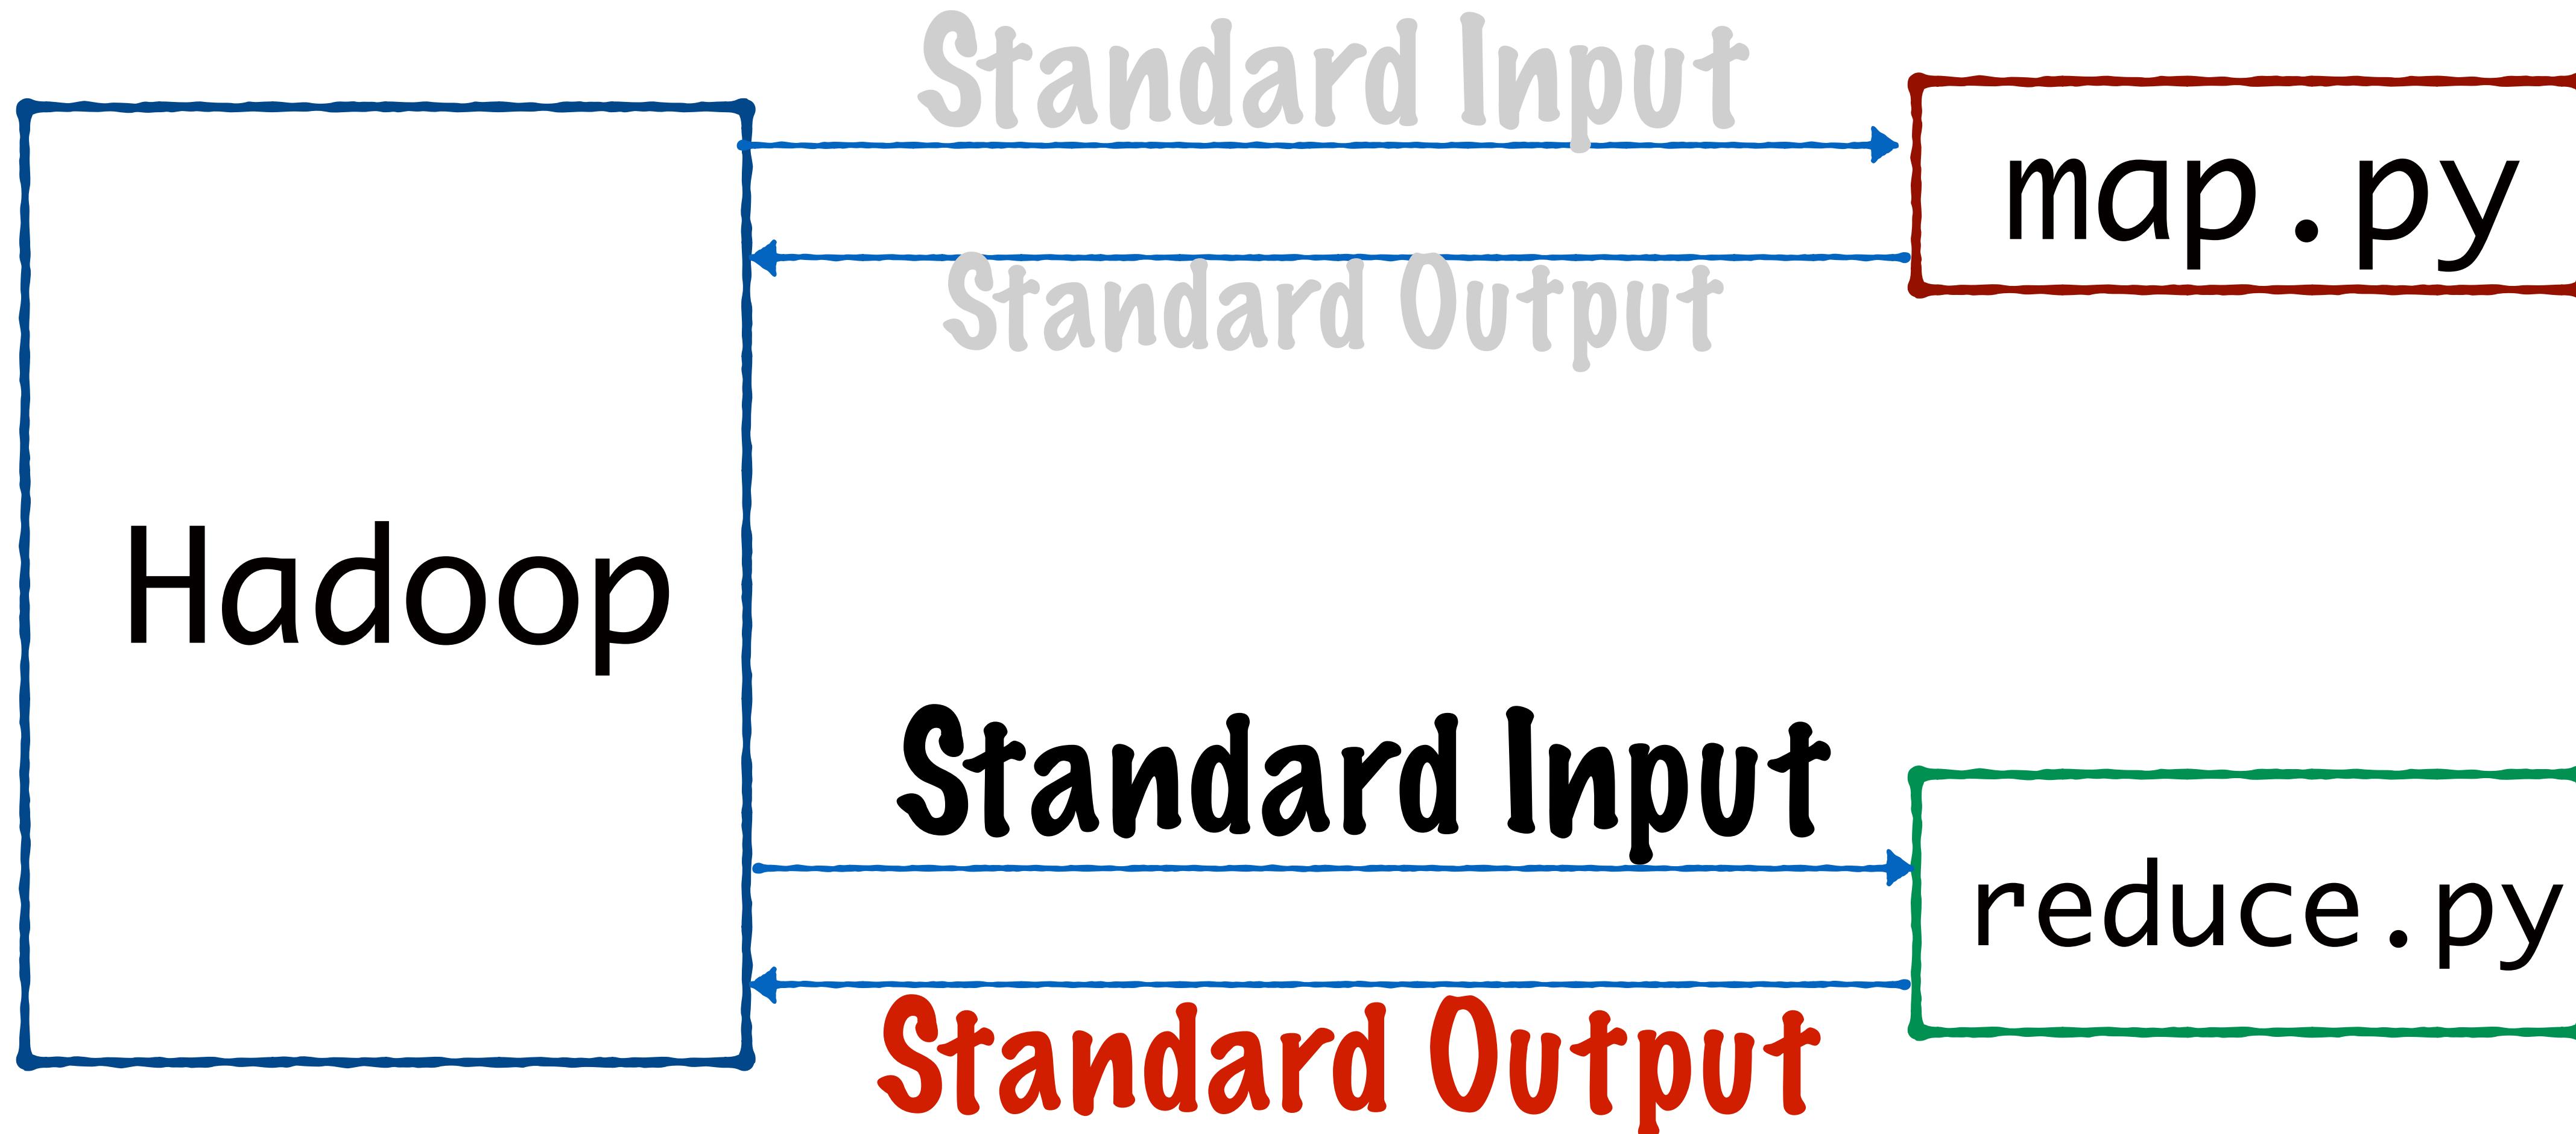
Hadoop Streaming API



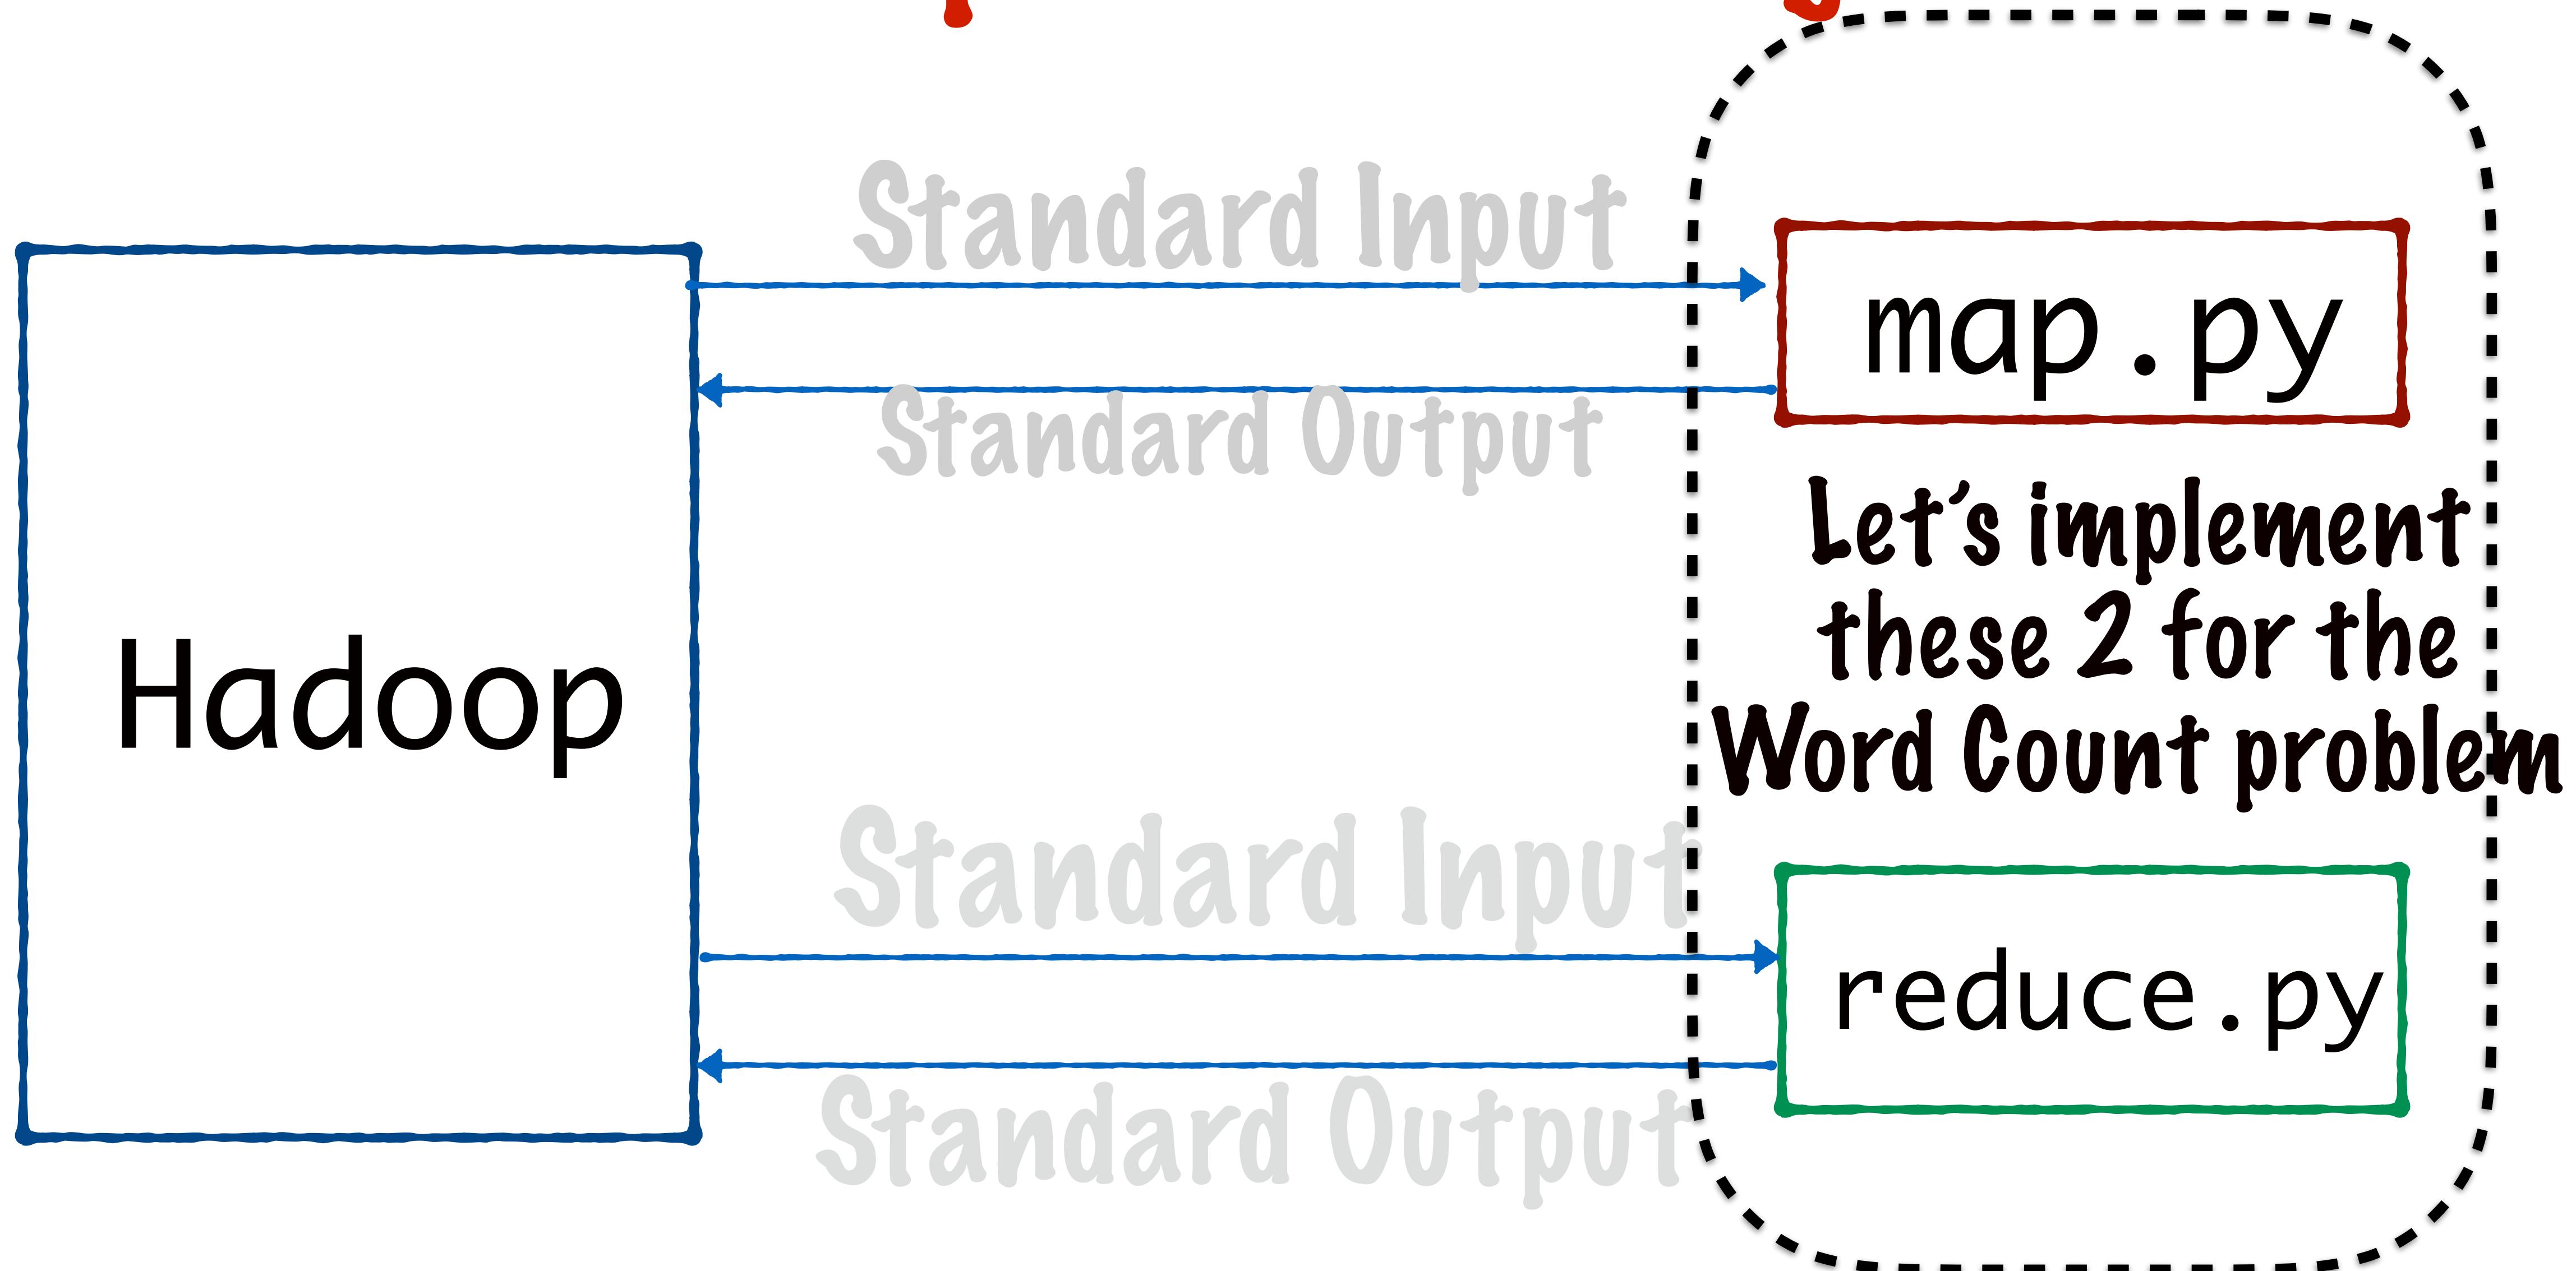
Hadoop Streaming API



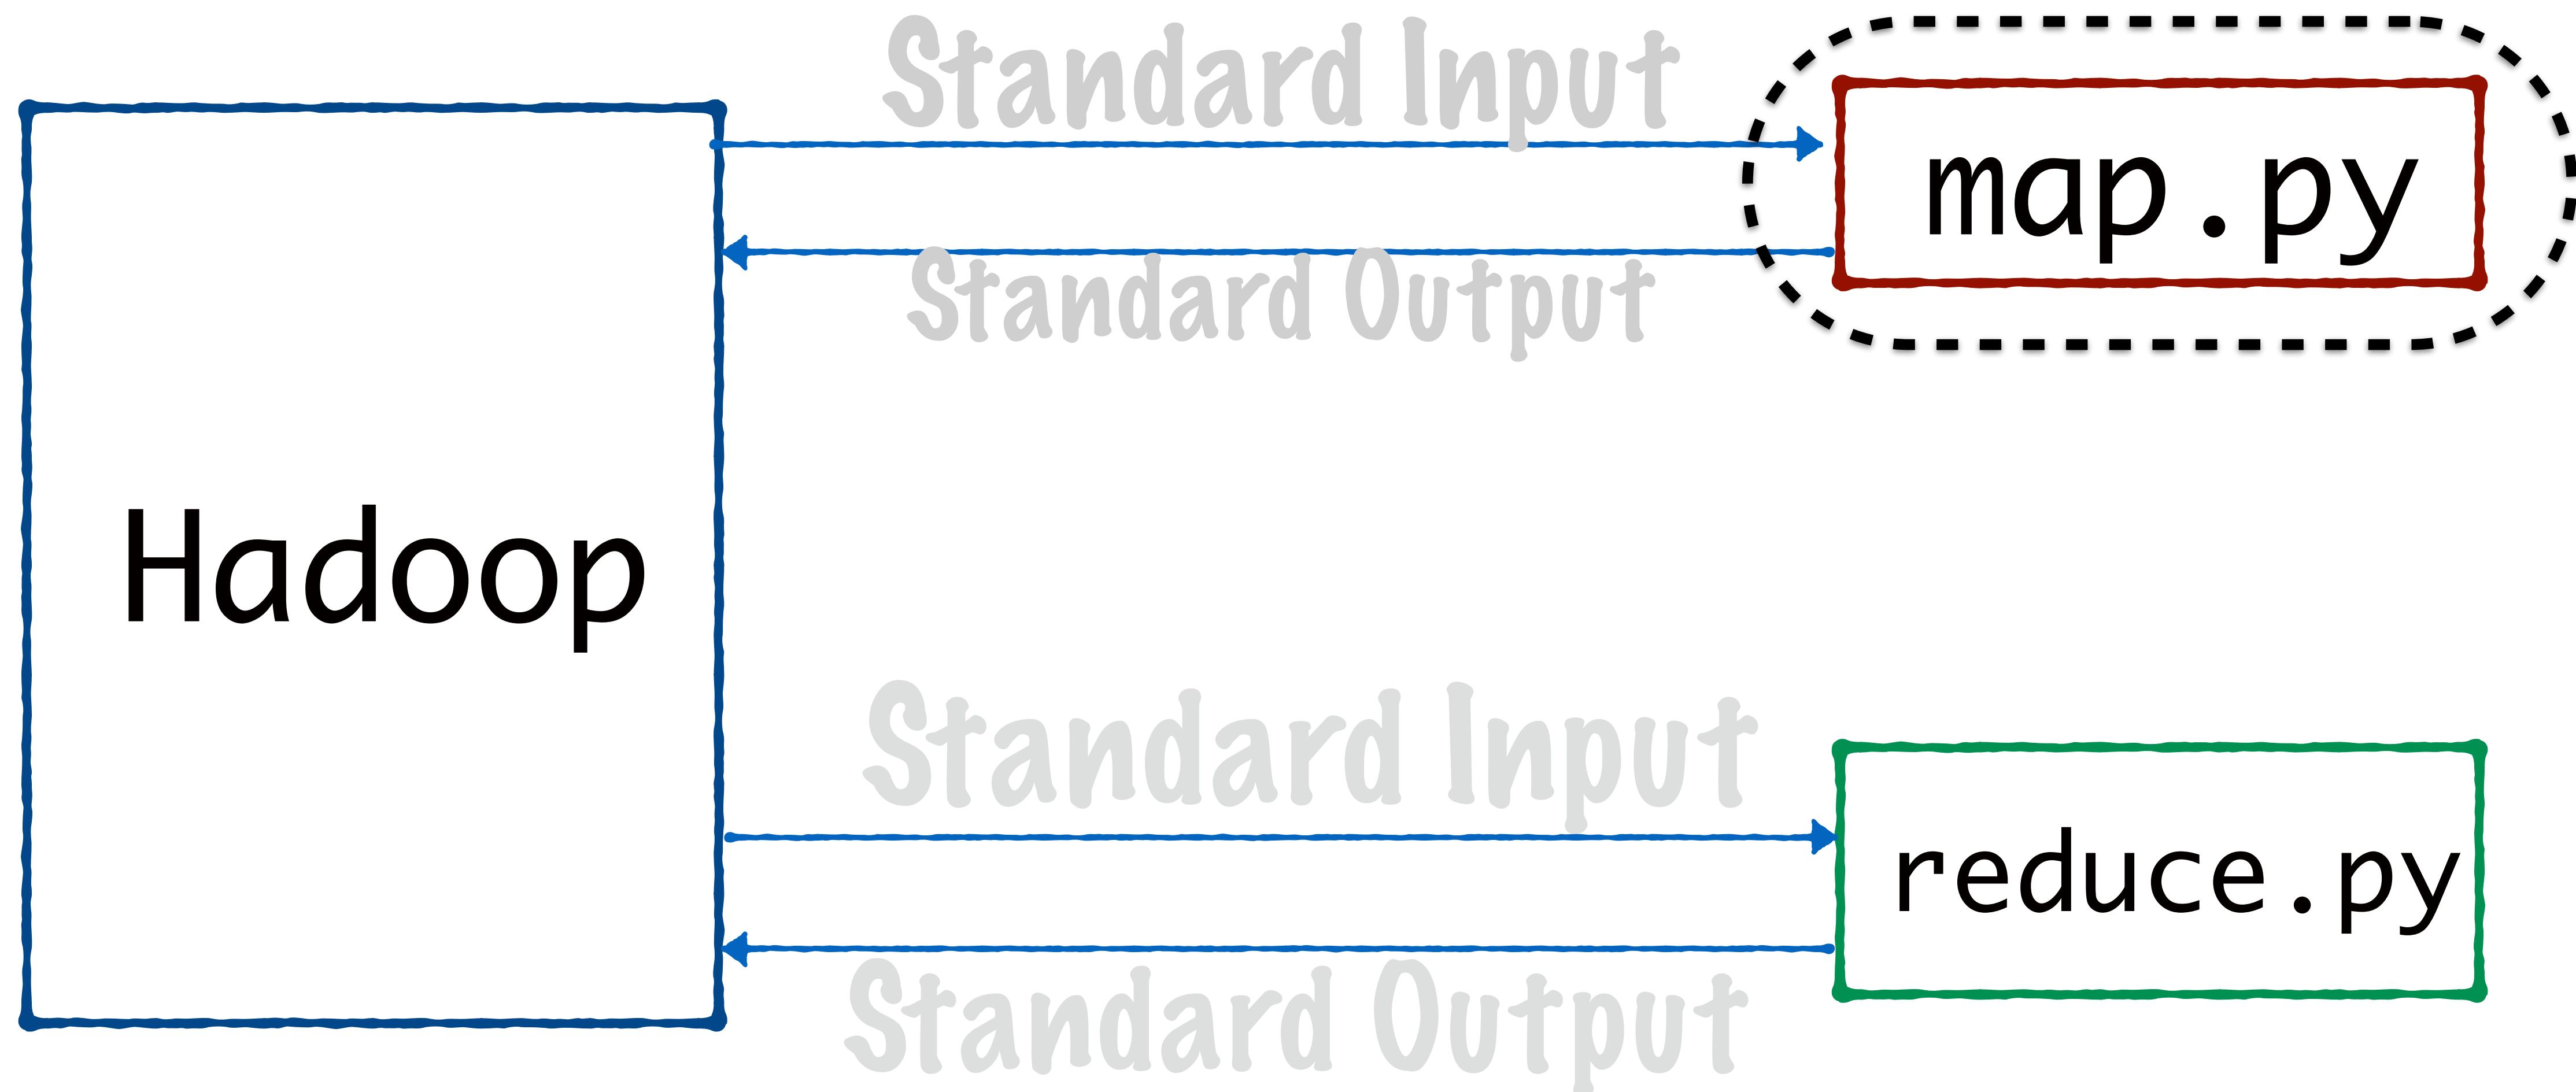
Hadoop Streaming API



Hadoop Streaming API



Hadoop Streaming API



Hadoop Streaming API

map.py

```
import sys
for line in sys.stdin:
    for word in line.split():
        print "%s\t%s" % (word, 1)
```

Hadoop Streaming API

map.py

This represents the code for the step

```
import  
for  
<1,Hey diddle diddle>
```

<linenum, Line of text>

map()

```
<Hey,1>  
<Diddle,1>  
<Diddle,1>  
<Word,1>
```

Hadoop Streaming API

map.py

The lines of the text file are sent from
Hadoop over Standard Input

```
import sys
for line in sys.stdin:
    for word in line.split():
        print "%s\t%s" % (word,1)
```

Hadoop Streaming API

map.py

For each word, we print
(word,1) to Standard Output

```
import sys
for line in sys.stdin:
    for word in line.split():
        print "%s\t%s" % (word, 1)
```

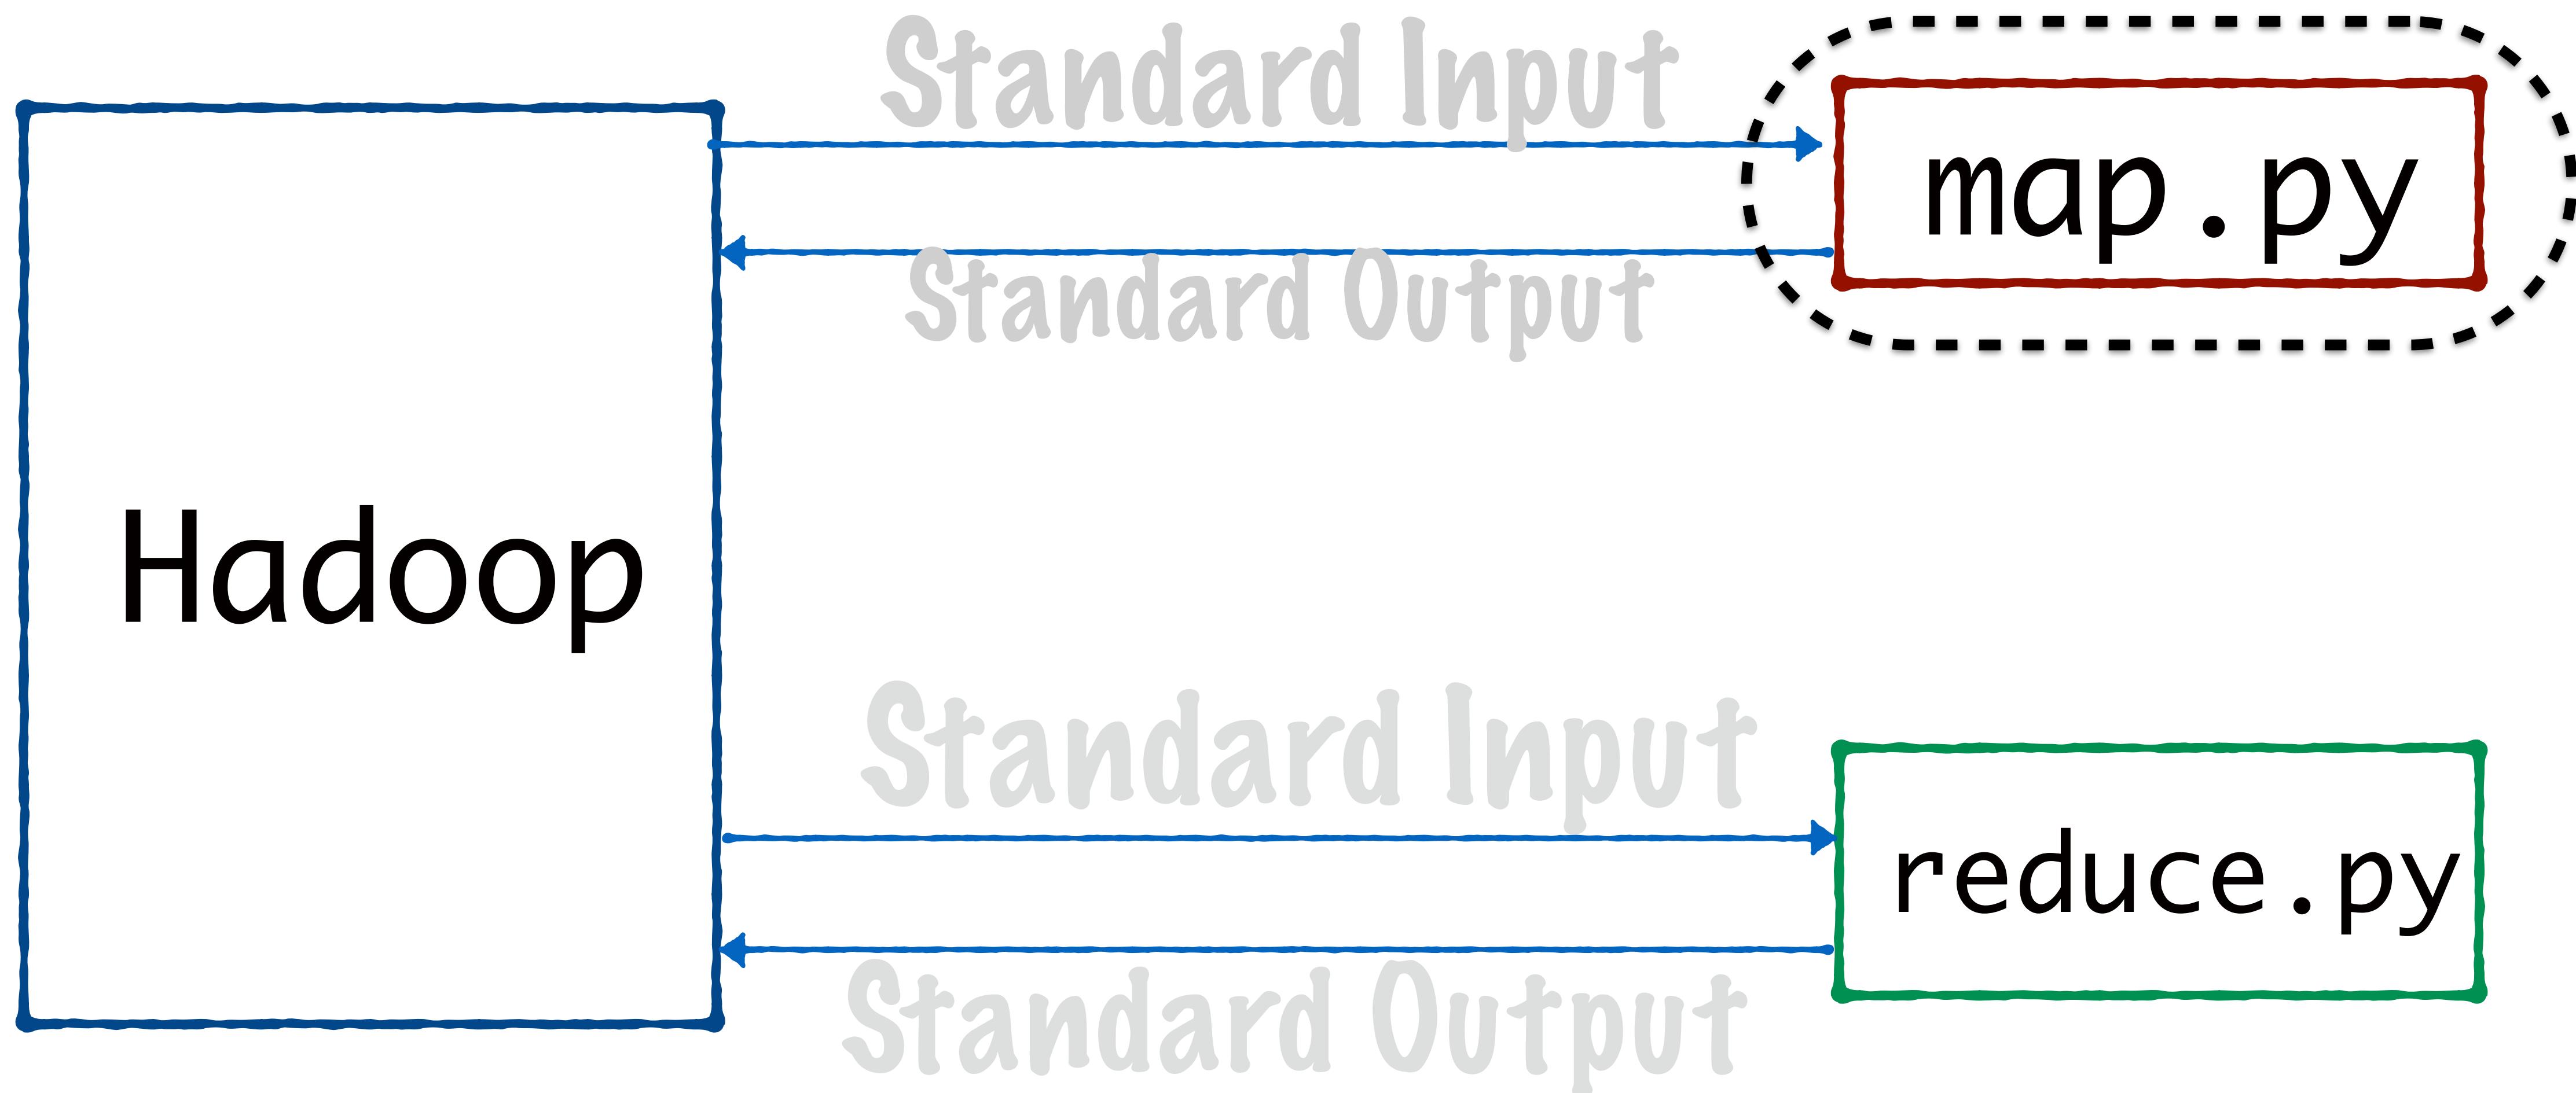
Hadoop Streaming API

map.py

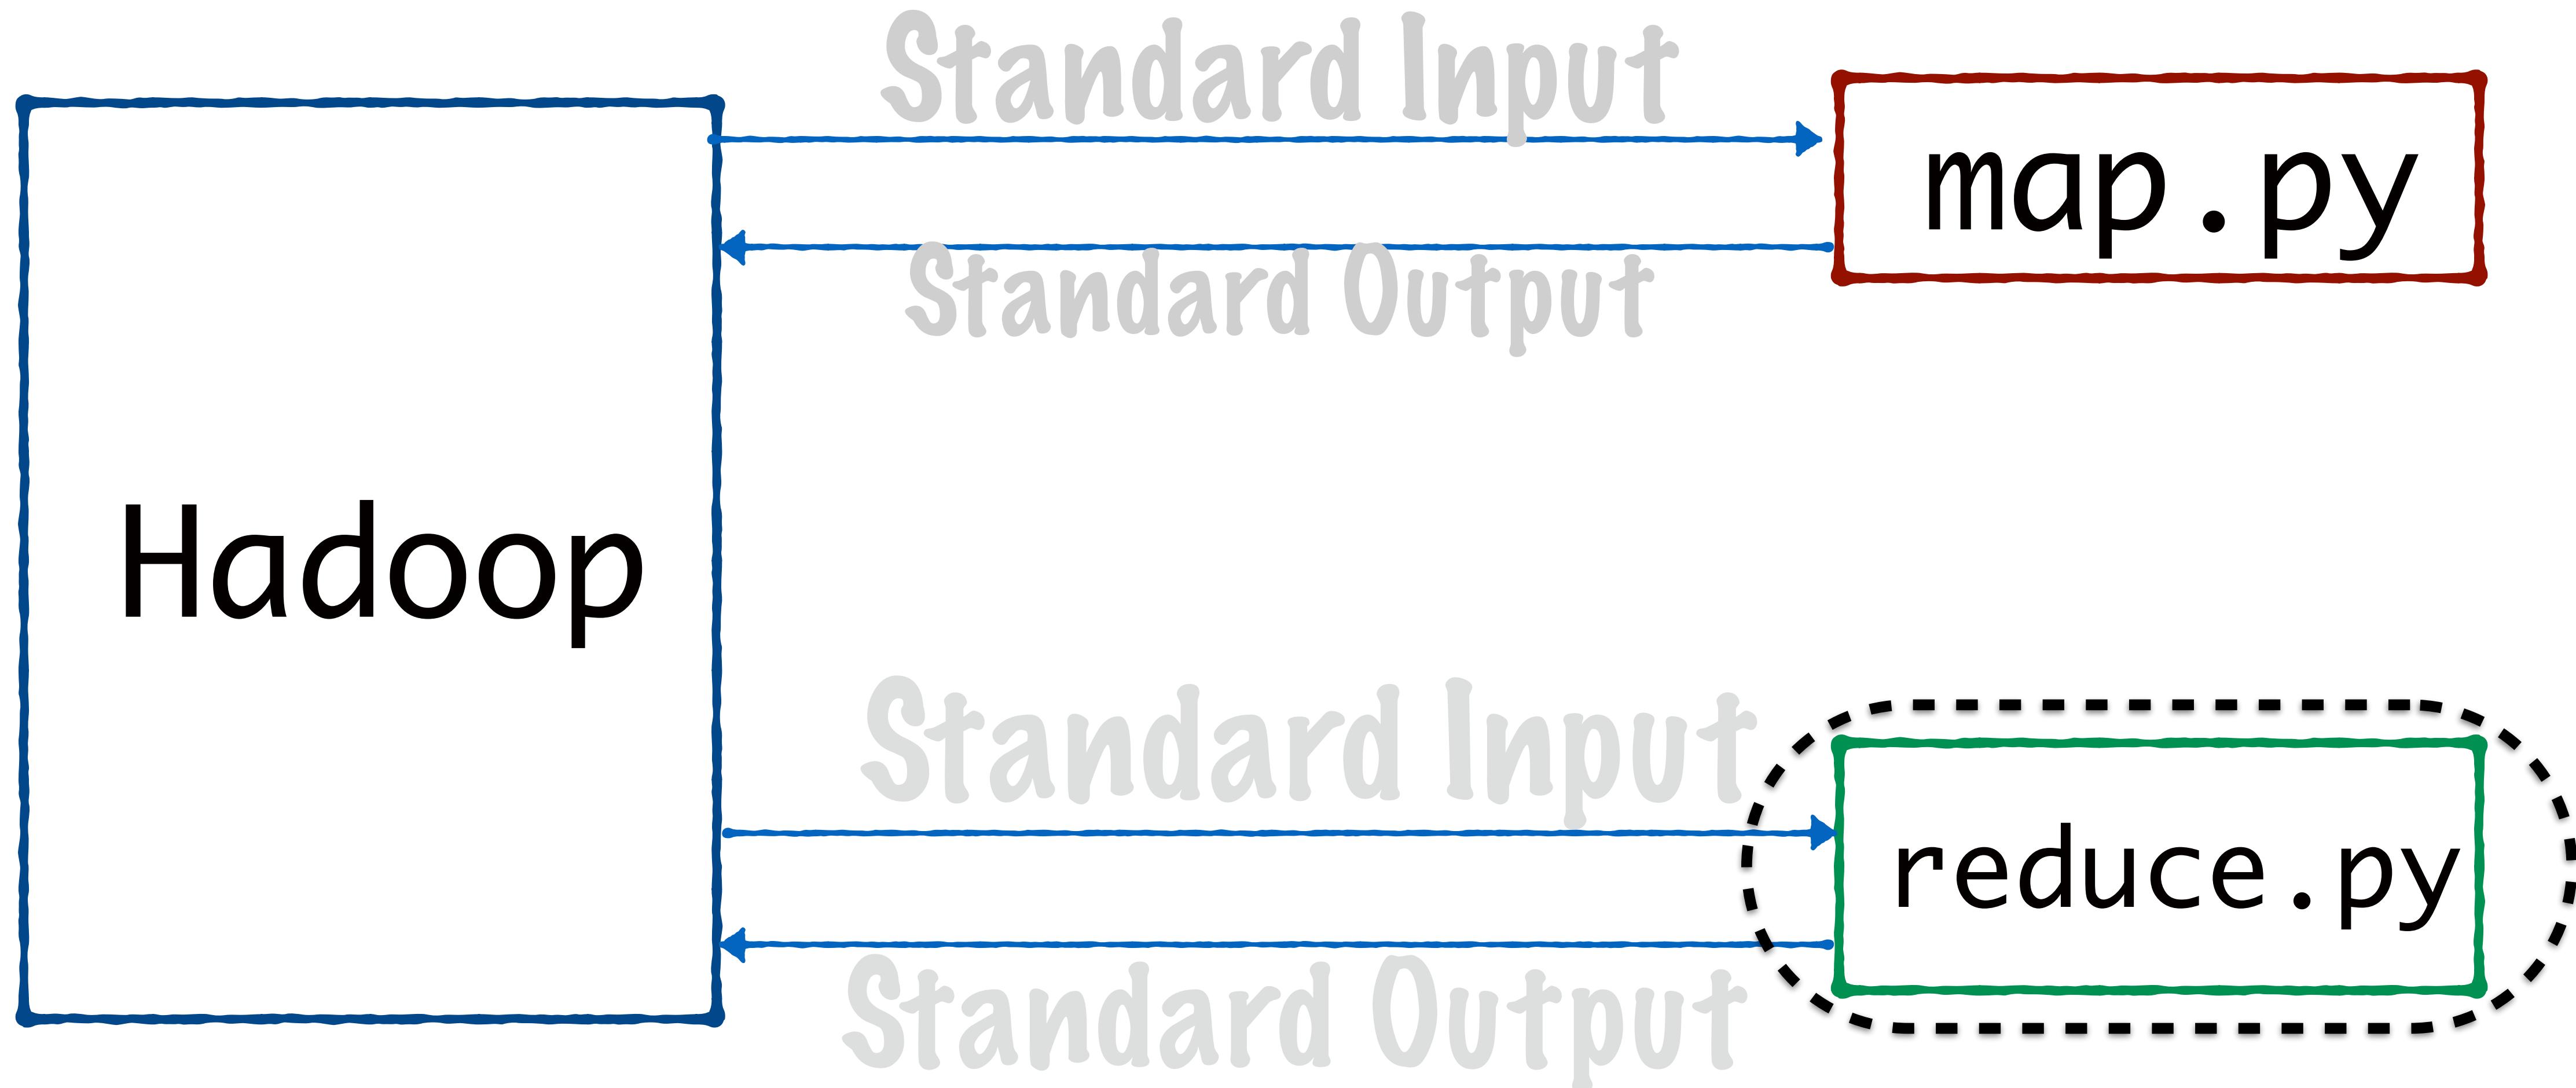
Hadoop expects the key and value to be separated by a tab

```
import sys
for line in sys.stdin:
    for word in line.split():
        print "%s\t%s" % (word, 1)
```

Hadoop Streaming API



Hadoop Streaming API



Hadoop Streaming API

reduce.py

The map output is sorted and sent to the reduce function

Unlike in Java, the sorted output is not merged

Hadoop Streaming API

reduce.py

```
import sys

(last_key,count) = (None, 0)
for line in sys.stdin:
    (key, value) = line.strip().split("\t")
    if last_key and last_key!=key:
        print "%s\t%s" % (last_key, count)
        (last_key, count) = (key, int(value))
    else:
        (last_key, count) = (key, count + int(value))
```

Hadoop Streaming API

reduce.py

```
import  
for  
    (last_key, count) = (key, 1)  
    for word, value in line.strip().split():  
        if word == last_key:  
            count += int(value)  
        else:  
            print(last_key, count)  
            (last_key, count) = (word, 1)  
    print(last_key, count)
```

This function represents
the code for the step

reduce()

<word, count>

Hadoop Streaming API

reduce.py

```
import sys
```

```
(last_key, count) = (None, 0)
for line in sys.stdin:
    (key, value) = line.strip().split("\t")
    if last_key and last_key!=key:
        print "%s\t%s" % (last_key, count)
        (last_key, count) = (key, int(value))
    else:
        (last_key, count) = (key, count + int(value))
```

Since the map output is coming
in like a stream, we need to keep
track of what the last key was

Hadoop Streaming API

reduce.py

As long as we see the same key,
we accumulate the count, else
we start over

```
import sys
```

```
(last_key, count) = (None, 0)
for line in sys.stdin:
    (key, value) = line.strip().split("\t")
    if last_key and last_key!=key:
        print "%s\t%s" % (last_key, count)
        (last_key, count) = (key, int(value))
    else:
        (last_key, count) = (key, count + int(value))
```

Hadoop Streaming API

reduce.py

In the beginning, we initialize
a value for the last key

```
import sys

(last_key, count) = (None, 0)
for line in sys.stdin:
    (key, value) = line.strip().split("\t")
    if last_key and last_key!=key:
        print "%s\t%s" % (last_key, count)
        (last_key, count) = (key, int(value))
    else:
        (last_key, count) = (key, count + int(value))
```

Hadoop Streaming API

reduce.py

Each pair is in 1 line of the
Standard Input, key and value
separated by tab

```
import sys

(last_key, count) = (None, 0)
for line in sys.stdin:
    (key, value) = line.strip().split("\t")
    if last_key and last_key!=key:
        print "%s\t%s" % (last_key, count)
        (last_key, count) = (key, int(value))
    else:
        (last_key, count) = (key, count + int(value))
```

Hadoop Streaming API

reduce.py

```
import sys
```

```
(last_key, count) = (None, 0)
for line in sys.stdin:
    (key, value) = line.strip().split("\t")
    if last_key and last_key != key:
        print "%s\t%s" % (last_key, count)
        (last_key, count) = (key, int(value))
    else:
        (last_key, count) = (key, count + int(value))
```

If we see a new key, we print the count for the last key and start over

Hadoop Streaming API

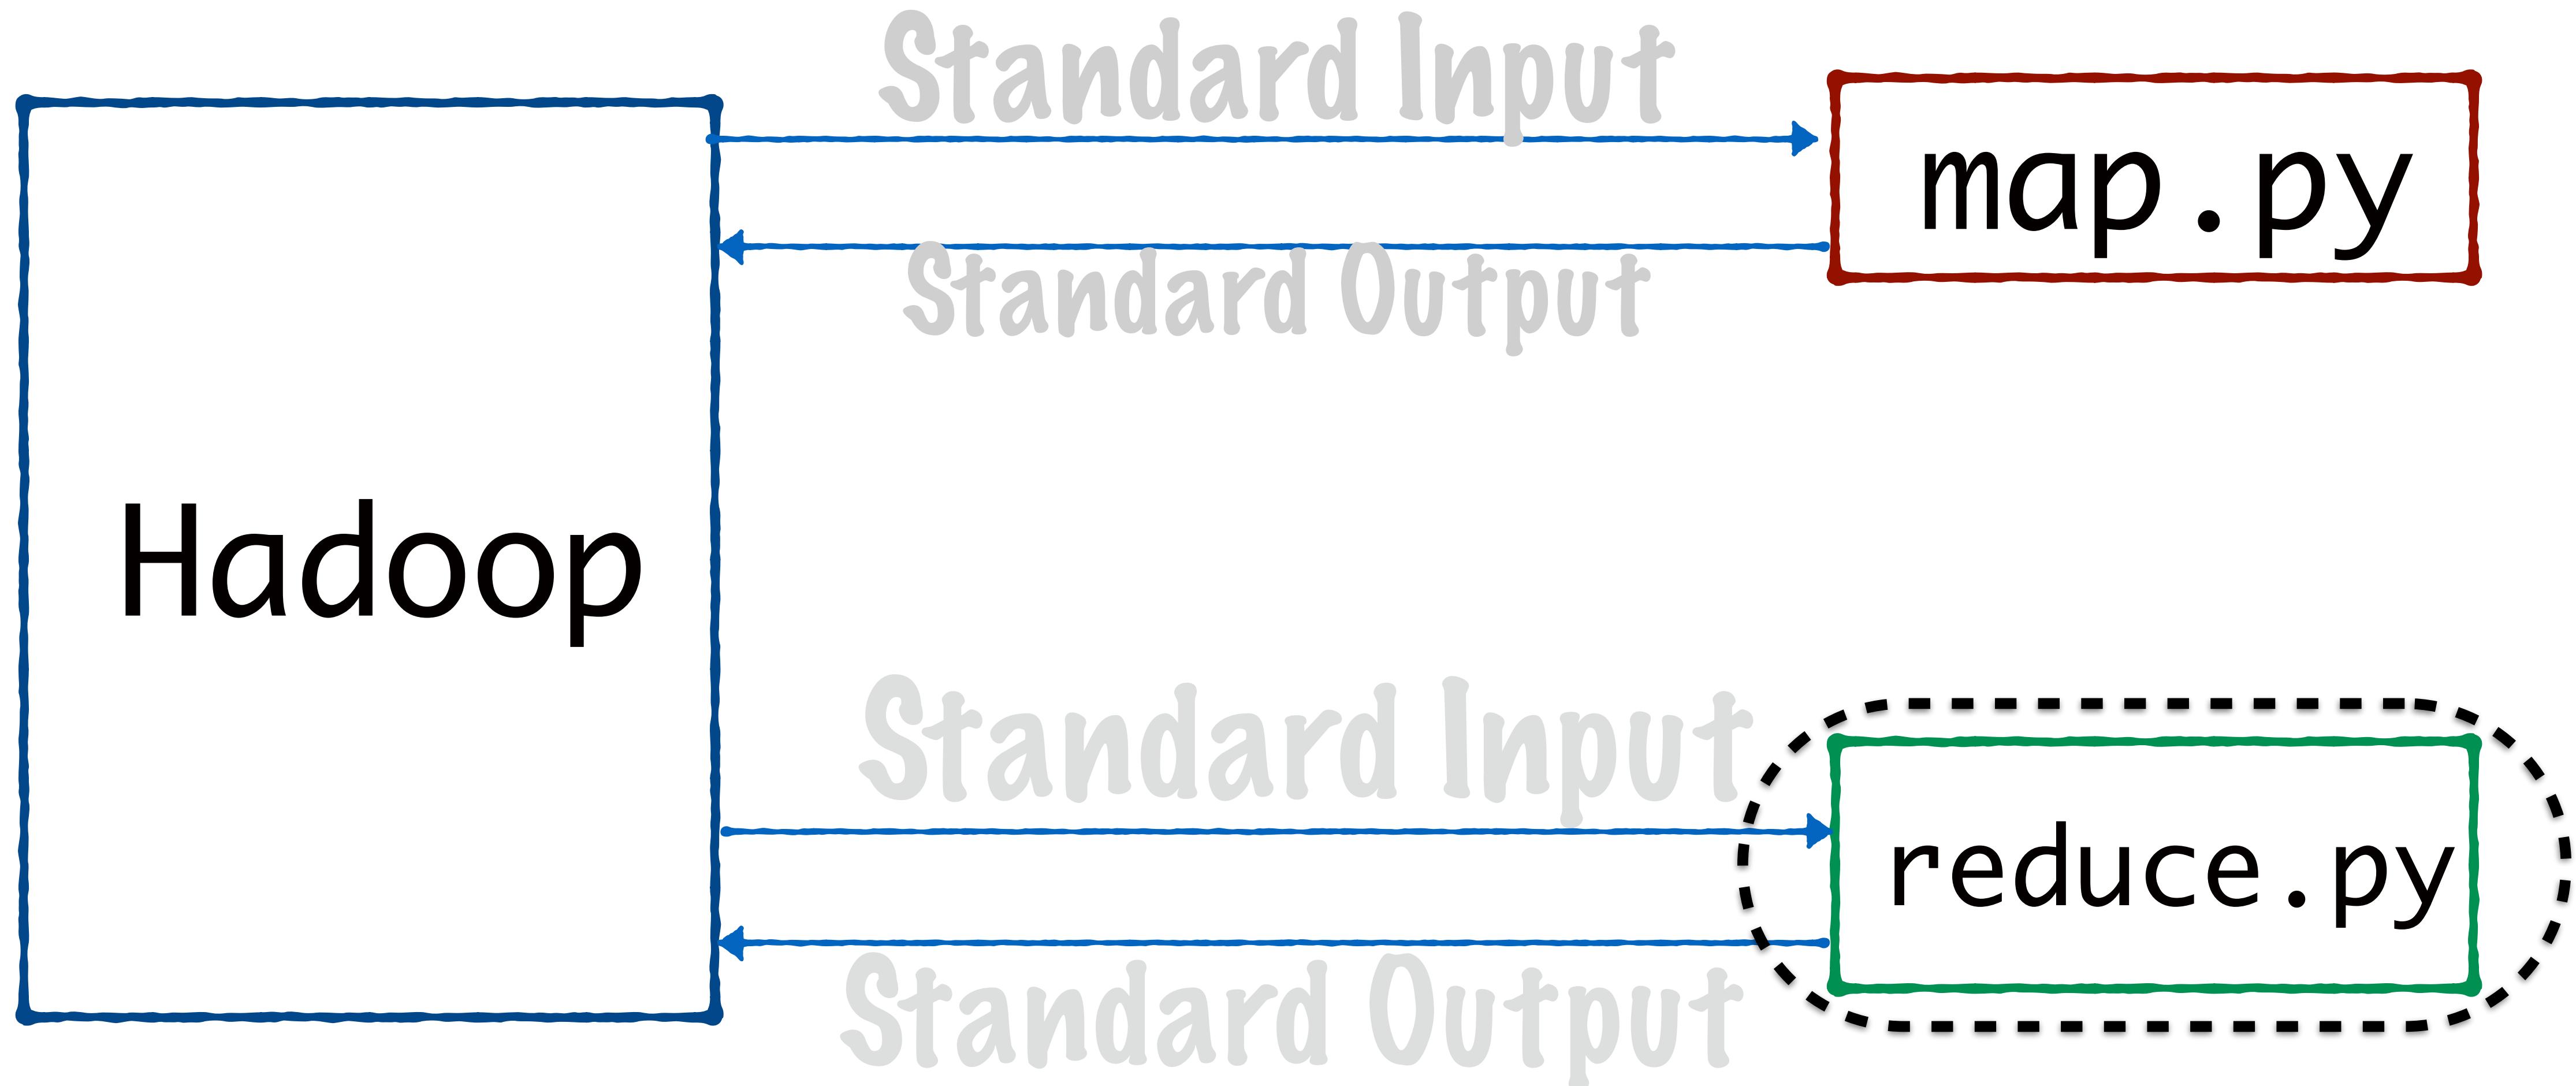
reduce.py

```
import sys
```

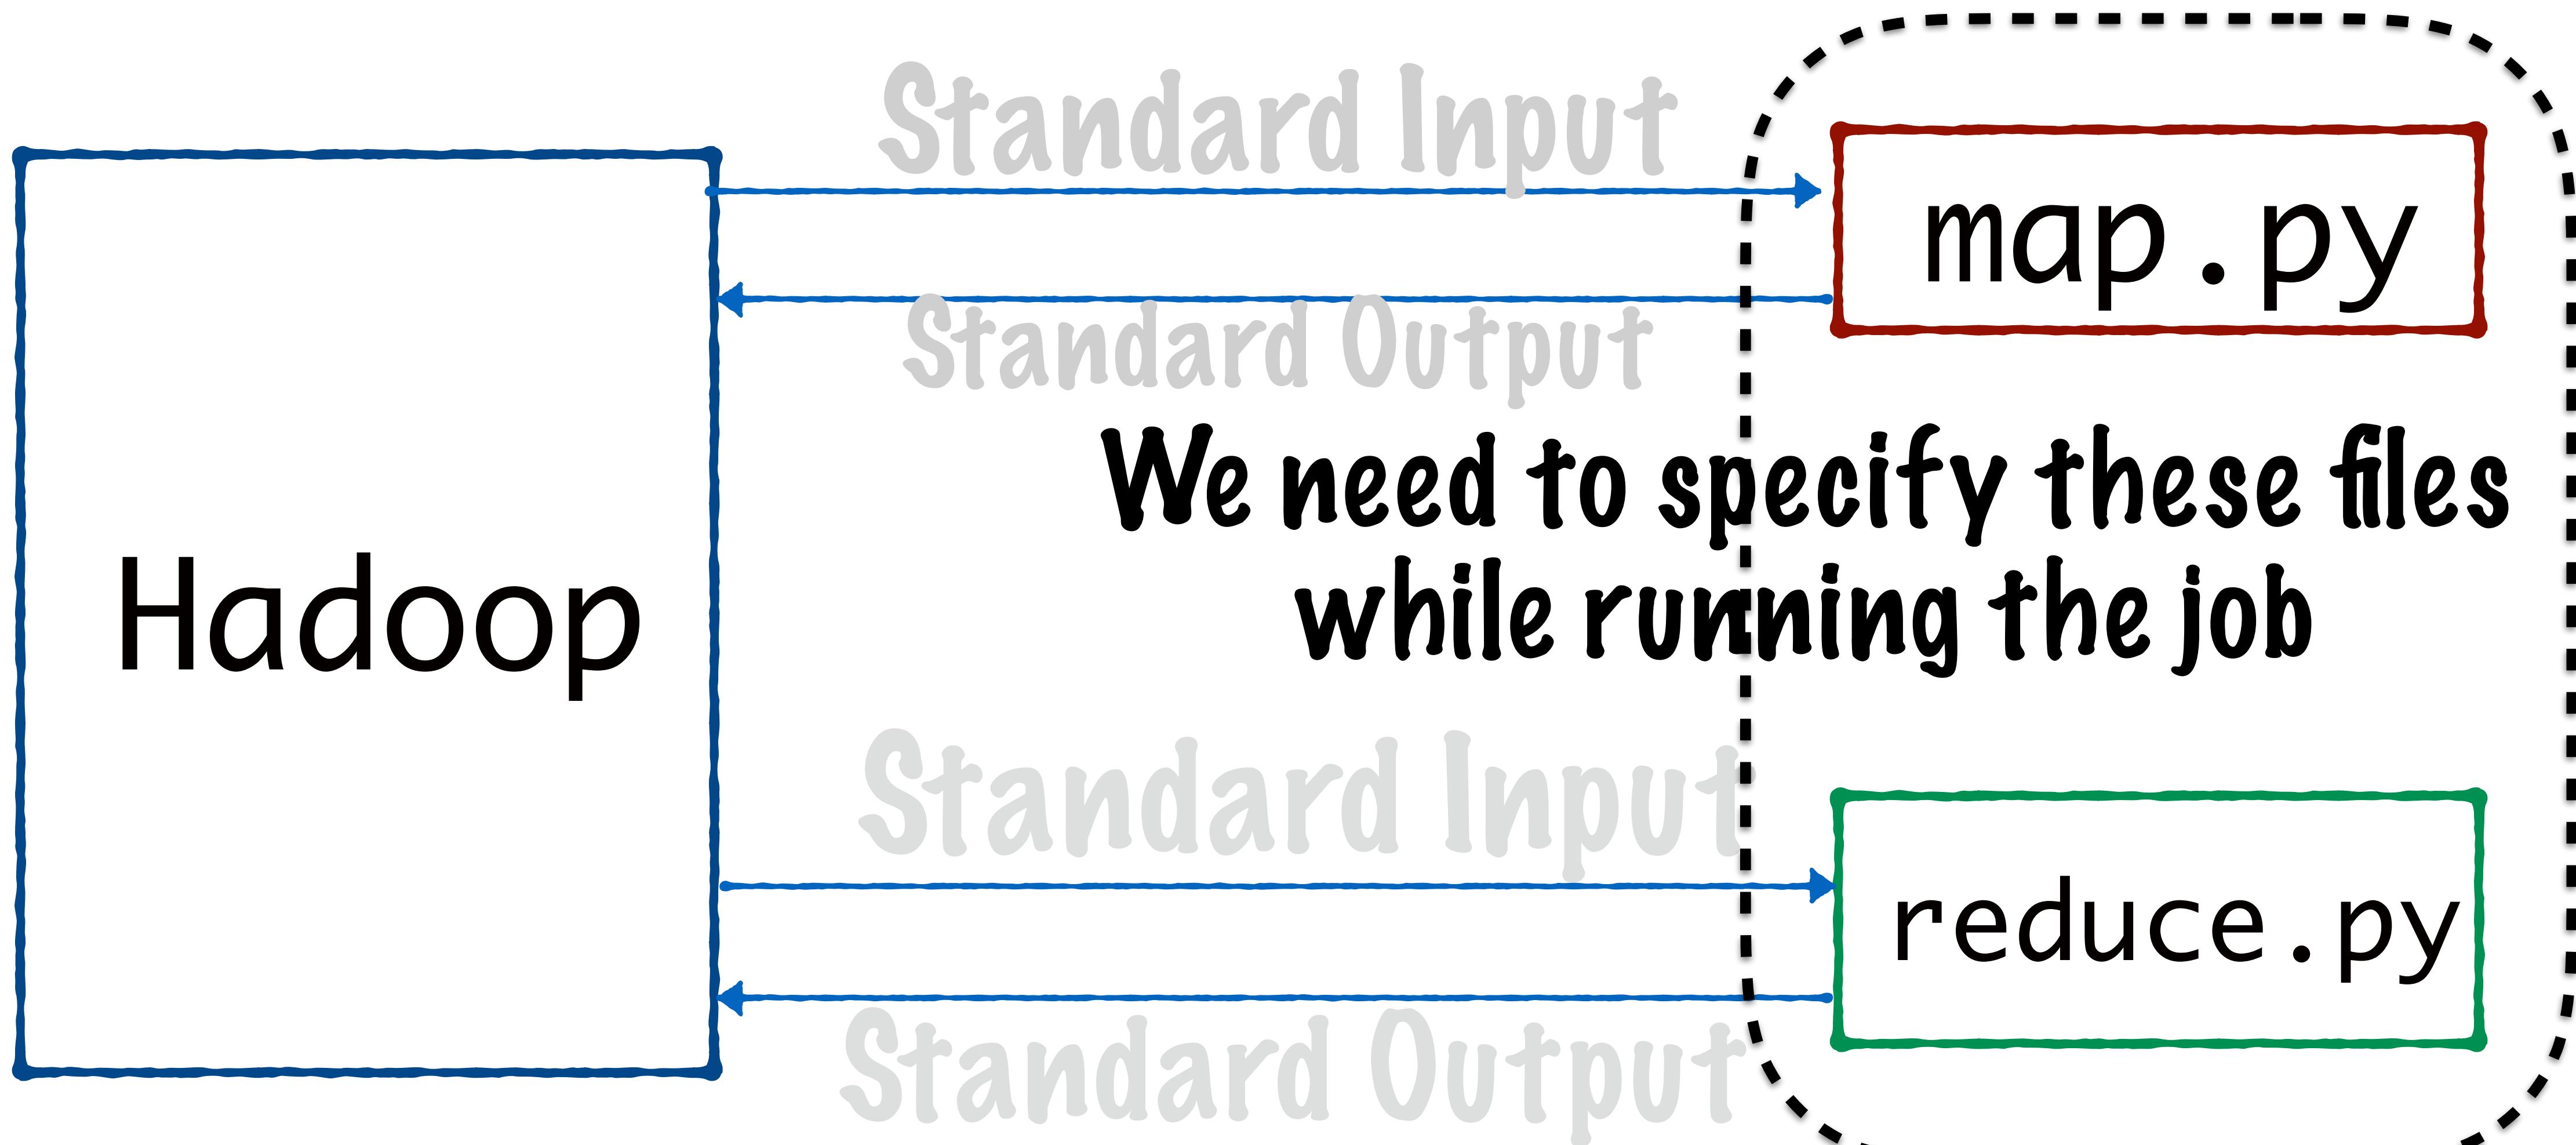
```
(last_key,count) = (None, 0)
for line in sys.stdin:
    (key, value) = line.strip().split("\t")
    if last_key and last_key!=key:
        print "%s\t%s" % (last_key,count)
        (last_key,count) = (key,int(value))
    else:
        (last_key,count) = (key, count + int(value))
```

If it's still the same key, then
we add to the word count

Hadoop Streaming API



Hadoop Streaming API



Hadoop Streaming API

Running a job with Hadoop Streaming

```
$ hadoop jar $HADOOP_HOME/share/hadoop/tools/lib/hadoop-streaming*.jar \
-files <mapCodePath>,<reduceCodePath> \
-input <inputFilePath> \
-output <outputFilePath> \
-mapper "python <mapCodePath>" \
-combiner "python <reduceCodePath>" \
-reducer "python <reduceCodePath>"
```

The jars for Hadoop
Streaming

Hadoop Streaming API

Running a job with Hadoop Streaming

```
$ hadoop jar $HADOOP_HOME/share/hadoop/tools/lib/hadoop-streaming*.jar \
  -files <mapCodePath>,<reduceCodePath> \
  -input <inputFilePath> \
  -output <outputFilePath> \
  -mapper "python <mapCodePath>" \
  -combiner "python <reduceCodePath>" \
  -reducer "python <reduceCodePath>"
```

The file paths for map.py and reduce.py

These files will be copied over to all the nodes
where computation is happening

Hadoop Streaming API

Running a job with Hadoop Streaming

```
$ hadoop jar $HADOOP_HOME/share/hadoop/tools/lib/hadoop-streaming*.jar \
-files <mapCodePath>,<reduceCodePath> \
  
-input <inputFilePath> \
-output <outputFilePath> \
-mapper "python <mapCodePath>" \
-combiner "python <reduceCodePath>" \
-reducer "python <reduceCodePath>"
```

The input text file path and output directory

Hadoop Streaming API

Running a job with Hadoop Streaming

```
$ hadoop jar $HADOOP_HOME/share/hadoop/tools/lib/hadoop-streaming*.jar \
-files <mapCodePath>,<reduceCodePath> \
-input <inputFilePath> \
-output <outputFilePath> \
-mapper "python <mapCodePath>" \
-combiner "python <reduceCodePath>" \
-reducer "python <reduceCodePath>"
```

We specify the file in which
the mapper is implemented

Hadoop Streaming API

Running a job with Hadoop Streaming

```
$ hadoop jar $HADOOP_HOME/share/hadoop/tools/lib/hadoop-streaming*.jar \
-files <mapCodePath>,<reduceCodePath> \
-input <inputFilePath> \
-output <outputFilePath> \
-mapper "python <mapCodePath>" \
-combiner "python <reduceCodePath>" \
-reducer "python <reduceCodePath>"
```

We specify the path for the files where combiner and reducer functions are implemented

Hadoop Streaming API

Running a job with Hadoop Streaming

```
$ hadoop jar $HADOOP_HOME/share/hadoop/tools/lib/hadoop-streaming*.jar \
-files <mapCodePath>,<reduceCodePath> \
-input <inputFilePath> \
-output <outputFilePath> \
-mapper "python <mapCodePath>" \
-combiner "python <reduceCodePath>" \
-reducer "python <reduceCodePath>"
```

For Word count problem, the combiner is the same as the reducer