

Say you have a **large**  
**amount of data**

You want to **explore it** and  
see if any **patterns** emerge

One way to do this is  
using a technique called

**Clustering**

# Clustering

## Anything

and I mean Anything in the world

can be described using

a set of numbers

# Clustering

**A person** (Age, Height, Weight)

**A sale** (Time, Amount, Product Id)

**A webpage** (Length, Frequencies of words)


# Clustering

A set of  $N$  numbers  
represents a point in  
an  **$N$ -Dimensional  
Hypercube**



# N-Dimensional Hypercube

A LINE IS A 1  
DIMENSIONAL  
SHAPE




Any point on a line  
can be represented  
using 1 number

# N-Dimensional Hypercube

A SQUARE IS A 2  
DIMENSIONAL  
SHAPE


Any point in a square  
can be represented  
using 2 numbers



# N-Dimensional Hypercube

A CUBE IS A 3  
DIMENSIONAL  
SHAPE

Any point in a cube  
can be represented  
with 3 numbers






# Clustering

A set of  $N$  numbers represents a point in an  **$N$ -Dimensional Hypercube**

This is just to say that any data in the world can be represented as points in an  $N$ -Dimensional space

# Clustering

THE OBJECTIVE OF  
CLUSTERING IS TO  
DIVIDE UP THESE USERS  
INTO GROUPS




Here is a set of  
points - each  
representing a  
Facebook User

USERS IN A GROUP  
ARE SIMILAR  
TO ONE ANOTHER

USERS FROM DIFFERENT  
GROUPS ARE  
VERY DIFFERENT FROM  
EACH OTHER

# Clustering


HERE'S ONE  
WAY OF DOING  
THIS..




USERS IN A GROUP  
ARE SIMILAR  
TO ONE ANOTHER

USERS FROM DIFFERENT  
GROUPS ARE  
VERY DIFFERENT FROM  
EACH OTHER

# Clustering




USERS IN A GROUP  
ARE SIMILAR  
TO ONE ANOTHER



HERE'S ANOTHER WAY ..

USERS FROM DIFFERENT  
GROUPS ARE  
VERY DIFFERENT FROM  
EACH OTHER

# Clustering



**BASICALLY, THE DISTANCE  
BETWEEN THE POINTS IS  
USED TO INDICATE THE  
SIMILARITY BETWEEN  
USERS**

**IN REAL LIFE THIS DISTANCE  
MIGHT TRANSLATE TO  
SOMETHING MEANINGFUL**

- 1. MAYBE THEY LIKE/FOLLOW  
THE SAME THINGS**
- 2. MAYBE THEY ARE FROM THE  
SAME STATE**
- 3. OR BOTH OF THE ABOVE**



# Clustering

THE CLUSTERS SHOULD BE SUCH THAT WE

USERS IN A GROUP ARE SIMILAR  
TO ONE ANOTHER

USERS FROM DIFFERENT GROUPS ARE  
VERY DIFFERENT FROM EACH OTHER

MAXIMIZE  
INTRACLUSTER  
SIMILARITY

MINIMIZE  
INTERCLUSTER  
SIMILARITY

# Clustering


MAXIMIZE  
INTRACLUSTER  
SIMILARITY

MINIMIZE  
INTERCLUSTER  
SIMILARITY

K-MEANS CLUSTERING IS  
A VERY FAMOUS  
ALGORITHM TO ACHIEVE  
EXACTLY THIS

# K-MEANS CLUSTERING


1. INITIALIZE A SET OF  
POINTS AS THE "K" MEANS  
(CENTROIDS OF THE  
CLUSTERS YOU WANT TO  
FIND)



YOU WANT TO DIVIDE THIS DATA INTO K CLUSTERS



# K-MEANS CLUSTERING



1 . INITIALIZE A SET OF POINTS  
AS THE "K" MEANS  
(CENTROIDS OF THE CLUSTERS  
YOU WANT TO FIND)


2 . ASSIGN EACH  
POINT TO THE  
CLUSTER BELONGING  
TO THE NEAREST  
MEAN

★ 3 . FIND THE NEW  
MEANS /CENTROIDS  
OF THE CLUSTERS

# K-MEANS CLUSTERING

RINSE AND REPEAT  
STEPS 2, 3 UNTIL THE  
MEANS DON'T  
CHANGE ANY MORE

1 . INITIALIZE A SET OF POINTS AS THE "K" MEANS  
(CENTROIDS OF THE CLUSTERS YOU WANT TO FIND)



2 . ASSIGN EACH POINT TO THE  
CLUSTER BELONGING TO THE NEAREST  
MEAN


3 . FIND THE NEW MEANS /  
CENTROIDS OF THE CLUSTERS

# K-MEANS CLUSTERING

## CONVERGENCE

RINSE AND REPEAT  
STEPS 2, 3 UNTIL THE MEANS DON'T  
CHANGE ANY MORE

1 . INITIALIZE A SET OF POINTS AS THE "K" MEANS  
(CENTROIDS OF THE CLUSTERS YOU WANT TO FIND)



2 . ASSIGN EACH POINT TO THE  
CLUSTER BELONGING TO THE NEAREST  
MEAN

3 . FIND THE NEW MEANS /  
CENTROIDS OF THE CLUSTERS

# K-MEANS CLUSTERING

Let's see how to  
achieve this using  
MapReduce

## K-MEANS CLUSTERING

One MapReduce job  
will run for each  
iteration that we just  
described

# K-MEANS CLUSTERING

There are a few questions we need to answer

1) What is the input and output of each iteration?

2) What are the Map and Reduce transformations in an iteration?

3) How do we specify the stopping condition?



# K-MEANS CLUSTERING

1) What is the input and output of each iteration?

In each iteration 1 MapReduce job will run

## Input

- 1) The current set of cluster centers
- 2) All the data points

## Output

- 1) A new set of cluster centers
- 2) Mapping of data point to Cluster center

# K-MEANS CLUSTERING

There are a few questions we need to answer

1) What is the input and output of each iteration?

2) What are the Map and Reduce transformations in an iteration?

3) How do we specify the stopping condition?



# K-MEANS CLUSTERING

There are a few questions we need to answer

1) What is the input and output of each iteration?

**2) What are the Map and Reduce transformations in an iteration?**

**3) How do we specify the stopping condition?**

# K-MEANS CLUSTERING

# Map and Reduce

## Input

- 1) The current set of cluster centers
- 2) All the data points



For each data point  $i$   
<nearest Cluster Center,  
data point  $i$ >

# K-MEANS CLUSTERING

# Map and Reduce

For each data point  $i$   
<nearest Cluster Center,  
data point  $i$ >

Sort/Merge



For each Cluster  
Center  
<Cluster Center,  
<List of  
assigned data  
points >>

# K-MEANS CLUSTERING

# Map and Reduce

For each Cluster  
Center

<Cluster Center,  
<List of  
assigned data  
points >>



For each data point

<New Cluster  
Center,  
<Data Point>>

# K-MEANS CLUSTERING

There are a few questions we need to answer

1) What is the input and output of each iteration?

**2) What are the Map and Reduce transformations in an iteration?**

**3) How do we specify the stopping condition?**



# K-MEANS CLUSTERING

There are a few questions we need to answer

1) What is the input and output of each iteration?

2) What are the Map and Reduce transformations in an iteration?

**3) How do we specify the stopping condition?**

# MapReduce Counters

**MapReduce allows users to  
define Counters that are set  
by the Mapper / Reducer**

# MapReduce Counters

The value of this  
counter at any time is  
accessible to the Job



# MapReduce Counters

In each iteration, if the  
new Cluster centers  $\neq$  old Cluster centers

The Reducer can  
increment a counter

# MapReduce Counters

**The Job will keep  
checking this to say if it  
wants to move on to  
another iteration or not**